

# Project **OBERON**

the **Design**  
of  
an Operating System (**OS**),  
a **Compiler**,  
and  
a **Computer**

Niklaus Wirth

Jürg Gutknecht

2013 revised

ISBN 0-201-54428-8



# CONTENTS

---

## PO

- o Preface 7
  - o.1 1<sup>st</sup> Edition 7
  - o.2 2<sup>nd</sup> Edition 11
- 1 Overview 19
  - 1.1 History & Motivation 19
  - 1.2 Overview 26

## I OS

- 2 Basic Concepts & System Structure 39
  - 2.1 Concepts 41
    - 2.1.1 Viewers 41
    - 2.1.2 Commands 44
    - 2.1.3 Tasks 48
    - 2.1.4 Tool Texts as Configurable Menus 52
    - 2.1.5 Extensibility 54
    - 2.1.6 Dynamic Loading 55
  - 2.2 The system structure 56
- 3 Task 63
  - 3.1 The Concept of Task 64
    - 3.1.1 Interactive Tasks 64
    - 3.1.2 Background Tasks 67

3.2	The Task Scheduler	69
3.3	The Concept of Command	74
3.3.1	Atomic Actions	74
3.3.2	Generic Text Selection	79
3.3.3	Generic Copy Viewer	79
3.4	Toolboxes	80
4	Display	87
4.1	Screen Layout Model	88
4.2	Viewers as Objects	93
4.3	Frames as Basic Display Entities	97
4.4	Display management	101
4.4.1	Viewers	103
4.4.2	Menu Viewers	112
4.4.3	Cursor Management	120
4.5	Raster Operations	126
4.6	Standard display configs and toolbox	133
5	Text	137
5.1	Text as an abstract data type	141
5.1.1	Loading and Storing Text	143
5.1.2	Editing Text	144
5.1.3	Accessing Text	147
5.2	Text Management	154
5.3	Text Frames	170
5.4	The Font Machinery	195
5.5	The Edit toolbox	201
6	Module Loader	203
6.1	Linking and loading	203

6.2	Module representation	206
6.3	The linking loader	210
6.4	The toolbox of the loader	213
6.5	The Oberon object file format	214
7	File System(FS)	217
7.1	Files	217
7.2	Implementation on a random-access store	225
7.3	Implementation on a disk	236
7.4	The file directory	246
7.5	The toolbox of file utilities	267
8	Storage Layout and Management	269
8.1	Layout and run-time organization	269
8.2	Heap Management	274
8.3	Kernel	286
8.4	The storage management toolbox	288
9	Device Driver(DD)	291
9.1	Overview	291
9.2	Keyboard and mouse	294
9.3	The SD-card (disk)	296
9.4	Serial asynchronous interface (RS 232)	297
9.5	Serial communications controller (SCC)	298
10	Compiler	299
10.1	Introduction	299
10.2	Code patterns	304
10.2.1	Assignment of constants	306
10.2.2	Simple expressions	307
10.2.3	Indexed variables	309

10.2.4	RECORD fields and pointers	312
10.2.5	Boolean expressions, IF statements	313
10.2.6	WHILE and REPEAT statements	316
10.2.7	FOR statements	317
10.2.8	Proper procedures	318
10.2.9	Functions	319
10.2.10	Dynamic ARRAYS	320
10.2.11	SETS	322
10.2.12	Imported variables and procedures	324
10.2.13	RECORD extensions with pointers	326
10.2.14	RECORD extensions as VAR parameters	329
10.2.15	ARRAY assignments and strings	331
10.2.16	Predeclared procedures	332
10.2.17	Predeclared functions	335
10.3	Internal data structures and module interfaces	336
10.3.1	Data structures	336
10.3.2	Module interfaces	343
10.4	The Parser	351
10.5	The scanner	357
10.6	Searching the symbol table, and handling symbol files	358
10.6.1	The structure of the symbol table	358
10.6.2	Symbol files	362
10.7	The code generator	370
10.7.1	Expressions	372
10.7.2	Relations	375

10.7.3	Set operations	376
10.7.4	Assignments	378
10.7.5	Conditional and repetitive statements	378
10.7.6	Boolean expressions	381
10.7.7	Procedures	382
10.7.8	Type extension	385
10.7.9	Import and export, global variables	390
10.7.10	Traps	392
11	RISC Processor Implementation	395
11.1	Introduction	395
11.2	The arithmetic and logic unit	399
11.2.1	Shifters	404
11.2.2	Multiplication	405
11.2.3	Division	411
11.3	Floating-point arithmetic	414
11.3.1	Floating-point addition	415
11.3.2	Floating-point multiplication	417
11.3.3	Floating-point division	418
11.4	The Control Unit	418
12	Processor's Environment	425
12.1	The SRAM memory	426
12.2	Peripheral interfaces	427
12.2.1	The PS/2 interface for the keyboard	428
12.2.2	The Mouse	430
12.2.3	The SPI interface for the SD-card (disk) and the Net	434
12.2.4	The display controller	436

12.2.5	The RS-232 interface	440
--------	----------------------	-----

## Spec

A	Oberon Programming Language	447
A.1	Introduction	447
A.1.1	Syntax	448
A.1.2	Vocabulary	448
A.2	Declarations	451
A.2.1	Constant	453
A.2.2	Type	453
A.2.3	Variable	459
A.3	Expressions	460
A.3.1	Operands	460
A.3.2	Operators	462
A.4	Statements	466
A.4.1	Statement Sequences	467
A.4.2	Assignments	467
A.4.3	Procedure Calls	468
A.4.4	IF	469
A.4.5	CASE	470
A.4.6	WHILE	472
A.4.7	REPEAT	472
A.4.8	FOR	473
A.5	Procedures	474
A.5.1	Parameters	475
A.5.2	Predefined	478
A.6	Modules	480
A.6.1	SYSTEM	482



B	Syntax of Oberon	485
---	------------------	-----



PO



## PREFACE

---

### 0.1 1<sup>st</sup> EDITION

This book presents the results of Project Oberon, namely an entire software environment for a modern workstation. The project was undertaken by the authors in 1986-89, and its primary goal was to design and implement an entire system from scratch, and to structure it in such a way that it can be described, explained, and understood as a whole. In order to become confronted with all aspects, problems, design decisions and details, the authors not only conceived but also programmed the entire system described in this book, and more.

Although there exist numerous books explaining principles and structures of OSes, there is a lack of descriptions for one actually implemented and used. We wished not only to give advice on how a system might be built, but to demonstrate how one was built. Program listings therefore play a key role in this text, because they alone contain the ultimate explanations. The choice of a suitable formalism therefore

assumed great importance, and we designed the language Oberon as not only an effective vehicle for implementation, but also as a publication medium for algorithms in the spirit in which Algol 60 had been created three decades ago. Because of its structure, the language Oberon is equally well suited to exhibit global, modular structures of programmed systems.

In spite of the small number of man-years spent on realizing the Oberon, and in spite of its compactness letting whose description fit a single book, it is not an academic toy, but rather a versatile workstation system that has found many satisfied and even enthusiastic users in academia and industry. The core system described here, consisting of storage, file, display, text, and viewer managers, of program loader and device drivers, draws its major power from a suitably chosen, flexible set of basic facilities and, most importantly, of their effective extensibility in many directions and for many applications. The latter is particularly enhanced by the language Oberon on the one, and by the efficiency of its basic core on the other hand. It is rooted in the application of the object-oriented paradigm which is employed wherever extensibility appears advantageous.

In addition to the core system, we describe in full detail the compiler for language Oberon and a graphics system, which both may be regarded as applications. The former reveals how a compact compiler is designed to achieve both fast

compilation and efficient, dense code. The latter stands as an example of extensible design based on object-oriented techniques, and it shows how a proper integration with an existing text system is possible. A network module is another addition to the core system allowing many workstations to be interconnected. We also show how the Oberon serves conveniently as the basis for a multi-server station, accommodating a file distribution, a printing, and an electronic-mail facility.

Compactness and regular structure, and due attention to efficient implementation of important details appear to be the key to economical software engineering. With the Oberon, we wish to refute Reiser's Law, which has been confirmed by virtually all recent releases of operating systems: In spite of great leaps forward, hardware being becoming faster is more slowly than software being slower. The Oberon has required a tiny fraction of the manpower demanded for the construction of widely-used commercial OSes, and a small fraction of their demands on computing power and storage capacity, while providing equal power and flexibility to the user, albeit without certain bells and whistles. The reader is invited to study how this was possible.

But most importantly, we hope to present a worth-while case study of a substantial piece of programming in large for the benefit of all those who are eager to learn from the experiences of others.

We wish to thank the many anonymous contributors of suggestions, advice, and encouragement. In particular we wish to thank our colleagues H. Mössenböck, B. Sanders and our associates at the Institut für Computer systeme for reading all or parts of the book draft. We are grateful to M. Brandis, R. Crelier, A. Disteli, M. Franz, and J. Templ for their work in porting the Oberon System successfully to various commercially available computers, and thus making the study of this book more worth-while for many readers. And we gratefully acknowledge the contribution of our school, ETH, for providing the environment and support which made it possible for us to pursue and complete this project.

Zürich, February 1992

N.W. and J.G.



0.2 2<sup>nd</sup> EDITION

Comments about plans to prepare a second edition to this book varied widely. Some felt that this book is outdated, that nobody is interested in a system of this kind any longer. "Why bother"? Others felt that there is an urgent need for this type of text, which explains an entire system in detail rather than merely proposing strategies and approaches. "By all means"!

Very much has changed in these last 30 years. But even without this change, it would be preposterous to propose and construct a system competing with existing, worldwide "standards". Indeed, very few people would be interested in using it. The community at large seems to be stuck with these gigantic software systems, and helpless against their complexity, peculiarities, and occasional unreliability.

But surely new systems will emerge, perhaps for different, limited purposes, allowing for smaller systems. One wonders where their designers will study and learn their trade. There is little technical literature, and my conclusion is that understanding is generally gained by doing, that is, "on the job". However, this is a tedious and sub-optimal way to learn. Whereas sciences are governed by principles and laws to be learned and understood, in engineering experience and practice are indispensable. Does Computer Science teach laws that hold for (almost) ever? More than any other field of engineering, it would be predestined to base on rigorous

mathematical principles. Yet, its core hardly is. Instead, one must rely on experience, that is, on studying sound examples.

The main purpose of and the driving force behind this project is to provide a single book that serves as an example of a system that exists, is in actual use, and is explained in all detail. This task drove home the insight that it is hard to design a powerful and reliable system, but even much harder to make it so simple and clear that it can be studied and fully understood. Above everything else, it requires a stern concentration on what is essential, and the will to leave out the rest, all the popular "bells and whistles".

Recently, a growing number of people has become interested in designing new, smaller systems. The vast complexity of popular OSES makes them not only obscure, but also provides opportunities for "back doors". They allow external agents to introduce spies and devils unnoticed by the user, making the system attack-able and corruptible. The only safe remedy is to build a safe system anew from scratch.

Turning now to a practical aspect: The largest chapter in the 1992 edition of this book dealt with the compiler translating Oberon programs into code for the NS32032 processor. Which is now neither available nor is its architecture recommendable. Instead of writing a new compiler for some other commercially available architecture, I decided to design my own in order to extend the desire for simplicity and regular-

ity to the hardware. The ultimate benefit of this decision is not only the software, but also the hardware of Oberon is described completely and rigorously. The processor is called RISC. The hardware modules are described exclusively in the language Verilog.

The decision for a new processor was expedited by the possibility to implement it, that is, to make it concrete and available. This is due to the advent of programmable gate arrays (FPGA), allowing to turn a design into a real, functioning processor on a single chip. As a result, the described system can be realized using a low-cost development board. Which, Xilinx Spartan-3 by Digilent, features a 1MB static memory, which easily accommodates the entire Oberon, including its compiler. It is shown, together with a display, keyboard and mouse in the photo below. In the lower, right corner, is the board.



The decision to develop our own processor required that the compiler and linking loader chapters had to be completely rewritten. However, it also provided the welcome chance to improve their clarity considerably. The new processor indeed allowed to simplify and straighten out the entire compiler.

For the description of a system to be comprehensible, the key element is the notation, formalism, or language in which it is defined. Algol 60, published 50 years ago, was proposed as a publication language, as a formalism in which algorithms could be defined without reference to particular computers, or to any mechanism at all. This was a great goal, but so far it was hardly achieved. Yet, it emphasized the importance of abstraction to be achieved by the notation with a mathematically rigorous foundation. At least, Algol was the first language based on a formally defined syntax. Algol was the result of early recognition that programs must never be written just to feed computers, but always to be understood and instructive to people.

In all my past work, I have tried to design a successor to Algol, that improves its rigor and at the same time extends its applicability from numerical algorithms to software systems. From a long sequence, starting with Algol, through Pascal, Modula, and Oberon, we have come closer to this goal than ever before, and than any other language in existence. The key lay in a continued struggle for sensible simplification.

The Oberon language, defined in 1988, underwent a revision in 2007, mostly discarding features that were either duplicated or not essential. Adaptation of the system's source code to the revised language was, besides the change of processor, the second important reason for numerous, local changes in this text. We summarize the various deletions of features:

1. The data types LONGINT, SHORTINT, and LONGREAL have been discarded, and with them the concept of type inclusion.
2. The LOOP and EXIT statements (repetitions with multiple exit points) have been discarded.
3. The WITH statement (regional type guard) has been discarded.
4. The RETURN statement has been removed and is now syntactically merged with the ending of function procedure declarations.
5. Objects declared in procedures are not accessible within their local counterparts. That is, objects must be either strictly local or global in order to be accessible.
6. Assignments to imported variables are not permitted (read-only export).
7. Forward procedure declarations have been discarded.

In contrast to these removals, there is a single addition (made in 2012):

- The data type BYTE has been added. Its values are integers  $x$  satisfying  $0 \leq x < 256$ . This addition prevents the frequent abuse of the type CHAR. BYTE is mainly used for elements of arrays and records in low level modules in order to economize the use of memory.

In spite of these two reasons for changes – one at the highest level, language, the other at the lowest, hardware – the remainder of this book proved to be pretty stable and still valid. It has been my desire to present the system essentially as it existed 25 years ago, without embellishments. The chapters 3 - 5 on tasking, the display and text, originally written by J. Gutknecht, have been carried over virtually unchanged. Significant changes, however, were necessary mainly in the descriptions of device drivers for keyboard and mouse. They now use the PS-2 interface standard. The disk has been replaced by a single SD-card (flash memory) with a standard SPI interface. The net interface no longer uses RS-485, but is also based on the SPI standard. Chapters on the compiler and linker are completely new.

Mostly thanks to the regularity of RISC instruction set, the compiler size could be reduced significantly. It now measures less than 2900 lines of program and compiles itself in about 3 seconds, which is proof of its efficiency. The entire system compiles itself in less than 10 seconds.

Considered extravagant and hardly necessary only years ago, run-time checks are generated automatically. In particular, they cover index range checks and access to NIL-pointers. Due to their efficiency they hardly affect run-time speed, but are a great benefit to programmers.

A welcome consequence of the language and processor simplifications is the fact that all parts had been written in assembler code in 1992 – and therefore were not included in the book – have now been expressed in Oberon as well. Vindicating my perennial efforts to obtain a high-level language which is powerful, flexible, and also efficient enough to express parts such as device drivers and raster operations, this was the necessary and final step to make this book comprehensive and complete.

#### REFERENCES

1. [Oberon 07 Report pdf](#)
2. [RISC Arch pdf](#)

## ACKNOWLEDGMENTS

I gratefully acknowledge the valuable contributions of Paul Reed. He designed the interfaces to various devices, such as the PS-2 and SPI, including the SD-card, acting as disk store. He suggested many improvements and simplifications. He originally decisively suggested a re-edition of this book of 30 year ago, and was the key impetus to do all this work. My thanks go to him.

Niklaus Wirth, September 2013



## OVERVIEW

---

### 1.1 HISTORY & MOTIVATION

How could anyone diligently concentrate on his work in an afternoon with such warmth, splendid sunshine, and blue sky. This rhetorical question was one I asked many times while spending a sabbatical leave in California in 1985. Back home everyone would feel compelled to profit from the sunny spells to enjoy life at leisure in the country-side, wandering or engaging in one's favourite sport. But here, every day was like that, and giving in to such temptations would have meant the end of all work. And, had I not chosen this location in the world because of its inviting, enjoyable climate?

Fortunately, my work was also enticing, making it easier to buckle down. I had the privilege of sitting in front of the most advanced and powerful workstation anywhere, learning the secrets of perhaps the newest fad in our fast developing trade, pushing colored rectangles from one place of the screen to another. This all had to happen under strict

observance of rules imposed by physical laws and by the newest technology. Fortunately, the advanced computer would complain immediately if such a rule was violated, it was a rule checker and acted like your big brother, preventing you from making steps towards disaster. And it did what would have been impossible for oneself, keeping track of thousands of constraints among the thousands of rectangles laid out. This was called computer-aided design. "Aided" is rather a euphemism, but the computer did not complain about the degradation of its role.

While my eyes were glued to the colorful display, and while I was confronted with the evidence of my latest inadequacy, in through the always open door stepped my colleague (JG). He also happened to spend a leave from duties at home at the same laboratory, yet his face did not exactly express happiness, but rather frustration. The chocolate bar in his hand did for him what the coffee cup or the pipe does for others, providing temporary relaxation and distraction. It was not the first time he appeared in this mood, and without words I guessed its cause. And the episode would reoccur many times.

His days were not filled with the great fun of rectangle-pushing; he had an assignment. He was charged with the design of a compiler for the same advanced computer. Therefore, he was forced to deal much more closely, if not intimately, with the underlying software system. Its rather

frequent failures had to be understood in his case, for he was programming, whereas I was only using it through an application; in short, I was an end-user! These failures had to be understood not for purposes of correction, but in order to find ways to avoid them. How was the necessary insight to be obtained? I realized at this moment that I had so far avoided this question; I had limited familiarization with this novel system to the bare necessities which sufficed for the task on my mind.

It soon became clear that a study of the system was nearly impossible. Its dimensions were simply awesome, and documentation accordingly sparse. Answers to questions that were momentarily pressing could best be obtained by interviewing the system's designers, who all were in-house. In doing so, we made the shocking discovery that often we could not understand their language. Explanations were fraught with jargon and references to other parts of the system which had remained equally enigmatic to us.

So, our frustration-triggered breaks from compiler construction and chip design became devoted to attempts to identify the essence, the foundations of the system's novel aspects. What made it different from conventional OSes? Which of these concepts were essential, which ones could be improved, simplified, or even discarded? And where were they rooted? Could the system's essence be distilled and extracted, like in a chemical process?

During the ensuing discussions, the idea emerged slowly to undertake our own design. And suddenly it had become concrete. "Crazy" was my first reaction, and "impossible". The sheer amount of work appeared as overwhelming. After all, we both had to carry our share of teaching duties back home. But the thought was implanted and continued to occupy our minds.

Sometime thereafter, events back home suggested that I should take over the important course about System Software. As it was the unwritten rule that it should primarily deal with operating system principles, I hesitated. My scruples were easily justified: After all I had never designed such a system nor a part of it. And how can one teach an engineering subject without first-hand experience?

Impossible? Had we not designed compilers, OSes, and document editors in small teams? And had I not repeatedly experienced that an inadequate and frustrating program could be programmed from scratch in a fraction of source code used by the original design? Our brainstorming continued, with many intermissions, over several weeks, and certain shapes of a system structure slowly emerged through the haze. After some time, the preposterous decision was made: we would embark on the design of an OS for our workstation (which happened to be much less powerful than the one used for my rectangle-pushing) from scratch.

The primary goal, to personally obtain first-hand experience, and to reach full understanding of every detail, inherently determined our manpower: two part-time programmers. We tentatively set our time-limit for completion to three years. As it later turned out, this had been a good estimate; programming was begun in early 1986, and a first version of the system was released in the fall of 1988.

Although the search for an appropriate name for a project is usually a minor problem and often left to chance and whim of the designers, this may be the place to recount how Oberon entered the picture in our case. It happened that around the time of the beginning of our effort, the space probe Voyager made headlines with a series of spectacular pictures taken of the planet Uranus and of its moons, the largest of which is named Oberon. Since its launch I had considered the Voyager project as a singularly well-planned and successful endeavor, and as a small tribute to it I picked the name of its latest object of investigation. There are indeed very few engineering projects whose products perform way beyond expectations and beyond their anticipated lifetime; mostly they fail much earlier, particularly in the domain of software. And, last but not least, we recall that Oberon is famous as the king of elfs.

The consciously planned shortage of manpower enforced a single, but healthy, guideline: Concentrate on essential functions and omit embellishments that merely cater to

established conventions and passing tastes. Of course, the essential core had first to be recognized and crystallized. But the basis had been laid. The ground rule became even more crucial when we decided that the result should be able to be used as teaching material. I remembered C.A.R. Hoare's plea that books should be written presenting actually operational systems rather than half baked, abstract principles. He had complained in the early 1970s that in our field engineers were told to constantly create new artifacts without being given the chance to study previous works that had proven their worth in the field. How right was he, even to the present day!

The emerging goal to publish the result with all its details let the choice of programming language (PL) appear in a new light: it became crucial. Modula-2 which we had planned to use, appeared as not quite satisfactory. Firstly, it lacked a facility to express extensibility in an adequate way. And we had put extensibility among the principal properties of the new system. By "adequate" we include machine-independence. Our programs should be expressed in a manner that makes no reference to machine peculiarities and low-level programming facilities, perhaps with the exception of device interfaces, where dependence is inherent.

Hence, Modula-2 was extended with a feature that is now known as type extension. We also recognized that Modula-2 contained several facilities that we would not need, that

do not genuinely contribute to its power of expression, but at the same time increase the complexity of the compiler. But the compiler would not only have to be implemented, but also to be described, studied, and understood. This led to the decision to start from a clean slate also in the domain of language design, and to apply the same principle to it: concentrate on the essential, purge the rest. The new language, which still bears much resemblance to Modula-2, was given the same name as the system: Oberon [1, 2]. In contrast to its ancestor it is terser and, above all, a significant step towards expressing programs on a high level of abstraction without reference to machine-specific features.

We started designing the system in late fall 1985, and programming in early 1986. As a vehicle we used our workstation Lilith and its language Modula-2. First, a cross-compiler was developed, then followed the modules of the inner core together with the necessary testing and down-loading facilities. The development of the display and the text system proceeded simultaneously, without the possibility of testing, of course. We learned how the absence of a debugger, and even more so the absence of a compiler, can contribute to careful programming.

Thereafter followed the translation of the compiler into Oberon. This was swiftly done, because the original had been written with anticipation of the later translation. After its availability on the target computer Ceres, together with

the operability of the text editing facility, the umbilical cord to Lilith could be cut off. The Oberon had become real, at least its draft version. This happened around the middle of 1987; its description was published thereafter [3], and a manual and guide followed in 1991 [5].

The system's completion took another year and concentrated on connecting the workstations in a network for file transfer [4], on a central printing facility, and on maintenance tools. The goal of completing the system within three years had been met. The system was introduced in the middle of 1988 to a wider user community, and work on applications could start. A service for electronic mail was developed, a graphics system was added, and various efforts for general document preparation systems proceeded. The display facility was extended to accommodate two screens, including color. At the same time, feedback from experience in its use was incorporated by improving existing parts. Since 1989, Oberon has replaced Modula-2 in our introductory programming courses.

## 1.2 OVERVIEW

This book consists of 3 main parts, excluding this preface and overview(PO) part:

1. OS
2. apps



### 3. hardware computer

First, OS. Implementation of a system proceeds bottom-up naturally, because higher level modules are clients of those lower level and cannot function without their imports. Description, on the contrary, should better be arranged in the reverse top-down way. This is because a system is designed with its expected applications and functions in mind. Decomposition into a hierarchy of modules is justified by the use of auxiliary functions and abstractions, or by postponing more detailed explanation later, when the need has been fully motivated. Thus, we will describe the OS part top-down essentially.

2 first explains the most important basic concepts like Viewers, Commands, Tasks and Texts, etc. which will be further elaborated in the following chapters. It ends with the system structure, to share readers a high-level overview understanding to the whole system.

Chapters 3 - 5 describe the outer core system:

3 focuses on the dynamic aspects. In particular, it introduces the fundamental operational units of task and command. Oberon's tasking model distinguishes the categories of interactive tasks and background tasks. The former are represented on the display screen by rectangular areas, so-called viewers. Yet the latter need not be connected with any displayed object. They are scheduled with low priority when

interactions are absent. A good example of a background task is the memory garbage collecting. Both of them are mapped to a single process by the task scheduler. Commands in Oberon are explicit, atomic units of interactive operations. They are realized in the form of exported parameterless procedures and replace the heavier-weight notion of program known from more conventional OSes. This chapter continues with a definition of a software toolbox as a logically connected collection of commands. It terminates with an outline of the system control toolbox.

4 explains Oberon's display system. It starts with a discussion of our choice of a hierarchical tiling strategy for the allocation of viewers. A detailed study of the exact role of Oberon viewers follows. Type Viewer is presented as an object class with an open message interface providing a conceptual basis for far-reaching extensibility. Viewers are then recognized as just a special case of so-called frames that may be nested. A category of standard viewers containing a menu frame and a frame of contents is investigated. The next topic is cursor handling. A cursor in Oberon is a marked path. Both viewer manager and cursor handler operate on an abstract logical display area rather than on individual physical monitors. This allows a unified handling of display requests, independent of number and types of monitors assigned. For example, smooth transitions of the cursor across screen boundaries are conceptually guaranteed. The chapter continues with the presentation of a

concise and complete set of raster operations that is used to place textual and graphical elements in the display area. An overview of the system display toolbox concludes the chapter.

5 introduces Text. Oberon distinguishes itself by treating Text as an abstract data type that is integrated in the central system. Numerous fundamental consequences are discussed. For example, a text can be produced by one command, edited by a user, and then consumed by a next command. Commands themselves can be represented textually in the form M.P, followed by a textual parameter list. Consequently, any command can be called directly from within a text (so called tool) simply by pointing at it with the mouse. However, the core of this chapter is a presentation of Oberon's text system as a case study in program modularization. The concerns of managing a text and displaying it are nicely separated. Both the text manager and the text display feature an abstract public interface as well as an internally hidden data structure. Finally in this chapter, Oberon's type-font management and the toolbox for editing are discussed.

Chapters 6 - 9 describe the inner core system:

6 explains the loader of modules and motivates the intro of data type *Module*. The chapter includes the management of the memory part holding program code and defines the format in which compiled modules are stored as object files. Furthermore, it discusses the problems of binding

separately compiled modules together and of referencing objects defined in other modules.

7 is devoted to the file system(FS), a part of crucial importance, because files are involved in almost every program and computation. The chapter consist of two distinct parts, the first introducing the type File and describing the structure of files, i.e. their representation on disk storage with its sequential characteristics, the second describing the directory of file names and its organisation as a B-tree for obtaining fast searches.

8 the memory management. A single, central storage management was one of the key design decisions, guaranteeing an efficient and economical use of storage. The chapter explains the store's partitioning into specific areas. Its central concern, however, is the discussion of dynamic storage management in the partition called the heap. The algorithm for allocation (corresponding to the intrinsic procedure NEW) and for retrieval (called garbage collection) are explained in detail.

9 describes the lowest level of the module hierarchy: device drivers, which contains drivers for some widely accepted interface standards. The first is PS-2, a serial transmission with synchronous clock. This is used for the keyboard and for the Mouse. The second is SPI, a standard for bi-directional, serial transmission with synchronous clock. This is used for the "disk", represented by an SDI-card (flash

memory), and for the network. And the third standard is RS-232 typically used for simple and slow data links. It is bidirectional and asynchronous.

The second part, consisting of Chapters 10 - 15, is devoted to what may be called first applications of the basic Oberon. These chapters are therefore independent to each other, only making reference to the upper Chapters 3 - 9.

Although the Oberon is well-suited for operating stand-alone workstations, a facility for connecting a set of computers should be considered as fundamental. Module *Net*, which makes transmission of files among workstations connected by a bus-like network possible, is the subject of ?? . It presents not only the problems of network access, of transmission failures and collisions, but also those of naming partners. The solutions are implemented in a surprisingly compact module which uses a network driver presented in

9.

When a set of workstations is connected in a network, the desire for a central server appears. A central facility serving as a file distribution service, as a printing station, and as a storage for electronic mail is presented in ?? . It emerges by extending the *Net* module of ?? , and is a convincing application of the tasking facilities explained in Section 2.2. In passing we note that the server operates on a machine that is not under observation by a user. This circumstance requires an increased degree of robustness, not only against trans-

mission failures, but also against data that do not conform to defined formats.

The presented system of servers demonstrates that Oberon's single-thread scheme need not be restricted to single-user systems. The fact that every command or request, once accepted, is processed until completion, is acceptable if the request does not occupy the processor for too long, which is mostly the case in the presented server applications. Requests arriving when the processor is engaged are queued. Hence, the processor handles requests one at a time instead of interleaving them which, in general, results in faster overall performance due to the absence of frequent task switching.

[10](#) describes the Oberon compiler. It translates source text in Oberon into target code, i.e. instruction sequences of some target computer. Its principles and techniques are explained in [6]. Both, source language and target architecture must be understood before studying a compiler. Both source language and the target computer's RISC architecture are presented in the Appendix.

Although here the compiler appears as an application module, it naturally plays a distinguished role, because the system (and the compiler itself) is formulated in the language which the compiler translates into code. Together with the text editor it was the principal tool in the system's development. The use of straight-forward algorithms for parsing

and symbol table organization led to a reasonably compact piece of software. A main contributor to this result is the language's definition: the language is devoid of complicated structures and rarely used embellishments.

The compiler and thereby the chapter is partitioned into two main parts. The first is language specific, but does not refer to any particular target computer. It consist of the scanner and the parser. This part is therefore of most general interest to the readership. The second part is, essentially, language-independent, but is specifically tailored to the instruction set of the target computer. It is called the code generator.

Texts play a predominant role in the Oberon. Their preparation is supported by the system's major tool, the editor. In ?? we describe another one, which handles graphic objects. At first, only horizontal or vertical lines and short captions are introduced as objects. The major difference to texts lies in the fact that their coordinates in the drawing plane do not follow from those of their predecessor automatically, because they form a set rather than a sequence. Each object carries its own, independent coordinates. The influence of this seemingly small difference upon an editor are far-reaching and permeate the entire design. There exist hardly any similarities between a text and a graphics editor. Perhaps one should be mentioned: the partitioning into three parts. The bottom module defines the respective abstract data structure for texts or graphics, together with, of course, the proce-

dures handling the structure, such as searches, insertions, and deletions. The middle module in the hierarchy defines a respective frame and contains all procedures concerned with displaying the respective objects including the frame handler defining interpretation of mouse and keyboard events. The top modules are the respective tool modules (Edit, Draw). The presented graphics editor is particularly interesting in so far as it constitutes a convincing example of Oberon's extensibility. The graphics editor is integrated into the entire system; it embeds its graphic frames into menu-viewers and uses the facilities of the text system for its caption elements. And lastly, new kinds of elements can be incorporated by the mere addition of new modules, i.e. without expanding, even without recompiling the existing ones. Two examples are shown in ?? itself: rectangles and circles.

The Draw System has been extensively used for the preparation of diagrams of electronic circuits. This application suggests a concept that is useful elsewhere too, namely a recursive definition of the notion of object. A set of objects may be regarded as an object itself and be given a name. Such an object is called a macro. It is a challenge to the designer to implement a macro facility such that it is also extensible, i.e. in no way refers to the type of its elements, not even in its input operations of files on which macros are stored.



?? presents two other tools, namely one used for installing an Oberon on a bare machine, and one used to recover from failures of the file store. Although rarely employed, the first was indispensable for the development of the system. The maintenance or recovery tools are invaluable assets when failures occur. And they do! ?? covers material that is rarely presented in the literature.

?? is devoted to tools that are not used by the Oberon presented so far, but may be essential in some applications. The first is a data link with a protocol based on the RS-232 standard shown in 9. Another is a standard set of basic mathematical functions. And the third is a tool for creating new macros for the Draw System.

The last part is a detailed description of the hardware:

11 defines the processor, for which the compiler generates code. The target computer is a truly simple and regular processor called RISC with only 14 instructions, represented not by a commercial processor, but implemented with an FPGA, a Field Programmable Gate Array. It allows its structure to be described in full detail. It is a straight-forward, von Neumann type device consisting of a register bank, an arithmetic-logic unit, including a floating-point unit. Typical optimization facilities, like pipelining and cache memory, have been omitted for the sake of transparency and simplicity. The processor circuit is described in the language Verilog.

12 describes the environment in which the processor is embedded. This environment consists of the interfaces to main memory and to all external devices.

#### REFERENCES

1. N. Wirth. The PL Oberon. *Software - Practice and Experience* 18, 7, (July 1988) 671-690.
2. M. Reiser and N. Wirth. *Programming in Oberon - Steps beyond Pascal and Modula*. AddisonWesley, 1992.
3. N. Wirth and J. Gutknecht. The Oberon. *Software - Practice and Experience*, 19, 9 (Sept. 1989), 857-893.
4. N. Wirth. Ceres-Net: A low-cost computer network. *Software - Practice and Experience*, 20, 1 (Jan. 1990), 13-24.
5. M. Reiser. *The Oberon - User Guide and Programmer's Manual*. Addison-Wesley, 1991.
6. N. Wirth. *Compiler Construction*. Addison-Wesley, Reading, 1996. ISBN 0-201-40353-6

Part I

OS



## BASIC CONCEPTS & SYSTEM STRUCTURE

---

In order to warrant the sizeable effort designing and constructing an entire OS from scratch, basic concepts need be novel. The principal concepts underlying Oberon and the dominant design decisions are 1st discussed. Upon this, the system structure was then presented, restricted to the coarsest, the composition and interdependence of the largest building blocks, *modules*. Finally, an overview, helps understanding the place, role, and significance of each part (chapter).

The fundamental OS objective is to present the computer at a certain abstraction level to the user and programmer. For example,

STORE requestable *variables* of a specified data type,

DISK sequences of characters or bytes called *files*,

DISPLAY rectangular areas called *viewers*,

KEYBOARD an *input stream* of characters, and

MOUSE a pair of *coordinates* and a set of *key states*.

Every abstraction is characterized by certain *properties* and governed by a set of *operations*. It is the *task* of OS to implement these operations and to manage them, constrained by the available resources of the underlying computer. This is commonly called *resource management* (RM).

Every abstraction inherently hides details, from which it abstracts. Hiding occurs at different levels. For example, the

- computer may make certain store parts or devices inaccessible according to operation modes (user/supervisor);
- PL may make certain parts inaccessible through a facility inherent in visibility rules.

The latter is of course much more flexible and powerful, while the former indeed plays an almost negligible role in our system. Hiding is important because it allows maintenance of certain properties (called *invariants*) of an abstraction to be guaranteed. Abstraction is indeed the key of any modularization, and without it every hope of being able to guarantee reliability and correctness vanishes. Clearly, Oberon was designed with the goal of establishing a modular structure on the basis of purpose-oriented abstractions. The availability of an appropriate PL is an indispensable prerequisite, and the importance of prudent choice cannot be over-emphasized.

## 2.1 CONCEPTS

### 2.1.1 *Viewers*

Whereas the abstractions of individual variables representing parts of the primary store, and of files representing parts of the disk store are well established notions and have significance in every computer system, abstractions regarding input and output devices became important with the advent of high interactivity between user and computer. High interactivity requires high bandwidth, and the only channel of human users with high bandwidth is the eye. Consequently, the computer's visual output unit must be properly matched with human eyes. This occurred with the advent of the *high-resolution display* (HRD) in the mid 1970s, which in turn had become feasible due to faster and cheaper electronic memory components. The HRD marked one of the few very significant break-throughs in the history of computer development. The typical bandwidth of a modern display is in the order of 100 MHz. Primarily the HRD made visual output a subject of abstraction and RM. In Oberon, the display is partitioned into *viewers*, also called windows, or more precisely, *frames*, rectangular areas of the screen(s). A viewer typically consists of:

**TITLE BAR** contains a subject name and commands menu,

**MAIN FRAME** some text, graphic, video, or other object.

Viewers (frames) can be nested, in principle to any depth.

System provides routines for generating, moving, and closing a viewer. It allocates a new one at a specified place, and upon request delivers hints as to where it might best be placed. It keeps track of their opened set. These are called *viewer management*, in contrast to contents handling.

High interactivity requires not only high visual output bandwidth but also input flexibility. Surely, there is no need for an equally high bandwidth, but a **keyboard** limited by the speed of typing to about 100 Hz is not good enough. The break-through on this front was achieved by a pointing device, **mouse**, appeared roughly the same time as HRD.

This was by no means just a lucky coincidence. The mouse comes to fruition only through HRD and appropriate software. It is itself a conceptually very simple device delivering signals when moved on the table. These signals allow the computer to update the position of a mark - the cursor - on the display. Since feedback occurs through human eyes, no great precision is required. For example, when the user wishes to identify a certain object, e.g. a letter on the screen, one moves the mouse as far as required until the mapped cursor reaches the object. This stands in marked contrast to a digitizer which is supposed to deliver exact coordinates. Oberon relies highly on it.



Perhaps the cleverest idea was to equip mice with buttons. By being able to signal a request with the same hand that determines the cursor position, the user obtains the direct impression of issuing position-dependent requests. Position-dependence is realized in software by delegating interpretation of the signal to a procedure - a so-called handler or interpreter - which is local to the viewer in whose area the cursor momentarily appears. A surprising flexibility of command activation can be achieved in this manner by appropriate software. Various techniques have emerged in this connection, e.g. pop-up menus, pull-down-menus, etc. which are powerful even under the presence of a single button only. For many applications, a mouse with several keys is far superior, and Oberon basically assumes 3 buttons to be available. The assignment of different functions to the keys may of course easily lead to confusion when every application prescribes different key assignment. This is, however, easily avoided by the adherence to certain "global" conventions. In Oberon, button of the:

LEFT primarily used for marking a position (setting a caret),

MIDDLE for issuing general commands (see below), and

RIGHT for selecting displayed objects.

Recently, it has become fashionable to use overlapping windows mirroring documents being piled up on the desktop. We have found this metaphor not entirely convincing. Par-

tially hidden windows are typically brought to the top and made fully visible before any operation is applied to their contents. In contrast to the insignificant advantage stands the substantial effort necessary to implement this scheme. It is a good example of a case where the benefit of a complication is incommensurate with its cost. Therefore, we have chosen a solution that is much simpler to realize yet has no genuine disadvantages: *tiled viewers* as shown in Fig 1:

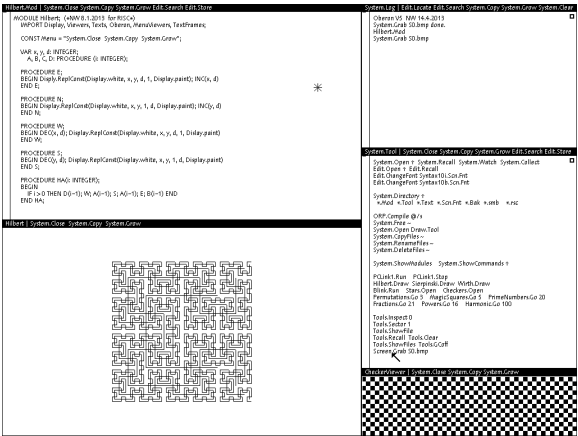


Figure 1: Oberon display with tiled viewers

2.1.2 Commands

Position-dependent commands with meaning fixed for each type of viewer must be supplemented by general commands. Conventionally, such commands are issued through the keyboard by typing the program's name that is to be executed

into a special command text. In this respect, Oberon offers a novel and much more flexible solution:

1<sup>st</sup> of all we remark that a program in the common sense of a text compiled as a unit is mostly a far too large action unit to serve as a command. Compare it with, for example, the insertion of a piece of text through a mouse command. In Oberon, the notion of a action unit is separated from the notion of compilation unit. The former is a command represented by a (exported) procedure, the latter a module. Hence, a module may, and typically does, define several, even many commands. Such a (general) command may be invoked at any time by pointing at its name in any text visible in any viewer on the display, and by clicking the middle mouse button. The command name has the form  $M.P$ , where  $P$  is the procedure's identifier and  $M$  the module in which  $P$  is declared. As a consequence, any command click may cause the loading of one or several modules, if  $M$  is not already present in main store. The next invocation of  $M.P$  occurs instantaneously, since  $M$  is already loaded. A further consequence is that modules are never (automatically) removed, because a next command may well refer to the same module.

Every command has the purpose to alter the state of some operands. Typically, they are denoted by text following the command identification, and Oberon fol-

lows this convention. Strictly speaking, commands are denoted as parameterless procedures; but the system provides a way for the procedure to identify the text position of its origin, and hence to read and interpret the text following the command, i.e. the actual parameters. Both reading and interpretation must, however, be programmed explicitly.

The parameter text must refer to objects that exist before command execution starts and are quite likely the result of a previous command interpretation. In most OSes, these objects are files registered in the directory, and they act as interfaces between commands. Oberon broadens this notion; the links between consecutive commands are not restricted to files, but can be any global variable, because modules do not disappear from storage after command termination, as mentioned above.

This tremendous flexibility seems to open Pandora's box, and indeed it does when misused. The reason is that global variables' states may completely determine and alter the effect of a command. The variables represent hidden states, hidden in the sense that the user is in general unaware of them and has no easy way to determine their value. The positive aspect of using global variables as interfaces between commands is that some of them may well be visible on the display.

All viewers - and with them also their contents - are organized in a data structure that is rooted in a global variable (in `Viewers`). Parts of this variable therefore constitute visible states, and it is highly appropriate to refer to them as command parameters.

One of the rules of what may be called the Oberon Programming Style is therefore to avoid hidden states, and to reduce the introduction of global variables. We do not, however, raise this rule to the rank of a dogma. There exist genuinely useful exceptions, even if the variables have no visible parts.

There remains the question of how to denote visible objects as command parameters. An obvious case is the use of the most recent selection as parameter. A procedure for locating that selection is provided by `Oberon`. (It is restricted to text selections). Another possibility is the use of the caret position in a text. This is used in the case of inserting new text; the pressing of a key on the keyboard is also considered to be a command, and it causes the character's insertion at the caret position.

A special facility is introduced for designating viewers as operands: the star marker. It is placed at the cursor position when the keyboard's mark key (`SETUP`) is pressed. The procedure `Oberon.MarkedViewer` identifies the viewer in whose area the star lies. Commands which take it as their parameter are typically followed

by an asterisk in the text. Whether the text in a text viewer, a graph in a graphic viewer, or any other part of the marked viewer is taken as the actual parameter depends on how the command is programmed.

Finally a most welcome property of the system should not remain unmentioned. It is a direct consequence of the persistent nature of global variables and becomes manifest when a command fails. Detected failures result in a trap that should be regarded as an abnormal command termination. In the worst case, global data may be left in an inconsistent state, but they are not lost, and a next command can be initiated based on their current state. A trap opens a small viewer and lists the sequence of invoked procedures with their local variables and current values. This information helps a programmer to identify its cause.

### 2.1.3 *Tasks*

From the presentations above it follows that Oberon is distinguished by a highly flexible scheme of command activation. The notion of a command extends from the insertion of a single character and the setting of a marker to computations that may take hours or days. It is moreover distinguished by a highly flexible notion of operand selection not restricted to registered, named files. And most importantly, by the

virtual absence of hidden states. The state of the system is practically determined by what is visible to the user.

This makes it unnecessary to remember a long history of previously activated commands, started programs, entered modes, etc. Modes are in our view the hallmark of user-unfriendly systems. It should at this point have become obvious that the system allows a user to pursue several different tasks concurrently. They are manifest in the form of viewers containing texts, graphics, or other displayable objects. The user switches between tasks implicitly when choosing a different viewer as operand for the next command. The characteristic of this concept is that task switching is under explicit control of the user, and the atomic units of action are the commands.

At the same time, we classify Oberon as a single-process (or single-thread) system. How is this apparent paradox to be understood? Perhaps it is best explained by considering the basic mode of operation. Unless engaged in the interpretation of a command, the processor is engaged in a loop continuously polling event sources. This loop is called the *central loop*; it is contained in Oberon which may be regarded as the system heart. The 2 fixed event sources are the mouse and the keyboard. If a keyboard event is sensed, control is dispatched to the handler installed in the so-called *focus viewer*, designated as the one holding the caret. If a mouse event (key) is sensed, control is dispatched to the handler

in which the cursor currently lies. This is all possible under the paradigm of a single, uninterruptible process.

The notion of a single process implies non-interruptability, and therefore also that commands cannot interact with the user. Interaction is confined to the selection of commands before their execution. Hence, there exists no input statement in typical Oberon programs. Inputs are given by parameters supplied and designated before command invocation.

This scheme at first appears as gravely restrictive. In practice it is not, if one considers single-user operation. It is this single user who carries out a dialog with the computer. A human might be capable of engaging in simultaneous dialogs with several processes only if the commands issued are very time-consuming. We suggest that execution of time-consuming computations might better be delegated to loosely coupled servers in a distributed system.

The primary advantage of a system dealing with a single process is that task switches occur at user-defined points only, where no local process state has to be preserved until resumption. Furthermore, because the switches are user-chosen, the tasks cannot interfere in unexpected and uncontrollable ways by accessing common variables. The system designer can therefore omit all kinds of protection mechanisms excluding such interference. This is significant simplification.



The essential difference between Oberon and multiprocess-systems is that in Oberon task switches occur between commands only, whereas in the latter switches may be invoked after any single instruction. Evidently, the difference is one of action granularity. Oberon's coarse granularity is entirely acceptable for a single-user system.

The system offers the possibility to insert further polling commands in the central loop. This is necessary if additional event sources are to be introduced. The prominent example is a network, where commands may be sent from other workstations. The central loop scans a list of so-called *task descriptors*. Each descriptor refers to a command procedure. The 2 standard events are selected only if their guard permits, i.e. if either keyboard input is present, or if a mouse event occurs. Inserted tasks must provide their own guard in the beginning of the installed procedure.

The example of a network inserting commands, called requests, raises a question: what happens if the processor is engaged in the execution of another command when the request arrives? Evidently, the request would be lost unless measures are taken. The problem is easily remedied by buffering the input. This is done in every input device driver, including the keyboard as well as the network drivers. The incoming signal triggers an interrupt, and the invoked interrupt handler accepts the input and buffers it. We emphasize that such interrupt handling is confined to drivers, system

components at the lowest level. An interrupt does not evoke a task selection and a task switch. Control simply returns to the point of interruption, and the interrupt remains unnoticeable to programs. There exists, as with every rule, an exception: an interrupt due to keyboard input of the abort character returns control to the central loop.

#### 2.1.4 *Tool Texts as Configurable Menus*

Certainly, the concepts of viewers specifying their own interpretation of mouse clicks, of commands invokable from any text on the display, of any displayed object being selectable as an interface between commands, and of commands being dialog-free, uninterruptible units of action, have considerable influence on the style of programming in Oberon, and they thoroughly change the style of using the computer. The ease and flexibility in the way pieces of text can be selected, moved, copied, and designated as command and parameters, drastically reduces the need for typing. The mouse becomes the dominant input device: the keyboard merely serves to input textual data. This is accentuated by the use of so-called *tool texts*, compositions of frequently used commands, which are typically displayed in the narrower system track of viewers. One simply doesn't type commands! They are usually visible somewhere already. Typically, the user composes a tool text for every project pursued. Tool texts can be regarded as individually configurable private menus.

The rarity of issuing commands by typing them has the most agreeable benefit that their names can be meaningful words. For example, the copy operation is denoted by `Copy` instead of `cp`, rename by `Rename` instead of `rn`, the call for a file directory excerpt is named `Directory` instead of `ls`. The need for memorizing an infinite list of cryptic abbreviations, which is another hallmark of user unfriendly systems, vanishes.

But the influence of the Oberon concept is not restricted to the style in which the computer is used. It extends also to the way programs are designed to communicate with the environment. The definition of the abstract type `Text` in the system's core suggests the replacement of files by texts as carrier of input and output data in very many cases. The advantage to be gained lies in the text's immediate editability. For example, the output of the command `System.Directory` produces the desired excerpt of the file directory in the form of a (displayed) text. Parts of it or the whole may be selected and copied into other texts by regular editing commands (mouse clicks). Or, the compiler accepts texts as input. It is therefore possible to compile a text, execute the program, and to recompile the re-edited text without storing it on disk between compilations and tests. The ubiquitous editability of text together with the persistence of global data (in particular viewers) allows many steps that do not contribute to the progress of the task actually pursued to be avoided.

### 2.1.5 *Extensibility*

An important objective in the Oberon design was extensibility. It should be easy to extend the system with new facilities by adding modules that make use of the already existing resources. Equally important, it should also reduce the system to those facilities that are currently and actually used. For example, a document editor processing documents free of graphics should not require the loading of an extensive graphics editor, a workstation operating as a stand-alone system should not require the loading of extensive network software, and a system used for clerical purposes need include neither compiler nor assembler. Also, a system introducing a new kind of display frame should not include procedures for managing viewers containing such frames. Instead, it should make use of existing viewer management. The staggering consumption of memory space by many widely used systems is due to violation of such fundamental rules of engineering. The requirement of many megabytes of store for an OS is, albeit commonly tolerated, absurd and another hallmark of user-unfriendliness, or perhaps manufacturer friendliness. Its reason is none other than inadequate extensibility.

We do not restrict this notion to procedural extensibility, which is easy to realize. The important point is that extensions may not only add further procedures and functions, but introduce their own data types built on the basis of those

provided by the system: data extensibility. For example, a graphics system should be able to define its graphics frames based on frames provided by the basic display module and by extending them with attributes appropriate for graphics.

This requires an adequate language feature. The Oberon language provides precisely this facility in the form of type extensions. The language was designed for this reason; Modula-2 would have been the choice, had it not been for the lack of a type extension feature. Its influence on system structure was profound, and the results have been most encouraging. In the meantime, many additions have been created with surprising ease. One of them is described at the end of this book. The basic system is nevertheless quite modest in its resource requirements (see Table 1).

#### 2.1.6 *Dynamic Loading*

Activation of commands residing in modules that are not present in the store implies the loading of the modules and, of course, all their imports. Invoking the loader is, however, not restricted to command activation; it may also occur through programmed procedure calls. This facility is indispensable for a successful realization of genuine extensibility. Modules must be loadable on demand. For example, a document editor loads a graphics package when a graphic element appears in the processed document, but not otherwise.

Oberon features no separate linker. A module is linked with its imports when it is loaded, and never before. As a consequence, every module is present only once, in main store (linked) as well as on backing store (unlinked, as file). Avoiding the generation of multiple copies in different, linked object files is the key to storage economy. Prelinked mega-files do not occur in Oberon, and every module is freely reusable.

## 2.2 THE SYSTEM STRUCTURE

The largest identifiable units of the system are its modules. It is therefore most appropriate to describe its structure in terms of modules. As their interfaces are explicitly declared, it is also easy to exhibit their interdependence in a directed graph. The edges indicate imports.

The module graph of a system programmed in Oberon is hierarchical, i.e. has no cycles. The lowest members of the hierarchy effectively import hardware only. We refer here to modules which contain device drivers. But module `Kernel` also belongs to this class; it "imports memory" and includes the disk driver. The modules on the top of the hierarchy effectively export to the user. As the user has direct access to command procedures, we call these top members command modules or tool modules.

The hierarchy of the basic system is shown in a table of direct imports and as a graph in Fig 2. The picture is simplified

by omitting direct import edges if an indirect path also leads from the source to the destination. For example, Files imports Kernel, the direct import is not shown, because a path from Kernel leads to Files via FileDir.

	Text- Frame	Menu- View	Ober- on	Texts	Fonts	Input	View	Disp	Modul	Files	FileDir	Kernel
System	x	x	x	x	x	x	x	x	x	x	x	x
Edit	x	x	x	x	x							
TextFrames		x	x	x	x	x	x	x	x			
MenuViewers			x			x	x	x				
Oberon				x	x	x	x	x	x	x		
Texts					x					x		
Fonts										x		
Viewers								x				
Display												
Modules										x		x
Files											x	x
FileDir												x




Figure 2: The structure of the Oberon core

Module names in the plural form typically indicate the definition of an abstract data type in the module. The type is ex-

ported together with the pertinent operations. Examples are Files, Modules, Fonts, Texts, Viewers, MenuViewers, and TextFrames. Modules whose names are in singular form typically denote a resource that the module manages, be it a global variable or a device. The variable or the device is itself hidden (not exported) and becomes accessible through the module's exported procedures. Examples are all device drivers, Input for keyboard and mouse, Kernel for memory and disk, and Display. Exceptions are the command modules whose name is mostly chosen according to the activity they primarily represent, like System, and Edit.

Oberon is, as already mentioned, the heart of the system containing the central loop to which control returns after each command interpretation, independent of whether it terminates normally or abnormally. Oberon exports several procedures of auxiliary nature, but primarily also the one allowing the invocation of commands (Call) and access to the command's parameter text through variable Oberon.Par. Furthermore, it contains global, exported variables: the log text. The log text typically serves to issue prompts and short failure reports of commands. The text is displayed in a log viewer that is automatically opened when module System is initialized. Oberon furthermore contains the 2 markers used globally on the display, the **mouse cursor** and the **star pointer**. It exports procedures to draw and to erase them, and allows the installation of different patterns for them.



The system shown in Fig 2 basically contains facilities for generating and editing texts, and for storing them in the FS. All other functions are performed by modules that must be added in the usual way by module loading on demand. This includes, notably, the compiler, network communication, document editors, and all sorts of programs designed by users. The high priority given in the system's conception to modularity, to avoid unnecessary frills, and to concentrate on the indispensable in the core, has resulted in a system of remarkable compactness. Although this property may be regarded as of little importance in this era of falling costs of large memories, we consider it to be highly essential. We merely should like to draw the reader's attention to the correlation between a systems' size and its reliability. Also, we do not consider it as good engineering practice to consume a resource lavishly just because it happens to be cheap. The following table lists the core's modules and the major application modules, and it indicates the size of code (in words) and static variables (in bytes) and, the number of source program lines.

## REFERENCES

1. N. Wirth. The PL Oberon. *Software - Practice and Experience* 18, 7, (July 1988) 671-690.

module	code	data	lines
Kernel	1123	8244	263
FileDir	1963	60	352
Files	2360	148	505
Modules	1214	112	226
Input	186	32	79
Fonts	628	56	115
Display	1033	84	190
Viewers	1324	104	206
Texts	2906	204	537
Oberon	1679	288	410
MenuViewers	1513	56	208
TextFrames	5786	292	874
System	2134	72	418
Edit	1096	1104	232
	24945	10856	4615
ORS	1762	992	319
ORB	2348	408	437
ORG	6699	34976	1125
ORP	5994	144	974
	16803	36520	2855
Graphics	3484	564	685
GraphicFrames	2832	288	498
Draw	690	40	164
Rectangles	649	40	118
Curves	1765	72	241
	9420	1004	1706

Table 1: Size of Oberon Modules

2. M. Reiser and N. Wirth. Programming in Oberon - Steps beyond Pascal and Modula. AddisonWesley, 1992. ISBN 0-201-56543-9
3. N. Wirth and J. Gutknecht. The Oberon. Software - Practice and Experience, 19, 9 (Sept. 1989), 857-893.
4. N. Wirth. Ceres-Net: A low-cost computer network. Software - Practice and Experience, 20, 1 (Jan. 1990), 13-24.
5. M. Reiser. The Oberon - User Guide and Programmer's Manual. Addison-Wesley, 1991. ISBN 0-201-54422-9



## TASK

---

Eventually, it is the generic ability to perform every conceivable task that turns a computing device into a versatile universal tool. Consequently, the issues of modeling and orchestrating of tasks are fundamental in the design of any OS. Of course, we cannot expect a single fixed tasking metaphor to be the ideal solution for all possible kinds of systems and modes of use. For example, different metaphors are probably appropriate in the cases of a closed mainframe system serving a large set of users in time-sharing mode, on one hand, and of a personal workstation operated by a single user at a high degree of interactivity, on the other.

In the case of Oberon, we have consciously concentrated on the domain of personal workstations. More precisely, we have directed Oberon's tasking facilities towards a single-user interactive personal workstation that is possibly integrated into a local area network.

We start in [3.1](#) with a clarification of the technical notion of task. In [3.2](#) we continue with a detailed explanation of the scheduling strategy. Then, in [3.3](#) we introduce the concept of

command. Finally, 3.4 provides an overview of predefined system-oriented toolboxes, i.e. coherent collections of commands devoted to some specific topic, e.g. system control and diagnosis, display management, and file management.

### 3.1 THE CONCEPT OF TASK

In principle, we distinguish 2 categories of tasks in Oberon:

INTERACTIVE	bound to local regions on the display and interactions with their contents,
BACKGROUND	system-wide and not necessarily related to any specific displayed entity.

#### 3.1.1 *Interactive Tasks*

Every interactive task is represented by a so-called viewer. Viewers constitute the interface to Oberon's display-system. They embody a variety of roles that are collected in an abstract data type *Viewer*. We shall give a deeper insight into the display system in 4. For the moment it suffices to know that viewers are represented graphically as rectangles on the display screen and that they are implicit carriers of interactive tasks. Fig 3 shows a typical Oberon display screen that is divided up into 6 viewers corresponding to 6 simultaneously active interactive tasks.




Figure 3: Typical display configuration with tool track on the right

In order to get firmer ground under our feet, we now present the programmed declaration of type `Viewer` in a slightly abstracted form:

```
Viewer = POINTER TO ViewerDesc;
ViewerDesc = RECORD X,Y,W,H, state: INT;
                handle: Handler END;
```

`X`, `Y`, `W`, `H` define the viewer's rectangle on the screen:

location `X`, `Y` of the lower left corner relative to the display origin,

width `W` and height `H`.

state informs about the current state of visibility

```
visible,
closed,
covered
```

while `handle` represents the functional interface of viewers:

```
Handler = PROC(V: Viewer; VAR M: ViewerMsg);
```

where `ViewerMsg` is some base type of messages whose exact declaration is of minor importance for now:

```
ViewerMsg = RECORD (*basic parameter fields*) END;
```

However, we should point out the use of object-oriented (OO) terminology. It is justified because `handle` is a procedure variable (a handler) whose identity depends on the specific viewer. A call `V.handle(V, M)` can therefore be interpreted as the sending of a message `M` to be handled by the method of the receiving viewer `V`.

We recognize an important difference between the standard OO model and our handler paradigm. The standard model is closed in the sense that only a fixed set of messages is understood by a given class of objects. In contrast, the handler paradigm is open because it defines just the root (`ViewerMsg`) of a potentially unlimited tree of extending



message types. For example, a concrete handler might be able to handle messages of type:

```
MyViewerMsg = RECORD (ViewerMsg)
                mypar: MyParameters
            END;
```

which is an extended type of ViewerMsg.

It is worth noting that our open OO model is extremely flexible. Notably, extending the set of message types that are handled by an object is a mere implementation issue, that is, it has no effect at all on the object's compile-time interface and on the system integrity. It is fair to mention though that such a high degree of extensibility does not come for free. The price to pay is the obligation of explicit message dispatching at runtime. The following chapters will capitalize on this property.

Coming back to the perspective of tasks, we note that each sending of a message to a viewer corresponds to an activation or reactivation of the interactive task that it represents.

### 3.1.2 *Background Tasks*

Oberon background tasks are not connected a priori with any specific aggregate in the system. Seen technically, they are instances of an abstract data type consisting of type

declarations `Task` and `TaskDesc` together with intrinsic operations `NewTask`, `Install` and `Remove`:

```
Task = POINTER TO TaskDesc;  
TaskDesc = RECORD state: INT; handle: PROC END;  
PROC NewTask(h: PROC; period: INT): Task;  
PROC Install(T: Task);  
PROC Remove (T: Task);
```

The procedures `Install` and `Remove` are called explicitly in order to transfer the state of the specified task from offline to idle and from idle to offline respectively. Installed tasks take their turns in becoming active, that is, in being executed. The installed handlers are simple, parameterless procedures specifying their own actions and conditions for execution, with one exception: Resumption may be delayed until a certain period of time has elapsed. This period is specified in milliseconds when a task is created.

The following 2 examples of concrete background tasks may serve a better understanding of our explanations. The 1st is a system-wide garbage collector collecting unused memory. The 2nd is a network monitor accepting incoming data on a local area network. In both examples the state of the task is captured entirely by global system variables. We shall come back to these topics in ?? and ?? respectively.

We should not end without drawing an important conclusion. Transfers of control between tasks are implemented in

Oberon as ordinary calls and ordinary procedures returns (procedure variables, actually). Preemption is not possible. From that we conclude, active tasks periods are sequentially ordered and controlled by a single thread. This simplification pays well: Locks of common resources are completely dispensable and deadlocks are not a topic.

### 3.2 THE TASK SCHEDULER

We start from the general assumption that, at any given time, a number of well-determined tasks are ready in the system to be serviced. Remember that 2 categories of tasks exist: interactive and background ones. They differ substantially in the criteria of activation or reactivation and in the priority of dispatching. Interactive tasks are (re)activated exclusively upon interactions by the user and are dispatched with high priority. In contrast, background tasks are polled with low priority.

We already know that interactive tasks are activated by sending messages. The types of messages used for this purpose are `InputMsg` and `ControlMsg` reporting keyboard events and mouse events respectively. Slightly simplified, they are declared as

```
InputMsg = RECORD (ViewerMsg) id, X, Y: INT;
               keys: SET; ch: CHAR END;
```

```
ControlMsg = RECORD (ViewerMsg) id,X,Y: INT END;
```

The field `id` specifies the exact request transmitted with this specific reactivation. In the case of `InputMsg` the possible requests are `consume` (the character specified by field `ch`) and `track` (mouse, starting from state given by `keys` and `X, Y`). In case of `ControlMsg` the choice is `mark` (the viewer at position `X, Y`) or `neutralize`. `Mark` means moving the global system pointer (typically represented as a star-shaped mark) to the current position of the mouse. `Neutralizing` a viewer is equivalent to removing all marks and graphical attributes from this viewer.

All tasking facilities are collected in 1 module, called `Oberon`. In particular, the module's definition exposes the declarations of the abstract data type `Task` and of the message types `InputMsg` and `ControlMsg`. The module's most important contribution, however, is the task scheduler (often referred to as "Oberon loop") that can be regarded as the system's dynamic center.

Before studying the scheduler in detail we need some more preparation. We start with the institution of the focus viewer. By definition, this is a distinguished viewer that by convention consumes subsequent keyboard input. Note that we identify the focus viewer with the focus task, hereby making use of the one-to-one correspondence between viewers and tasks.

Module Oberon provides the following facilities in connection with the focus viewer: A global variable `FocusViewer`, a procedure `PassFocus` for transferring the role of focus to a new viewer, and a defocus variant of `ControlMsg` for notifying the old focus viewer of such a transfer.

The implementation details of the abstract data type `Task` are hidden from the clients. It is sufficient to know that all task descriptors are organized in a ring and that a pointer points to the previously activated task. The ring is guaranteed never to be empty because the above mentioned garbage collector is installed as a permanent sentinel task at system loading time.

The following is a slightly abstracted version of the actual scheduler code operating on the task ring. It should be associated with procedure `Loop` in the module `Oberon`.

```

get mouse position and state of keys;
REPEAT
  IF keyboard input available THEN read character
    IF character is escape THEN
      broadcast neutralize message to viewers
    ELSIF character is mark THEN
      send mark message to viewer containing mouse
    ELSE send consume message to focus viewer
  END;
get mouse position and state of keys

```

```

    ELSIF at least one key pressed THEN
        REPEAT
            send track message to viewer containing mouse;
            get mouse position and state of keys
        UNTIL all keys released
    ELSE (*no key pressed*)
        send track message to viewer containing mouse;
        take next task in ring as current task;
        call its handler
            (if specified time period has elapsed)
        get mouse position and state of keys
    END
UNTIL FALSE

```

The system executes a sequence of uninterrupted procedures (tasks). Interactive tasks are triggered by input data being present, either from the keyboard, mouse, or other input sources. Background tasks are taken up in a round-robin manner. Interactive tasks have priority.

Having consciously excluded exceptional program behavior in our explanations so far, some comments about the way of runtime continuation in the case of a failing task or, in other words, in the case of a trap are in order here. On the (abstract) level of tasks, we can identify 3 sequential actions of recovery taken after a program failure:

recovery after program failure =

```

BEGIN save current system state;
      call installed trap handler;
      roll back to start of task scheduler
END

```

Essentially, the system state is determined by the values of all global and local variables at a given time. The trap handler typically opens an extra viewer displaying the cause of the trap and the saved system state. Notice in the program fragment above that background tasks are removed from the ring after failing. This is an effective precaution against cascades of repeated failures. Obviously, no such precaution is necessary in the case of interactive tasks because their reactivation is under control of the user of the system.

Summarizing the essence of the Oberon tasking system:

- A multitasking system based on a 2-category model
  - `INTERACTIVE` interfacing with the display system, high-priority scheduled upon user interactions;
  - `BACKGROUND` stand-alone, low-priority scheduled.
- Task activations are modeled as message passing, eventually, procedures calls assigned to variables. They're sequentially ordered, controlled by a single thread.

### 3.3 THE CONCEPT OF COMMAND

An OS constitutes a general purpose platform on which applications can build upon. To software designers the platform appears as interface to "the system" and (in particular) to the underlying hardware. Unfortunately, interfaces defined by conventional OSes often suffer from an all too primitive access mechanism that is based solely on the concept of "software interrupt" or "supervisor call" and on files taking the role of "connecting pipes". The situation is especially ironic when compared with the development of high-level PLs towards extreme abstraction.

We have put greatest emphasis in Oberon on closing the semantic gap between applications and the system platform. The result of our efforts is a highly expressive and consistent *application programming interface* (API) in the form of an explicit hierarchy of module definitions. Perhaps the most significant and notable outcome of this approach is a collection of very powerful and system-wide abstract data types such as Task, Frame, Viewer, File, Font, Text, Module, Reader, Scanner, Writer, etc.

#### 3.3.1 *Atomic Actions*

The most important generic function of any OS is executing programs. A clarification of the term program as it is used in Oberon comprises 2 views: a static and a dynamic one.



STATICALLY, program means a software package with an *entry point*. More formally, it is a pair  $(M^*, P)$ , where  $M$  is a module,  $P$  is an *exported parameterless procedure* of  $M$ , and  $M^*$  denotes the hierarchy consisting of  $M$  itself and all imported modules, directly or indirectly.

Note: 2 different hierarchies  $M^*$  and  $N^*$  are not generally disjoint, rather, intersect a superset of the OS.

DYNAMICALLY VIEWED, it is an *atomic action* called *command*, operating on the global system state. Here *atomic* means *no user interaction*. This definition is just a necessary consequence of our *non-preemptive task model* scheduling with the benefit of a single *carrier thread*. We can argue like this:

When a traditional interactive program requires input from the user, the current task is normally preempted in favor of another task that produces the required input data. Therefore, a traditional interactive program can be viewed as a sequence of atomic actions interrupted by actions that possibly belong to other programs. Whereas interruption in traditional systems may occur at any time, yet in Oberon it can occur only after task (command) completion.

Quintessentially, programs are represented as commands: exported parameterless procedures not interacting with the user.

Returning to programs calling and execution we now arrive at the following refined version:

```
call program(M*, P) = BEGIN
    load module hierarchy M*;
    call command P
END
```

The system interface to the command mechanism itself is again provided by Oberon. Its primary operation is to "call command by name and pass actual parameters":

```
PROC Call(name: ARRAY OF CHAR;
          pars: ParList; VAR res: INT);
```

name    the desired command name in form of M.P,  
 pars    the actual parameters list (APL), and  
 res    the result code.

But in fact we have separated the setting of parameters:

```
PROC SetPar(F: Display.Frame;
            T: Texts.Text; pos: INT);
```

        F    indicates the calling viewer,  
 Pair (T, pos)    specifies the starting position of a text.

from the actual call:

```
PROC Call(name: ARRAY OF CHAR; VAR res: INT);
```

Notice the occurrence of yet another abstract data type `Text`, exported by `Texts`. We shall devote 5 to a thorough discussion of Oberon's text system. For the moment we can simply look at a text as a sequence of characters.

The APL is handed over to the command by Oberon as an exported global variable:

```
Par: RECORD vwr: Viewers.Viewer;  
      frame: Display.Frame;  
      text: Texts.Text; pos: INT  
END
```

In principle, commands operate on the entire system and can access the current global state via the system's powerful abstract modular interface, of which the APL is just one component. Another one is the so-called *system log* which is a system-wide protocol reporting on the progress of command execution and on exceptional events in chronological order. The log is represented as a global variable:

```
Log: Texts.Text;
```

It should have become clear by now that implementers of commands may rely on a rich arsenal of abstract global facilities that reflect the current system state and make it

accessible. In other words, they may rely on a high degree of system integration. Therefore, Oberon features an extraordinarily broad spectrum of mutually integrated facilities. For example, the system distinguishes itself by a complete integration of the abstract data types Viewer (4) and Text (5) that we encountered above.

Oberon assists the integration of these types with the following conceptual features, including the 1st 2 we have just introduced:

1. Standard parameter list for commands,
2. system log,
3. generic text selection, and
4. generic copy viewer.

At this point we should add a word of clarification to our use of the term *generic*. It is synonymous with "interpretable individually by any viewer (interactive task)" and is typically used in connection with messages or orders whose receiver's exact identity is unknown.

Let us now go into a brief discussion of the generic facilities without, however, leaving the level of our current abstraction and understanding.

### 3.3.2 *Generic Text Selection*

Textual selections are characterized by

- a text,
- a stretch of characters within that text, and
- a time stamp.

Without further qualification "the text selection" always means "the most recent text selection". It can be obtained programmatically by calling:

```
PROC GetSelection(VAR text: Texts.Text;  
                  VAR beg, end, time: LONGINT);
```

The parameters specify the desired stretch of text starting at position beg and ending at end-1 as well as the associated time stamp. The procedure is implemented in form of a broadcast of a so-called selection message to all viewers. The declaration of this message is

```
SelectionMsg = RECORD (ViewerMsg) text: Texts.Text;  
                      beg, end, time: INT  END;
```

### 3.3.3 *Generic Copy Viewer*

Generic copying is synonymous with reproducing and cloning. It is the most elementary generic operation possible. Again,

a variant of type `ViewerMsg` is used for the purpose of transmitting requests of the desired type:

```
CopyMsg = RECORD (ViewerMsg) vwr:Viewers.Viewer END;
```

Receivers of a copy message typically generate a clone of themselves and return it to the sender via field `vwr`.

Let us now summarize the concept of command in Oberon:

- As an OS Oberon presents to clients a highly expressive *modular interface* that exports many powerful abstract data types like `Viewer` and `Text`. A rich arsenal of global data types and generic facilities achieve *system integration* at a high degree.
- Programs are modeled as *commands*, *exported parameterless procedures* not interacting with the user.
- The collection of commands provided by a module appears as its *user interface*. Parameters are passed to commands via a global parameter list, registered by the calling task in the central Oberon.
- Commands operate on the global state of the system.

### 3.4 TOOLBOXES

Modules typically appear in 3 different forms.

- 1<sup>st</sup> encapsulates some data, which can be accessed only through exported procedures and functions. `FileDir`, encapsulating the file directory and protecting it from disruptive access, is a good example.
- 2<sup>nd</sup> represent an abstract data type, exporting it and its associated operators. `Files`, `Modules`, `Viewers`, and `Texts` are typical examples.
- 3<sup>rd</sup> collection of procedures pertaining to the same topic, like RS232 handling communication over a serial line.

Oberon adds a 4<sup>th</sup> form:

*toolbox* A pure collection of commands on the *top* of the modular hierarchy. Toolbox modules are "imported" by system users at run-time. In other words, their definitions define the user interface. Typical examples are `System` and `Edit`.

As a rule of thumb there exists a toolbox for every topic or application. As an example of a toolbox definition we quote an annotated version of `System`:

DEFINITION `System`;

```
(*System management, Ch. 3 and 8*)  
PROC SetUser; (*identification*)  
PROC SetFont; (*for typed text*)  
PROC SetColor; (*for typed text&graphics*)
```

```

PROC SetOffset;(*for typed text*)
PROC Date;      (*set/get date&time*)
PROC Collect;  (*garbage*)
(*Display management, Ch. 4*)
PROC Open;    (*viewer*)
PROC Close;   (*viewer*)
PROC CloseTrack;
PROC Recall;(*most recently closed viewer*)
PROC Copy;    (*viewer*)
PROC Grow;    (*viewer*)
PROC Clear;   (*log*)
(*Module management, Ch. 6*)
PROC Free;      (*specified modules*)
PROC ShowCommands;(*of specified module*)
PROC ShowModules; (*show loaded modules*)
(*File management, Ch. 7*)
PROC Directory;
PROC CopyFiles;
PROC RenameFiles;
PROC DeleteFiles;
(*System inspection, Ch. 8*)
PROC Watch;(*tasks, memory&disk storage*)
END System;

```

An important consequence of our integrated systems approach is the possibility of constructing a universal, interactive command interpreter bound to viewers of textual



contents. If the text obeys the following syntax (specified in EBNF), we call it command tool:

```
CommandTool = {[Comment]CommandName[ParameterList]}
```

If present, the parameter list is made available to the called command via fields `text` and `pos` in the global variable `Par` that is exported from `Oberon`. Because this parameter list is interpreted individually by each command, its format is completely open. However, we postulate some conventions and rules for the purpose of a standardized user interface:

1. The elements of a textual parameter list (TPL) are universal syntactical tokens like name, literal string, integer, real number, and special character.
2. An arrow "`^`" in the TPL refers to the current text selection for continuation. In the special case of the arrow following the command name immediately, the entire parameter list is represented by the text selection.
3. An asterisk "`*`" in the TPL refers to the currently marked viewer. Typically, the asterisk replaces the name of a file. In such a case the contents of the viewer marked by the system pointer (star) is processed by the command interpreter instead of the contents of a file.

4. An at-character "@" in the TPL indicates that the selection marks the (beginning of the) text which is taken as operand.
5. A terminator-character "~" terminates the TPL in case of a variable number of parameters.

Because command tools are ordinary, editable texts (in contrast to menus in conventional systems), they can be customized "on the fly", which makes the system highly flexible. We refer again to Fig 3 that shows a typical Oberon screen layout consisting of 2 vertical tracks,

1. a wider user track on the left, and  
3 documents are displayed in the user track:
  - a) a text,
  - b) a graphic, and
  - c) a picture.
2. a narrow system track on the right,  
where we find 1 logviewer displaying the system log,  
2 tool-viewers making available
  - a) the standard system tool, and
  - b) a customized private tool respectively.

In concluding this chapter, let us exemplify the concepts of command and tool by the system control section of the System toolbox. Consisting of the commands

SetUser	installing the user's identification,
Date	displaying or setting the system date and time,
SetFont	presetting the system type-font for typed text,
SetColor	setting the system color, and
Collect	activating the garbage collector.

They are used to control system-wide facilities (detail each).

In summary,

- a toolbox is a special module, defined as a *collection of commands*;
- appearing at the top of the modular hierarchy, the entire toolboxes fix the *user interface* for the system;
- command tools are sequences of textually represented command calls, editable and customizable;
- in a typical screen layout, tools are displayed in the viewers within the *system track*.



## DISPLAY

---

Of the interface a PC presents to a user, the display screen is the most important. At the 1st sight, it simply represents a rectangular output area. However, in combination with a mouse, it quickly develops into a sophisticated interactive I/O platform of almost unlimited flexibility. It is mainly this Janus-faced characteristic that makes it stand out from ordinary external devices to OS-managed core system.

In this chapter we shall give more detailed insight into why

- central position a display system takes within an OS, and
- its determining influence on entire system architecture.

In particular, we shall show that it is a natural basis or anchor for functional extensibility.

4.1    SCREEN LAYOUT MODEL

In the early 70s, Xerox PARC in California launched the Smalltalk project with the goal of conceiving and developing new more natural ways to communicate with PCs [Goldberg]. Perhaps the most conspicuous among several significant achievements of this endeavor is the idea of applying the desktop metaphor to the display screen. This metaphor comprises a desktop and the collection of possibly mutually overlapping pages of paper laid out on it. By projecting such a configuration onto the screen surface we get the familiar picture of Fig 4 showing a collection of partially or totally visible rectangular areas on a background, so-called *windows* or *viewers*.




Figure 4: Desktop showing partially overlapping viewers

The desktop metaphor is used by many modern OSes and user interface shells both as

a	for	to
natural	the	separate displayed data
model	system	belonging to different tasks, and
power		organize the display screen
-ful	users	interactively, according to
tool		individual taste and preference.

However, in this metaphor there are inherent drawbacks, primarily connected with overlapping.

1<sup>st</sup>, any efficient management of overlapping viewers must rely on

- a subordinate management of (arbitrary) sub-rectangles, and
- sophisticated clipping operations.

This is so because partially overlapped viewers must be partially restored under control of the *viewer manager*. For example, in Fig 4, rectangles *a*, *b*, and *c* in viewer *B* ought to be restored individually after closing of *A*.

2<sup>nd</sup>, there is a significant danger of covering viewers completely and losing them forever.

3<sup>rd</sup>, no canonical heuristic algorithms exist for automatic allocation of screen space to newly opened viewers.

Experience has shown that partial overlapping is desirable and beneficial in rare cases only, so the additional complexity of its management is hard to justify. [Binding; Wille] Therefore, alternate strategies to structure a display screen have been looked for. An interesting class of established solutions can be titled as *tiling*. There are at least 2 variants. [Cohen]

Perhaps the most *unconstrained* (hence obvious) one is based on iterated horizontal or vertical splitting of existing viewers. Starting with the full screen and successively opening *A*, *B*, *C*, *D*, *E*, and *F* we get to a configuration as in Fig 5.




Figure 5: Viewer configuration resulting from unconstrained tiling



The 2<sup>nd</sup> variant is *hierarchic tiling*. Again, the hierarchy starts with a full screen that is now decomposed into a number of *vertical tracks*, each of which is further decomposed into a number of horizontal viewers. We decided in favor of this kind of tiling in Oberon, mainly because the algorithm of reusing the area of a closed viewer is simpler and more uniform. For example, assume that in Fig 5 viewer *F* has been closed. Then, it is straightforward to reverse the previous opening operation by extending viewer *E* at its bottom end. However, if the closed viewer is *B*, no such simple procedure exists. For example, the freed area can be shared between viewers *C* and *D* by making them extend to their left. Clearly, no such complicated situations can occur in the case of hierarchic tiling.

It is also used in Xerox PARC's Cedar system [Teitelman]. However, Oberon differs in:

- 1<sup>st</sup>, It supports quick temporary context switching by overlaying one track or any contiguous sequence of tracks with new layers. In Fig 6 a snapshot of a standard display screen is graphically represented. It suggests 2 original tracks and 2 levels of overlay, where the top layer is screen-filling.
- 2<sup>nd</sup>, Oberon displays do not provide reserved areas for system-wide facilities, while standard Cedar screens feature a command row at the top and an icon row at the bottom. And




Figure 6: Overlay of tracks and sequences of tracks

$3^{rd}$ , It is based on a different heuristic strategy for the automatic placement of new viewers. As a Cedar default invariant, the area of every track is divided up evenly among the viewers in this track. When a new viewer is to be placed, the existing viewers in the track are requested to reduce their size and move up appropriately. The newly opened is then allocated in the freed spot at the bottom. In contrast, Oberon normally splits the largest existing viewer in a given track into 2 halves of equal size. As an advantage of this latter allocation strategy we note that existing contents are kept stable.

## 4.2 VIEWERS AS OBJECTS

Although everybody seems to agree on the meaning of the term *viewer*, no 2 different system designers actually do. The original role of a viewer as merely a separate display area has meanwhile become heavily overloaded with additional functionality. Depending on the underlying system are viewers' individual views on a certain configuration of objects, carriers of tasks, processes, applications, etc. Therefore, we first need to define our own precise understanding of the concept of viewer.

The best guide to this aim is the abstract data type *Viewer* we introduced in 3. We recapitulate: *Viewer* serves as a template describing viewers abstractly as “black boxes” in terms of a state of visibility, a rectangle on the display screen, and a message handler. The exact functional interface provided by a given variant of viewer is determined by the set of messages accepted. This set is structured as a customized hierarchy of type extensions.

We can now obtain a more concrete specification of the role of viewer by identifying some basic categories of universal messages that are expected to be accepted by all variants of viewer. For example, we know that messages reporting about user interactions as well as messages defining a generic operation are universal. These 2 categories of universal messages document the roles of viewers as interactive tasks and as parts of an integrated system respectively. In to-

tal, there are 4 such categories. They are here listed together with the corresponding topic and message dispatchers:

Dispatcher	Topic	Message
task scheduler	dispatching tasks	report user interaction
cmd interpreter	processing command	define generic operation
view manager	organizing display area	location/size change
doc manager	operating on document	content/format change

These topics essentially define the role of viewers. In short, we may look at an viewer as a non-overlapped rectangular box on the screen both

- acting as an *integrated display area* for some objects of a document, and
- *representing an interactive task* in the form of a sensitive editing area.

Shifting emphasis a little and regarding the various message dispatchers as subsystems, we recognize immediately the role of viewers as *integrator of the different subsystems via message-based interfaces*. In this light, Viewer appears as a common object-oriented basis of Oberon’s subsystems.

The topics listed above constitute some kind of contents backbone of the 3, 4 and 5. Task scheduling and command interpreting are already familiar to us from §3.2 and 3.3. Viewer and text management will be the topics of §?? and 5.2, respectively. Thereby, the built-in type Text will serve

as a prime example of a document type. The activities that a viewer performs are basically controlled by events or, more precisely, by messages representing event notices. We shall explain this in §?? and 5.3 in detail cases of an abstract class of standard viewers and a class of viewers displaying standard text, respectively.

Here is a preliminary overview of some archetypal kinds of message:

- After each stroke a keyboard message with the typed character is sent to the focus viewer and after each click a mouse message of the new state is sent to the viewer containing the mouse.
- The message often representing some generic operation is to be interpreted individually by recipients. Obvious examples are "return current textual selection", "copy-over stretch of text", and "produce a copy (clone)". Notice that generic operation is the key to extensibility.
- In a tiling viewer environment, every opening of a new viewer and every change of size or location of an existing viewer has an obvious effect on adjacent viewers. The viewer manager therefore issues a message to every affected viewer requesting to adjust its size appropriately.

- Whenever the contents or the format of a document has changed, a message notifying all visible viewers of the change is broadcast. Notice that broadcasting messages by a model (document) to the entirety of its potential views (viewers) is an interesting implementation of the famous model-view-controller (MVC) pattern that dispenses models from “knowing” (registering) their views.

## 4.3 FRAMES AS BASIC DISPLAY ENTITIES

When we introduced viewers in 3 and 4.2, we simplified the abstraction aim. We know already that viewers appear as elements of 2nd order in the tiling hierarchy. Having treated them as black boxes so far we have not revealed anything about the hierarchy continuation. As a matter of fact, viewers are neither elementary display entities nor atoms. They are just a special case of so-called (*display*) *frames*. They are arbitrary rectangles displaying a collection of objects or a document excerpt. In particular, they may recursively contain other frames, a capability that makes them an extremely powerful tool for any display organizer.

Type Frame is declared as

```
Frame = POINTER TO FrameDesc;
FrameDesc = RECORD next, dsc: Frame;
                X, Y, W, H: INT;
                handle: Handler END;
```

`next` and `dsc` are connections to further frames, whose names suggest a recursive hierarchical structure:

	points to
<i>next</i>	the next frame on the same level, while
<i>dsc</i>	the (1st) descendant, i.e. the next lower level of the nested frames hierarchy.

X, Y, W, H, and the handler handle serve the original purpose to which we introduced them. In particular, the handler allows frames to react individually on the receipt of messages:

```
Handler = PROC(F: Frame; VAR M: FrameMsg);
```

where FrameMsg represents the root of a potentially unlimited tree hierarchy of possible messages to frames:

```
FrameMsg = RECORD END;
```

Having now introduced the concept of frames, we can reveal the whole truth about viewers. As a matter of fact, Viewer is a derived type, an extension of Frame:

```
Viewer = POINTER TO ViewerDesc;
```

```
ViewerDesc = RECORD (FrameDesc) state: INT END;
```

These declarations formally express the fact that viewers are nothing but a special case (or variant or subclass) of general frames, additionally featuring a state of visibility. In particular, viewers inherit the hierarchical structure of frames. This is an extremely useful property immediately opening an unlimited spectrum of possibilities for designers of a specific subclass of viewers to organize the representing rectangular area. For example, the area of viewers of, say, class Desktop may take the role of a background being covered



by an arbitrary collection of possibly mutually overlapping frames. In other words, our decision of using a tiling viewer scheme globally can easily be overwritten locally.

An even more important example of a predefined structure is provided by the abstract class, called *menu viewers*, whose shape is familiar from most snapshots taken of the standard Oberon display screen. A menu viewer consists of a thin rectangular boundary line and an interior area being vertically decomposed into a menu region at the top and a contents region at the bottom (see Fig 7).




Figure 7: The compositional structure of a menu viewer

In terms of data structures, the class of menu viewers is defined as a type extension of *Viewer* with an additional component *menuH* specifying the menu frame height:

```
MenuViewer = POINTER TO MenuViewerDesc;
MenuViewerDesc = RECORD (ViewerDesc) menuH: INT END;
```

Each menu viewer  $V$  specifies exactly 2 descendants:

- the menu frame  $V.dsc$ , and
- the frame of main contents or main frame  $V.dsc.next$ .

Absolutely nothing is fixed about the 2 descendant frames contents. The standard menu frame is a text frame, displaying a line of commands in inverse video mode. By definition, the nature of the main frame specifies its type:

- If it is a text frame, we call the viewer a text one;
- If a graphics frame, we call it a graphics viewer, etc.

## 4.4 DISPLAY MANAGEMENT

Oberon's display system comprises 2 main topics:

1. *viewer management* (VM), and
2. *cursor handling*.

Let us 1st turn to the much more involved topic of 1 and postpone 2 to the end of this section. Before we can actually begin our explanations we need to introduce the concept of the *logical display area*. It is modeled as a 2-dimensional Cartesian plane housing the totality of objects to be displayed. The essential point of this abstraction is a rigorous decoupling of any aspects of physical display devices. As a matter of fact, any concrete assignment of display monitors to certain finite regions of the display area is a pure matter of configuring the system.

Being a subsystem of OS with well-defined modular structure, the display system appears in the form of a small hierarchy of modules. Its core is a linearly ordered set consisting of 3 modules:

Display,  
Viewers, and  
MenuViewers,

the latter building upon the former. Conceptually, each

module contributes an associated class of display-oriented objects and a collection of related service routines.

The following is an overview of the subsystem VM.

Module	Type	Service
MenuViewer	Viewer	Message handling for menu viewers
Viewers	Viewer	Tiling VM
Display	Frame	Block-oriented raster operations

Modules on upper lines import lower ones, and  
types on upper lines extend those on lower.

Inspecting the "Type" column we recognize precisely our familiar types

Frame,  
Viewer, and  
MenuViewer respectively,

where the last is an abbreviation of `MenuViewers.Viewer`. In addition to the core modules of the display system a section in Oberon provides a specialized API that simplifies the use of the VM package by applications in the case of standard Oberon display configurations. We shall come back to this topic in §4.6.

For this moment let us concentrate on the VM core and in particular the Viewers and MenuViewers, saving the Display

for the next section. Typically, we start a module presentation by listing and commenting its definition, and refer to subsequent listings for its implementation.

#### 4.4.1 Viewers

Focusing 1st on module Viewers we can roughly define the domain of its responsibility as *initializing and maintaining the global layout of the display area*. From the previous discussion we are well acquainted already with the structure of the global display space as well as its building blocks:

The display area is hierarchically tiled with frames, where the first two levels in the frame hierarchy correspond to *tracks* and *viewers* respectively.

The formal definition:

```

DEFINITION Viewers;
  IMPORT Display;                                (*message ids*)
  CONST restore = 0; modify = 1; suspend = 2;
  TYPE Viewer = POINTER TO ViewerDesc;
    ViewerDesc = RECORD (Display.FrameDesc)
      state: INT END;
    ViewerMsg = RECORD (Display.FrameMsg)
      id, X, Y, W, H: INT;
      state: INT END;
  VAR curW: INT; (*currently configured width*)

```

```

PROC InitTrack(W, H: INT; Filler: Viewer);
PROC OpenTrack(X, W: INT; Filler: Viewer);
PROC CloseTrack(X: INT);           (*track handling*)

PROC Open(V: Viewer; X, Y: INT);
PROC Change(V: Viewer; Y: INT);
PROC Close(V: Viewer);           (*viewer handling*)

PROC This(X, Y: INT): Viewer;
PROC Next(V: Viewer): Viewer;
PROC Recall(VAR V: Viewer);      (*miscellaneous*)
PROC Locate(X,H: INT; VAR fil,bot,alt,max: Viewer);
PROC Broadcast(VAR M: Display.FrameMsg);
END Viewers.

```

The 1st 3 support the track structure of the display area.

- InitTrack creates a new track of width W and height H by partitioning off a vertical strip of width W from the display area.

In addition, it initializes the newly created one with a 3<sup>rd</sup> parameter, a filler viewer. The filler viewer essentially serves as background filling up the track at its top end. It reduces to height 0 if the track is covered completely by productive viewers.

Configuring the display area is part of system initialization. It amounts to executing a sequence of steps:

```
NEW(Filler);
Filler.handle := HandleFiller;
InitTrack(W, H, Filler)
```

where `HandleFiller` is supposed to handle messages that require modifications of size and cursor drawing.

- The global variable `curW` indicates the already configured part width of the display area. Note that configuring starts with  $x = 0$  and is non-reversible in sense that the grid defined by the initialized tracks cannot be refined later. However, remember that it can be coarsened at any time by overlaying a contiguous sequence of existing tracks by a single new track.

`OpenTrack` serves exactly this purpose. The track (or sequence of tracks) to be overlaid in the display area must be spanned by the segment  $[X, X + W)$ .

- `CloseTrack`, inverse to `OpenTrack`, is called to
  - close the (topmost) track located at  $X$  in the display area, and
  - restore the previously covered track (or sequence of tracks).

The 2nd 3 are to organize viewers within individual tracks.

- Open allocates a viewer at given position. More precisely,
  1. locates the viewer containing point  $(X, Y)$ ,
  2. splits it horizontally at height  $Y$ , and
  3. opens the new one  $V$  in the lower part of area.

In the special case of  $Y$  coinciding with the upper boundary line of viewer in 1, it is closed automatically.

- Change allows to change the height of a given viewer  $V$  by moving its upper boundary line to a new location  $Y$  (within the limits of its neighbors).
- Close removes the given  $V$  from the display area.

Fig 8 makes these operations clear.




Figure 8: Basic operations on viewers



The last group provides miscellaneous services.

- This identifies the viewer displayed at  $(X, Y)$ .
- Next returns the next upper neighbor of a given displayed viewer  $V$ .
- Recall allows recalling and restoring the most recently closed viewer.
- Locate assists heuristic allocation of new viewers. For any given track and desired minimum height, it offers a choice of some distinguished viewers in this track:
  - the filler viewer,
  - the one at bottom,
  - an alternative choice, and
  - the viewer of maximum height.
- Finally, Broadcast broadcasts a message to the display area, that is, sends the given message to all currently displayed viewers.

It is now a good time to throw a glance behind the scenes. Let us start with revealing Viewer internal data structure. Remember that according to the principle of information hiding an internal data structure is fully private to the containing module and accessible only through the module's

procedural interface. Fig 9 shows a data structure view of the display snapshot taken in Fig 6. Note that the overlaid tracks and viewers are still part of the internal data structure.

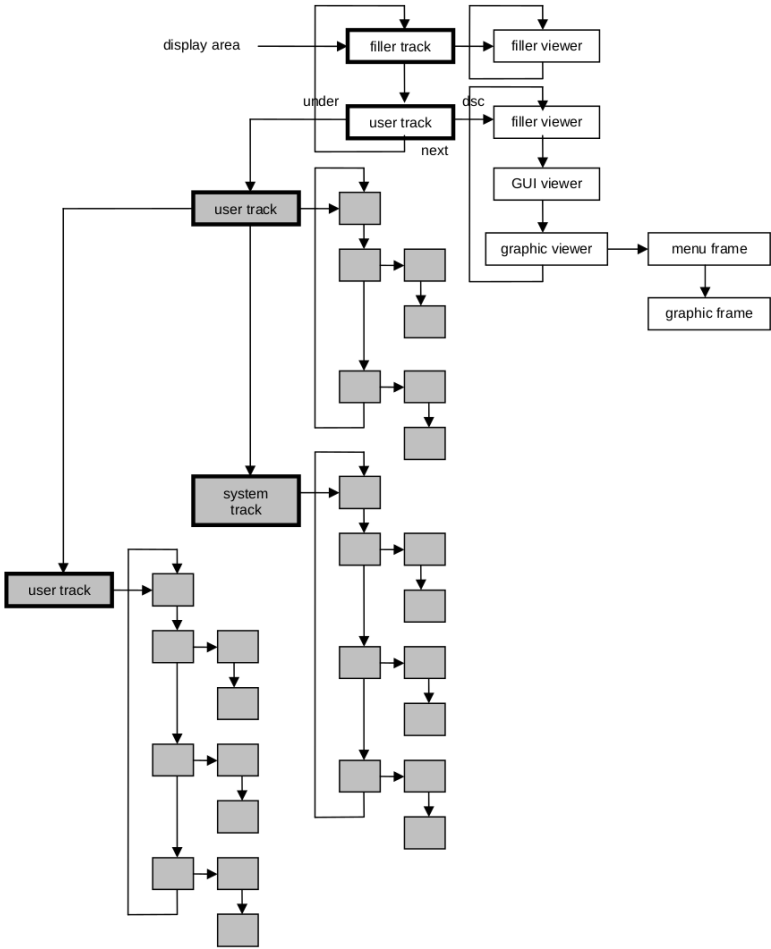


Figure 9: Internal data structure snapshot of Fig 6

In the data structure we recognize an anchor that represents the display area and points to a list of tracks, each track pointing to a list of viewers, each viewer in turn pointing to a list of arbitrary sub-frames. Both the list of tracks and the list of viewers are closed to a ring, where the filler track (filling up the display area) and the filler viewers (filling up the tracks) act as anchors. Additionally, each track points to a (possibly empty) list of tracks lying underneath. These frames are invisible on the display, and shaded in Fig 9.

Technically, the track descriptor `TrackDesc` is a private extension of the viewer descriptor `ViewerDesc`. Repeating the declarations of viewer descriptors and frame descriptors, we get to this hierarchy of types:

```
TrackDesc = RECORD (ViewerDesc)
               under: Display.Frame   END;

ViewerDesc = RECORD (FrameDesc) state: INT END;

FrameDesc = RECORD  next, dsc: Frame;
                   X, Y, W, H: INT;
                   handle: Handler   END;
```

It is noteworthy that the data structure of the VM is heterogeneous with `Frame` as base type. It provides a nice example of a nested hierarchy of frames with the additional property

that the 1st 2 levels correspond to the 1st 2 levels in the type hierarchy defined by *Track*, *Viewer*, and *Frame*.

In an object-oriented environment objects are autonomous entities in principle. However, they may be bound to some higher instance (other than the system) temporarily. For example, we can look at the objects belonging to a module's private data structure as bound to this module. Deciding if an object is currently bound then becomes a fundamental problem. In the case of viewers, this information is contained in an extra instance variable called *state*.

As a system invariant, we have for every viewer *V*

$$V \text{ is bound to module } \textit{Viewers} \Leftrightarrow V.state \neq 0$$

If we call any displayed viewer *visible* and each covered by an overlaying track *suspended*, we can refine this invariant to

$$\{V \text{ is } \textit{visible} \Leftrightarrow V.state > 0\} \text{ and } \\ \{V \text{ is } \textit{suspended} \Leftrightarrow V.state < 0\}$$

In addition, more detailed information about the kind of viewer *V* is given by the magnitude  $|V.state|$ :

$ V.state $	kind of viewer
0	closed
1	filler
-1	productive

The magnitude  $|V.state|$  is kept invariant by *Viewers*. It

could be used, for example, to distinguish different levels of importance or preference with the aim of supporting a smarter algorithm for heuristic new viewers allocation. *state* is read-only to modules other than Viewers.

We are now sufficiently prepared to understand how the exported Viewers procedures work behind. They all operate on the internal dynamic data structure just explained.

- This, Next, Locate, Change use it as a reference only or operate on individual elements
- InitTrack, OpenTrack, Open add new elements, and
- CloseTrack, Close remove elements.

Most have side-effects on existing elements (*size* or *state*).

Let us now change perspective and look at Viewers as a general low-level viewer manager (VMer) whose exact contents are unknown to it (and whose controlling software might have been developed years later). In short, let us look at Viewers as a manager of black boxes. Such an abstraction immediately makes it impossible for the implementation to call fixed procedures for, say, changing a viewer's size or state. The facility needed is a *message-oriented interface*.

```
TYPE ViewerMsg = RECORD (Display.FrameMsg) id,
                        X, Y, W, H, state: INT END;
```

There’re 3 variants of Viewer messages, discriminated by id:

- restore contents,
- modify height (extend or reduce at bottom), and
- suspend (close temporarily or permanently).

The additional components of the message inform about the desired new location, size, and state. The following table lists senders, messages, and recipients of viewer messages.

Originator	Message	Recipients	
OpenTrack	suspend temporarily	viewers covered by opening track	
CloseTrack	suspend permanently	viewers in closing track	
Open	modify or suspend	upper	opening viewer
Change	modify	neighbor of	changing viewer
Close	suspend permanently		closing viewer

4.4.2 Menu Viewers

So far, we have treated viewers abstractly. Next step we focus on a special class called *menu viewers*. From the earlier definition we know they’re characterized by a structure consisting of 2 vertically tiled *descendant* frames,

a frame of    at the

---

menu    top, and

contents    bottom.

Because the nature and contents of these frames are typically unknown by their “ancestor” (or “parent”) viewer, a collection of abstract messages is again a postulating form of interface. As net effect, the handling of menu viewers boils down to a combination of preprocessing, transforming and forwarding messages to the descendant frames. In short, the display space is hierarchically organized and message passing within it obeys the pattern of strict parental control.

Again, we start detailed discussion with module interface:

```

DEFINITION MenuViewers;
  IMPORT Viewers, Display;          (*message ids*)
  CONST extend = 0; reduce = 1; move = 2;
  TYPE Viewer = POINTER TO ViewerDesc;
      ViewerDesc = RECORD (Viewers.ViewerDesc)
                      menuH: INT      END;
      ModifyMsg = RECORD (Display.FrameMsg)
                      id, dY, Y, H: INT  END;

  PROC Handle(      V: Display.Frame;
                  VAR M: Display.FrameMsg);

  PROC New(Menu , Main: Display.Frame;
           menuH, X, Y: INT      ): Viewer;
END MenuViewers.

```

The interface represented by this definition is conspicuously narrow. There are just 2 procedures:

a generator procedure `New`, and

Returns a newly created menu viewer displaying the 2 (arbitrary) frames passed as parameters.

a standard message handler `Handle`.

Implements the entire “behavior” of an object and in particular the above message dispatching functionality.

Message handlers in Oberon are implemented in the form of procedure variables that obviously must be initialized properly at object creation time. In other words, some concrete behavior must explicitly be bound to each object, where different instances of the same object type could potentially have a different behavior and/or the same instance could change its behavior during its lifetime. Our object model is therefore *instance-centered*.

Conceptually, the creation of an object is an atomic action consisting of 3 basic steps:

1. Allocate memory block;
2. Install message handler;
3. Initialize state variables.



In the case of a standard menu *Viewer* creation:

```
NEW(V); V.handle := Handle; V.dsc      := Menu;
      V.menuH    := menuH ; V.dsc.next := Main
```

New equals to create V; open V at X,Y. Opening V needs Viewers' assistance.

Implementing Handle embodies the standard message handling strategy. This is a coarse-grained view:

```
IF message reports about user interaction THEN
  IF variant is mouse tracking THEN
    IF mouse is in menu region THEN
      IF mouse is in upper menu region and
        left key is pressed THEN
        handle changing of viewer
      ELSE delegate handling to menu-frame
    END
  ELSE
    IF mouse is in main-frame THEN
      delegate handling to main-frame
    END
  END
ELSIF variant is keyboard input THEN
  delegate handling to menu frame;
  delegate handling to main frame
END
```

```
ELSIF message defines generic operation THEN
  IF message requests copy (clone) THEN
    send copy-message to menu frame to get a copy;
    send copy-message to main frame to get a copy;
    create menu viewer clone from copies
  ELSE
    delegate handling to menu frame;
    delegate handling to main frame
  END

ELSIF message reports about change of contents THEN
  delegate handling to menu frame;
  delegate handling to main frame

ELSIF message requests change of location/size THEN
  IF operation is restore THEN
    draw viewer area and border;
    send menu frm modify-msg to make extend from H 0;
    send main frm modify-msg to make extend from H 0
  ELSIF operation is modify THEN
    IF operation is extend THEN
      extend viewer area and border;
      send modify-msg to menu frm to make it extend;
      send modify-msg to main frm to make it extend
    ELSE (*reduce*)
      send modify-msg to main frm to make it reduce;
      send modify-msg to menu frm to make it reduce;
```

```

        reduce viewer area and border
    END
    ELSIF operation is suspend THEN
        send main frm modify-msg to make reduce to H 0;
        send menu frm modify-msg to make reduce to H 0
    END
END

```

In principle, the handler acts as a message dispatcher that either processes a message directly and/or delegates its processing to the descendant frames. Note that the handler's main alternative statement discriminates precisely among the 4 basic categories of messages.

From the above outlined algorithm, handling copy messages, that is, requests for generating a copy or clone of a menu viewer, we can derive a general recursive scheme for the creation of a clone of an arbitrary frame:

1. Send copy message to each element in the list of descendants;
2. Generate copy of the original frame descriptor;
3. Attach copies of descendants to the copy of descriptor.

The essential point here is the use of new outgoing messages in order to process a given incoming message. We can regard message processing as a transformation mapping incoming

messages into a set of outgoing ones, with possible side-effects. The simplest one, the input message being simply passed on to descendant(s), is called *delegation*.

As a fine point we clarify that the above algorithm is designed to create a deep copy of a composite object (a menu viewer in our case). If a shallow copy would be desired, the descendants would not have to be copied, and the original descendants instead of their copies would be attached to the copy of the composite object.

Another example of message handling is provided by mouse tracking. Assume that a mouse message is received by a menu viewer while the mouse is located in the upper part of its menu frame and the left mouse key is kept down. This means "change viewer's height by moving its top line vertically". No message to express the required transformation of the sub-frames yet exists. Consequently, module `MenuView` takes advantage of our open (extensible) message model and simply introduces one called:

```
ModifyMsg = RECORD (Display.FrameMsg)
    id, dY, Y, H: INT
END;
```

Field `id` specifies one of the following 2 variants:

1. *extend*, or

Requests the frame to move by the vertical translation vector  $dY$  and then extend to height  $H$  at bottom.

2. *reduce*.

Requests the frame to reduce to height  $H$  at bottom and then move by  $dY$ .

In both cases  $Y$  indicates the Y-coordinate of the new lower-left corner. Fig 10 summarizes this graphically.

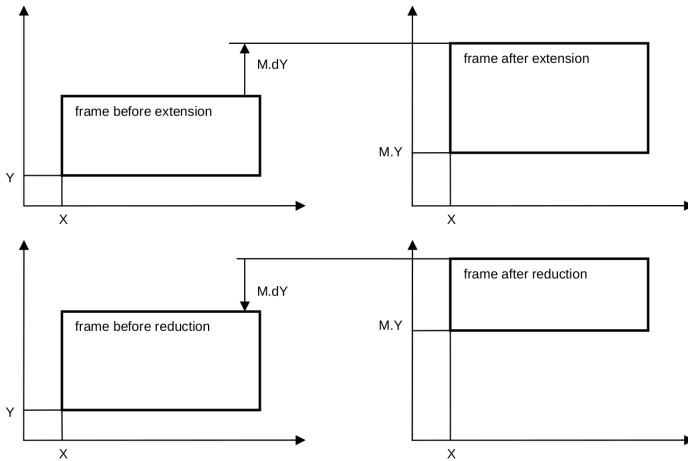


Figure 10: The modify frame operation

Messages arriving from the VMer requesting the receiving viewer to extend or reduce at its bottom are also mapped into `ModifyMsgs`. Of course, no translation,  $dY = 0$ .

The attentive reader might perhaps have asked why the standard handler is exported by `MenuViewers` at all. The

thought behind is code reusability. For example, a message handler for a subclass of menu viewers could be implemented effectively by reusing menu viewer's standard one. After having first handled all new or differing cases it would simply (super-)call the standard handler subsequently.

#### 4.4.3 *Cursor Management*

Traditionally, a cursor indicates and visualizes on the screen the current location of the caret in a text or, more generally, the current *focus of attention*. A small arrow or similar graphic symbol is typically used for this purpose. In Oberon, we have slightly generalized and abstracted this concept. A cursor is a path in the logical display area whose current position can be made visible by a *marker*.

The VMer and the cursor handler are 2 concurrent users of the same display area. Actually, we should imagine 2 parallel planes,

	displaying
one	viewers, and
the other	cursors.

If there's just 1 physical plane we take care of painting markers non-destructively, for example in inverse-video mode. Then, no precondition must be established before drawing a marker. However, in the case of a viewer task

painting destructively in its viewer's area, the area must be locked first after turning invisible all markers in the area.

The technical support of cursor management is also contained in Oberon. The corresponding API:

DEFINITION Oberon;

```
TYPE Marker = RECORD Fade, Draw: PROC(x,y: INT) END;
      Cursor = RECORD marker: Marker; X,Y: INT;
                                     on: BOOL END;
```

```
VAR Arrow, Star    : Marker;
    Mouse, Pointer: Cursor;
```

```
PROC OpenCursor(VAR c: Cursor);
PROC FadeCursor(VAR c: Cursor);
PROC DrawCursor(VAR c: Cursor;
                VAR m: Marker; X, Y: INT);
PROC MarkedViewer(): Viewers.Viewer;
PROC RemoveMarks(X, Y, W, H: INT);
...
```

END Oberon.

The state of a cursor is given by

on    its mode of visibility,  
(X, Y)    its position in the display area, and  
marker    the current marker.

Marker is an abstract data type with an interface consist-

ing of operations Fade and Draw. The main benefit of this abstraction is once more conceptual independence of the underlying hardware. For example, they can

- adapt to a given monitor hardware with built-in cursor support or, in case of absence of such support, simply
- be implemented as identical procedures (an involution) drawing the marker pattern in inverse video mode.

The functional interface to cursors consists of 3 operations:

-Cursor	to
Open	open a new cursor,
Fade	switch off the marker of an open cursor, and
Draw	extend the path of a cursor to a new position and mark it with the given marker.

We emphasize that the marker representing a given cursor can change its shape dynamically on the fly.

2 cursors are predefined:

cursor	represents	built-in marker typically
Mouse	the mouse	a small NW-pointing Arrow
Pointer	a global system pointer	a Star symbol

The pointer can be used to mark any displayed object. It serves primarily as an implicit parameter of commands.



2 assisting service procedures are added in connection with the predefined cursors:

Marked -Viewer	returns the viewer currently marked by the pointer, equivalent to <code>Viewers.This(Pointer.X, Pointer.Y)</code> .
Remove -Marks	turns invisible within a given rectangle in display area, used to lock the rectangle for its caller.

Summary the essential concept points of cursor handling:

- By virtue of the use of
  - abstract markers, and
  - the logical display area,

any potential hardware dependence is encapsulated in system modules and is therefore hidden from the application programmer. Cursors are moving uniformly within the whole display area, even across screen boundaries.

- Cursor handling is decentralized by delegating it to the individual handlers that are installed in viewers. Typically, a handler reacts on the receipt of a mouse tracking message by drawing the mouse cursor at the indicated new position. The benefit of such individualized handling is flexibility. For example, a smart local handler might choose the shape of the visualizing

marker depending on the exact location, or it might force the cursor onto a grid point.

- Even though cursor handling is decentralized, there is some intrinsic support for cursor drawing built into the Cursor declaration. Cursors are full value objects and, as such, can "memorize" their current state. Consequently, the interface operations FadeCursor and DrawCursor need to refer to the desired future state only.
- Looking at the VMer as one user of the display area, the cursor handler is the 2nd (and logically concurrent) user of the same resource. If there is just one physical plane implementing the display area, any region must be locked by a current user before destructive painting. Therefore, markers are usually painted non-destructively in inverse-video mode.

Let us now recapitulate the entire section.

- The central resource managed by the display subsystem is the *logical display area* whose purpose is abstraction from the underlying display monitor hardware.

The display area is primarily used by the VMer for the accommodation of tracks and viewers, which are merely the 1st 2 levels of a potentially unlimited nested

hierarchy of display frames. For example, standard menu viewers contain 2 subordinate frames:

- a menu frame, and
  - a main frame of contents.
- Viewers are treated as black boxes by the VMer and are addressed via messages.

Viewers and, more generally frames, are used as elements of *message-based interfaces* connecting the display subsystem with other subsystems like

- the task scheduler, and
  - the various document managers.
- Finally, the display area is also the living room of cursors. In Oberon, a cursor is a marked path, 2 standard cursors Mouse and Pointer are predefined.

#### 4.5 RASTER OPERATIONS

In ?? we introduced the display area as an abstract concept, modeled as a 2-dimensional Cartesian plane. So far, this view of the display space was sufficient because we were interested in its global structure only and ignored contents completely. However, if we are interested in what is displayed, we need to reveal more details about the model.

The Cartesian plane representing the display area is discrete. We consider points in the display area as grid points or picture elements (*pixel*), and we assume contents to be generated by assigning colors to the pixels. For the moment, the number of possible colors a pixel can attain is irrelevant. In the binary case of 2 colors we think of one representing background color and the other foreground color.

The most elementary operation generating contents in a discrete plane is "set color of pixel" or "set pixel" for short. While a few drawing algorithms directly build on this atomic operation, block-oriented functionality (traditionally called *raster operations*) plays a much more important role in practice. A *block* is a rectangular area of pixels whose bounding lines are parallel to the axes of the coordinate system.

Raster operations are based on a common principle:

- A source block of width  $SW$  and height  $SH$  is placed at a given destination point  $(DX, DY)$  in the display area.

- In the simplest case, the destination block (DX, DY, SW, SH) is plainly overwritten by the source one.
- In general, the new value of each pixel in the destination block is a combination of its old value and the value of the corresponding source pixel:

$$d := F(s, d)$$

$F$  is sometimes called the combination mode of the raster operation. The raster is stored as an array of values of type SET, each set representing 32 black/white pixels. The modes of combining source and destination is implemented by the following set operations:

mode	operation
replace	$s$
paint	$s + d$ (or)
invert	$s / d$ (xor)

Note that invert is equivalent with inverse video mode if  $s$  is TRUE for all pixels. There are many different variants of raster operations. Some refer to a source block in the display area, others specify a constant pattern to be taken as source block. Some variants require replication of the source block within a given destination block (DX, DY, DW, DH) rather than simple placement.

The challenge when designing a raster interface is finding a unified, small and complete set of raster operations that covers all needs, in particular including the need of placing character glyphs. The amazingly compact resulting set of Oberon raster operations is exported by module Display:

DEFINITION Display;

```

    CONST black = 0; white = 1;           (*colors*)
          replace = 0; paint = 1; invert = 2;
                                     (*operation modes*)
    PROC Dot      (col,      x,y,      mode: INT);
    PROC ReplConst (col,      x,y,w,h, mode: INT);
    PROC CopyBlock (sx,sy, w,h, dx,dy,  mode: INT);
    PROC CopyPattern(col, patadr, x,y,  mode: INT);
    PROC ReplPattern(col, patadr, x,y,w,h, mode: INT);

```

END Display.

In the parameter lists of the above raster operations, mode is the mode of combination (replace, paint, or invert). CopyBlock copies the source block (sx, sy, w, h) to position (dx, dy) and uses mode to combine new contents in the destination block (dx, dy, w, h). It is assumed tacitly that the numbers of colors per pixel in the source block and in the destination area are identical. It is perhaps informative to know that CopyBlock is essentially equivalent with the famous BitBlt (bit block transfer) in the SmallTalk project [Goldberg]. In Oberon, it is used primarily for scrolling contents within a viewer.

The remaining raster operations use a constant pattern. Patterns are implemented as arrays of bytes, and the parameter `patadr` is the address of the relevant pattern. The 1st 2 bytes indicate width  $w$  and height  $h$  of the pattern. Pattern data are given as a sequence of bytes to be placed into the destination block from left to right and from bottom to top. Each line takes an integral number of bytes. Hence, the number of data bytes is  $((w + 7)/8) * h$ . Fig 11 shows an example:




Figure 11: A pattern and its encoding as an array of bytes (in hex)

Some standard patterns are included in module `Display` and exported as global variables. Among them

patterns	intended to represent
arrow,	the cursor,
hook, and	the caret, and
star	the marker.

A group of predefined patterns supports drawing graphics.

- `col` in the pattern-oriented raster operations specifies the pattern's foreground color. Colors black (background) and white are predefined.
- `CopyPattern` copies the pattern to location `x, y` in the display area, using the given combination mode. It is probably the most frequently used operation of all because it is needed to write text.
- `ReplPattern` replicates the given pattern to the given destination block. It starts at bottom left and proceeds from left to right and from bottom to top.
- `Dot` and `ReplConst` are special cases of `CopyPattern` and `ReplPattern` respectively, taking a fixed implicit pattern consisting of a single foreground pixel.
  - `Dot` is exactly our previously mentioned "set pixel".
  - `ReplConst` is used to draw horizontal and vertical lines of various widths.

The raster operations are a prominent example of the use of Oberon's data type `SET`. Formally, variables are sets of integers between 0 and 31. Here, they are taken as sets of bits numbered from 0 to 31. We consider the replication of 1's (mode = replace or paint) in the rectangle with origin `x, y`, width `w`, and height `h`. Every line consists of 1024 pixels, or 32 words. `a1, ar, a0, a1` are addresses.



```

VAR al, ar, a0, a1: INT;
    left, right, pixl, pixr: SET;

al := base + y*128;
ar := ((x+w-1) DIV 32)*4 + al;
al := (x DIV 32)*4 + al;
left := {(x MOD 32) .. 31};
right := {0 .. ((x+w-1) MOD 32)};
FOR a0 := al TO al + (h-1)*128 BY 128 DO
    SYSTEM.GET(a0, pixl);
    SYSTEM.GET(ar, pixr);
    SYSTEM.PUT(a0, pixl + left);
    FOR a1 := a0+4 TO ar-4 BY 4 DO
        SYSTEM.PUT(a1, {0 .. 31});
    END
    SYSTEM.PUT(ar, pixr + right)
END

```

The definition (and even more so the implementation) of module `Display` provides support for a restricted class of possible hardware configurations only. Any number of display monitors is theoretically possible. However, they must be mapped to a regular horizontal array of predefined cells in the display area. Each cell is vertically split into 2 congruent regions, where the corresponding monitor is supposed to be able to select and display one of the 2 regions alternatively. Finally, it is assumed that all cells hosting black-and-white monitors are allocated to the left of all cells

hosting color monitors. Fig 12 gives an impression of such a configuration.




Figure 12: General, regular cell structure of display area

Under these restrictions any concrete configuration can be parameterized by the variables of the definition above. Unit, Width, and Height specify the extent of a displayed region, where Width and Height are width and height in pixel units, and Unit is the size of a pixel in units of  $1/36'000$  mm. This unit is a common divisor of all of the standard metric units used by the typesetting community, like mm, inch, Pica point and point size of usual printing devices. Bottom and UBottom specify the bottom y-coordinate of the primary region and the secondary region respectively. Finally, Left and ColLeft give the left x-coordinate of the area of black-and-white monitors and of color monitors respectively.

## 4.6 STANDARD DISPLAY CONFIGS AND TOOLBOX

Let us now take up again our earlier topic of configuring the display area. We have seen that no specific layout of the display area is distinguished by the general viewer management itself. However, some support of the familiar standard Oberon display look is provided by Oberon.

In the terminology of this module, a standard configuration consists of one or several horizontally adjacent displays, where a display is a pair consisting of 2 equal height tracks,

*a user track on the left, and*  
*a system track on the right.*

Note that even though no reference to any physical monitor is made, a display is typically associated with a monitor in reality. This is the relevant excerpt of the definition:

DEFINITION Oberon;

```
PROC OpenDisplay(UW, SW, H: INT);
PROC OpenTrack(X, W: INT);
PROC DisplayWidth (X: INT): INT;
PROC DisplayHeight(X: INT): INT;
PROC UserTrack    (X: INT): INT;
PROC SystemTrack  (X: INT): INT;
PROC AllocateUserViewer (DX: INT; VAR X,Y: INT);
PROC AllocateSystemViewer(DX: INT; VAR X,Y: INT);
```

END Oberon.

- `OpenDisplay` initializes and opens a new display of

parameter	dimension
<code>UW</code>	User track Width,
<code>SW</code>	System track Width,
<code>H</code>	Height.

- `OpenTrack` overlays the sequence of existing tracks spanned by the segment  $[X, X + W)$  by a new track.
- Both `OpenDisplay` and `OpenTrack` take from the client the burden of creating a filler viewer.
- `DisplayWidth`, `DisplayHeight`, `UserTrack` and `SystemTrack` return width or height of the respective structural entity located at position  $X$  in the display area.
- `AllocateUserViewer` and `AllocateSystemViewer` make proposals for the allocation of a new viewer in the desired track of the display located at  $DX$ .

In 1st priority, the location is determined by the system pointer that can be set manually. If the pointer is not set, a location is calculated on the basis of some heuristics whose strategies rely on different splitting fractions that are applied in the user track and in the system track respectively, with the aim of generating aesthetically satisfactory layouts.

In addition to the programming interface provided by Oberon for the case of standard display layouts, the display management section in the System toolbox provides a user interface:

```
DEFINITION System; (*Display management*)
  PROC Open;    (*viewer*)
  PROC Close;   (*viewer*)
  PROC CloseTrack;
  PROC Recall;  (*most recently closed viewer*)
  PROC Copy;   (*viewer*)
  PROC Grow;   (*viewer*)
  PROC Clear;  (*system log*)
END System.
```

In turn, these commands are called to

- open a text viewer in the system track,
- close a viewer,
- close a track,
- recall (and reopen) the most recently closed viewer,
- copy a viewer,
- grow a viewer, and
- clear the system log.

Close, Copy, Grow, CloseTrack, and Recall are generic commands. The 1st 3 are typically included in the title bar of a menu viewer. Their detailed implementations follow subsequently.



TEXT

---

At the beginning of the computing era, text was the only medium mediating information between users and computers. Not only was a textual notation used to denote all kinds of data and objects via names and numbers (represented by sequences of characters and digits respectively), but also for the specification of programs (based on the notions of formal language and syntax) and tasks. Actually, not even the most modern and most sophisticated computing environments have been able to make falter the dominating role of text substantially. At most, they have introduced alternative models like graphical user interfaces (GUI) as a graphical replacement for *command lines*.

There are many reasons for the popularity of text in general and in connection with computers in particular. To name but a few:

- Text containing any arbitrary amount of information can be built from a small alphabet of widely standardized elements (characters),

- their building pattern is extremely simple (lining up elements), and
- the resulting structure is most elementary (a sequence).

And perhaps most importantly, *syntactically structured text can be parsed and interpreted by a machine.*

In computing terminology, sequences of elements are called *files* and, in particular, sequences of characters are known as *text files*. Looking at their binary representation, we find text files excellently suited to be stored in computer memories and on external media. Remember that individual characters are usually encoded in 1 byte each (ASCII). We can therefore identify the binary structure of text files with sequences of bytes, matching perfectly the structure of any underlying computer storage. We should recall at this point that, with the possible exception of line-break control characters, rendering information is not part of ordinary text files. For example, the choices of character style and of paragraph formatting parameters are entirely left to the rendering interpreter.

Unfortunately, in conventional computing environments, text is merely used for input/output, and its potential is not nearly exploited optimally. Input texts are typically read from the keyboard under control of some text editor, interpreted and then discarded. Output text is volatile. Once displayed on the screen it is no longer available to any other



parts of the program. The root of the problem is easily located: Conventional OSes neither feature an integrated management nor an abstract programming interface (PI) for texts.

Of course, such poor support of text on the level of programming must reflect itself on the user surface. More often than not, users are forced to retype a certain piece of text instead of simply copy/pasting it from elsewhere on the screen. Investigations have shown that, in average, up to 80% of required input text is already displayed somewhere.

Motivated by our positive experience with integrated text in the Cedar system [Teitelman] we decided to provide a central text management in Oberon at a sufficiently low system level. However, this is not enough. We actually need an abstract PI for text, that is, an abstract data type `Text`, together with a complete set of operations. We shall devote 5.1 to the explanation of this data type. In 5.2, we take a closer look at the basic text management, including data structures and algorithms used for the implementation of type `Text`.

Text frames are a special class of display frames. They appear typically (but not necessarily) as frames within a menu viewer (see 4.4.2). Their role is double-faced:

- a) Rendering text on the display screen, and

- b) interpreting interactive editing commands.

The details will be discussed in [5.3](#).

With the aim of exploiting the power of modern bitmap-displays and also of reusing the results of earlier projects in the field of digital font design, we decided in favor of supporting “rich texts” in Oberon, including graphical attributes and in particular font specification. In [5.4](#) we shall explain the font machinery, starting from an abstract level and proceeding down to the level of raster data.

## 5.1 TEXT AS AN ABSTRACT DATA TYPE

The concept of abstraction is arguably the most important achievement of PL development. It provides a powerful tool to create simplified views of complicated things and connections. 2 prominent examples of program abstractions

	embodying simplified views on
abstract data types	a certain kind of data, and
definitions (interfaces)	a certain piece of program.

We shall now give a precise definition of the notion of text in Oberon by presenting it as an abstract data type. It is important not to confuse this type with the far less powerful one `String` as it is often supported by advanced PLs. Here we carefully avoid revealing any implementation aspects of the abstract type `Text`. Our viewpoint is that of an application program operating on text abstractly or using it as a medium of communication.

Nevertheless, let us first use a symbolic looking glass to get a refined understanding of the concept of character in the context of rich texts. We know that each character represents a textual element of information. If displayed, it also refers to some specific graphical pattern, often called *glyph*. In Oberon, we do justice to both aspects by thinking of the ASCII as an index into a font that is into a set of glyphs of the

same style. Representing characters as pairs (*font*, *ref*), where *font* designates a font and *ref* the character's ASCII code and adding 2 more attributes *color* and *vertical offset*, we get to a quadruple representation (*font*, *ref*, *col*, *voff*) of characters. The components *font*, *color*, and *vertical offset* together are often referred to as *looks*. With that, we can now define a (rich) text as a *sequence of characters with looks*. We shall treat the topic of fonts and glyphs thoroughly in 5.4.

For the moment, however, let us continue our discussion of the abstract data type *Text*. Formally, we define it as

```
Text = POINTER TO TextDesc;
TextDesc = RECORD len: INT; notify: Notifier END;
```

There is only 1

- state variable *len*, and

Represents the current length of the described text (i.e. the number of characters in the sequence).

- method *notify*.

Occasionally called *after-method*. Notify interested clients of state changes.

By definition, each abstract data type comes with a complete set of operations. In the case of *Text*, 3 groups corresponding to 3 different topics need to be considered,

1. loading (from file), storing (to file),
2. editing, and
3. accessing (reading and writing) respectively.

#### 5.1.1 *Loading and Storing Text*

The 1<sup>st</sup> file group introduces a pair of mutually inverse operations, *internalize* and *externalize*, respectively meaning

- load from file and build up an internal data structure, and
- serialize the internal data structure and store it in file.

There are 3 corresponding procedures:

```
PROC Open (T: Text; name: ARRAY OF CHAR);  
PROC Load (T: Text; f: Files.File; pos: INT;  
           VAR len: INT);  
PROC Store(T: Text; f: Files.File; pos: INT;  
           VAR len: INT);
```

Logical entities like texts are stored on external media in form of *section*, addressed by pair (file, pos) consisting of

file descriptor, and starting position.

In general, the structure of section obeys:

```
section = identification type length contents.
```

Open	internalizes a named text file (consisting of a single text section)
Load	internalizes an arbitrary text section starting at (f, pos)
Store	externalizes a text section to (f, pos)
T	designates the internalized text
len	returns the length of the section

Note that in case of Load the identification of section must have been read and consumed before the loader is called.

### 5.1.2 *Editing Text*

Next group operations support text editing, comprising:

```
PROC Append      (T: Text;          B: Buffer);
PROC Insert      (T: Text; pos: INT; B: Buffer);
PROC Delete      (T: Text; beg, end: INT);
PROC ChangeLooks(T: Text; beg, end: INT; sel: SET;
                  fnt: Fonts.Font; col, voff: INT);
```

Again, we should 1st explain the types of parameters:

- Delete and ChangeLooks each take a stretch of text as an argument which, by definition, is an interval [beg, end) within the given text.
- In the parameter lists of Insert and Append we recognize a new data type Buffer.

Buffers are a facility to hold anonymous sequences of characters. Buffer presents itself again as an abstract data type:

```
Buffer  = POINTER TO BufDesc;  
BufDesc = RECORD len: INT END;
```

len specifies the current length of the buffered sequence.

The following represent the intrinsic operations on buffers:

```
PROC OpenBuf(B: Buffer);  
PROC Copy(SB, DB: Buffer);  
PROC Save(T: Text; beg, end: INT; B: Buffer);
```

Their function, in turn:

- opening a given buffer B,
- copying buffer SB to DB, and
- saving a stretch [beg, end) of text in a given buffer, recalling the most recently deleted stretch of text and putting it into B.

Buffer is used as an auxiliary data type in editing.

- Delete d~s the given stretch [beg, end) within text T,
- Insert i~s the buffer's contents at position within T, and
- Append(T, B) is a shorthand of Insert(T, T.len, B).

Note that, as a side-effect of Insert and Append, the buffer involved is emptied.

Finally, ChangeLooks allows to change selected looks within the given stretch [beg, end) of T. sel is a mask selecting a subset of the looks set {font, color, vertical offset}.

It is time now to come back to the notifier concept. Recapitulate that notify is an *after-method*. It must be installed by the client when opening the text and is called at the end of every editing operation. Its signature

```
Notifier = PROC(T: Text; op, beg, end: INT);
```

op, beg, and end report

- about the operation that calls the notifier, and
- on the affected stretch [beg, end) of the text.

There're 3 possible ops corresponding to different editing operations:



op	operation
delete	Delete
insert	Insert(, Append)
replace	ChangeLooks

By far the most important application of the notifier is updating the display, i.e. adjusting all affected views of the text that are currently displayed to the new state of the text (the model). We shall come back to this important matter when discussing text frames in 5.3.

In concluding this part it is worth noting that the operations just discussed have been designed to be equally useful for

- interactive text editors, as well as
- programmed text generators/manipulators.

### 5.1.3 Accessing Text

The 3rd and last group of operations on texts: Accessing, i.e. *reading* and *writing*.

According to the principle of separation of concerns, one of our guiding principles, the access mechanism operates on extra aggregates (called *readers* and *writers*) rather than on texts themselves.

Readers are used to read texts sequentially. Their type:

```
Reader = RECORD eot: BOOL; (*end of text*)  
              fnt: Fonts.Font; col, voff: INT  
            END;
```

A reader must 1st be opened at the desired position in the text before it can then be moved forward incrementally by reading character-by-character. Its state variables indicate end-of-text and expose the looks of the character last read. The corresponding operators are

```
PROC OpenReader(VAR R: Reader; T: Text; pos: INT);  
PROC Read      (VAR R: Reader; VAR ch: CHAR);
```

- OpenReader sets up a reader R at position pos in text T.
- Read returns the character at the current position of R and makes R move to the next position.

The current position of R is returned by a call to the function:

```
PROC Pos(VAR R: Reader): INT;
```

In 3 we learned that commands plus parameter lists are often embedded in ordinary texts. When interpreting such commands, the underlying text appears as a sequence of tokens like name, number, special symbol etc. much rather than as a sequence of characters. Therefore, we have adopted the well-known concepts of syntax and scanning from the

discipline of compiler construction, including functional support. The Oberon scanner recognizes tokens of some universal classes. They are name, string, integer, real, longreal, and special character. The exact syntax of universal Oberon tokens is:

```

token = name | string | integer | real | spexchar
name  = ident { .ident }
ident = letter { letter | digit }
string = ' ' { char } ' '
integer = [+|-] number
real = [+|-] number . number [E [+|-] number]
number = digit { digit }
spexchar = any character except letters, digits,
           space, tab, or carriage-return

```

Scanner is defined correspondingly as

```

Scanner = RECORD (Reader) nextCh: CHAR;
           line, class, i, len: INT;
           x: REAL; c: CHAR;
           s: ARRAY 32 OF CHAR
END;

```

This type is actually a variant record type with class as discriminating tag. Depending on its class the value of the current token is stored in one of the fields *i*, *x*, *c*, or *s*.

- `len` gives the length of `s`,
- `nextCh` typically exposes the character terminating the current token, and
- `line` counts the number of lines scanned.

The operations on scanners:

```
PROC OpenScanner(VAR S: Scanner; T: Text; pos: INT);
PROC Scan      (VAR S: Scanner);
```

They correspond exactly to their counterparts `OpenReader` and `Read` respectively. Writers are dual to readers. They serve the purpose of creating and extending texts. However, again, they do not operate on texts directly. Rather, they act as self-contained aggregates, continuously consuming and buffering textual data.

The formal declaration of `Writer` resembles `Reader`:

```
Writer = RECORD buf: Buffer;
               fnt: Fonts.Font; col, voff: INT
            END;
```

`buf` is an internal buffer containing the consumed data. `fnt`, `col`, and `voff` specify the current looks for the next character consumed by this writer.

The following procedures constitute the `Writer` API:

```
PROC OpenWriter(VAR W: Writer);  
PROC SetFont   (VAR W: Writer;  fnt: Fonts.Font);  
PROC SetColor  (VAR W: Writer;  col: INT);  
PROC SetOffset (VAR W: Writer;  voff: INT);
```

OpenWriter opens a new writer with an empty buffer.

SetFont, SetColor, and SetOffset set the respective current look. For example, SetFont(W, fnt) is equivalent with `W.fnt := fnt`. These procedures are included because fnt, col, and voff are read-only for clients.

The question may arise how data is produced and transferred to writers. The answer is a set of writer procedures, each of them handling an individual data type:

```
PROC Write      (VAR W: Writer; ch: CHAR);  
PROC WriteLn    (VAR W: Writer);  
PROC WriteString(VAR W: Writer; s: ARRAY OF CHAR);  
PROC WriteInt   (VAR W: Writer; x, n: INT);  
PROC WriteHex   (VAR W: Writer; x: INT);  
PROC WriteReal  (VAR W: Writer; x: REAL; n: INT);  
PROC WriteRealFix(VAR W: Writer; x: REAL; n, k: INT);  
PROC WriteClock (VAR W: Writer; d: INT);
```

The following is schematic fragment of a client program that creates textual output:

```
open writer; set desired font;
```

```
REPEAT
    process;
    write result to writer;
    append writer buffer to output text
UNTIL ended
```

Of course, writers can be reused. For example, a single global writer is typically shared by all of the procedures within a module. In this case, the writer needs to be opened just once at module loading time.

Typically, however, accessing aggregates are of a transient nature and are bound to a certain activity, which manifests itself in their allocation on the stack without any possibility of referencing them from the outside of the activity, in contrast to the underlying texts that are allocated on the system heap and have a much longer life time.

Let us summarize:

- Text in Oberon is a powerful abstract data type with intrinsic operations from 3 areas:
  1. Loading/storing,
  2. editing, and
  3. accessing (reading/writing).

- The latter 2 areas on their part introduce further abstract types Buffer, Reader, Scanner, and Writer.

In combination they guarantee a clean separation of very different concerns. The benefits of such a rigorous decoupling are numerous. For example, it makes it possible to freely choose (and vary) the granularity at which a text and its views are updated.

- Finally, an after-method is used to allow context-dependent post-processing of editing operations. It is used primarily for preserving consistency between text models and their views.

## 5.2 TEXT MANAGEMENT

The art and challenge of modularization lie in finding an effective decomposition of a topic into modules with relatively thin interfaces or, in other words, into modules with a great potential for information hiding. Text systems provide a welcome opportunity of an exercise. A closer analysis immediately leads to the following separate concerns corresponding to the components Model, View and Controller of the MVC scheme: Text management, text rendering, and text editing. If we combine View and Controller and add an auxiliary font handling module Fonts, we arrive at the following three-module import hierarchy:

Module	Object type	Service
TextFrames	Frame	Text rendering and editing
Texts	Text	Text management
Fonts	Font	Font management

Note that, in contrast to the display-subsystem, the associated object types are not connected hierarchically here.

Separate Sections 5.3 and 5.4 will be devoted to modules TextFrames and Fonts respectively. In the current Section we focus on module Texts. Regarding it as a model of the abstract data type Text presented in the previous Section, its definition is congruent with the specification of the abstract data type itself, and we need not repeat it here.



The main topics of this Section are internal representation and file representation of texts. We first emphasize that the internal representation of a text is a completely private matter of module Texts that is encapsulated and hidden from clients. In particular, the representation could be changed at any time without invalidating any single client. In principle, the same is true for the file representation. However, stability is of paramount importance here because files serve the additional purposes of backing up text on external media and of porting text to other environments. Our choice of an internal representation of text was determined by a catalogue of requirements and desired properties. The wish list looks like this:

1. lean data structure
2. closed under editing operations
3. efficient editing operations
4. efficient sequential reading
5. efficient direct positioning
6. super efficient internalizing
7. preserving file representations

With the exception of 5, we found these requirements met perfectly by an adequately generalized variant of the piece

list technique that was originally used for Xerox PARC's Bravo text editor and also for ETH's former document editors Dyna and Lara [Gutknecht]. The original piece list is able to describe a vanilla text without looks. It is based on:

- A text is regarded as a sequence of pieces, where a piece is a section of a text file consisting of a sequence of contiguous characters.
- Each piece is represented by a descriptor (*f*, *pos*, *len*), designating *file*, starting *position*, and *length* respectively. The whole text is represented as a list of piece descriptors (in short: *piece list*). The editing operations operate on the list rather than on pieces themselves.

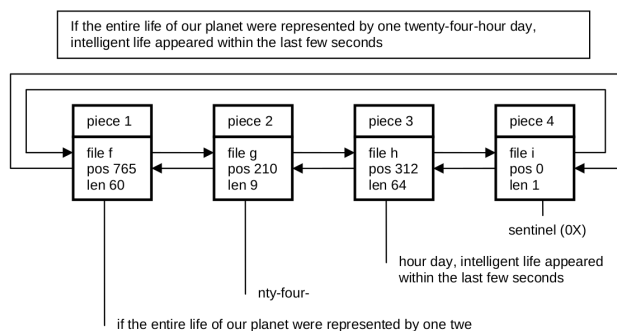


Figure 13: Piece chain representing a text

Fig 13 shows a typical piece list representing (the current state of) a text. Investigating the effects of the basic editing operations *delete* and *insert* on this piece list, we end up with these algorithms:

```

delete stretch [beg, end) of text = BEGIN
  split pieces at both beg and end;
  remove p-descriptors from beg to end from the chain
END

```

```

insert stretch of text at pos = BEGIN
  split piece at pos;
  insert p-descriptors representing the stretch there
END

```

Splitting is superfluous if the desired splitting point happens to coincide with the piece beginning. Fig ?? and 16 show the resulting list after a delete, and an insert respectively:

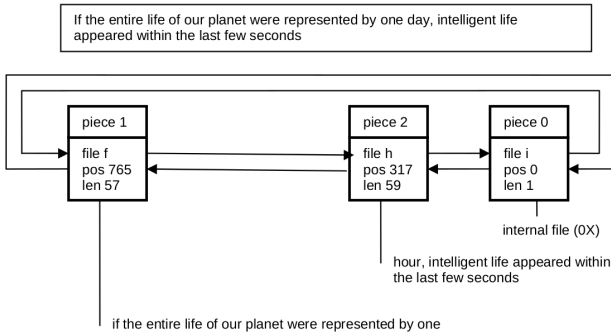


Figure 14: Piece chain after delete operation

Checking our above wish list we immediately recognize the requirements 1, 2, and 3 as met. Requirement 4 is also met assuming an efficient mechanism for direct positioning in files. Requirement 6 can be checked off because the piece

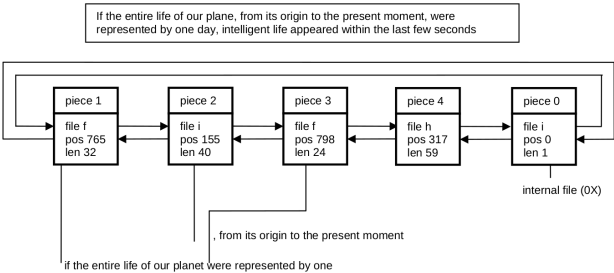


Figure 15: Piece chain after insert operation

list initially consists of a single piece spanning the entire text file. Finally, requirement 7 is met simply because the operations do not affect file representations at all.

In Oberon we adapted the piece list technique to texts with looks (*rich texts*). Formally, we 1st define a run as a stretch of text whose characters show identical looks. Now, we require the piece list to subordinate itself to the run structure. This obviously means that every piece needs to be contained within a single run. Fig 16 visualizes such a compliant piece list representing a text with varying looks. There are only 2 new aspects compared to the original version of the piece list discussed above: An additional operation to change looks and the initial state of the piece chain.

```
change looks in a stretch [beg, end) of text = BEGIN
    split pieces at both beg and end;
    change looks in piece descriptors from beg to end
END
```

This shows that wish list requirements 2 and 3 still satisfied.




Figure 16: Generalized piece chain representing a rich text

Initially, the pieces are identical with runs, and the number of elements in the piece list is equal to the one of runs. Because this number is typically small in comparison with the total characters number in a text, requirement 6 is still met.

We conclude that the new aspects do not invalidate the positive rating given above to the piece technique with regard to wish list requirements 1-4, 6 and 7. However, requirement of efficient direct positioning remains. The problem is the necessity to scan through the list sequentially to locate the piece that contains the desired position. We investigated different solutions of this efficiency problem. They are based on different data structures connecting the piece descriptors, among them a piece tree and a variant of the piece list featuring an additional long-distance link like in a skip-list.

Eventually, we decided in favor of a simpler solution that we can easily justify by pointing out that the typical editing scenario is zooming into a local region of text, i.e. positioning at an arbitrary location once and subsequently positioning at locations in its immediate neighborhood many times. Therefore, an appropriate solution is caching the most recently calculated values (*pos*, *piece*) of the translation map. Of course, this does not solve the problem of cache misses. Notice, however, that this problem is acute only in the case of extremely long piece lists that do not occur in ordinary texts and editing sessions.

We shall now illustrate the piece technique in detail at the example of 2 important but basic operations: Insert and read. Let us start with an overview of the data types involved. Apart from some auxiliary private variables marked with an arrow, the types *Text*, *Buffer*, and *Reader* are already familiar to us from the previous Section. *Piece* is completely private and hidden from the clients.

```
Text = POINTER TO TextDesc;
Notifier = PROC(T: Text; op, beg, end: INT);

TextDesc = RECORD len: INT; notify: Notifier;
               --> trailer: Piece;
               --> org: INT;
               --> pce: Piece
            END;
```

```

Buffer = POINTER TO BufDesc;
BufDesc = RECORD len: INT;
               --> header, last: Piece
            END;

Reader = RECORD eot: BOOL; fnt: Fonts.Font;
               col, voff: INT;
               --> ref: Piece;
               --> org, off: INT;
               rider: Files.Rider
            END;

--> Piece = POINTER TO PieceDesc;
--> PieceDesc = RECORD f: Files.File;
                  off, len: INT;
                  fnt: Fonts.Font;
                  col, voff: INT;
                  prev, next: Piece
                END;

```

As depicted in Fig 13, the piece list is implemented as a doubly linked list with a sentinel piece closing it to a ring. The field trailer in TextDesc points to the sentinel piece. Fields org and pce implement a translation cache consisting of merely one entry (org, pce). It links a position org with a piece pce. The fields header and last in Buffer refer to the implementation of buffers as piece lists. They point to

the first and last piece descriptors respectively. Finally, the fields *ref*, *org*, and *off* in *Reader* memorize the current piece, its origin, and the current offset within this piece. The fields *f*, *off*, and *len* in *Piece* specify the underlying file, starting position in the file, and length of the piece. *fnt*, *col*, and *voff* are its looks. Finally *prev* and *next* are pointers to the previous and next piece in the list respectively.

*FindPiece* and *SplitPiece* are auxiliary procedures that are used by almost all piece-oriented operations.

```
PROC FindPiece(T: Text; pos: INT; VAR org: INT; VAR p: Piece)
  VAR p: Piece; porg: INT;
  BEGIN p := T.pce; porg := T.org;
  1) IF pos >= porg THEN
      WHILE pos >= porg + p.len DO INC(porg, p.len); p := p.next;
  2) ELSE p := p.prev; DEC(porg, p.len);
      WHILE p < porg DO p := p.prev; DEC(porg, p.len) END
  END;
  3) T.pce := p; R.org := porg; (*update cache*)
      pce := p; org := porg
  END FindPiece;
```

Explanations (referring to the line numbers in the above code excerpt) 1) search to the right (next) 2) search to the left (prev) 3) update cache if more than 50 pieces traversed

```
1) PROC SplitPiece (p: Piece; off: INT; VAR pr: Piece);
```



```

    VAR q: Piece;
BEGIN
2)  IF off > 0 THEN NEW(q);
    q.fnt := p.fnt; q.col := p.col; q.voff := p.voff;
    q.len := p.len - off;
    q.f := p.f; q.off := p.off + off;
    p.len := off;
3)  q.next := p.next; p.next := q;
4)  q.prev := p; q.next.prev := q;
    pr := q
    ELSE pr := p
    END
END SplitPiece;

```

Explanations: 1) return right part piece pr after split 2) generate new piece only if remaining length > 0 3) insert new piece in forward chain 4) insert new piece in backward chain

Procedure Insert handles text insertion. It operates on a buffer that contains the stretch of text to be inserted:

```

PROC Insert (T: Text; pos: INT; B: Buffer);
    VAR pl, pr, p, qb, qe: Piece; org, end: INT;
BEGIN
1) FindPiece(T, pos, org, p); SplitPiece(p, pos - org, pr);
2) IF T.org >= org THEN
    T.org := org - p.prev.len; T.pce := p.prev

```

```

    END;
    pl := pr.prev; qb := B.header.next;
3) IF (qb # NIL) & (qb.f = pl.f) & (qb.off = pl.off + pl.len
    & (qb.fnt = pl.fnt) & (qb.col = pl.col) & (qb.voff = pl
    pl.len := pl.len + qb.len; qb := qb.next
    END;
    IF qb # NIL THEN qe := B.last;
4)   qb.prev := pl; pl.next := qb; qe.next := pr; pr.prev :=
    END;
5) T.len := T.len + B.len; end := pos + B.len;
6) B.last := B.header; B.last.next := NIL; B.len := 0;
7) T.notify(T, insert, pos, end)
    END Insert;

```

Explanations: 1) split piece to isolate point of insertion 2) adjust cache if necessary 3) merge pieces if possible 4) insert buffer 5) update text length 6) empty buffer 7) notify

Procedure Read implements sequential reading of characters in texts. It operates on a text reader:

```

PROC Read (VAR R: Reader; VAR ch: CHAR);
BEGIN
1) Files.Read(R.rider, ch); R.fnt := R.ref.fnt; R.col := R.
    INC(R.off);
2) IF R.off = R.ref.len THEN
3)   IF R.ref.f = WFile THEN R.eot := TRUE END;
    R.org := R.org + R.off; R.off := 0;

```

```

4)   R.ref := R.ref.next; R.org := R.org + R.off; R.off := 0;
5)   Files.Set(R.rider, R.ref.f, R.ref.off)
      END
      END Read;

```

Explanations: 1) read character from file and update looks in reader 2) if piece boundary reached 3) check if sentinel piece reached 4) move reader to next piece 5) position file rider

Procedure Read is typically used as a primitive by text scanners and in particular by the built-in scanner Scan for the recognition of universal tokens, as they were defined in the previous section. Scanning is a rather complex operation that, for example, includes the conversion of a sequence of digits into an internal floating-point representation and vice-versa. Scanning a real number involves recognizing  $m$  and  $d$ , and computing  $x = m \cdot 10^d$ . This is done using procedure Ten( $d$ ) computing  $10^d$  by repeated multiplication maintaining the invariant  $t \cdot p^n = 10n_0$ , where  $n_0$  is the initial value of  $n$ .

```

PROC Ten(n: INT): REAL;
VAR t, p: REAL;
BEGIN t := 1.0; p := 10.0;
WHILE n > 0 DO
  IF ODD(n) THEN t := p * t END ;
  p := p*p; n := n DIV 2

```

```
END ;  
RETURN t  
END Ten;
```

Writing a real number in decimal form is more complicated. It involves the computation of  $m$  and  $d$  from  $x = n \cdot 2^e$  so that  $x = m \cdot 10^d$  with  $1.0 \leq m < 10.0$ . First,  $e$  is obtained with the standard function  $\text{UNPK}(x, e)$ , then  $d$  is computed (from the relationship  $10^k = 2^k \cdot \log_2(10)$ ) as  $d = e / \log_2(10)$ . In order to avoid a real division for obtaining  $d$ , we use the approximation  $1.0 / \log_2(10) = 77 \text{ DIV } 256$ , and then compute  $x := x / \text{Ten}(e)$  or  $x := x * \text{Ten}(-e)$ . Further details are to be taken from the listings of `WriteReal` and `WriteRealFix`.

In spite of its apparent simplicity the piece list technique interoperates with other system components in quite a subtle way. For example, after a while of editing, there are typically numerous cross references between the documents involved. In other words, pieces of one document may point to foreign files that is to files that were originally associated with other documents. As a consequence, the FS must either employ some smart garbage collection algorithm or not recycle file pages at all, even if a new version of a file of the same name has been created in the meantime.

A problem of another kind, again affecting the FS, arises if, say, a single text line is composed of several small pieces. Then, reading this line sequentially may necessitate several

jumps to different positions in different files at a high pace. Depending on the quality of the file buffering mechanism, this may lead to significantly hesitant mouse tracking.

And finally, typed characters that are supposed to be inserted into a text need to be stored on the so-called keyboard file. For this (continuously growing) file, several readers and one writer must be allowed to coexist concurrently.

As a consequence, the following qualities of the underlying FS are mandatory for the piece technique to work properly:

1. Once a file page is allocated it must not be reused (until system restart).
2. A versatile file buffering mechanism supporting multiple buffers per file is required.
3. Files must be allowed to be open in read mode and in write mode simultaneously.

The format of text sections in files obeys a set of syntactical rules (productions) that can easily be specified in EBNF-notation:

```
TextSection = ident header {char}.  
header = type offset run {run} null length.  
run = font [name] color offset length.
```

In the TextSection production `ident` is an identifier for text blocks. In the header production `type` is a type-discriminator, `offset` is the offset to the character part, `run` is a run-descriptor, `null` is a nullcharacter, and `length` is the length of the character sequence. In the run production `font`, `color`, and `offset` are specifications of looks, and `length` is the run-length. In order to save space, font names are coded as ordinal numbers within a text section. If and only if a font appears for the first time in a text block it is followed by the actual font name.

Let us conclude this Section with 2 side-remarks and a summary. Remarks: For compatibility reasons, plain ASCII-files are accepted as text files as well. They are mapped to texts consisting of a single run with standard looks.

Internalizing a text section from a file is extremely efficient because it is obviously sufficient to read the header and translate it into the initial state of the piece list.

Summary: The mechanism used for the implementation of the abstract data type Text is completely hidden from clients. It is a generalized version of the original piece list technique, adapted to texts with looks. The piece list technique in turn is based on the principle of indirection: Operations operate on descriptors of texts rather than on texts themselves. The benefits are efficiency and non-destructive operations. However, the technique works properly only in combination with

an efficient (and reliable) garbage collector and a suitable FS.

### 5.3 TEXT FRAMES

The tasks of text frames are text rendering and user interaction. A text frame represents a text view and a controller in the form of an interactive text editor. Technically, text frames are a subclass of display frames and, as such, are objects with an open message interface of the kind explained in Chapter 4.

The geometric layout of text frames is determined by 2 areas: A rectangle of contents and a vertical scroll-bar along the left borderline. The type of text frames is a direct extension of `Display.Frame`:

```
Frame = POINTER TO FrameDesc;
FrameDesc = RECORD (Display.FrameDesc)
    text: Texts.Text;
    org, col, lsp: INT;
    left, right, top, bot: INT;
    markH, time: INT;
    hasCar, hasSel, hasMark: BOOL;
    carloc: Location;
    selbeg, selend: Location
END;
```

Fields `text` and `org` specify the text part to be displayed, the former referring to the underlying text and the latter designating the starting position of the displayed part. Fields `col`



and `lsp` are rendering parameters. They specify the frame's background color and the line spacing. Fields `left`, `right`, `top`, and `bot` are margins. They determine the rectangle of contents. `mark` indicates whether there is a position marker, which is a small horizontal bar indicating the position of the displayed part relative to the whole text. `markH` represents its location within the text frame.

Caret and selection are two important features associated with a text frame. The caret indicates a focus, and it serves as an implicit "point of insertion" for placing consumed characters (for example from the keyboard). The selection is a stretch of displayed text. Additionally it serves as a parameter for various operations and commands, among them delete and change looks. The state and location of the caret is given by the variables `car` and `carloc` respectively. Analogously, the state of the selection and its begin and end are reflected by the fields `sel`, `selbeg`, and `selend` in the frame descriptor. Field `time` is a time stamp on the current selection.

In principle, caret and selection could be regarded as ingredients of the underlying text (the model) equally well. However, we deliberately decided to associate these features with frames (views) in order to get increased flexibility. For example, two different selections in adjacent viewers displaying the same text are normally interpreted as one extensive selection across their span. The auxiliary type

Location summarizes information about a location in a text frame. Its definition is:

```
Location = RECORD org, pos, dx, x, y: INT END;
```

$x$ ,  $y$  specify the envisioned location relative to the text frame's origin, and  $dx$  is the width of the character at this location.  $pos$  is the corresponding position in the text and  $org$  is the origin position of the corresponding text line.

The following is a simplified version of the message handler employed by text frames. It fully determines the behavior and capabilities of text frames.

```
PROC Handle* (F: Display.Frame; VAR M: Display.FrameMsg);
  VAR F1: Frame; buf: Texts.Buffer;
BEGIN
  CASE M OF
    Oberon.InputMsg:
      1)   IF M.id = Oberon.track THEN
            Edit(F(Frame), M.X, M.Y, M.keys)
          ELSIF M.id = Oberon.consume THEN
      2)   IF F(Frame).hasCar THEN
            Write(F(Frame), M.ch, M.fnt, M.col, M.voff)
          END
        END |
        Oberon.ControlMsg:
      3)   IF M.id = Oberon.defocus THEN Defocus(F(Frame))
```

```

4)    ELSIF M.id = Oberon.neutralize THEN Neutralize(F(Frame))
      END |
5)    Oberon.SelectionMsg:
      GetSelection(F(Frame), M.text, M.beg, M.end, M.time) |
7)    Oberon.CopyMsg: Copy(F(Frame), F1); M.F := F1 |
      MenuViewers.ModifyMsg: Modify(F(Frame), M.id, M.dY, M.Y, M.H) |
8)    CopyOverMsg: CopyOver(F(Frame), M.text, M.beg, M.end) |
9)    UpdateMsg: IF F(Frame).text = M.text THEN Update(F(Frame), M) E
      END
      END Handle;

```

#### Explanations:

1. Mouse tracking message: Call built-in editor immediately
2. Consume message: In case of valid caret insert character
3. Defocus message: Remove caret
4. Neutralize message: Remove caret and selection
5. Selection message: Return current selection with time stamp
6. Copy message: Create a copy (clone)
7. Modify message: Translate and change size
8. Copyover message: Copy given stretch of text to caret

9. Update message: If text was changed then update display

We recognize again our categories of universal messages introduced in Chapter 4, Table 4.6: Messages in lines 1) and 2) report about user interactions. Messages in 3), 4), 5), 6), and 7) specify generic operations. Messages in 8) require a change of location or size. Messages of the latter kind arrive from the ancestor menu viewer via delegation. They are generated by the interaction handler and preprocessed by the original viewer message handler. Finally, messages in line 9) report about changes of contents.

The text frame handler is encapsulated in a module called `TextFrames`. This module exports the above introduced types `Frame` (text frame) and `Location`, as well as the procedure `Handle`. Furthermore, it exports type `UpdateMsg` to report on changes made to a displayable text.

```
UpdateMsg = RECORD (Display.FrameMsg)
  id: INT;
  text: Texts.Text;
  beg, end: INT
END;
```

Field `id` names one of the operators `replace`, `insert`, or `delete`. The remaining fields `text`, `beg`, and `end` restrict the change to a range. Additional procedures generate a new standard menu text frame and contents text frame respectively:

```
PROC NewMenu(name, commands: ARRAY OF CHAR): Frame;
PROC NewText(text: Texts.Text; pos: INT): Frame;
```

This completes the minimum definition of module TextFrames. In addition, this module exports a set of useful service procedures supporting the composition of custom handlers from elements of the standard handler:

```
PROC Edit (F: Frame; X, Y: INT; Keys: SET);
PROC Write (F: Frame; ch: CHAR; fnt: Fonts.Font; col, voff: INT);
PROC Defocus (F: Frame);
PROC Neutralize (F: Frame);
PROC GetSelection (F: Frame; VAR text: Texts.Text;
VAR beg, end, time: INT);
PROC CopyOver (F: Frame; text: Texts.Text; beg, end: INT);
PROC Copy (F: Frame; VAR F1: Frame);
PROC Modify (F: Frame; id, dY, Y, H: INT);
PROC Update (F: Frame; VAR M: UpdateMsg);
```

The module also supports mouse tracking inside text frames:

```
PROC TrackCaret (F: Frame; X, Y: INT; VAR keysum: SET);
PROC TrackSelection (F: Frame; X, Y: INT; VAR keysum: SET);
PROC TrackLine (F: Frame; X, Y: INT; VAR org: INT; VAR keysum: SET);
PROC TrackWord (F: Frame; X, Y: INT; VAR pos: INT; VAR keysum: SET);
```

Let us now take a closer look at the implementation of some selected operations. For this purpose, we must first explain

the notion of line descriptor that is used to optimize the operation of locating positions within text frames.

```

Line = POINTER TO LineDesc;
LineDesc = RECORD
    len, wid: INT;
    eot: BOOL;
    next: Line
END;

```

Each line descriptor provides detailed information about a single line of text that is currently displayed: len is the number of characters on the line, wid is the line width, eot indicates terminating line, and next points to the next line descriptor.

Text frames maintain a private data structure called line list that describes the list of text lines displayed:

```

Frame = POINTER TO FrameDesc;
FrameDesc = RECORD (Display.FrameDesc)
    text: Texts.Text;
    org, col, lsp: INT;
    left, right, top, bot: INT;
    markH, time: INT;
    hasChar, hasSel, hasMark: BOOL;
    carloc, selbeg, selend: Location;
--> trailer: Line

```

```
END;
```

Field trailer represents a sentinel element that closes the line list to a ring. The line list contains useful summary information about the current contents of the text frame. It can be used beneficially by some related data types, for example by type `Location` that was introduced earlier:

```
Location = RECORD
  org, pos, dx, x, y: INT;
  --> lin: Line
END;
```

The built-in editor procedure `Edit` is a worthwhile part to look at in more detail. It is called by the task scheduler to handle mouse events within a text frame. The following code excerpt shows nicely how the different components of the text system interoperate.

```
PROC Edit* (F: Frame; X, Y: INT; Keys: SET);
VAR M: CopyOverMsg;
    text: Texts.Text;
    buf: Texts.Buffer;
    v: Viewers.Viewer;
    loc0, loc1: Location;
    beg, end, time, pos: INT;
    keysum: SET;
```

```

    fnt: Fonts.Font;
    col, voff: INT;
BEGIN
    IF X < F.X + Min(F.left, barW) THEN (*cursor is in scroll
    Oberon.DrawMouse(ScrollMarker, X, Y); keysum := Keys;
    IF Keys = {2} THEN (*ML, scroll up*)
        TrackLine(F, X, Y, pos, keysum);
        IF (pos >= 0) & (keysum = {2}) THEN
            RemoveMarks(F); Oberon.RemoveMarks(F.X, F.Y, F.W, F
            Show(F, pos)
        END
    ELSIF Keys = {1} THEN (*MM*) keysum := Keys;
        REPEAT Input.Mouse(Keys, X, Y); keysum := keysum + Key
            Oberon.DrawMouse(ScrollMarker, X, Y)
        UNTIL Keys = {};
        IF ~(keysum = {0, 1, 2}) THEN
            IF 0 IN keysum THEN pos := 0
            ELSIF 2 IN keysum THEN pos := F.text.len - 100
            ELSE pos := (F.Y + F.H - Y) * (F.text.len) DIV F.H
            END ;
            RemoveMarks(F); Oberon.RemoveMarks(F.X, F.Y, F.W, F
            Show(F, pos)
        END
    ELSIF Keys = {0} THEN (*MR, scroll down*)
        TrackLine(F, X, Y, pos, keysum);
        IF keysum = {0} THEN
            LocateLine(F, Y, loc0); LocateLine(F, F.Y, loc1);
            pos := F.org - loc1.org + loc0.org;

```



```

    IF pos < 0 THEN pos := 0 END ;
    RemoveMarks(F); Oberon.RemoveMarks(F.X, F.Y, F.W, F.H);
    Show(F, pos)
  END
END
ELSE (*cursor is in text area*)
  Oberon.DrawMouseArrow(X, Y);
  IF 0 IN Keys THEN (*MR: select*)
    TrackSelection(F, X, Y, keysum);
    IF F.hasSel THEN
      IF keysum = {0, 2} THEN (*MR, ML: delete text*)
        Oberon.GetSelection(text, beg, end, time);
        Texts.Delete(text, beg, end, TBuf);
        Oberon.PassFocus(Viewers.This(F.X, F.Y)); SetCaret(F, beg)
      ELSIF keysum = {0, 1} THEN (*MR, MM: copy to caret*)
        Oberon.GetSelection(text, beg, end, time);
        M.text := text; M.beg := beg; M.end := end;
        Oberon.FocusViewer.handle(Oberon.FocusViewer, M)
      END
    END
  ELSIF 1 IN Keys THEN (*MM: call*)
    TrackWord(F, X, Y, pos, keysum);
    IF (pos >= 0) & ~(0 IN keysum) THEN Call(F, pos, 2 IN keysum) END
  ELSIF 2 IN Keys THEN (*ML: set caret*)
    Oberon.PassFocus(Viewers.This(F.X, F.Y));
    TrackCaret(F, X, Y, keysum);
    IF keysum = {2, 1} THEN (*ML, MM: copy from selection to caret*)
      Oberon.GetSelection(text, beg, end, time);

```

```

    IF time >= 0 THEN
        NEW(TBuf); Texts.OpenBuf(TBuf);
        Texts.Save(text, beg, end, TBuf); Texts.Insert(F.text,
        SetSelection(F, F.carloc.pos, F.carloc.pos + (end -
        SetCaret(F, F.carloc.pos + (end - beg))
    ELSIF TBuf # NIL THEN
        NEW(buf); Texts.OpenBuf(buf);
        Texts.Copy(TBuf, buf); Texts.Insert(F.text, F.carloc
        SetCaret(F, F.carloc.pos + buf.len)
    END
    ELSIF keysum = {2, 0} THEN (*ML, MR: copy looks*)
        Oberon.GetSelection(text, beg, end, time);
        IF time >= 0 THEN
            Texts.Attributes(F.text, F.carloc.pos, fnt, col, vo
            IF fnt # NIL THEN Texts.ChangeLooks(text, beg, end,
        END
    END
END
END Edit;
```

We see that the editing operation is determined by the first key pressed (primary key) and can then be varied by “interclicking” that is, by clicking a secondary key while holding down the primary key. As a convention, (inter)clicking all keys means cancelling the operation. Mouse clicks and subsequent actions can now be summarized as follows: 1. In the scroll bar

primary key	secondary key	action
ML		scroll designated line to the top
MM		scroll proportional to mouse position
MM	ML	scroll to the end of the text
MM	MR	scroll to the beginning of the text

## 2. In the text area

primary key	secondary key	action
ML		set caret
ML	MM	copy selection to caret
ML	MT	copy looks
MM		call selected procedure
MR		select
MR	ML	delete selection
MR	MM	copy selection to caret

In the text area the keys are interpreted according to their generic semantics:

left key    =   point key  
middle key   =   execute key  
right key    =   select key

Let us “zoom into” one of the editing operations, for example into TrackCaret.

```
PROC TrackCaret (F: Frame; X, Y: INT; VAR keysum: SET);
```

```

        VAR loc: Location; keys: SET;
    BEGIN
1)  IF F.trailer.next # F.trailer THEN
2)      LocateChar(F, X - F.X, Y - F.Y, F.carloc);
3)      FlipCaret(F);
4)      keysum := {};
        REPEAT
            Input.Mouse(keys, X, Y); keysum := keysum + keys;
            Oberon.DrawMouseArrow(X, Y);
            LocateChar(F, X - F.X, Y - F.Y, loc);
            IF loc.pos # F.carloc.pos THEN FlipCaret(F); F.carloc := loc;
5)      UNTIL keys = {};
6)      F.hascar := TRUE
        END
    END TrackCaret;

```

Explanations:

1. guard guarantees non-empty line list
2. locates the character pointed at
3. drags caret to new location
4. -
5. tracks mouse and drags caret accordingly
6. set caret state

TrackCaret makes use of two auxiliary procedures FlipCaret and LocateChar. FlipCaret is used to turn off/on the pattern of the caret. LocateChar is an important operation that is used to locate the character at a given Cartesian position (x,y) within the frame.

```
PROC FlipCaret (F: Frame);
BEGIN
1) IF F.carloc.x < F.W THEN
2)   IF (F.carloc.y >= 10) & (F.carloc.x + 12 < F.W) THEN
3)     Display.CopyPattern(Display.white, Display.hook,
        F.X + F.carloc.x, F.Y + F.carloc.y - 8, Display.invert)
    END
  END
END FlipCaret;
```

Explanations:

1. -
2. if there is room for drawing the caret
3. copy standard hook-shaped pattern to caret location  
in inverse video mode

```
PROC LocateChar (F: Frame; x, y: INT; VAR loc: Location);
VAR R: Texts.Reader;
    patadr, pos, lim: INT;
    ox, dx, u, v, w, h: INT;
```

```

1) BEGIN LocateLine(F, y, loc);
2)   lim := loc.org + loc.lin.len - 1;
3)   pos := loc.org; ox := F.left; dx := eolW;
4)   Texts.OpenReader(R, F.text, loc.org);
5)   WHILE pos # lim DO
6)     Texts.Read(R, nextCh);
7)     Fonts.GetPat(R.fnt.raster, nextCh, dx, u, v, w, h, pa
      IF ox + dx <= x THEN
        INC(pos); ox := ox + dx;
        IF pos = lim THEN dx := eolW END
      ELSE lim := pos
      END
    END ;
8)   loc.pos := pos; loc.dx := dx; loc.x := ox
   END LocateChar;

```

Explanations:

1. locate text line corresponding to at y
2. set limit to the last actual character on this line
3. start locating loop with first character on this line
4. setup reader and read first character of this line
5. -
6. scan through characters of this line until limit or x is reached

7. get character width dx of current character
8. return location found

Note that the need to read characters from the text (again) in `LocateChar` has its roots in the so-called proportional fonts in which our rich texts are represented. We found that keeping character widths is an unnecessary optimization thanks to the buffering capabilities of the underlying file system. In the case of fixed-pitch fonts a simple division by the character width would be sufficient, of course.

Finally, procedure `LocateLine` uses the line list to locate the desired text line without reading text at all.

```
PROC LocateLine (F: Frame; y: INT; VAR loc: Location);
  VAR L: Line; org, cury: INT;
  BEGIN org := F.org;
1)   org := F.org; L := F.trailer.next; cury := F.H - F.top - asr;
2)   WHILE (L.next # F.trailer) & (cury > y + dsr) DO
      org := org + L.len; L := L.next; cury := cury - lsp
3)   END;
4)   loc.org := org; loc.lin := L; loc.y := cury
  END LocateLine;
```

Explanations:

1. start with first line in the frame

2. -
3. traverse line chain until last line or y is reached
4. return found line

After text editing text, rendering is our next topic. Let us pursue the case in that a user pressed the point-key and then interclicked the middle key, corresponding to line 56) in procedure Edit. Remember that notifier is called at the end of every editing operation and in particular at the end of Texts.Insert. In case of standard text frames, the notifier simply broadcasts an update message into the display space:

```
PROC NotifyDisplay (T: Texts.Text; op, beg, end: INT);
  VAR M: UpdateMsg;
BEGIN M.id := op; M.text := T; M.beg := beg; M.end := end
  Viewers.Broadcast(M)
END NotifyDisplay;
```

Let us now take the perspective of a text frame receiving an update message. Looking at line 9) in the text frame handler, we see that procedure Update is called, which in turn calls procedure Insert in TextFrames:

```
PROC Insert (F: Frame; beg, end: INT);
  VAR R: Texts.Reader; L, L0, l: Line;
      org, len, curY, botY, Y0, Y1, Y2, dY, wid: INT;
BEGIN
```



```

IF beg < F.org THEN F.org := F.org + (end - beg)
ELSE
1)   org := F.org; L := F.trailer.next; curY := F.Y + F.H - F.top
    WHILE (L # F.trailer) & (org + L.len <= beg) DO
        org := org + L.len; L := L.next; curY := curY - lsp
2)   END;
3)   IF L # F.trailer THEN
        botY := F.Y + F.bot + dsr;
4)   Texts.OpenReader(R, F.text, org); Texts.Read(R, nextCh);
5)   len := beg - org; wid := Width(R, len);
6)   ReplConst (F.col, F, F.X + F.left + wid, curY - dsr, L.wid
7)   DisplayLine(F, L, R, F.X + F.left + wid, curY, len);
8)   org := org + L.len; curY := curY - lsp;
    Y0 := curY; L0 := L.next;
    WHILE (org <= end) & (curY >= botY) DO
        NEW(l);
        Display.ReplConst(F.col, F.X + F.left, curY - dsr, F.W - l
        DisplayLine(F, l, R, F.X + F.left, curY, 0);
        L.next := l; L := l;
        org := org + L.len; curY := curY - lsp
9)   END;
10)  IF L0 # L.next THEN Y1 := curY;
11)  L.next := L0;
    WHILE (L.next # F.trailer) & (curY >= botY) DO
        L := L.next; curY := curY - lsp
12)  END;
    L.next := F.trailer;
    dY := Y0 - Y1;

```

```

        IF Y1 > curY + dY THEN
13)      Display.CopyBlock (F.X + F.left, curY + dY + lsp
                                F.W - F.left, Y1 - curY - dY
                                F.X + F.left, curY + lsp - dY);
        Y2 := Y1 - dY
        ELSE Y2 := curY
        END;
14)      curY := Y1; L := L0;
        WHILE curY # Y2 DO
            Display.ReplConst(F.col, F.X + F.left, curY - dY,
                               DisplayLine(F, L, R, F.X + F.left, curY, 0));
            L := L.next; curY := curY - lsp
15)      END
        END
        END
        END;
16) UpdateMark(F)
    END Insert;

```

Some explanations:

1. -
2. search line where inserted part starts
3. if it is displayed in this viewer
4. setup reader on this line

5. get width of unaffected part of line (avoid touching it)
6. clear remaining part of line
7. display new remaining part of line
8. -
9. display newly inserted text lines
10. if it was not a one line update
11. -
12. skip overwritten text lines
13. use fast block move to adjust reusable lines
14. -
15. redisplay previously overwritten text lines
16. adjust position marker

Special care needs to be exercised in the implementation to avoid "flickering" and to minimize processing time. Concretely, the following measures are taken for this purpose:

1. Avoid writing the same data again.
2. Keep the number of newly rendered text lines at a minimum.

3. Use raster operations (block moves) to adjust reusable displayed lines.

Of course, the rules governing the rendering and formatting process crucially influence the complexity of procedures like Insert. For text frames we have consciously chosen the simplest possible set of formatting rules. They can be summarized as:

1. For a given text frame the distance between lines is constant.
2. There are no implicit line breaks.

It is exactly this set of rules that makes it possible to display a text line in one pass. Two passes are inevitable if line distances have to adjust to font sizes or if lines must be broken implicitly.

Updating algorithms make use of the following one-pass rendering procedures Width and DisplayLine:

```
PROC Width (VAR R: Texts.Reader; len: INT): INT;
  VAR patadr, pos, ox, dx, x, y, w, h: INT;
1) BEGIN pos := 0; ox := 0;
  WHILE pos < len DO
    Fonts.GetPat(R.fnt, nextCh, dx, x, y, w, h, pat);
    ox := ox + dx; INC(pos); Texts.Read(R, nextCh)
2) END;
```

```

3) RETURN ox
   END Width;

```

Explanations:

1. -
2. scan through len characters of this line
3. return accumulated width

Procedures Width and LocateChar are similar. Therefore the above comment about relying on the buffering capabilities of the underlying FS applies to procedure Width equally well.

```

PROC DisplayLine (F: Frame; L: Line; VAR R: Texts.Reader; X, Y, L
    VAR patadr; NX, Xlim, dx, x, y, w, h: INT;
1) BEGIN NX := F.X + F.W; Xlim := NX - 40;
2) WHILE (nextCh # CR) & ((nextCh > " ") OR (X < Xlim)) & (R.fnt #
4)   Fonts.GetPat(R.fnt, nextCh, dx, x, y, w, h, patadr);
5)   IF (X + x + w <= NX) & (h # 0) THEN
6)     Display.CopyPattern(R.col, patadr, X + x, Y + y, Display.in
       END;
7)   X := X + dx; INC(len); Texts.Read(R, nextCh)
3) END;
8) L.len := len + 1; L.wid := X + eolW - (F.X + F.left);
9) L.eot := R.fnt = NIL; Texts.Read(R, nextCh)
   END DisplayLine;

```

## Explanations:

1. set right margin
2. -
3. display characters of this line
4. get width dx, box x, y, w, h, and pattern address of next character
5. if there is enough space in the rectangle of contents
6. display pattern
7. jump to location of next character; read next character
8. -
9. setup line descriptor

Procedure `DisplayLine` is again similar to `LocateChar`, and the comment about relying on the file system's buffering capabilities applies once more. The principal difference between `LocateChar` and `Width` on one hand and `DisplayLine` on the other hand is the fact that the latter accesses the display screen physically. Therefore, possession of the screen lock is a tacit precondition for calling `DisplayLine`.

A quick look at an auxiliary procedure that updates the position marker concludes our tour behind the scenes of the text system:

```
PROC UpdateMark (F: Frame);
  VAR oldH: INT;
BEGIN
1)  oldH := F.markH; F.markH := F.org * F.H DIV (F.text.len + 1));
    IF (F.mark > 0) & (F.left >= barW) & (F.markH # oldH) THEN
2)    Display.ReplConst(Display.white, F.X + 1, F.Y + F.H - 1 - oldH);
3)    Display.ReplConst(Display.white, F.X + 1, F.Y + F.H - 1 - F.markH);
    END
END UpdateMark;
```

### Explanations

1. shows how the marker's position is calculated.  
Loosely spoken, the invariant is:  
distance from top of frame / frame height =  
text position of first character in frame / text length
2. erase the old marker
3. draw the new marker

And this in turn concludes our section on text frames. Recapitulating the most important points:

- The tasks of text editing (input oriented) and text rendering (output oriented) are combined in the concept

of text frames. Text frames constitute a subclass of display frames, and they are implemented in a separate module called `TextFrames`.

- The implementation of `TextFrames` accesses the displayed text exclusively via the “official” abstract interface of module `Texts` discussed in Section 5.2. It maintains a private data structure of line lists to accelerate locating requests.
- Text frames use simple formatting rules that allow super-efficient rendering of text in a single pass. In particular, line spacing is fixed for every text frame. Therefore, different styles of a base font are possible within a given text frame while different sizes are not.

Putting into relation the different extensions of type `Display.Frame` that we came across in Chapters 4 and 5, we obtain the type hierarchy as shown in Fig ??.




Figure 17: Extensions of type *Display.Frame*



## 5.4 THE FONT MACHINERY

We saw in the previous Sections that Oberon texts support attribute specifications (“looks”) for characters. Three different attributes are supported: font, color, and vertical offset. Let us first focus on the font attribute. A font can be regarded as a style the standard character set is designed in. Typically, an entire text is typeset in a single style, that is, there is one font per text. However, sometimes, an author wants to emphasize titles or words by changing the size of the font or by varying it to bold face or italics. In special texts, special characters like mathematical symbols or other kinds of icons may occur. In even more complex documents, mathematical or chemical formulae might flow within the text.

This generalized view leads us to a different interpretation of the notion of font. We can regard a font as an indexed library of (graphical) objects, mostly but not necessarily glyphs. In the case of ordinary characters it is natural to use the ASCII-code as an index, ending up with an interpretation of text as sequence of pairs (library, index). Note that this is a very general view indeed that, in principle, is equivalent with defining text as sequence of arbitrary objects.

The imaging model of characters provides 2 levels of abstraction. On the first level, characters are black boxes specified by a set of metric data  $x$ ,  $y$ ,  $w$ ,  $h$ , and  $dx$ .  $(x, y)$  is a vector from the current point of reference on the base line to

the origin of the box.  $w$  and  $h$  are width and height of the box, and  $dx$  is the distance to the point of reference of the next character on the same base line. On the second level of abstraction, a character is defined by a digital pattern or glyph that is to be rendered into the box. Fig ?? visualizes this model of characters.

The additional 2 character attributes color and vertical offset appear now as parameters for the character model. The vertical offset allows translating the glyph vertically and the color attribute specifies the foreground color of the pattern.




Figure 18: The geometric character model

Good examples of procedures operating on the first level of abstraction are procedures `LocateChar` and `Width` that we discussed in the previous section, as well as text formatters for a remote printer. In contrast, procedure `DisplayLine` operates on the second level.

The representation of characters as digital patterns is merely the last step in a complex font design and rendering process. At the beginning is a generic description of the shape of each character in the form of outlines and hints. Outlines are typically composed of straight lines and spline-curves. Hints are included to assist the digitizer in its effort to faithfully map the filled character outlines into the device raster. For example, hinting can guarantee consistency of serif shapes and stem widths across an entire font within a text, independent of the relative positions of the characters with respect to the grid lines. Automatic digitization produces digital patterns of sufficiently high quality for printing media resolutions. For screen resolutions, however, we prefer to add a hand-tuning step. This is the reason why digital patterns are not produced "on the fly" in Oberon.

Oberon's font management is encapsulated in module `Fonts`, with a low-level extension into the module `Display` that we already know from Chapter 4. The interface to module `Fonts` is very simple and narrow:

```
MODULE Fonts;
  TYPE Font = POINTER TO FontDesc;
  FontDesc = RECORD
    name: ARRAY 32 OF CHAR;;
    height, minX, maxX, minY, maxY: INT;
    next: Font
  END;
```

```

VAR Default: Font;
PROC GetPat(fnt: Font; ch: CHAR; VAR dx, x, y, w, h, pa
PROC This (name: ARRAY OF CHAR): Font;
PROC Free;
END Fonts.

```

Variable name in type `Font` is the name of the underlying file. The variables `height`, `minX`, `maxX`, `minY`, and `maxY` designate line height and summary metric data. `Default` is a system-wide default font. It is installed at system loading time. `GetPat` delivers the geometric data for a given character in a given font (see Fig ??). This is a procedure to internalize (load) a font from a file given by its name. `Free` releases from storage fonts that are no longer needed.

Type `Font` should again be regarded as an abstract data type with two intrinsic operations `This` and `GetPat`. Thinking of the immutable nature of fonts, multiple internal copies of the same font are certainly undesirable. Therefore, internalized fonts are cached in a private list that manifests itself in a private field next in type `FontDesc`. The cache is maintained by the internalizing procedure `This` according to the following scheme:

```

search font in cache;
IF found THEN return cached internalization
ELSE internalize font; cache it END

```

The implementation of type Font did not raise many challenges. One, however, is an undesirable side-effect of caching. The problem arises if a font is used for a limited time only. Because it is referenced by the cache it will never be collected by the system's garbage collector. Two possible solutions offer themselves: a) provide an explicit freeing operation and b) enforce some special handling by the garbage collector based on a concept of "weak" pointers.

We conclude this section with a formal specification of the font file format. Note that on the one hand, the file format is completely private to the managing Fonts module and on the other hand, it should be ultimately stable because it is probably used for long-term backup and for wide-range data exchange across multi-system platforms.

This is an EBNF specification of the Oberon font file format:

```
FontFile = ident header contents.
header = abstraction family variant height minX maxX minY maxY.
contents = nofRuns { beg end } { dx x y w h } { rasterByte }.
```

ident, abstraction, family, and variant are one-byte values indicating file identification, abstraction (first level without raster bytes, second level with raster bytes), font family (Times Roman, Oberon, etc.), and variant (bold face, italics etc.). The values height, minX, maxX, minY and maxY are 2 bytes long each. They define in turn line height, minimum x-coordinate (of a box), maximum x-coordinate, minimum

y-coordinate, and maximum y-coordinate. All values in production contents are 2 bytes long. `noRuns` specifies the number of runs within the ASCII-code range (intervals occupied without gaps) and every pair `[beg, end]` describes one run. The tuples `(dx, x, y, w, h)` are the metric data of the corresponding characters (in their ASCII-code order), and the sequence of `rasterByte` gives the total of raster information.

In summary, fonts in Oberon are indexed libraries of objects. The objects are descriptions of character images in 2 levels of abstraction: As metric data of black boxes and as binary patterns (glyphs). Type `Font` is an abstract data type with intrinsic operations to internalize and to get character object data. Internalized fonts are cached in a private list.

## 5.5 THE EDIT TOOLBOX

We have seen that every text frame integrates an interactive text editor that we can regard as an interpreter of a set of built-in commands (intrinsic commands). Of course, we would like to be able to extend this set by custom editing commands (extrinsic commands). Adding additional editing commands was indeed a worthwhile stress test for the underlying texts API. Module Edit is the result of this effort. It is a toolbox of consisting of some standard extrinsic editing commands.

### DEFINITION Edit;

```
PROC Open; (*text viewer*)
PROC Show; (*text*)
PROC Locate; (*position*)
PROC Search; (*pattern*)
PROC Store; (*text*)
PROC Recall; (*deleted text*)
PROC CopyFont;
PROC ChangeFont;
PROC ChangeColor;
PROC ChangeOffset;
END Edit.
```

The first group of commands in Edit is used to display, locate, and store texts or parts of texts. In turn they open a text file and display it, open a program text and show the

declaration of a given object, locate a given position in a displayed text (main application: locating an error found by the compiler), search a pattern, and store the current state of a displayed text. Commands in the next group are related with editing. They allow restoring of the previously deleted part of text, copying a font attribute to the current text selection, and change attributes of the current text selection. Note that the commands `CopyFont`, `ChangeFont`, `ChangeColor`, and `ChangeOffset` are extrinsic variations of the intrinsic copy-look operation. The implementations of the toolbox commands are given in the Appendix.

#### REFERENCES

- GUTKNECHT J. Gutknecht, "Concept of the Text Editor Lara", Communications of the ACM, Sept. 1985, Vol.28, No.9.
- TEITELMAN W. Teitelman, "A tour through Cedar", IEEE Software, 1, (2), 44-73 (1984).



## MODULE LOADER

---

### 6.1 LINKING AND LOADING

When the execution of a command  $M.P$  is requested, module  $M$  containing procedure  $P$  must be loaded, unless it

- already loaded for its another command had earlier been executed, or
- had been imported by another one before.

Modules are available in the form of so-called *object files*, generated by the compiler. The term *loading* refers to the transfer of module code from file into main memory, from where processors fetch individual instructions. It involves also a certain amount of transformation as required by

- the object file format on the one hand, and
- the storage layout, on the other.

A system typically consists of many modules, and hence loading modules also involves linking them together, in particular linking with already loaded ones. Before loading,

- references to another module's objects are relative to the base address of this one;
- the linking or binding process converts them into absolute addresses.

This process may require a significant amount of address computations. But they're simple enough and can be executed very swiftly, if the data are organized in an appropriate way. Nevertheless and surprisingly, in many OSes it needs more time than compilation. The remedy, which system designers offer, is to separate linking from loading:

- A set of compiled modules is 1st linked into a object file with absolute addresses.
- The loader then merely transfers it into main store.

We consider it an unfortunate solution. Instead of trying to cure it with aid of an additional processing stage and tool, it is wiser to remove the malady at its core, namely to speed up linking itself.

Indeed, there is no separate linker. In Oberon, they are integrated and fast enough to avoid any desire for pre-linking. Furthermore, the system extensibility crucially depends on

the possibility to link additional modules to already loaded ones by calls from any module. This is called *dynamic loading*, with prelinked object files which is impossible.

- Newly loaded simply refer to already loaded, whereas
- prelinked files lead to multiple copies of same module.

Evidently, the essence of linking is the conversion of

- relative addresses as generated by the compiler for all external references, into
- absolute ones as required during program execution.

To proceed, we must consider an additional complication.

Assume that a module M1 is to be compiled which is a client of (that is, it imports) M0. The interface of M1 - in the form of a symbol file - does not specify the entry addresses of its exported procedures, but merely specifies a unique number (pno) for each one of them. The reason are:

- In this way the implementation of M0 may be modified, causing a change of entry addresses without affecting its interface specification. And
- a crucial property of the scheme of separate modules compilation: Changes of the implementation of M0 must not necessitate the recompilation of clients (M1).

The consequence is

- the binding of entry addresses to pnos must be performed by the linker. To make this possible,
- the object file must contain a list (table) of its entry addresses, one for each pno used as index to the table.

Similarly, the object file must contain a table of imported modules, containing their names. An external reference in the program code then appears in the form of a pair consisting of

- a module number (mno) - used as index to the import table (of modules), and
- a procedure number (pno) - used as index to the entry table of this module.

Certain linkage information must not only be provided in each object file, but also be present along with each loaded module's program code, because a module to be loaded must be linkable with modules loaded at any earlier time without reading their object files again.

## 6.2 MODULE REPRESENTATION

The primary requirement is

a system must be represented in a form that allows to add new modules quickly.

What is a sensible representation for this purpose? The simplest solution that comes to mind is a list of module blocks containing sections for

- the global data,
- the program code, and perhaps
- meta data for the linking process.

The list is rooted in a variable global to the loader module, here called `Modules`.




Figure 19: System of 4 modules

The 1st part, containing the link to the next module, is called *module descriptor*. In Oberon, it contains further links to the various sections of a module. The type `Module` is defined as:

```

TYPE Module = POINTER TO ModDesc;
      ModuleName = ARRAY 32 OF CHAR;
      ModDesc = RECORD name: ModuleName;
  
```

```
next: Module;  
key, num, size, refcnt, data, code,  
imp, cmd, ent, ptr: INT (*addresses*)  
END;
```

*key* the module's key used for version consistency checking. Changes, if and only if, the module's interface and thereby its symbol file changes.

*num* the module's number, index of the module's entry in a global module table, referenced by the processor's MT register. The invariant relationship is

$$\text{ModTable}[\text{mod.mno}] = \text{mod.data}$$

for all *mod* in the module list.

*size* the entire module block's size excluding the descriptor, and

*refcnt* the number of other modules importing this module. Used to check whether a module can be released by procedure `Modules.Free`.

The section with meta data follows the data and code areas and consists of several parts.

- Imports is an array of the modules imported by this module, each entry being the address of the respective module descriptor.

- Commands is a sequence of procedure identifiers followed by their offset in the code section. This section is used when activating a command.
- Entries is an array of offsets of all exported entities (including commands). This section is used by the loader itself for linking.
- Pointer refs is an array of offsets of global pointer variables in the data section. These are used by the garbage collector as the roots of graphs of heap objects in use.

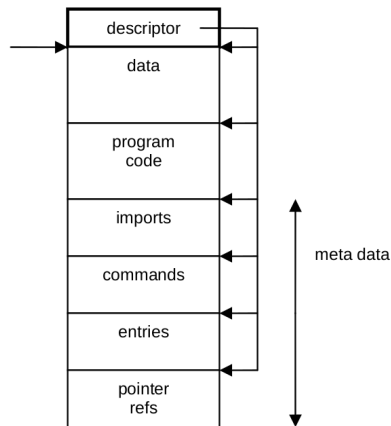


Figure 20: Module block headed by descriptor

### 6.3 THE LINKING LOADER

The purpose of the loader is to read object files, and to transform the file representation of modules into their internal image.

The loader is represented by procedure Load in module Modules. It accepts a name and returns a pointer to the specified module's descriptor. It 1st scans the list searching for the named module. Only if it is not present, the module is loaded and added to the list. Duplications therefore cannot occur.

```
mod := root;
WHILE (mod # NIL) & (name # mod.name) DO
  mod := mod.next
END;

IF mod = NIL THEN (*load*)
  F := ThisFile(name);
  Files.Set(R, F, 0);
  ...
```

First, the header of the respective object file is read. It specifies the required size of the block which is allocated in the module area at the position indicated by the global variable AllocPtr. Then the list of imports of the module being imported is read, and these module are imported.



Evidently procedure Load is used recursively. Because cyclic imports are excluded, recursion always terminates.

```
Files.ReadString(R, impname); (*imports*)
WHILE (impname[0] # 0X) & (res = 0) DO
  Files.ReadInt(R, impkey);
  Load(impname, impmod);
  import[nofimps] := impmod;
  importing := name1;
  IF res = 0 THEN
    IF impmod.key = impkey THEN
      INC(impmod.refcnt);
      INC(nofimps)
    ELSE
      error(3, name);
      imported := impname
    END
  END;
  Files.ReadString(R, impname)
END
```

The loading process stops, if a key mismatch is detected (err = 3). After successful loading of all imports, the loading of the actual module proceeds by allocating a descriptor and then reading the remaining sections of the file. The data is allocated (and cleared) and the code section is read in a straight-forward way without alteration.

At the very end of the file 3 integers called `fixorgP`, `fixorgD`, and `fixorgT` are read. They are the anchors of linked lists in the program code of instructions that need fixups. These fixups are performed only after the entire file had been read. Traversing the P-list, the pairs `mno-pno` are replaced by computed offsets in BL instructions (procedure calls). Traversing the D-list, addresses of LDR instructions and instruction pairs are fixed up, and traversing the T-list, addresses of type descriptors are computed and inserted. This low-level piece of code is shown below for call instructions (BL). Those for the D-List and the T-list are analogous.

```
adr := mod.code + fixorgP*4;
WHILE adr # mod.code DO
  SYSTEM.GET(adr, inst);
  mno := inst DIV 100000H MOD 10H; (*decompose*)
  pno := inst DIV 1000H MOD 100H;
  disp := inst MOD 1000H;
  SYSTEM.GET(mod.imp + (mno-1)*4, impmod);
  SYSTEM.GET(impmod.ent + pno*4, dest);
  dest := dest + impmod.code;
  offset := (dest - adr - 4) DIV 4;    (*compose*)
  SYSTEM.PUT(adr, (offset MOD 10000000H)+0F7000000H);
  adr := adr - disp*4
END;
```

After the module has been loaded successfully, its initialization body is executed. Apart from `Load`, module `Modules` also contains the procedures

```
PROC ThisCommand(mod: Module; name: ARRAY OF CHAR):
                                Command;
PROC Free      (              name: ARRAY OF CHAR);
```

- The former yields the procedure named `name` from module `mod`. It is used in `TextFrames.Call` for activating command procedures.
- The latter unloads the named module, i.e. removes it from the list of loaded modules.

The frequent use of the low-level procedures `SYSTEM.GET` and `SYSTEM.PUT` is easily justified in base modules such as the loader or device drivers. After all, here data are transferred into untyped main storage.

## 6.4 THE TOOLBOX OF THE LOADER

User commands directed to the loader are contained in module `System`. The toolbox offers the following 3 commands:

```
System.ShowModules
System.ShowCommands modname
System.Free {modname}
```

- The 1st command opens a viewer and provides a list of all loaded modules. The list indicates the block length and the number of clients importing a module (the reference count).
- ShowCommands opens a viewer and lists the commands provided by the specified module. The commands are prefixed by the module name, and hence can immediately be activated by a mouse click.
- Free is called in order to remove modules either to regain storage space or to replace a module by a newly compiled version. A module can be dispensed only if
  1. it has no clients, and
  2. does not declare any record types which are extensions of imported types.

## 6.5 THE OBERON OBJECT FILE FORMAT

The name extension of object files is `.rsc`. Their syntax:

```
CodeFile = name key version size
imports typedesc varsize strings code commands
          entries ptrrefs fixP fixD fixT body "0".
imports = {modname key} 0X.
typedesc = nof {byte}.
strings = nof {char}.
```

```
code = nof {word}.  
commands = {comname offset} 0X.  
entries = nof {word}.  
ptrrefs = {word} 0.
```

- fixP, fixD, fixT are the origins of chains of instructions to be updated (fixed up).
- body is the entry point offset of the module body.



## FILE SYSTEM(FS)

---

### 7.1 FILES

It is essential that a computer system has a facility for storing data over longer periods of time and for retrieving the stored data. Such a facility is called a FS. Evidently, a FS cannot accommodate all possible data types and structures that will be programmed in the future. Hence, it is necessary to provide a simple, yet flexible enough base structure that allows any data structure to be mapped onto this base structure (and vice-versa) in a reasonably straight-forward and efficient way. This base structure, called file, is a sequence of bytes. As a consequence, any given structure to be transformed into a file must be sequentialized. The notion of sequence is indeed fundamental, and it requires no further explanation and theory. We recall that texts are sequences of characters, and that characters are typically represented as bytes.

The sequence is also the natural abstraction of all physically moving storage media. Among them are magnetic tapes and

disks. Magnetic media have the welcome property of non-volatility and are therefore the primary choices for storing data over longer periods of time, especially over periods where the equipment is switched off. Sequential access is also necessary for media that allow access only by large blocks, such as flash-RAMs and SD-cards.

A further advantage of the sequence is that its transmission between media is simple too. The reason is that its structural information is inherent and need not be encoded and transmitted in addition to the actual data. This implicitness of structural information is particularly convenient in the case of moving storage media, because they impose strict timing constraints on transmission of consecutive elements. Therefore, the process which generates (or consumes) the data must be effectively decoupled from the transmission process that observes the timing constraints. In the case of sequences, this decoupling is simple to achieve by dividing a sequence into subsequences which are buffered. A sequence is output to the storage medium by alternately generating data (and filling the buffer holding the current subsequence) and transmitting data (fetching elements from the buffer and transmitting them). The size of the subsequences (and the buffer) depends on the storage medium under consideration: there must be no timing constraints between accesses to consecutive subsequences.



The file is not a static data structure like the array or the record, because the length may increase dynamically, i.e. during program execution. On the other hand, the sequence is less flexible than general dynamic structures, because it cannot change its form, but only its length, since elements can only be appended but not inserted. It might therefore be called a semi-dynamic structure. The discipline of purely sequential access to a file is enforced by restricting access to calls of specific procedures, typically read and write procedures for scanning and generating a file. In the jargon of data processing, a file must be opened before reading or writing is possible. The opening implies the initialization of a reading and writing mechanism, and in particular the fixing of its initial position. Hence each (opened) file not only has a value and a length, but also a position attributed to it. If reading must occur from several positions (still sequentially) alternately, the file is "multiply opened"; it implies that the same file is represented by several variables, each denoting a different position.

This widespread view of files is conceptually unappealing, and the Oberon FS therefore departs from it by introducing the notion of a rider. A file simply has a value, the sequence of bytes, and a length, the number of bytes in the sequence. Reading and writing occurs through a rider, which denotes a position. "Multiple opening" is achieved by simply using several riders riding on the same file. Thereby the two

concepts of data structure (file) and access mechanism (rider) are clearly distinct and properly disentangled.

Given a file *f*, a rider *r* is placed on a file by the call `Files.Set(r, f, pos)`, where *pos* indicates the position from which reading or writing is to start. Calls of `Files.Read(r, x)` and `Files.Write(r, x)` implicitly increment the position beyond the element read or written, and the file is implicitly denoted via the explicit parameter *r*, which denotes a rider. The rider has two (visible) attributes, namely *r.eof* and *r.res*. The former is set to `FALSE` by `Files.Set`, and to `TRUE` when a read operation could not be performed, because the end of the file had been reached. *r.res* serves as a result variable in procedures `ReadBytes` and `WriteBytes` allowing one to check for correct termination.

A FS must not only provide the concept of a sequence with its accessing mechanism, but also a registry. This implies that files be identified, that they can be given a name by which they are registered and retrieved. The registry or collection of registered names is called the file system's directory. Here we wish to emphasize that the concepts of files as data structure with associated access facilities on the one hand, and the concept of file naming and directory management on the other hand must also be considered separately and as independent notions. In fact, in the Oberon their implementation underscores this separation by the existence of two modules: `Files` and `FileDir`. The following proce-

dures are available. They are summarized by the interface specification (definition) of module Files.

```

DEFINITION Files;
TYPE File = POINTER TO FileDesc;
FileDesc = RECORD END ;
Rider = RECORD eof: BOOL; res: INT END ;
PROC Old(name: ARRAY OF CHAR): File;
PROC New(name: ARRAY OF CHAR): File;
PROC Register(f: File);
PROC Close(f: File);
PROC Purge(f: File);
PROC Length(f: File): INT;
PROC Date(f: File): INT);
PROC Set(VAR r: Rider; f: File; pos: INT);
PROC ReadByte(VAR r: Rider; VAR x: BYTE);
PROC ReadBytes(VAR r: Rider; VAR x: ARRAY OF BYTE; n: INT);
PROC Read(VAR r: Rider; VAR ch: CHAR);
PROC ReadInt(VAR r: Rider; VAR n: INT);
PROC ReadSet(VAR r: Rider; VAR s: SET);
PROC ReadReal(VAR r: Rider; VAR x: REAL);
PROC ReadString(VAR r: Rider; VAR s: ARRAY OF CHAR);
PROC ReadNum(VAR r: Rider; VAR n: INT);
PROC WriteByte(VAR r: Rider; x: BYTE);
PROC WriteBytes(VAR r: Rider; x: ARRAY OF BYTE; n: INT);
PROC WriteInt(VAR r: Rider; n: INT);
PROC WriteSet(VAR r: Rider; s: SET);
PROC WriteReal(VAR r: Rider; x: REAL);

```

```

PROC WriteString(VAR r: Rider; x: ARRAY OF CHAR);
PROC WriteNum(VAR r: Rider; n: INT);
PROC Pos(VAR r: Rider): INT;
PROC Base(VAR r: Rider): File;
PROC Rename(old, new: ARRAY OF CHAR; VAR res: INT);
PROC Delete(name: ARRAY OF CHAR; VAR res: INT);
END Files.

```

New(name) yields a new (empty) file without registering it in the directory. Old(name) retrieves the file with the specified name, or yields NIL, if it is not found in the directory. Register(f) inserts the name of f (specified in the call of New) in the directory. An already existing entry with this name is replaced. Close(f) must be called after writing is completed and the file is not to be registered. Close actually stands for "close buffers", and is implied in the procedure Register. Procedure Purge will be explained at the end of section 7.2.

The sequential scan of a file f (reading characters) is programmed as shown in the following template:

```

VAR f: Files.File; r: Files.Rider;
f := Files.Old(name);
IF f $\neq$ NIL THEN
Files.Set (r, f, 0); Files.Read (r, x);
WHILE ~ r.eof DO ... x ...; Files.Read(r, x) END
END

```

The analogous template for a purely sequential writing is:

```
f := Files.New(name); Files.Set(r, f, 0);  
WHILE ... DO Files.Write (r, x); ... END  
Files.Register(f)
```

There exist two further procedures; they do not change any files, but only affect the directory. `Delete(name, res)` causes the removal of the named entry from the directory. `Rename(old, new, res)` causes the replacement of the directory entry old by new.

It may surprise the reader that these two procedures, which affect the directory only, are exported from module `Files` instead of `FileDir`. The reason is that the presence of the two modules, together forming the FS, is also used for separating the interface into a public and a private (or semi-public) part. The definition (in the form of a symbol file) of `FileDir` is not intended to be freely available, but restricted to use by system programmers. This allows the export of certain sensitive data, (such as file headers) and sensitive procedures (such as `Enumerate`) without the danger of misuse by inadvertent users.

Module `Files` constitutes a most important interface whose stability is utterly essential, because it is used by almost every module programmed. During the entire time span of development of the Oberon, this interface had changed only once. We also note that this interface is very terse, a

factor contributing to its stability. Yet, the offered facilities have in practice over years proved to be both necessary and sufficient.

## 7.2 IMPLEMENTATION ON A RANDOM-ACCESS STORE

A file cannot be allocated as a block of contiguous storage locations, because its length is not fixed. Neither can it be represented as a linked list of individual elements, because this would lead to inefficient use of storage - more might be used for the links than the elements themselves. The solution generally adopted is a compromise between the two extremes: files are represented as lists of blocks (subsequently called sectors) of fixed length. A block is appended when the last one is filled. On the average, each file therefore wastes half of a sector. Typical sector sizes are 0.5, 1, 2, or 4 Kbytes, depending on the device used as store.

It immediately follows that access to an element is not as simple as in the case of an array. The primary concern in the design of a FS and access scheme must be the efficiency of access to individual elements while scanning the sequence, at least in the case when the next element lies within the same sector. This access must be no more complicated than a comparison of two variables followed by an indexed access to the file element and the incrementing of an address pointing to the element's successor. If the successor lies in another sector, the procedure may be more involved, as transitions to the next sector occur much less frequently.

The second most crucial design decision concerns the data structure in which sectors are organized; it determines how a succeeding sector is located. The simplest solution is

to link sectors in a list. This is acceptable if access is to be restricted to purely sequential scans. Although this would be sufficient for most applications, it is unnecessarily restrictive for media other than purely sequential ones (tapes). After all, it is sometimes practical to position a rider at an arbitrary point in the file rather than always at its beginning. This is made possible by the use of an indexed sector table, typically stored as a header in the file. The table is an array of the addresses of the file's data sectors. Unfortunately, the length of the table needed is unknown. Choosing a fixed length for all files is controversial, because it inevitably leads to either a limitation of file length (when chosen too small) that is unacceptable in some applications, or to a large waste of file space (when chosen too large). Experience shows that in practice most files are quite short, i.e. in the order of a few thousand bytes. The dilemma is avoided by a two-level table, i.e. by using a table of tables.

The scheme chosen in Oberon is slightly more complex in order to favor short files (< 64 K bytes): Each file header contains a table of 64 entries, each pointing to a 1K byte sector. Additionally, it contains a table of 12 entries, the so-called extensions, each pointing to an index sector containing 256 further sector pointers. The file length is thereby limited to  $64 + 12 \cdot 256$  sectors, or  $3'211'264$  bytes (minus the length of the header). The chosen structure is illustrated in Fig. ??.

sec[o] always points to the sector containing the file header.



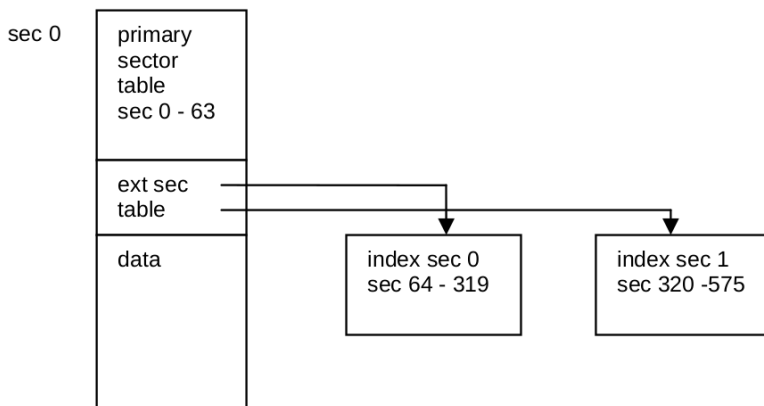


Figure 21: File header and extension sectors

The header contains some additional data, namely the length of the file (in bytes), its name, and date and time of its creation. The size of the header is 352 bytes; the remaining 672 bytes of the first sector are used for data. Hence, truly short files occupy a single sector only. The declaration of the file header is contained in the definition of module FileDir. An abbreviated version containing the fields relevant so far is:

```

FileHeader = RECORD
  leng: INT;
  ext: ARRAY 12 OF SectorPointer;
  sec: ARRAY 64 OF SectorPointer
END
  
```

We now turn our attention to the implementation of file access, and first present a system that uses main storage for the file data instead of a disk and therefore avoids the problems introduced by sector buffering. The key data structure in this connection is the Rider, represented as a record.

```
Rider = RECORD
eof: B00L; res, pos, adr: INT;
file: File
END
```

A rider is initialised by a call `Set(r, f, pos)`, which places the rider `r` on file `f` at position `pos`. From this it is clear that the rider record must contain fields denoting the attached file and the rider's position on it. We note that they are not exported. However, their values can be obtained by the function procedures `Pos(r)` and `Base(r)`. This allows a (hidden) representation most appropriate for an efficient implementation of `Read` and `Write` without being unsafe.

Consider now the call `Read(r, x)`; its task is to assign the value of the byte designated by the rider's position to `x` and to advance the position to the next byte. Considering the structure by which file data are represented, we easily obtain the following program, assuming that the position is legal, i.e. non-negative and less than the file's length. `a`, `b`, `c` are local variables, `HS` is the size of the header (in sector `o`), `SS`

is the sector size, typically a power of 2 in order to make division efficient.

```
a := (r.pos + HS) DIV SS; b := (r.pos + HS) MOD SS;  
IF a < 64 THEN c := r.file.sec[a]  
ELSE c := r.file.ext[(a - 64) DIV 256].sec[(a - 64) MOD 256]  
END ;  
SYSTEM.GET(c + b, x) ; INC (r.pos)
```

In order to gain efficiency, we use the low-level procedure GET that assigns the value at address  $c+b$  to  $x$ . This program is reasonably short, but involves considerable address computations at every access, and in particular at positions larger than  $64 * SS$ . Fortunately, there exists an easy remedy, namely that of caching the address of the current position. This explains the presence of the field *adr* in the rider record. The resulting program is shown below; note that in order to avoid the addition of *HS*, *pos* is defined to denote the genuine position, i.e. the abstract position augmented by *HS*.

```
SYSTEM.GET(r.adr, x); INC(r.adr); INC(r.pos);  
IF r.pos MOD SS = 0 THEN  
  m := r.pos DIV SS;  
  IF m < 64 THEN r.adr := r.file.sec[m]  
  ELSE r.adr := r.file.ext[(m - 64) DIV 256].sec[(m - 64) MOD 256]  
END  
END
```

We emphasize that in all but one out of 1024 cases only three instructions and a single test are to be executed. This improvement therefore is crucial to the efficiency of file access, and to that of the entire Oberon. We now present the entire file module (for files on a random-access store).

```

MODULE MFiles; (*NW 24.8.90 / 12.10.90 / 20.6.2013*)
IMPORT SYSTEM, Kernel, FileDir;

(*A file consists of a sequence of sectors. The first sector
Part of the header is the sector table, an array of addresses
A file is referenced through riders each of which indicates
CONST
HS = FileDir.HeaderSize;
SS = FileDir.SectorSize;
STS = FileDir.SecTabSize;
XS = FileDir.IndexSize;
TYPE File* = POINTER TO FileDesc;
Index = POINTER TO IndexRecord;
IndexRecord = RECORD sec: FileDir.IndexSector END ;
Rider* =
RECORD eof*: BOOL;
res*, pos, adr: INT;
file: File
END ;
FileDesc =
RECORD mark: INT;
name: FileDir.FileName;
len, date: INT;
```

```

ext: ARRAY FileDir.ExTabSize OF Index;
sec: FileDir.SectorTable
END ;
PROC Old*(name: ARRAY OF CHAR): File;
VAR head: INT;
namebuf: FileDir.FileName;
BEGIN
FileDir.Search(name, head); RETURN SYSTEM.VAL(File, head)
END Old;
PROC New*(name: ARRAY OF CHAR): File;
VAR f: File; head: INT;
BEGIN f := NIL; Kernel.AllocSector(0, head);
IF head $\neq$ 0 THEN
f := SYSTEM.VAL(File, head); f.mark := FileDir.HeaderMark;
f.len := HS; f.name := name;
f.date := Kernel.Clock(); f.sec[0] := head
END ;
RETURN f
END New;
PROC Register*(f: File);
BEGIN
IF (f $\neq$ NIL) & (f.name[0] > 0X) THEN FileDir.Insert(f.name, f.
END Register;
PROC Length*(f: File): INT;
BEGIN RETURN f.len - HS
END Length;
PROC Date*(f: File): INT;
BEGIN RETURN f.date

```

```
END Date;
```

```
PROC Set*(VAR r: Rider; f: File; pos: LONGINT);
```

```
VAR m, n: INT;
```

```
BEGIN r.eof := FALSE; r.res := 0;
```

```
IF f $\neq$ NIL THEN
```

```
IF pos < 0 THEN r.pos := HS
```

```
ELSIF pos > f.len - HS THEN r.pos := f.len
```

```
ELSE r.pos := pos + HS
```

```
END ;
```

```
r.file := f; m := r.pos DIV SS; n := r.pos MOD SS;
```

```
IF m < STS THEN r.adr := f.sec[m] + n
```

```
ELSE r.adr := f.ext[(m-STs) DIV XS].sec[(m-STs) MOD XS] + n
```

```
END
```

```
END
```

```
END Set;
```

```
PROC ReadByte*(VAR r: Rider; VAR x: BYTE);
```

```
VAR m: INT;
```

```
BEGIN
```

```
IF r.pos < r.file.len THEN
```

```
SYSTEM.GET(r.adr, x); INC(r.adr); INC(r.pos);
```

```
IF r.adr MOD SS = 0 THEN
```

```
m := r.pos DIV SS;
```

```
IF m < STS THEN r.adr := r.file.sec[m]
```

```
ELSE r.adr := r.file.ext[(m-STs) DIV XS].sec[(m-STs) MOD XS]
```

```
END
```

```
END
```

```
ELSE x := 0; r.eof := TRUE
```

```
END
```

```

END ReadByte;
PROC WriteByte*(VAR r: Rider; x: BYTE);
VAR k, m, n, ix: INT;
BEGIN
  IF r.pos < r.file.len THEN
    m := r.pos DIV SS; INC(r.pos);
    IF m < STS THEN r.adr := r.file.sec[m]
    ELSE r.adr := r.file.ext[(m-STS) DIV XS].sec[(m-STS) MOD XS]
    END
  ELSE
    IF r.adr MOD SS = 0 THEN
      m := r.pos DIV SS;
      IF m < STS THEN Kernel.AllocSector(0, r.adr); r.file.sec[m] := r.adr
      ELSE n := (m - STS) DIV XS; k := (m - STS) MOD XS;
      IF k = 0 THEN (*new index*)
        Kernel.AllocSector(0, ix); r.file.ext[n] := SYSTEM.VAL(Index, ix)
      END ;
      Kernel.AllocSector(0, r.adr); r.file.ext[n].sec[k] := r.adr
    END
  END ;
  INC(r.pos); r.file.len := r.pos
END ;
SYSTEM.PUT(r.adr, x); INC(r.adr)
END WriteByte;
PROC Pos*(VAR r: Rider): INT;
BEGIN RETURN r.pos - HS
END Pos;
PROC Base*(VAR r: Rider): File;

```

```
BEGIN RETURN r.file  
END Base;  
END MFiles.
```

Allocation of a new sector occurs upon creating a file (Files.New), and when writing at the end of a file after the current sector had been filled. Procedure AllocSector yields the address of the allocated sector. It is determined by a search in the sector reservation table for a free sector. In this table, every sector is represented by a single bit indicating whether or not the sector is allocated. Although conceptually belonging to the FS, this table resides within module Kernel.

Deallocation of a file's sectors could occur as soon as the file is no longer accessible, neither through a variable of any loaded module nor from the file directory. However, this moment is difficult to determine. Therefore, the method of garbage collection is used in Oberon for the deallocation of file space. In consideration of the fact that file space is large and the collection of unused sectors relatively time-consuming, we confine this process to system initialization. It is represented by procedure FileDir.Init. At that time, the only referenced files are those registered in the directory. Init therefore scans the entire directory and records the sectors referenced in each file in the sector reservation table (see Sect. 7.4).



For applications where system startup and initialization is supposed to occur very infrequently, such as for server systems, a procedure `Files.Purge` is provided. Its effect is to return the sectors used by the specified file to the pool of free sectors. Evidently, the programmer then bears the responsibility to guarantee that no references to the purged file continue to exist. This may be possible in a closed server system, but files should not be purged under normal circumstances, as a violation of said precondition will lead to unpredictable disaster. The following procedures used for allocating, deallocating, and marking sectors in the sector reservation table are defined in module `Kernel`:

```
PROC AllocSector(hint: INT; VAR sec: INT); (*used in WriteByte*)
PROC MarkSector(sec: INT); (*used in Init*)
PROC FreeSector(sec: INT); (*used in Purge*)
```

### 7.3 IMPLEMENTATION ON A DISK

First we recall that the organization of files as sets of individually allocated blocks (sectors) is inherently required by the allocation considerations of dynamically growing sequences. However, if the storage medium is a tape, a disk, or a flash-RAM, there exists an additional reason for the use of blocks. They constitute the subsequences to be individually buffered for transmission in order to overcome the timing constraints imposed by the medium. If an adequate space utilization is to be achieved, the blocks must not be too long. A typical size is 1, 2, or 4K bytes. This necessity of buffering has a profound influence on the implementation of file access. The complication arises because the abstraction of the sequence of individual bytes needs to be maintained. The increase in complexity of file access is considerable, as can be seen by comparing the program listings of the two respective implementations.

The first, obvious measure is to copy the file's sector table into primary store when a file is "opened" through a call of `New()` or `Old()`. The record holding this copy is the file descriptor, and the value `f` denoting the file points to this handle (instead of the actual header on disk). The descriptor also contains the remaining information stored in the header, in particular the file's length.

If a file is read (or written) in purely sequential manner, a single buffer is appropriate for the transfer of data. For

reading, the buffer is filled by reading a sector from the disk, and bytes are picked up individually from the buffer. For writing, bytes are deposited individually, and the buffer is written onto disk as a whole when full. The buffer is associated with the file, and a pointer to it is contained in the descriptor.

However, we recall that several riders may be placed on a file and be moved independently. It might be appealing to associate a buffer with each rider. But this proposal must quickly be rejected when we realize that several riders may be active at neighbouring positions. If these positions refer to the same sector, which is duplicated in the riders' distinct buffers, the buffers may easily become inconsistent. Obviously, buffers must not be associated with riders, but with the file itself. The descriptor therefore contains the head of a list of linked buffers. Each buffer is identified by its position in the file. An invariant of the system is that no two buffers represent the same sector.

Even with the presence of a single rider, the possibility of having several buffers associated with a file can be advantageous, if a rider is frequently repositioned. It becomes a question of strategy and heuristics when to allocate a new buffer. In the Oberon, we have adopted the following solution:

1. The first buffer is created when the file is opened (New, Old).

2. Additional buffers may be allocated when a rider is placed (or repositioned) on the file.
3. At most four buffers are connected to the same file.
4. Purely sequential movements of riders do not cause allocation of buffers.
5. Separate buffers are generated when extensions of the file's sector table need be accessed (rider position > 64K). Each buffers the 256 sector addresses of the respective index sector.

The outlined scheme requires and is based upon the following data structures and types:

File =

Buffer =

Index =

POINTER TO FileDesc;

POINTER TO BufferRecord;

POINTER TO IndexRecord;

FileDesc = RECORD next: File;

aleng, bleng: INT;

(\*file length\*)

nofbufs: INT; (\*no. of buffers allocated\*)

modH, registered: BOOL; (\*header has been modified\*)

```

firstbuf: Buffer:
(*head of buffer chain*)
sechint: DiskAdr;
(*sector hint*)
name: FileDir.FileName;
date: INT;
ext: ARRAY FileDir.ExTabSize OF Index;
sec: ARRAY 64 OF DiskAdr
END;
BufferRecord = RECORD apos, lim: INT; (*lim = no. of bytes*)
mod: BOOL;
(*buffer has been modified*)
next: Buffer;
(*buffer chain*)
data: FileDir.DataSector
END;
IndexRecord = RECORD adr: DiskAdr;
mod: BOOL;
(*index record has been modified*)
sec: FileDir.IndexSector
END;
Rider =

RECORD eof: BOOL;
res: INT;
file: File;
apos, bpos: INT;
buf: Buffer

```

END ;

(\*end of file reached\*)  
(\*no. of unread bytes\*)  
(\*position\*)  
(\*hint: likely buffer\*)

In order to increase efficiency of access, riders have been provided with a field containing the address of the element of the rider's position. From the conditions stated above for the allocation of buffers, it is evident that the value of this field can be a hint only. This implies that there can be no reliance on its information. Whenever it is used, its validity has to be checked. The check consists in a comparison of the riders' position `r.apos` with the hinted buffer's actual position `r.buf.apos`. If they differ, a buffer with the desired position must be searched and, if not present, allocated. The advantage of the hint lies in the fact that the hint is correct with a very high probability. The check is included in procedures `Read`, `ReadByte`, `Write`, and `WriteByte`. Some fields of the record types require additional explanations:

1. The length is stored in a "preprocessed" form, namely by the two integers `aleng` and `bleng` such that `aleng` is a sector number and

$$\begin{aligned}\text{length} &= (\text{aleng} * \text{SS}) + \text{bleng} - \text{HS} \\ \text{aleng} &= (\text{length} + \text{HS}) \text{ DIV } \text{SS}\end{aligned}$$

$$\text{bleng} = (\text{length} + \text{HS}) \text{ MOD } \text{SS}$$

The same holds for the form of the position in riders (apos, bpos).

2. The field `nofbufs` indicates the number of buffers in the list headed by `firstbuf`:

$$1 \leq \text{nofbufs} \leq \text{Maxbufs}.$$

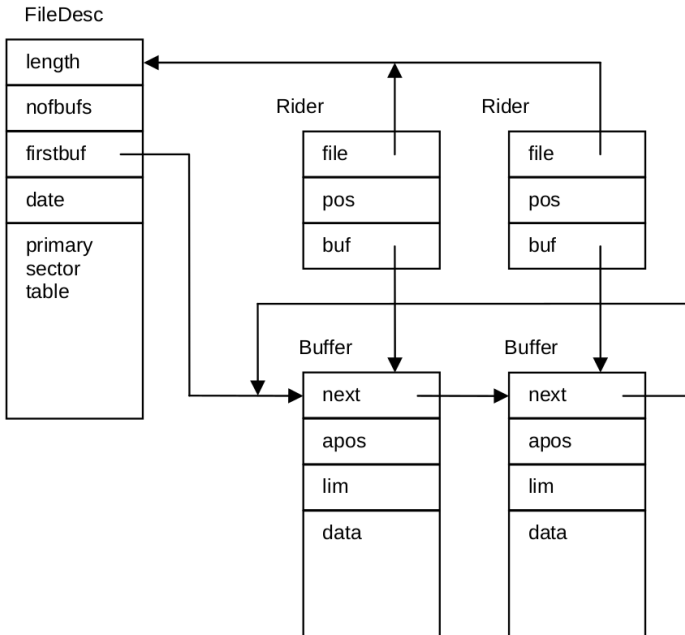
3. Whenever data are written into a buffer, the file becomes inconsistent, i.e. the data on the disk are outdated. The file is updated, i.e. the buffer is copied into the corresponding disk sector, whenever the buffer is reallocated, e.g. during sequential writing after the buffer is full and is "advanced". During sequential reading, a buffer is also advanced and reused, but needs not be copied onto disk, because it is still consistent. Whether a buffer is consistent or not is indicated by its state variable `mod` (modified). Similarly, the field `modH` in the file descriptor indicates whether or not the header had been modified.
4. The field `sechint` records the number of the last sector allocated to the file and serves as a hint to the kernel's allocation procedure, which allocates a next sector with an address larger than the hint. This is a measure to gain speed in sequential scans.

5. The buffer's position is specified by its field `apos`. Used as index in the file header's sector table, it yields the sector corresponding to the current buffer contents. The field `lim` specifies the number of bytes `s` stored in the buffer. Reading cannot proceed beyond this limiting index; writing beyond it implies an increase in the file's length. All buffers except the one for the last sector are filled and specify `lim = SS`.
6. The hidden rider field `buf` is merely a hint to speed up localization of the concerned buffer. A hint is likely, but not guaranteed to be correct. Its validity must be checked before use. The buffer hint is invalidated when a buffer is reallocated and/or a rider is repositioned. The structure of riders remains practically the same as for files using main store. The hidden field `adr` is merely replaced by a pointer to the buffer covering the rider's position. A configuration of a file `f` with two riders is shown in Fig ??.

Some comments concerning module Files follow.

1. After the writing of a file has been completed, its name is usually registered in the directory. Register invokes procedure `Unbuffer`. It inspects the associated buffers and copies those onto disk which had been modified. During this process, new index sectors may have to be transferred as well. If a file is to remain anonymous and local to a module or command, i.e. is not to be



Figure 22: File  $f$  with two riders and two buffers

registered, but merely to be read, the release of buffers must be specified by an explicit call to Close (meaning "close buffers"), which also invokes Unbuffer.

2. Procedure Old (and for reasons of consistency also New) deviates from the general Oberon programming rule that an object be allocated by the calling (instead of the called) module. This rule would suggest the statements

$$\text{New}(f); \text{Files.Open}(f, \text{name})$$

instead of

$$f := \text{Files.Old}(\text{name}).$$

The justification for the rule is that any extension of the type of  $f$  could be allocated, providing for more flexibility. And the reason for our deviation in the case of files is that, upon closer inspection, not a new file, but only a new descriptor is to be allocated. The distinction becomes evident when we consider that several statements  $f := \text{Files.Old}(\text{name})$  with different  $f$  and identical name may occur, probably in different modules. In this case, it is necessary that the same descriptor is referenced by the delivered pointers in order to avoid file inconsistency. Each (opened) file must have exactly one descriptor. When a file is opened, the first action is therefore to inspect whether a descriptor of this file already exists. For this purpose, all descriptors are linked together in a list anchored by the

global variable root and linked by the descriptor field next. This measure may seem to solve the problem of avoiding inconsistencies smoothly. However, there exists a pitfall that is easily overlooked: All opened files would permanently remain accessible via root, and the garbage collector could never remove a file descriptor nor its associated buffers. This would be unacceptable. In order to hide this list from the garbage collector, it is represented by integers (addresses) instead of pointers.

3. Sector pointers are represented by sector numbers of type *INT*. Actually, we use the numbers multiplied by 29. This implies that any single-bit error leads to a number which is not a multiple of 29, and hence can easily be detected. Thereby the crucial sector addresses are software parity checked and are safe (against single-bit errors) even on computers without hardware parity check. The check is performed by procedures *Kernel.GetSector* and *Kernel.PutSector*.

#### 7.4 THE FILE DIRECTORY

A directory is a set of pairs, each pair consisting of a name (key) and an object (here: file). It serves to retrieve objects by their name. If efficiency matters, the directory is organized as an ordered set, ordered according to the keys. The most frequently used structures for ordered sets are trees and hash tables. The latter have disadvantages when the size of the set is unknown, particularly when its order of magnitude is unknown, and when deletions occur. The Oberon system therefore uses a tree structure for its file directory, more specifically a B-tree, which was developed especially for cases where not individual pairs, but only sets of pairs as a whole (placed on a disk sector) can be accessed.

For a thorough study of B-trees we refer the reader to the literature [1, 2]. Here it must suffice to specify the B-tree's principal characteristics:

1. In a B-tree of order  $N$ , each node (called page) contains  $m$  elements (pairs), where  $N \leq m \leq 2N$ , except the root, where  $0 \leq m \leq 2N$ .
2. A page with  $m$  elements has either 0 descendants, in which case it is called a leaf page, or  $m + 1$  descendants.
3. All leaf pages are on the same (bottom) level.

height	minimum	maximum
1	0	24
2	25	624
3	625	15624
4	15625	390624

From 3, it follows that the B-tree is a balanced tree. Its height, and with it the longest path's length, has an upper bound of, roughly,  $2 * \log k$ , where  $k$  is the number of elements and the logarithm is taken to the base  $N$  and rounded up to the next larger integer. Its minimal height is  $\log k$  taken to the base  $2N$ .

On each page, space must be available for  $2N$  elements and for  $2N + 1$  references to descendants. Hence,  $N$  is immediately determined by the size of a page and the size of elements. In the case of Oberon, names are limited to 32 characters (bytes), and the object is a reference to the associated file (4 bytes). Each descendant pointer takes 4 bytes, and the page size is given by the sector size (1024) minus the number of bytes needed to store  $m$  (2 bytes). Hence

$$N = ((1024 - 2 - 4) \text{ DIV } (32 + 4 + 4)) \text{ DIV } 2 = 12$$

A B-tree of height  $h$  and order 12 may contain the following minimal and maximal number of elements:

It follows that the height of the B-tree will never be larger than 4, if the disk has a capacity of less than about 400 Mbyte, and assuming that each file occupies a single 1K sector. It is rarely larger than 3 in practice.

The definition of module FileDir shows the available directory operations. Apart from the procedures Search, Insert, Delete, and Enumerate, it contains some data definitions, and it should be considered as the non-public part of the FS's interface.

```
DEFINITION FileDir;
IMPORT SYSTEM, Kernel;
CONST
  FnLength = 32;
  (*max length of file name*)
  SecTabSize = 64; (*no. of entries in primary table*)
  ExTabSize = 12;
  SectorSize = 1024;
  IndexSize = SectorSize DIV 4;
  (*no. of entries in index sector*)
  HeaderSize = 352;
  DirRootAdr = 29;
  DirPgSize = 24;
  (*max no. of elements on page*)
  TYPE DiskAdr = INT;
  FileName = ARRAY FnLength OF CHAR;
  SectorTable = ARRAY SecTabSize OF DiskAdr;
```

```

ExtensionTable = ARRAY ExTabSize OF DiskAdr;
EntryHandler = PROC (name: FileName; sec: DiskAdr; VAR continue: BOOLEAN);
FileHeader = RECORD (*first page of each file on disk*)
  mark: INT;
  name: FileName;
  aleng, bleng, date: INT;
  ext: ExtensionTable;
  sec: SectorTable
END ;
IndexSector = RECORD (Kernel.Sector)
  x: ARRAY IndexSize OF LONGINT;
END ;
DataSector = ARRAY SectorSize OF BYTE;
DirEntry = RECORD
  name: FileName;
  adr, p: DiskAdr
END ;
DirPage = RECORD
  mark: INT;
  m: INT; (*no. of elements on page*)
  p0: DiskAdr;
  e: ARRAY DirPgSize OF DirEntry;
END ;
PROC Search(name: FileName; VAR fad: DiskAdr);
PROC Insert(name: FileName; fad: DiskAdr);
PROC Delete(name: FileName; VAR fad: DiskAdr);
PROC Enumerate(prefix: ARRAY OF CHAR; proc: EntryHandler);
END FileDir.

```

Procedures Search, Insert, and Delete represent the typical operations performed on a directory. Efficiency of the first operation is of primary importance. But the B-tree structure also guarantees efficient insertion and deletion, although the code for these operations is complex. Procedure Enumerate is used to obtain excerpts of the directory. The programmer must guarantee that no directory changes are performed by the parametric procedure of Enumerate.

As in the presentation of module Files, we first discuss a version that uses main storage rather than a disk for the directory. This allows us to concentrate on the algorithms for handling the directory, leaving out the additional complications due to the necessity to read pages (sectors) into main store for selective updating and of restoring them onto disk. In particular, we point out the definitions of the data types for B-tree nodes, called DirPage, and elements, called DirEntry. The component E.p of an entry E points to the page in which all elements (with index k) have keys  $E.p.e[k].name > E.name$ . The pointer p.po points to a page in which all elements have keys  $p.po.e[k].name < p.e[o].name$ . We can visualize these conditions by Fig. ??, where names have been replaced by integers as keys.

Procedure *Search* starts by inspecting the root page. It performs a binary search among its elements, according to the following algorithm. Let  $e[0 \dots m-1]$  be the ordered keys and  $x$  the search argument.






Figure 23: Example of a B-tree of order 2

```

L := 0; R := m;
WHILE L < R DO
  i := (L+R) DIV 2;
  IF x <= e[i] THEN R := i ELSE L := i + 1 END
END;
IF (R < m) & (x = e[R]) THEN found END

```

The invariant is

$$e[L - 1] < x \leq e[R]$$

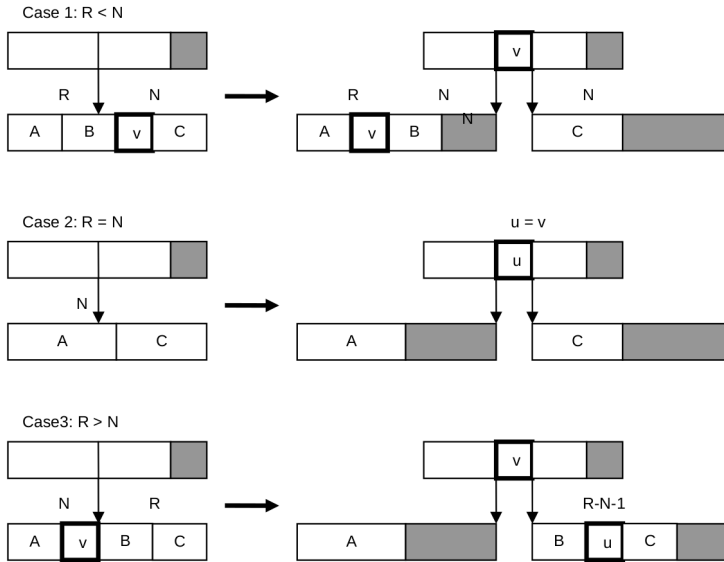
If the desired element is not found, the search continues on the appropriate descendant page, if there is one. Otherwise the element is not contained in the tree.

Procedures insert and delete use the same algorithm for searching an element within a page. However, they use recursion instead of iteration to proceed along the search path of pages. We recall that the depth of recursion is at most four. The reason for the use of recursion is that it facilitates the formulation of structural changes, which are performed during the "unwinding" of recursion, i.e. on

the return path. First, the insertion point (respectively the position of the element to be deleted) is searched, and then the element is inserted (deleted).

Upon insertion, the number of elements on the insertion page may become larger than  $2N$ , violating B-tree condition 1. This situation is called page overflow. The invariant must be reestablished immediately. It could be achieved by moving one element from either end of the array  $e$  onto a neighbouring page. However, we choose not to do this, and instead to split the overflowing page into two pages immediately. The process of a page split is visualized by Fig ??, in which we distinguish between three cases, namely  $R < N$ ,  $R = N$ , and  $R > N$ , where  $R$  marks the insertion point.  $a$  denotes the overflowing,  $b$  the new page, and  $u$  the inserted element. The  $2N + 1$  elements ( $2N$  from the full page  $a$ , plus the one element  $u$  to be inserted) are equally distributed onto pages  $a$  and  $b$ . One element  $v$  is pushed up in the tree. It must be inserted in the ancestor page of  $a$ . Since that page obtains an additional descendant, it must also obtain an additional element in order to maintain B-tree rule 2.

A page split may thus propagate, because the insertion of element  $v$  in the ancestor page may require a split once again. If the root page is full, it is split too, and the emerging element  $v$  is inserted in a new root page containing a single

Figure 24: Page split when inserting element  $u$ 

element. This is the only way in which the height of a B-tree can increase.

When an element is to be deleted, it cannot simply be removed, if it resides on an internal page. In this case, it is first replaced by another element, namely one of the two neighbouring elements on a leaf page, i.e. the next smaller (or next larger) element, which is always on a leaf page. In the presented solution, the replacing element is the largest on the left subtree (see procedure `del`). Hence, the actual deletion always occurs on a leaf page.

Upon deletion, the number of elements in a page may become less than  $N$ , violating invariant 1. This event is called page underflow. Since restructuring the tree is a relatively complicated operation, we first try to reestablish the invariant by borrowing an element from a neighbouring page. In fact, it is reasonable to borrow several elements, and thereby to decrease the likelihood of an underflow on the same page upon further deletions. The number of elements that could be taken from the neighbouring page  $b$  is  $b.m - N$ . Hence we will borrow

$$k = (b.m - N + 1) \text{ DIV } 2$$

elements. The process of page balancing then distributes the elements of the underflowing and its neighbouring page equally to both pages (see procedure underflow). However, if (and only if) the neighbouring page has no elements to spare, the two pages can and must be united. This action, called page merge, places the  $N - 1$  elements from the underflowing page, the  $N$  elements from the neighbouring page, plus one element from the ancestor page onto a single page of size  $2N$ . One element must be taken from the ancestor page, because that page loses one descendant and invariant rule 2 must be maintained. The events of page balancing and merging are illustrated in Fig ?? .  $a$  is the underflowing page,  $b$  its neighbouring page, and  $c$  their ancestor;  $s$  is the position in the ancestor page of (the pointer to the) underflowing page  $a$ . Two cases are distinguished,

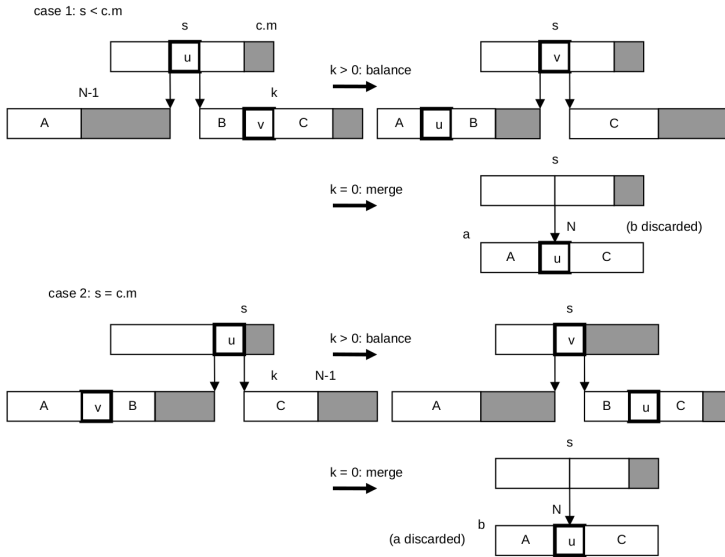


Figure 25: Page balancing and merging when deleting element

namely whether the underflowing page is the rightmost element ( $s = c.m$ ) or not (see procedure underflow).

Similarly to the splitting process, merging may propagate, because the removal of an element from the ancestor page may again cause an underflow, and perhaps a merge. The root page underflows only if its last element is removed. This is the only way in which the B-tree's height can decrease.

```
MODULE BTree;
IMPORT Texts, Oberon;
CONST N = 3;
TYPE Page = POINTER TO PageRec;
```

```

Entry = RECORD
key, data: INT;
p: Page
END ;
PageRec = RECORD
m: INT; (*no. of entries on page*)
p0: Page;
e: ARRAY 2*N OF Entry
END ;
VAR root: Page; W: Texts.Writer;
PROC search(x: INT; VAR p: Page; VAR k: INT);
VAR i, L, R: INT; found: BOOL; a: Page;
BEGIN a := root; found := FALSE;
WHILE (a $\neq$ NIL) & ~found DO
L := 0; R := a.m; (*binary search*)
WHILE L < R DO
i := (L+R) DIV 2;
IF x <= a.e[i].key THEN R := i ELSE L := i+1 END
END ;
IF (R < a.m) & (a.e[R].key = x) THEN found := TRUE
ELIF R = 0 THEN a := a.p0 ELSE a := a.e[R-1].p
END
END ;
p := a; k := R
END search;
PROC insert(x: INT; a: Page; VAR h: BOOL; VAR v: Entry);
(*a $\neq$ NIL. Search key x in B-tree with root a; if found
Otherwise insert new item with key x. If an entry is to be p

```

```

assign it to v. h := "tree has become higher"*)
VAR i, L, R: INT;
b: Page; u: Entry;
BEGIN (*a $\neq$ NIL & ~h*)
L := 0; R := a.m; (*binary search*)
WHILE L < R DO
i := (L+R) DIV 2;
IF x <= a.e[i].key THEN R := i ELSE L := i+1 END
END ;
IF (R < a.m) & (a.e[R].key = x) THEN (*found*) INC(a.e[R].data)
ELSE (*item not on this page*)
IF R = 0 THEN b := a.p0 ELSE b := a.e[R-1].p END ;
IF b = NIL THEN (*not in tree, insert*)
u.p := NIL; h := TRUE; u.key := x
ELSE insert(x, b, h, u)
END ;
IF h THEN (*insert u to the left of a.e[R]*)
IF a.m < 2*N THEN
h := FALSE; i := a.m;
WHILE i > R DO DEC(i); a.e[i+1] := a.e[i] END ;
a.e[R] := u; INC(a.m)
ELSE NEW(b); (*overflow; split a into a,b and assign the middle ent
IF R < N THEN (*insert in left page a*)
i := N-1; v := a.e[i];
WHILE i > R DO DEC(i); a.e[i+1] := a.e[i] END ;
a.e[R] := u; i := 0;
WHILE i < N DO b.e[i] := a.e[i+N]; INC(i) END
ELSE (*insert in right page b*)

```

```

DEC(R, N); i := 0;
IF R = 0 THEN v := u
ELSE v := a.e[N];
WHILE i < R-1 DO b.e[i] := a.e[i+N+1]; INC(i) END ;
b.e[i] := u; INC(i)
END ;
WHILE i < N DO b.e[i] := a.e[i+N]; INC(i) END
END ;
a.m := N; b.m := N; b.p0 := v.p; v.p := b
END
END
END
END insert;
PROC underflow(c, a: Page; s: INT; VAR h: BOOL);
(*a = underflowing page, c = ancestor page,
s = index of deleted entry in c*)
VAR b: Page;
i, k: INT;
BEGIN (*h & (a.m = N-1) & (c.e[s-1].p = a) *)
IF s < c.m THEN (*b := page to the right of a*)
b := c.e[s].p; k := (b.m-N+1) DIV 2; (*k = nof items available*)
a.e[N-1] := c.e[s]; a.e[N-1].p := b.p0;
IF k > 0 THEN (*balance by moving k-1 items from b to a*) i := 0;
WHILE i < k-1 DO a.e[i+N] := b.e[i]; INC(i) END ;
c.e[s] := b.e[k-1]; b.p0 := c.e[s].p;
c.e[s].p := b; DEC(b.m, k); i := 0;
WHILE i < b.m DO b.e[i] := b.e[i+k]; INC(i) END ;
a.m := N-1+k; h := FALSE

```



```

ELSE (*merge pages a and b, discard b*) i := 0;
WHILE i < N DO a.e[i+N] := b.e[i]; INC(i) END ;
i := s; DEC(c.m);
WHILE i < c.m DO c.e[i] := c.e[i+1]; INC(i) END ;
a.m := 2*N; h := c.m < N
END

ELSE (*b := page to the left of a*) DEC(s);
IF s = 0 THEN b := c.p0 ELSE b := c.e[s-1].p END ;
k := (b.m-N+1) DIV 2; (*k = nof items available on page b*)
IF k > 0 THEN i := N-1;
WHILE i > 0 DO DEC(i); a.e[i+k] := a.e[i] END ;
i := k-1; a.e[i] := c.e[s]; a.e[i].p := a.p0;
(*move k-1 items from b to a, one to c*) DEC(b.m, k);
WHILE i > 0 DO DEC(i); a.e[i] := b.e[i+b.m+1] END ;
c.e[s] := b.e[b.m]; a.p0 := c.e[s].p;
c.e[s].p := a; a.m := N-1+k; h := FALSE
ELSE (*merge pages a and b, discard a*)
c.e[s].p := a.p0; b.e[N] := c.e[s]; i := 0;
WHILE i < N-1 DO b.e[i+N+1] := a.e[i]; INC(i) END ;
b.m := 2*N; DEC(c.m); h := c.m < N
END
END
END underflow;

PROC delete(x: INT; a: Page; VAR h: BOOL);
(*search and delete key x in B-tree a; if a page underflow arises,
balance with adjacent page or merge; h := "page a is undersize"*)
VAR i, L, R: INT; q: Page;
PROC del(p: Page; VAR h: BOOL);

```

```

VAR k: INT; q: Page; (*global a, R*)
BEGIN k := p.m-1; q := p.e[k].p;
IF q $\neq$ NIL THEN del(q, h);
IF h THEN underflow(p, q, p.m, h) END
ELSE p.e[k].p := a.e[R].p; a.e[R] := p.e[k];
DEC(p.m); h := p.m < N
END
END del;
BEGIN (*a $\neq$ NIL*)
L := 0; R := a.m; (*binary search*)
WHILE L < R DO
i := (L+R) DIV 2;
IF x <= a.e[i].key THEN R := i ELSE L := i+1 END
END ;
IF R = 0 THEN q := a.p0 ELSE q := a.e[R-1].p END ;
IF (R < a.m) & (a.e[R].key = x) THEN (*found*)
IF q = NIL THEN (*a is leaf page*)
DEC(a.m); h := a.m < N; i := R;
WHILE i < a.m DO a.e[i] := a.e[i+1]; INC(i) END
ELSE del(q, h);
IF h THEN underflow(a, q, R, h) END
END
ELSE delete(x, q, h);
IF h THEN underflow(a, q, R, h) END
END
END delete;
PROC Search*(key: INT; VAR data: INT);
BEGIN search(key, root, data)

```

```

END Search;
PROC Insert*(key: INT; VAR data: INT);
VAR h: BOOL; u: Entry; q: Page;
BEGIN h := FALSE; u.data := data; insert(key, root, h, u);
IF h THEN (*insert new base page*)
q := root; NEW(root);
root.m := 1; root.p0 := q; root.e[0] := u
END
END Insert;
PROC Delete*(key: INT);
VAR h: BOOL;
BEGIN h := FALSE; delete(key, root, h);
IF h THEN (*base page size underflow*)
IF root.m = 0 THEN root := root.p0 END
END
END Delete;
BEGIN NEW(root); root.m := 0
END BTree.

```

The B-tree is also a highly appropriate structure for enumerating its elements, because during the traversal of the tree each page is visited exactly once, and hence needs to be read (from disk) exactly once too. The traversal is programmed by the procedure Enumerate and uses recursion. It calls the parametric procedure proc for each element of the tree. The type of proc specifies as parameters the name and the (address of) the enumerated element. The third parameter

continue is a Boolean VAR-parameter. If the procedure sets it to FALSE, the process of enumeration will be aborted.

Enumerate is used for obtaining listings of the names of registered files. For this purpose, the actual procedure substituted for proc merely enters the given name in a text and ignores the address (sector number) of the file, unless it requires special file information such as the file's size or creation date.

The set of visited elements can be restricted by specifying a string which is to be a prefix to all enumerated names. The least name with the specified prefix is directly searched and is the name (key) of the first element enumerated. The process then proceeds up to the first element whose name does not have the given prefix. Thereby, the process of obtaining all elements whose key has a given prefix avoids traversal of the whole tree, resulting in a significant speedup. If the prefix is the empty string, the entire tree is traversed.

The principle behind procedure Enumerate is shown by the following sketch, where we abstract from the B-tree structure and omit consideration of prefixes:

```
PROC Enumerate(  
  proc: PROC (name: FileName; adr: INT; VAR continue: BOOL));  
  VAR continue: BOOL; this: DirEntry;  
  BEGIN continue := TRUE; this := FirstElement;  
  WHILE continue & (this $\neq$ NIL) DO
```

```

proc(this.name, this.adr, continue); this := NextEntry(this)
END
END Enumerate

```

From this sketch we may conclude that during the process of traversal the tree structure must not change, because the function `NextEntry` quite evidently relies on the structural information stored in the elements of structure itself. Hence, the actions of the parametric procedure must not affect the tree structure. Enumeration must not be used, for example, to delete a given set of files. In order to prevent the misuse of the indispensable facility of element enumeration, the interface of `FileDir` is not available to users in general.

The handling of the directory stored on disk follows exactly the same algorithms. The accessed pages are fetched from the disk as a whole (each page fits onto a single disk sector) and stored in buffers of type `DirPage`, from where individual elements can be accessed. In principle, these buffers can be local to procedures `insert` and `delete`. A single buffer is allocated globally, namely the one used by procedure `Search`. The reason for this exception is not only that iterative searching requires one buffer only, but because procedure `Files.Old` and in turn `Search` may be called when the processor is in the supervisor mode and hence uses the system- (instead of the user-) stack, which is small and would not accommodate sector buffers.

Naturally, an updated page needs to be stored back onto disk. Omission of sector restoration is a programming error that is very hard to diagnose, because some parts of the program are executed very rarely, and hence the error may look sporadic and mistakenly be attributed to malfunctioning hardware.

Oberon's file directory represents a single, ordered set of name-file pairs. It is therefore also called a flat directory. Its internal tree structure is not visible to the outside. In contrast, some file systems use a directory with a visible tree structure, notably UNIX. In a search, the name (key) guides the search path; the name itself displays structure, in fact, it is a sequence of names (usually separated by slashes or periods). The first name is then searched in the root directory, whose descendants are not files but subdirectories. The process is repeated, until the last name in the sequence has been used (and hopefully denotes a file).

Since the search path length in a tree increases with the logarithm of the number of elements, any subdivision of the tree inherently decreases performance since  $\log(m + n) < \log(m) + \log(n)$  for any  $m, n > 1$ . It is justified only if there exist sets of elements with common properties. If these property values are stored once, namely in the subdirectory referencing all elements with common property values, instead of in every element, not only a gain in storage economy results, but possibly also in accesses which depend on those proper-

ties. The common properties are typically an owner's name, a password, and access rights (read or write protection), properties that primarily have significance in a multi-user environment. Since Oberon was conceived explicitly as a single-user system, there is little need for such facilities, and hence a flat directory offers the best performance with a simple implementation.

Every directory operation starts with an access to the root page. An obvious measure for improving efficiency is to store the root page "permanently" in main store. We have chosen not to do this for four reasons:

1. If the hardware fails, or if the computer is switched off before the root page is copied to disk, the file directory will be inconsistent with severe consequences.
2. The root page has to be treated differently from other pages, making the program more complex.
3. Directory accesses do not dominate the computing process; hence, any improvement would hardly be noticeable in overall system performance. The payoff for the added complexity would be small.
4. Procedure *Init* is called upon system initialization in order to construct the sector reservation table. Therefore, this procedure (and the module) must be allowed to refer to the structure of a file's sector table(s), which

is achieved by placing its definitions into the module *FileDir* (instead of *Files*). Unlike *Enumerate*, *Init* traverses the entire B-tree. The sector numbers of files delivered by *TraverseDir* are entered into a buffer. When full, the entries are sorted, whereafter each file's head sector is read and the sectors indicated in its sector table are marked as reserved. The sorting speeds up the reading of the header sectors considerably. Nevertheless, the initialization of the sector reservation table clearly dominates the start-up time of the computer. For a file system with 10,000 files it takes in the order of 15s to record all files.



## 7.5 THE TOOLBOX OF FILE UTILITIES

We conclude this chapter with a presentation of the commands which constitute the toolbox for file handling. These commands are contained in the tool module *System*, and they serve to copy, rename, and delete files, and to obtain excerpts of the file directory. Procedures *CopyFiles*, *RenameFiles*, and *DeleteFiles* all follow the same pattern. The parameter text is scanned for file names, and for each operation a corresponding procedure is called. If the parameter text contains an arrow, it is interpreted as a pointer to the most recent text selection which indicates the file name. In the cases of *CopyFiles* and *RenameFiles* which require two names for a single action, the names are separated by "⇒" indicating the direction of the copy or rename actions.

Procedure *Directory* serves to obtain excerpts of the file directory. It makes use of procedure *FileDir.Enumerate*. The parametric procedure *List* tests whether or not the delivered name matches the pattern specified by the parameter of the directory command. If it matches, the name is listed in the text of the viewer opened in the system track. Since the pattern may contain one or several asterisks (wild cards), the test consists of a sequence of searches of the pattern parts (separated by the asterisks) in the file name. In order to reduce the number of calls of *List*, *Enumerate* is called with the first part of the pattern as parameter prefix. Enumeration then starts with the least name having the specified prefix,

and terminates as soon as all names with this prefix have been scanned.

If the specified pattern is followed by an option directive "!", then not only file names are listed, but also the listed files' creation date and length. This requires that not only the directory sectors on the disk are traversed, but that additionally for each listed file its header sector must be read. The two procedures use the global variables *pat* and *diroption*.

#### REFERENCES

1. R. Bayer and E. M. McCreight. Organization and maintenance of large ordered indexes. *Acta Informatica*, 1, 3, (1972), 173-189.
2. D. Comer. The ubiquitous B-tree. *ACM Comp Surveys*, 11, 2, (June 1979), 121-137.

## STORAGE LAYOUT AND MANAGEMENT

---

Centralized RM (CRM) is a crucial property of Oberon.

ITS ADVANTAGE is that replication of management algorithms and a premature partitioning of resources are avoided.

THE DISADVANTAGE is that management algorithms are fixed once and forever and remain the same for all applications.

The success of a CRM therefore depends crucially on its flexibility and its efficient implementation. This chapter presents the scheme and the algorithms governing main storage in Oberon.

### 8.1 LAYOUT AND RUN-TIME ORGANIZATION

The storage layout of Oberon is determined by the structure of code and data typical in the use of modular, high-level PLs, and in particular of the language Oberon. It suggests the subdivision of storage into 3 areas:

**MODULE SPACE** Each module specifies procedures (code) and global (static) variables. Its initialization can be regarded as a procedure implicitly called upon loading. Space must be allocated for code and data in *blocks* by the loader. Typically, modules contain very few global variables, hence the space size is primarily determined by code.

**WORKSPACE (STACK)** Execution of every command invokes a sequence of procedures, each of which uses (possibly zero) parameters and local variables. Since procedure calls and completions follow a strict first-in last-out order, the stack is the uniquely suited strategy for local storage allocation. Deallocation upon completion of a procedure is achieved by merely resetting the pointer identifying the top of the stack. Since this operation is performed by a single instruction, it costs virtually no time. Because Oberon is a single-process system, a single stack suffices. Furthermore, after completion of a command, the stack is empty. This fact will be important in simplifying the scheme for reclamation of dynamically allocated space.

**DYNAMIC SPACE (HEAP)** Apart from global (static) variables, and local (stack-allocated) variables, a program may refer to anonymous variables referenced through pointers. Such variables are allocated truly dynamically through calls of an explicit **NEW** operation. These

variables are allocated in the so-called *heap*. Their deallocation is "automatic", when free storage is needed and they are no longer referenced from any of the loaded modules. This process is called *garbage collection* (GC). Every record allocated in the heap contains a (hidden) pointer to the type descriptor, called *type tag*, used by GC.

Unfortunately, the number of distinct spaces is larger than 2. If it were 2, no arbitrary size limitation would be necessary; merely the sum of their sizes would be inherently limited by the size of the store. In the case of 3 spaces, arbitrarily determined size limits are unavoidable. Address mapping hardware can alleviate (and delegate) this problem using a virtual address space which is so large that limits will hardly ever be reached.

Such a scheme is implemented by tables mapping virtual into physical addresses, requiring multiple memory accesses for every reference. Of course, the need for a double or a triple access for every memory reference is avoided by a translation cache in the (hardware) unit. Nevertheless, a decrease in performance is unavoidable for each cache miss. Furthermore, an additional subcycle is required for every access in order to look up the cached translation table. Without a virtual address scheme, each module block must consist of an integral number of physically adjacent pages. Holes generated by the release of modules must be reused. We

employ the simple scheme of marking the released space as a hole, and of allocating a new block in the 1st hole encountered that is large enough (1st-fit strategy). Considering the relative infrequency of module releases, efforts to improve the strategy are not worth the resulting added complexity.

It is remarkable that the code for module allocation and release without virtual addressing is only marginally more complicated than with it. The only remaining advantages of an MMU are a better storage utilization, because no holes occur (a negligible advantage), and that inadvertent references to unloaded modules, e.g. via installed procedures, lead to an invalid address trap.

It is worth recalling that the concept of address mapping was introduced as a requirement for virtual memory implemented with disks as backing store, where pages could be moved into the background in order to obtain space for newly required pages, and could then be retrieved from disk on demand, i.e. when access was requested. This scheme is called *demand paging*. It is not used in Oberon, and one may fairly state that demand paging has lost its significance with the availability of large, primary stores.

Experience in the use of the RISC predecessor Ceres leads to the conclusion that whereas address translation through an MMU was an essential feature for multi-user OSeS, it constitutes a dispensable overkill for single-user workstations. The fact that modern semiconductor technology made it possible

to integrate the entire translation and caching scheme into a single chip, or even into the processor itself, led to the hiding (and ignoring) of the scheme's considerable complexity. Its side effects on execution speed are essentially unpredictable. This makes systems with MMU virtually unusable for applications with tight real-time constraints. The RISC processor does indeed not feature an address mapping unit.

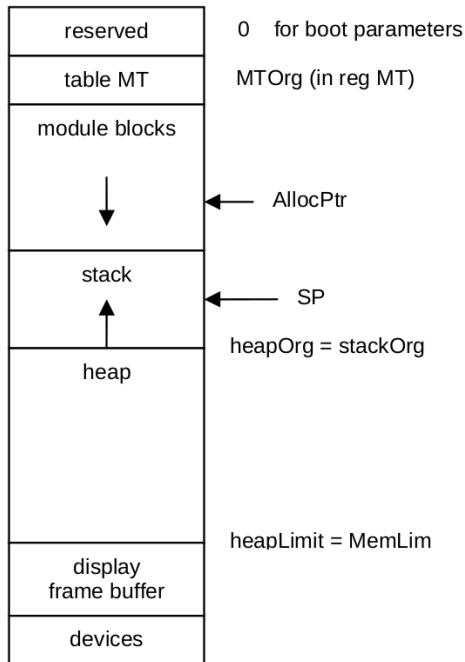


Figure 26: Storage layout

The RISC processor features 16 registers (of 32 bits). R0 - R11 are for expression evaluation. R12 - R15 have fixed, system-wide usage:

- R12 address of the module table MT (typically constant)
- R13 base address for variables in the current module SB (static base)
- R14 stack pointer SP
- R15 return address LNK (fixed by RISC's BL instruction)

The used memory layout is shown in Figure ??.

## 8.2 HEAP MANAGEMENT

The term *dynamic storage* is used here for all variables that are allocated neither statically (global variables) nor on the stack (local variables), but through invocation of the intrinsic procedure `NEW` in the heap. Such variables are anonymous and are referenced exclusively via pointers.

The space allocated to such dynamic variables becomes free and reusable as soon as the last reference to it vanishes. This event is hard, and in multiprocess systems even impossible to detect. The usual remedy is to ignore it and instead to determine the accessibility of all allocated variables (records, objects) only at the time when more storage space is needed. This process is then called *garbage collection*.

Oberon does not provide an explicit deallocation procedure allowing the programmer to signal that a variable will no longer be referenced. The 1st reason for this omission is that usually a programmer would not know when to call for



deallocation. And 2ndly, this "hint" could not be taken as trustworthy. An erroneous deallocation, i.e. one occurring when there still exist references to the object in question, could lead to a multiple allocation of the same space with disastrous consequences. Hence, it appears wise to fully rely on system management to determine which areas of the store are truly reusable.

Before discussing the scheme for storage reclamation, which is the primary subject of heap management, we turn our attention to the problem of allocation, i.e. the implementation of procedure NEW. The simplest solution is to maintain a list of free blocks and to pick the 1st one large enough. This strategy leads to a relatively large fragmentation of space and produces many small elements, particularly in the 1st part of the list. We therefore employ a somewhat more refined scheme and maintain 4 lists of available space. 3 of them contain pieces of fixed size, namely 32, 64, and 128 bytes. The 4th list contains pieces whose size is any multiple of 256. We note that the choice of the values permits the merging of any 2 contiguous elements into an element of the next list. This scheme keeps fragmentation, i.e. the emergence of small pieces in large numbers, reasonably low with minimal effort. The body of procedure NEW consists of relatively few instructions, and typically only a small fraction of them needs to be executed.

The statement `NEW(p)` is compiled into an instruction sequence assigning the address of pointer variable  $p$  to a fixed register (`R0`) and the type tag to another register (`R1`). The type tag is a pointer to a type descriptor containing information required by GC. This includes the size of the space occupied and now to be allocated. The effect of `NEW` is the assignment of the address of the allocated block to  $p$ , and the assignment of the tag to a prefix of the block. (see Fig. ??)




Figure 27: Heap allocation of dynamic variable  $p^$  by `NEW(p)`

In conclusion, we emphasize that this scheme makes the allocation of an object very efficient. Nevertheless, it is considerably more costly than that of a variable explicitly declared and therefore allocated either globally or on the stack. We now turn to the problem of storage reclamation or GC. There exist 2 essentially different schemes:

1<sup>st</sup>, the reference counting, and

2<sup>nd</sup>, the mark-scan schemes.

In the former, every object carries a (hidden) *reference count*, indicating the number of existing references to it. The scheme works as follows:

1. `NEW(p)` initializes the reference count of  $p^{\wedge}$  to 1.
2. An assignment  $q := p$  decrements the reference count of  $q^{\wedge}$  by 1, performs the assignment, then increments the reference count of  $p^{\wedge}$  by 1.

When a reference count reaches zero, the element is linked into the free list. There are 2 disadvantages inherent in this approach:

1. the non-negligible overhead in pointer assignments.
2. circular data structures never become recognized as free, even if no external references point to their elements.

Oberon employs the 2nd scheme which involves no hidden operations like the 1st one, but relies on a process initiated when free storage has become scarce and more is needed. It consists of 2 phases:

`MARK PHASE` all referenced and therefore still accessible elements are marked.

`SCAN PHASE` their unmarked complement is released.

Its primary disadvantage is that the process may be started at moments unpredictable to the system's user. During the process, the computer then appears to be blocked. It follows that an interactive system using mark-scan GC must guarantee that the process is sufficiently fast in order to be hardly noticeable. Modern processors make this possible, even with large main stores. Nevertheless, finding all accessible nodes in an entire computer system within, say, a second appears to be a formidable feat.

We recognize that the mark phase essentially is a tree traversal, or rather a forest traversal. The roots of the trees are all named pointer variables in existence. We shall postpone the question of how these roots are to be found, and 1st present a quick tutorial about tree traversal. In general, nodes of the traversed structure may contain many pointers (branches). We shall, however, 1st restrict our attention to a binary tree, because the essential problem and its solution can be explained better in this way.

The essential problem alluded to is that of storage utilization by the traversal algorithm itself. Typically, information about the nodes already visited must be retained, be it explicitly, or implicitly as in the case of use of recursion. Such a strategy is plainly unacceptable, because the amount of storage needed is unpredictable and may become very large, and because GC is typically initiated just when more storage is unavailable. The task may seem impossible, yet a solution

lies in the idea of inverting pointers along the path traversed, thus keeping the return path open. It is embodied in the following procedure, whose task is to traverse the tree given by the parameter *root*, and to mark every node. Mark values are assumed to be initially 0. Let the data structure be defined by the types

```
Ptr = POINTER TO Node;
Node = RECORD m: INT; L, R: Ptr END;
```

and the algorithm by the procedure

```
PROC traverse(root: Ptr);
  VAR p, q, r: Ptr;
BEGIN p := root; q := root;
  REPEAT (*p # NIL*) INC(p.m); (*mark*)
    IF p.L # NIL THEN (*pointer rotation*)
      r := p.L; p.L := p.R; p.R := q; q := p; p := r
    ELSE
      p.L := p.R; p.R := q; q := NIL
    END
  UNTIL p = q
END traverse
```

We note that only 3 local variables are required, independent of the size of the tree to be traversed. The 3rd, *r*, is in fact merely an auxiliary variable to perform the rotation of values




Figure 28: Rotation of four pointers

$p.L$ ,  $p.R$ ,  $q$ , and  $p$  as shown in Fig. ?? . A snapshot of a tree traversal is shown in Fig. ?? .




Figure 29: Tree traversal (original at left, snapshot at right)

The pair  $p, q$  of pointers marks the position of the process. The algorithm traverses the tree in a left to right, depth 1st fashion. When it returns to the root, all nodes have been marked. How are these claims convincingly supported? The best way is by analyzing the algorithm at an arbitrary node. We start with the hypothesis H that, given the initial state P, the algorithm will reach state Q, (see Fig ??). State Q differs from P by the node and its descendants B and C having been marked, and by an exchange of  $p$  and  $q$ . We now apply the algorithm to state P, assuming that B and C are not empty. The process is illustrated in Fig ?? . Po stands for P in Fig. ?? . Transitions  $P0 \rightarrow P1$ ,  $P2 \rightarrow P3$ , and  $P4 \rightarrow P5$  are the direct




Figure 30: Transition from state P to Q

results of applying the pointer rotation as specified by the sequence of five assignments in the algorithm. Transitions  $P1 \rightarrow P2$  and  $P3 \rightarrow P4$  follow from the hypothesis  $H$  being applied to the states  $P1$  and  $P3$ : subtrees are marked and  $p$ ,  $q$  interchanged. We note in passing that the node is visited 3 times. Progress is recorded by the mark value which is incremented from 0 to 3.

Fig. ?? demonstrates that, if  $H$  holds for steps  $P1 \rightarrow P2$  and  $P3 \rightarrow P4$ , then it also holds for step  $P1 \rightarrow P5$ , which visits the subtree  $p$ . Hence, it also holds for the step  $\text{root} \rightarrow \text{root}$ , which traverses the entire tree.

This proof by recursion relies on the algorithm performing correct transitions also in the case of  $p.L$  being NIL, i.e.  $B$  being the empty tree. In this case, state  $P1$  is skipped; the 1st transition is  $P0 \rightarrow P2$  (see Fig ??). If  $p.L$  is again NIL, i.e. also  $C$  is empty, the next transition is  $P2 \rightarrow P4$ . This concludes the demonstration of the algorithm's correctness.

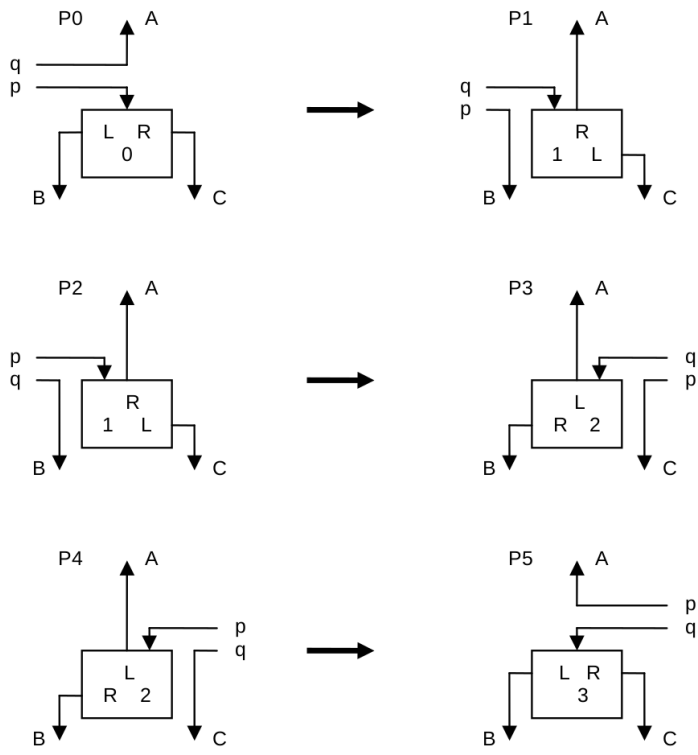


Figure 31: Transitions from  $P_0$  to  $P_5$ , visiting nodes 3 times




Figure 32: Direct transition from  $P_0$  to  $P_2$ , if  $p.L = NIL$

We now modify the algorithm of tree traversal to the case where the structure is not confined to a binary tree, but



may be a tree of any degree, i.e. each node may have any number  $n$  of descendants. For practical purposes, however, we restrict  $n$  to be in the range  $0 \leq n \leq N$ , and therefore may represent all nodes by the type

```
Node = RECORD m, n: INT;
        dsc: ARRAY N OF Node
    END
```

In principle, the binary tree traversal algorithm might be adopted almost without change, merely extending the rotation of pointers from  $p.L, p.R, q, p$  to  $p.dsc[0], \dots, p.dsc[n-1], q, p$ . However, this would be an unnecessarily inefficient solution. The following is a more effective variant. Moreover, it caters for the case of inhomogeneous graphs, where different nodes have different numbers of descendants. The key lies in associating with every node, in addition to the tag, a 2nd private field  $mk$ . It serves 2 purposes. The 1st is as a mark, with  $mk > 0$  indicating that the node had been visited. The 2nd is to store the address of the next descendant to be visited. The underlying data structure is shown in Fig ??.

Type descriptors consist of the following fields:

size    in bytes, of the described record type,  
 base    a table of pointers to the base types descriptors(3 elements only)  
 offsets of the descendant pointers in the described type(1 word each)




Figure 33: Record and its type descriptor

We note that the mark value, starting with zero (unmarked), is used as a counter of descendants already traversed, and hence as an index to the descendant field to be processed next. The algorithm can be applied not only to trees, but to arbitrary structures, including circular ones, if the continuation condition  $p \neq 0$  (actually  $p \geq \text{heapOrg}$ ) is extended to  $(p \geq \text{heapOrg}) \ \& \ (\text{offadr} = 0)$ . This causes a descendant that is already marked to be skipped. Here the array  $M$  stands for the entire memory.

```
PROC traverse(root: Ptr);
  VAR offadr, offset: INT; p, q, r: Ptr;
```

```

BEGIN p := root; q := root;
  REPEAT (*p # NIL*) offadr := p.mk; (*mark*)
    IF offadr = 0 THEN
      tag := p.tg; offadr := tag + 16
    ELSE INC(offadr, 4) END;
    p.mk := offadr; offset := M[offadr];
    IF offset # -1 THEN (*move down*)
      r := M[p+offset]; offadr := M[r-4];
      IF offadr = 0 THEN
        M[p+offset] := q; q := p; p := r
      END
    ELSE (*move up*)
      offadr := M[q-4]; offset := M[offadr]
      IF p # q THEN
        r := M[q+offset]; M[q+offset] := p;
        p := q; q := r
      END
    END
  UNTIL (p = q) & (offset = -1)
END traverse.

```

The mark is included in each record's hidden prefix. The prefix takes 2 words only; the first is used for the tag. The other is reserved for the garbage collector and used as mark and offset address. The end of the list of descendant pointers is marked by an entry with value -1. And finally, assignments involving M are expressed as

assignment	for
SYSTEM.GET(a, x)	$x := M[a]$
SYSTEM.PUT(a, x)	$M[a] := x$

The scan phase is performed by a relatively straight-forward algorithm. The heap, i.e. the storage area between HeapOrg and HeapLimit (the latter is a variable), is scanned element by element, starting at HeapOrg. Elements marked are unmarked, and unmarked elements are freed by linking them into the appropriate list of available space. As the heap may always contain free elements, the scan phase must be able to recognize them in order to skip them or merge them with an adjacent free element. For this purpose, the free elements are also considered as prefixed. The prefix serves to determine the element's size and to recognize it as free due to a special (negative) mark value. The encountered mark values and the action to be taken are:

<i>mk</i> value	state	action
= 0	unmarked	collect, mark free
> 0	marked	unmark
< 0	free	skip or merge

8.3    KERNEL

The kernel lies at the bottom of the module hierarchy. It contains the procedures for dynamic storage allocation and

retrieval as described before. The procedures are New, Mark, and Scan. Kernel also contains the driver routines for the disk. They are used by modules FileDir and Files. The "disk" is actually an SD-card, a high-volume flash-RAM. It is accessed purely sequentially, byte-wise, by a standard, serial peripheral interface (SPI). Within Kernel a table called SectorMap is allocated keeping track of blocks (sectors) occupied by files. A single bit indicates, whether a sector is allocated or not. This table is accessed by the procedures AllocSector, MarkSector, and FreeSector. Reading and writing is done sector-wise by procedures GetSector and PutSector. Sector numbers are always a multiple of 29 for the purpose of redundancy checks.

Furthermore, the kernel contains a timer counting milliseconds and, perhaps, a real time clock, showing date and time. Clock data are packed into a single word as follows:

6	4	5	5	6	6
year	month	day	hour	minute	second

Figure 34: Encoding of date and time (year starting with 2000)

DEFINITION Kernel;(\*NW/PR 11.4.86/27.12.95/15.5.2013\*)

```

CONST SectorLength = 1024;
TYPE Sector = ARRAY SectorLength OF BYTE;
VAR allocated, NofSectors, heapOrg, heapLim,
    stackOrg, MemSize: INT;
```

```
PROC New(VAR ptr: INT; tag: INT);
PROC Mark(pref: INT);
PROC Scan;
PROC ResetDisk;
PROC MarkSector(sec: INT);
PROC FreeSector(sec: INT);
PROC AllocSector(hint: INT; VAR sec: INT);
PROC GetSector(src: INT; VAR dest: Sector);
PROC PutSector(dest: INT; VAR src: Sector);
PROC Time(): INT; (*milliseconds*)
PROC Clock(): INT;
PROC SetClock(dt: INT);
PROC Install(adr, procadr: INT);
PROC Init;
END Kernel.
```

#### 8.4 THE STORAGE MANAGEMENT TOOLBOX

The user can obtain information about the system state and resources through its toolbox, a set of commands contained in the too module System. These commands are:

```
PROC Watch;
PROC Collect; / n
PROC SetClock; / year,month,day,hour,minute,second
```

Command Watch shows the amount of storage occupied in the heap, the number of disk sectors allocated on the

disk, and the number of tasks installed. The command `Collect` allows to control the frequency of GC. The number  $n$  indicates how many commands are executed before the next GC.





## DEVICE DRIVER(DD)

---

### 9.1 OVERVIEW

DDs are collections of procedures that constitute the immediate interface between hardware and software. They refer to those parts of the computer hardware that are usually called *peripheral*. Computers typically contain a system bus which transmits data among its different parts. Processor and memory are considered as its internal parts; the remaining parts, such as disk, keyboard, display, etc, are considered as external or peripheral, notwithstanding the fact that they are often contained in the same cabinet or board.

Such peripheral devices are connected to the system bus via special registers (data buffers) and transceivers (switches, buffers in the sense of digital electronics). These registers and transceivers are addressed by the processor in the same way as memory locations - they are said to be *memory-mapped* - and they constitute the hardware interface between processor bus and device. References to them are typically

confined to specific driver procedures which constitute the software interface.

Drivers are inherently hardware specific, and the justification of their existence is precisely that they encapsulate these specifics and present to their clients an appropriate abstraction of the device. Evidently, this abstraction must still reflect the essential characteristics of the device, but not the details (e.g. the addresses of its interface registers).

Our justification to present the drivers connecting the Oberon with the RISC computer in detail is on the one hand the desire for completeness. But on the other hand it is also in recognition of the fact that their design represents an essential part of the engineering task in building a system. This part may look trivial from a conceptual point of view; it certainly is not so in practice. In order to reduce the number of interface types, standards have been established. The RISC computer also uses such interface standards, and we will concentrate on them in the following presentations.

The following devices are presented:

**KEYBOARD** Considered as a serial device delivering one byte of input data per key stroke. It is connected by a serial line according to the PS/2 and ASCII standards. The software is contained in module Input (9.2), and the hardware is explained in 12.2.1.

**MOUSE** A pointing device delivering coordinates as well as key states. The software is also part of Input (9.2).

**DISPLAY** The interface to it is an area of memory contains the displayed information, exactly one bit per pixel for a monochrome display. This area is called *frame buffer* or *bitmap*. Here the size of the display area is 768 lines and 1024 dots per line, representing a raster. The software is module *Display*, which primarily consists of operations to draw frequently occurring patterns (called *raster-ops*, explained in 4.5). The actual display requires a hardware interface called a *display controller*, the connection between which and the display follows the VGA-Standard (see 12.2.4).

**DISK** Our RISC computer does not use a magnetic, rotating disk to store non-volatile data. Instead, it uses an SD-card (flash-RAM). The driver is contained in module *Kernel* (8.3). The hardware is discussed in 12.2.3.

**NET** In the original text, a network was presented consisting of a bus connecting many computers, based on the RS-485 standard. It was implemented by the serial communications controller Zilog 8530, operating at a frequency of 230 Kb/s. The name SCC has been retained as a generic interface, behind which the packet transport has now been re-implemented as a simple wireless network (Nordic nRF24L01 controller) in the

regulation-free 2.4GHz industrial/scientific/medical (ISM) frequency band.

In all driver modules the implementation-dependent procedures `SYSTEM.PUT`, `SYSTEM.GET`, and `SYSTEM.BIT` are used to access the registers of the device interface. Their 1st parameter is the address of the register, the 2nd an expression or variable.

## 9.2 KEYBOARD AND MOUSE

Their driver procedures are in `Input`. `Available() > 0` signals that a character has been typed on the keyboard. The character is read by calling `Read(c)`. `Input` is restricting the data to the ASCII character set Latin-1, i.e. the values lie in the range  $0X \leq c < 80X$  (7-bit values). `Mouse(keys, x, y)` yields the current keys state and coordinates of the mouse.

```
MODULE Input;
  PROC Available(): INT;
  PROC Read(VAR c: CHAR);
  PROC Mouse(VAR keys: SET; VAR x, y: INT);
  PROC SetMouseLimits(w, h: INT);
  PROC Init;
END Input.
```

The driver software accesses the keyboard via the Standard PS/2 interface represented by an 8-bit register for the re-

ceived data `kbdCode`, and a single-bit flag indicating whether a byte had been received.

The keyboard codes received from the keyboard via a PS/2 line are not identical with the character values delivered to the Read procedure. A conversion is necessary. This is so, because modern keyboards treat all keys in the same way, including the ones for upper case, control, alternative, etc. Separate codes are sent to signal the pushing down and the release of a key, followed by another code identifying which key had been pressed or released. This requires, besides a translation table from codes to characters, a set of state variables. They are the global, boolean variables `Recd`, `Up`, `Shift`, `Ctrl`, and `Ext`. Procedure `Peek` determines whether an actual character is present, or merely a code signalling a key shift. `Peek` controls the state.

Procedure `Mouse` fetches a word from the mouse interface register and decomposes it into its components (keys state and coordinates). (`kb` is the bit indicating whether a code had been received from the keyboard).




Figure 35: Format of the mouse register

### 9.3 THE SD-CARD (DISK)

SD-card are high-volume memory devices based on flash-store technology. They are typically organized as individually accessible blocks of 1K bytes. The driver for the SD-card is contained in module `Kernel`, which also handles allocation and reservation of blocks, here in analogy to rotating disks still called *sectors*.

```

TYPE Sector = ARRAY SectorLength OF BYTE;
PROC GetSector(src: INT; VAR dst: Sector);
PROC PutSector(dst: INT; VAR src: Sector);
PROC AllocSector(hint: INT; VAR sec: INT);
PROC MarkSector(sec: INT);
PROC FreeSector(sec: INT);

```

Data transfer is sequential and handled by procedures `ReadSD` and `WriteSD` by issuing commands. These are for transmitting a block address, for receiving, and for sending a block of data. Synchronous transmission of sequences of words follows the SPI standard, which uses 3 lines, one for data input, one for data output, and one for the clock (see also [12.2.3](#)). The hardware interface contains a 32-bit register. The bit-rate is 8.3 MB/s.

## 9.4 SERIAL ASYNCHRONOUS INTERFACE (RS 232)

The RS-232 standard serves to transmit sequences of bytes over a data line asynchronously. This implies that there is no separate clock line (see also [12.2.5](#)). The hardware interface contains a 10-bit register for the transmitter and one for the receiver. The data rate used here is 19,200 bit/s. A byte is sent and received over the line by the following programs.

```
CONST data = -56;
      stat = -52; (*device register addresses*)
PROC Send(x: BYTE);
BEGIN
    REPEAT UNTIL SYSTEM.BIT(stat, 1);
    SYSTEM.PUT(data, x)
END Send;

PROC Rec(VAR x: BYTE);
BEGIN
    REPEAT UNTIL SYSTEM.BIT(stat, 0);
    SYSTEM.GET(data, x)
END Rec;
```

These procedures are used in the driver module RS232 presented in `??`. This module itself is not used in the Oberon core, but it was instrumental in building the System on a host computer and downloading it. It is characterized by a very simple interface.

## 9.5    SERIAL COMMUNICATIONS CONTROLLER (SCC)

The interface of the driver for the network was taken over from the original design using a serial communications controller Zilog 8530. The implementation changed totally. It was designed by Paul Reed for the wireless controller Nordic nRF24L01.

DEFINITION SCC;

```

    TYPE Header = RECORD valid: BOOL;
                      dadr, sadr, typ: BYTE;
                      len: INT; (*of data following header*)
    END;

    PROC Start(filter: BOOL);
    PROC Send(VAR head: Header; buf: ARRAY OF BYTE);
    PROC Available(): INT;
    PROC ReceiveHead(VAR head: Header);
    PROC Receive(VAR x: BYTE);
    PROC Skip(m: INT);
    PROC Stop;
END SCC.
```



## COMPILER

---

### 10.1 INTRODUCTION

The compiler is the primary tool of the system builder. It therefore plays a prominent role in the Oberon System, although it is not part of the basic system. Instead, it constitutes a tool module - an application - with a single command: `Compile`. It translates program texts into machine code. Therefore, it is as a program inherently machine-dependent; it acts as the interface between source language and target computer.

In order to understand the process of compilation, the reader needs to be familiar with the source language Oberon defined in Appendix 1, and with the target computer RISC, defined in Appendix 2.

The language is defined as an infinite set of sequences of symbols taken from the language's vocabulary. It is described by a set of equations called syntax. Each equation defines a syntactic construct, or more precisely, the set of sequences of symbols belonging to that construct. It specifies

how that construct is composed of other syntactic constructs. The meaning of programs is defined in terms of semantic rules governing each such construct.

Compilation of a program text proceeds by analyzing the text and thereby decomposing it recursively into its constructs according to the syntax. When a construct is identified, code is generated according to the semantic rule associated with the construct. The components of the identified construct supply parameters for the generated code.

It follows that we distinguish between 2 kinds of actions: analyzing steps and code generating steps. In a rough approximation we may say that the former are source language dependent and target computer independent, whereas the latter are source language independent and target computer dependent. Although reality is somewhat more complex, the module structure of this compiler clearly reflects this division. The main module of the compiler is ORP (for Oberon to RISC Parser) It is primarily dedicated to syntactic analysis, parsing. Upon recognition of a syntactic construct, an appropriate procedure is called the code generator module ORG (for Oberon to RISC Generator). Apart from parsing, ORP checks for type consistency of operands, and it computes the attributes of objects identified in declarations.

Whereas ORP mirrors the source language and is independent of a target computer, ORG reflects the target computer, but is independent of the source language.

Oberon program texts are regarded as sequences of symbols rather than sequences of characters. Symbols themselves, however, are sequences of characters. We refrain from explaining the reasons for this distinction, but mention that apart from special characters and pairs such as +, &, <=, also identifiers, numbers, and strings are classified as symbols. Furthermore, certain capital letter sequences are symbols, such as IF, END, etc. Each time the syntax analyzer (parser) proceeds to read the next symbol, it does this by calling procedure Get, which constitutes the so-called scanner residing in module ORS (for Oberon to RISC Scanner). It reads from the source text as many characters as needed to recognize the next symbol.

In passing we note that the scanner alone reflects the definition of symbols in terms of characters, whereas the parser is based on the notion of symbols only. The scanner implements the abstraction of symbols. The recognition of symbols within a character sequence is called *lexical analysis*.

Ideally the recognition of any syntactic construct, say A, consisting of subconstructs, say B1, B2, ..., Bn, leads to the generation of code that depends only on

1. the semantic rules associated with A, and
2. on (attributes of) B1, B2, ..., Bn.

If this condition is satisfied, the construct is said to be *context-free*, and if all constructs of a language are context-free, then also the language is context-free. Syntax and semantics of Oberon adhere to this rule, although with a significant exception. This exception is embodied by the notion of declarations. The declaration of an identifier, say  $x$ , attaches permanent properties to  $x$ , such as the fact that  $x$  denotes a variable and that its type is  $T$ . These properties are "invisible" when parsing a statement containing  $x$ , because the declaration of  $x$  is not also part of the statement. The "meaning" of identifiers is thus inherently *context-dependent*.

Context-dependence due to declarations is the immediate reason for the use of a global data structure which represents the declared identifiers and their properties (attributes). Since this concept stems from early assemblers where identifiers (then called symbols) were registered in a linear table, the term symbol table tends to persist for this structure, although in this compiler it is considerably more complex than an array. Basically, it grows during the processing of declarations, and it is searched while expressions and statements are processed. Procedures for building and for searching are contained in module ORB.

A complication arises from the notion of exports and imports in Oberon. Its consequence is that the declaration of an identifier  $x$  may be in a module, say  $M$ , different from where  $x$  is referenced. If  $x$  is exported, the compiler includes  $x$

together with its attributes in the symbol file of the compiled module *M*. When compiling another module which imports *M*, that symbol file is read and its data are incorporated into the symbol table. Procedures for reading and writing symbol files are contained in module *ORB*, and no other module relies on information about the structure of symbol files.

The syntax is precisely and rigorously defined by a small set of syntactic equations. As a result, the parser is a reasonably perspicuous and short program. In spite of the high degree of regularity of the target computer, the process of code generation is more complicated, as shown by module *ORG*.

The resulting module structure of the compiler is shown in Fig. 36 in a slightly simplified manner. In reality *OCS* is imported by all other modules due to their need for procedure *OCS.Mark*. This, however, will be explained later.

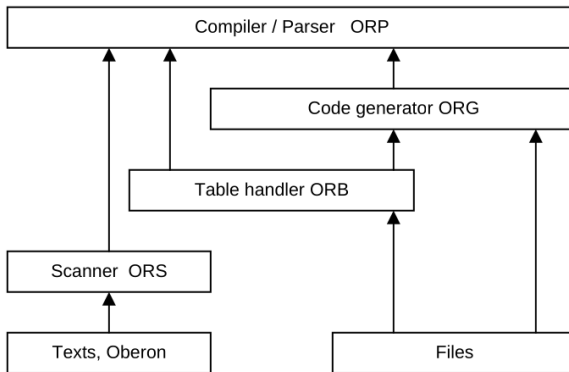


Figure 36: Compiler's module structure

## 10.2 CODE PATTERNS

Before it is possible to understand how code is generated, one needs to know which code is generated. In other words, we need to know the goal before we find the way leading to the goal. A fairly concise description of this goal is possible due to the structure of the language. As explained before, semantics are attached to each individual syntactic construct, independent of its context. Therefore, it suffices to list the expected code - instead of an abstract semantic rule - for each syntactic construct.

As a prerequisite to understanding the resulting instructions and in particular their parameters, we need to know where declared variables are stored, i.e. which are their addresses. This compiler uses the straight-forward scheme of sequential allocation of consecutively declared variables. An address is a pair consisting of a base address (in a register) and an offset. Global variables are allocated in the module's data section and the respective base address register is SB (Static Base, see Chapter 6). Local variables are allocated in a procedure activation record on the stack; the respective base register is SP (Stack Pointer). Offsets are positive integers.

The amount of storage needed for a variable (called its *size*) is determined by the variable's type. The sizes of basic types are prescribed by the target computer's data representation. The following holds for the RISC processor:

Types	Size
BYTE, CHAR, BOOL	1
INT, REAL, SET, POINTER, PROC	4

The size of an array is its element’s multiplied by the number of elements; The size of a record is the sum of all its fields’.

A complication arises due to so-called alignment. By alignment is meant the adjustment of an address to a multiple of the variable’s size. Alignment is performed for variable addresses as well as for record field offsets. The motivation for alignment is the avoidance of double memory references for variables being "distributed" over 2 adjacent words. Proper alignment enhances processing speed quite significantly. Variable allocation using alignment is shown by the example in Fig. 37.




Figure 37: Alignment of variables

We note in passing that a reordering of the four variables lessens the number of unused bytes, as shown in Fig. 38.

Memory instructions compute the address as the sum of a register (base) and an offset constant. Local variables




Figure 38: Improved order of variables

use the stack pointer SP (R14) as base, global variables the static base SB (R13) Every module has its own SB value, and therefore access to global (and imported) variables requires 2 instructions, one for fetching the base value, and one for loading or storing data. If the compiler can determine, whether the correct base value has already been loaded into the SB register, the former instruction is omitted.

The first 7 sample patterns contain global variables only, and their base SB is assumed to hold the appropriate value. Parameters of branch instructions denote jump distances from the instruction’s own location (PC-relative).

10.2.1 *Assignment of constants*

We begin with a simple example of assigning constants to variables. The variables used in this example are global; their base register is SB. Each assignment results in a single instruction. The constant is embedded within the instruction as a literal operand.

```
MODULE Pattern1;
  VAR ch: CHAR;      0
```



```

        k: INT ;      4
        x: REAL;      8
        s: SET ;      12

BEGIN                                     module entry code
    ch := "0";      40000030  MOV  R0 R0 30H
                                B0D00000  STR  R0 SB 0
    k := 10 ;      4000000A  MOV  R0 R0 10
                                A0D00004  STR  R0 SB 4
    x := 1.0;      60003F80  MOV' R0 R0 3F800000H
                                A0D00008  STR  R0 SB 8
    s := {0, 4, 8}  40000111  MOV  R0 R0 111H
                                A0D0000C  STR  R0 SB 12
END Pattern1.                               module exit code

```

### 10.2.2 *Simple expressions*

The result of an expression containing operators is always stored in a register before it is assigned to a variable or used in another operation.

Registers for intermediate results are allocated sequentially in ascending order R0, R1, ..., R11. Integer multiplication and division by powers of 2 are represented by shifts (LSL, ASR). Similarly, the modulus by a power of 2 is obtained by masking off leading bits. The operations of set union, difference, and intersection are represented by logical operations (OR, AND).

```

MODULE Pattern2;
  VAR i, j, k, n: INT ;    0,  4,  8, 12
      x, y: REAL;          16, 20
      s, t, u: SET ;       24, 28, 32

BEGIN i := (i + 1) * (i - 1);  LDR  R0 SB 0
                                ADD  R0 R0 1
                                LDR  R1 SB 0
                                SUB   R1 R1 1
                                MUL   R0 R0 R1
                                STR   R0 SB 0
      k := k DIV 17            ; LDR  R0 SB 8
                                DIV   R0 R0 17
                                STR   R0 SB 8
      k := 8 * n                ; LDR  R0 SB 12
                                LSL   R0 R0 3
                                STR   R0 SB 8
      k := n DIV 2              ; LDR  R0 SB 12
                                ASR   R0 R0 1
                                STR   R0 SB 8
      k := n MOD 16             ; LDR  R0 SB 12
                                AND   R0 R0 15
                                STR   R0 SB 8
      x := -y / (x - 1.0)      ; LDR  R0 SB 16
                                MOV'  R1 R0 3F80H
                                FSB   R0 R0 R1
                                LDR  R1 SB 20
                                FDV   R0 R1 R0

```

	MOV	R1	R0	0
	FSB	R0	R1	R0
	STR	R0	SB	16
s := s + t * u	LDR	R0	SB	28
	LDR	R1	SB	32
	AND	R0	R0	R1
	LDR	R1	SB	24
	OR	R0	R1	R0
	STR	R0	SB	24

END Pattern2.

### 10.2.3 Indexed variables

References to elements of arrays make use of the possibility to add an index value to an offset. The index must be present in a register and be multiplied by the size of the array elements. (For integers with size 4 this is done by a shift of 2 bits). Then this index is checked whether it lies within the bounds specified in the array's declaration. This is achieved by a comparison, actually a subtraction, and a subsequent branch instruction causing a trap, if the index is either negative or beyond the upper bound.

If the reference is to an element of a multi-dimensional array (matrix), its address computation involves several multiplications and additions. The address of an element  $A[i_{k-1}, \dots, i_1, i_0]$  of a  $k$ -dimensional array  $A$  with lengths  $n_{k-1}, \dots, n_1, n_0$  is

$$\text{adr}(A) + ((\dots((i_{k-1} * n_{k-2}) + i_{k-2}) * n_{k-3} + \dots) * n_1 + i_1) * n_0 + i_0$$

Note that for index checks CMP is written instead of SUB to mark that the subtraction is merely a comparison, that the result remains unused and only the condition flag registers hold the result.

```

MODULE Pattern3;
  VAR i, j, k, n:      INT;    0, 4, 8, 12
      a: ARRAY 10      OF INT;    16
      x: ARRAY 10, 10  OF INT;    56
      y: ARRAY 10, 10, 10 OF INT;  456

BEGIN
    k := a[i];      LDR R0 SB 0
                    CMP R1 R0 10
                    BLHI R12
                    LSL R0 R0 2
                    ADD R0 SB R0
                    LDR R0 R0 16
                    STR R0 SB 8
                    n := a[5];  LDR R0 SB 36
                    STR R0 SB 12
    x[i, j] := 2 ;   LDR R0 SB 0
                    CMP R1 R0 10
                    BLHI R12
                    MUL R0 R0 40

```

```

                                ADD R0 SB R0
                                LDR R1 SB 4
                                CMP R2 R1 10
                                BLHI R12
                                LSL R1 R1 2
                                ADD R0 R0 R1
                                MOV R1 R0 2
                                STR R1 R0 56
y[i, j, k] := 3 ;             LDR R0 SB 0
                                CMP R1 R0 10
                                BLHI R12
                                MUL R0 R0 400
                                ADD R0 SB R0
                                LDR R1 SB 4
                                CMP R2 R1 10
                                BLHI R12
                                MUL R1 R1 40
                                ADD R0 R0 R1
                                LDR R1 SB 8
                                CMP R2 R1 10
                                BLHI R12
                                LSL R1 R1 2
                                ADD R0 R0 R1
                                MOV R1 R0 3
                                STR R1 R0 456
                                MOV R0 R0 6
                                STR R0 SB 1836
y[3, 4, 5] := 6
END Pattern3.

```

10.2.4 *RECORD fields and pointers*

Fields of records are accessed by computing the sum of the record's (base) address and the field's offset. If the record variable is statically declared, the sum is computed by the compiler.

```

MODULE Pattern4;
  TYPE Ptr = POINTER TO Node;
    Node = RECORD num: INT;          0
               name: ARRAY 8 OF CHAR; 4
               next: Ptr              12
    END ;
  VAR p, q: Ptr ;                    12, 16
      r: Node;                        20
BEGIN
  r.num      := 10 ;  MOV R0 R0 10
                               STR R0 SB 20
  p.num      := 6 ;   LDR R0 SB 12 (p)
                               MOV R1 R0 6
                               STR R1 R0 0
  p.name[7]  := "0";  LDR R0 SB 12
                               MOV R1 R0 30H
                               STR R1 R0 11 (4+7)
  p.next     := q ;   LDR R0 SB 12
                               LDR R1 SB 16
                               STR R1 R0 12
  p.next.next := NIL  LDR R0 SB 12 (p)

```

```

LDR R0 R0 12 (p.next)
MOV R1 R0 0 (NIL)
STR R1 R0 12 (p.next.next)

END Pattern4.

```

### 10.2.5 *Boolean expressions, IF statements*

Conditional statements imply that parts of them are skipped. This is done by the use of branch instructions whose operand specifies the distance of the branch. The instructions refer to the condition-register as an implicit operand. Its value is determined by a preceding instruction, typically a compare or a bit-test instruction.

The Boolean operators & and OR are purposely not defined as total functions, but rather by the equations

```

p & q = if p then q else FALSE
p OR q = if p then TRUE else q

```

Consequently, Boolean operators must be translated into branches too. Evidently, branches stemming from if statements and branches stemming from Boolean operators should be merged, if possible. The resulting code therefore does not necessarily mirror the structure of the if statement directly, as can be seen from the code in Pattern5. We must conclude that code generation for Boolean expressions differs in some aspects from that for arithmetic expressions.





IF ODD(n) OR (n IN s) THEN	LDR SB R0 ...
	LDR R0 SB 0 (n)
	AND R0 R0 1
	BNE 5
	LDR R0 SB 4 (s)
	LDR R1 SB 0
	ADD R1 R1 1
	ROR R0 R0 R1
	BPL 2
n := -1000	MOV R0 R0 -1000
	STR R0 SB 0
END ;	
IF n < 0 THEN	LDR SB R0 ...
	LDR R0 SB 0
	CMP R0 R0 0
	BGE 3
s := {}	MOV R0 R0 0 {}
	STR R0 SB 4
	B 17
ELSIF n < 10 THEN	LDR SB R0 ...
	LDR R0 SB 0
	CMP R0 R0 10
	BGE 3
s := {0}	MOV R0 R0 1
	STR
	B 10
ELSIF n < 100 THEN	LDR SB R0 ...
	LDR R0 SB 0

```

                                CMP R0 R0 100
                                BGE 3
                                MOV R0 R0 2
                                STR R0 SB 4
                                B 3

ELSE
    s := {2}
                                MOV R0 R0 4
                                LDR SB R0 ...
                                STR R0 SB 4

END
END Pattern5.

```

#### 10.2.6 *WHILE and REPEAT statements*

```

MODULE Pattern6;
  VAR i: INT;
  BEGIN i := 0;
                                MOV R0 R0 0
                                STR R0 SB 0

    WHILE i < 10 DO
                                LDR SB R0 ...
                                LDR R0 SB 0
                                CMP R0 R0 10
                                BGE 4

                                i := i + 2
                                LDR R0 SB 0
                                ADD R0 R0 2
                                STR R0 SB 0

    END ;
                                B -8
    REPEAT i := i - 1
                                LDR SB R0 ...
                                LDR R0 SB 0

```

```

                                SUB R0 R0 1
                                STR R0 SB 0
UNTIL i = 0                    LDR R0 SB 0
                                CMP R0 R0 0
                                BNE -7

END Pattern6.

```

### 10.2.7 *FOR statements*

```

MODULE Pattern7;
  VAR i, m, n: INT;
BEGIN
  FOR i := 0 TO n-1 DO
    MOV R0 R0 0
    LDR R1 SB 8
    SUB R1 R1 1
    CMP LNK R0 R1
    BGT 7
    STR R0 SB 0
    LDR R0 SB 4
    LSL R0 R0 1
    STR R0 SB 4
    LDR R0 SB 0
    ADD R0 R0 1
    B -11
  END
END Pattern7.

```

### 10.2.8 *Proper procedures*

Procedure bodies are surrounded by a prolog (entry code) and an epilog (exit code). They reposition the stack pointer SP (see Chapter 6), which holds the address of the procedure activation record on the stack. The immediate value of the first instruction indicates the space taken by variables local to the procedure, rounded up to the next multiple of 4.

Procedure calls use a branch and link (BL) instruction. Parameters are loaded into registers prior to the call and pushed on the stack after the call. Every parameter occupies a multiple of 4 bytes. In the case of value parameters the value is loaded, and in the case of VAR-parameters, the variable's address is loaded.

```

MODULE Pattern8;
  VAR i: INT;
  PROC P(x: INT; VAR y: INT);
    VAR z: INT;
  BEGIN
    SUB SP   SP 16   adjust SP
    STR LNK SP  0   push ret adr
    STR R0   SP  4   push x
    STR R1   SP  8   push @y
    z := x;        LDR R0   SP  4   x
                    STR R0   SP 12   z
    y := z         LDR R0   SP 12   z
                    LDR R1   SP  8   @y
  
```

```

                                STR R0  R1  0   y
END P;                          LDR LNK SP  0   pop ret adr
                                ADD SP  SP 16
                                B    R15

BEGIN P(5, i)                   MOV R0  R0  5
                                ADD R1  SB  0   @i
                                BL  -14          call

END Pattern8.

```

### 10.2.9 *Functions*

They are handled in exactly the same manner as proper procedures, except that a result is returned in register R0. If the function is called in an expression at a place where intermediate results are held in registers, these values are put onto the stack before the call, and they are restored after return (not shown here).

```

MODULE Pattern9;
  VAR x: REAL;
  PROC F(x: REAL): REAL;
  BEGIN
                                SUB  SP SP 8
                                STR LNK SP 0   push ret adr
                                STR  R0 SP 4   push x
                                IF x >= 1.0 THEN
                                    LDR  R0 SP 4
                                    MOV' R1 R0 3F80H
                                    FSB  R0 R0 R1

```

```

                                BLT   4
                                LDR   R0 SP 4
                                BL    -9
                                BL    -10
                                STR   R0 SP 4

                                END ;
                                RETURN x
                                LDR   R0 SP 4
                                LDR   LNK SP 0      pop ret adr
                                ADD   SP SP 8
                                B     R15

                                END Pattern9.

```

#### 10.2.10 *Dynamic ARRAYS*

Dynamic array parameters are passed by loading a descriptor on the stack, regardless of whether they are value- or VAR- parameters. The descriptor consists of the actual variable's address and the array's length. (Only 1-dimensional dynamic arrays are handled).

Elements of dynamic arrays are accessed like those of static arrays. However, even when the index is a constant, the check cannot be performed by the compiler.

```

MODULE Pattern10;
  VAR a: ARRAY 12 OF INT;
  PROC P(x: ARRAY OF INT);
    VAR i, n: INT;

```

```

BEGIN                                SUB SP SP 20
                                     STR LNK SP 0
                                     STR R0 SP 4          x
                                     STR R1 SP 8 x.len
n := x[i];                          LDR R0 SP 12          i
                                     LDR R1 SP 8          x.len
                                     CMP R2 R0 R1
                                     BLHI R12
                                     LSL R0 R0 2
                                     LDR R1 SP 4          x
                                     ADD R0 R1 R0
                                     LDR R0 R0 0
                                     STR R0 SP 16
x[i+1] := n+5                       LDR R0 SP 12          i
                                     ADD R0 R0 1
                                     LDR R1 SP 8          x.len
                                     CMP R2 R0 R1
                                     BLHI R12
                                     LSL R0 R0 2
                                     LDR R1 SP 4          x
                                     ADD R0 R1 R0
                                     LDR R1 SP 16          n
                                     ADD R1 R1 5
                                     STR R1 R0 0
END P;                              LDR LNK SP 0
                                     ADD SP SP 20
                                     B R15

```

```

BEGIN P(a);          ADD R0 SB 0      a
                     MOV R1 R0 12     a.len
                     BL -29

END Pattern10.

```

#### 10.2.11 SETs

This code pattern exhibits the construction of sets. If the specified elements are constants, the set value is computed by the compiler. Otherwise, sequences of move and shift instructions are used. Since shift instructions do not check whether the shift count is within sensible bounds, the results are unpredictable, if elements outside the range 0 .. 31 are involved.

```

MODULE Pattern11;
  VAR s: SET; m, n: INT;
BEGIN
  s := {m};          LDR R0 SB 4      m
                     MOV R1 R0 1
                     LSL R0 R1 R0
                     STR R0 SB 0      s

  s := {0 .. n};     LDR R0 SB 8      n
                     MOV R1 R0 -2
                     LSL R0 R1 R0
                     XOR R0 R0 -1
                     STR R0 SB 0

  s := {m .. 31};    LDR R0 SB 4      m

```



```

                                MOV R1 R0 31
                                MOV R2 R0 -2
                                LSL R1 R2 R1
                                MOV R2 R0 -1
                                LSL R0 R2 R0
                                XOR R0 R0 R1
                                STR R0 SB 0      s
s := {m .. n};                LDR R0 SB 4      m
                                LDR R1 SB 8      n
                                MOV R2 R0 -2
                                LSL R1 R2 R1
                                MOV R2 R0 -1
                                LSL R0 R2 R0
                                XOR R0 R0 R1
                                STR R0 SB 0      s
IF n IN {2,3,5,7,11,13} THEN  MOV R0 R0 28ACH
                                LDR R1 SB 8
                                ADD R1 R1 1
                                ASR' R0 R0 R1
                                BPL 2
                                m := 1
                                MOV R0 R0 1
                                STR R0 SB 4      m
                                END
END Pattern11.

```

### 10.2.12 *Imported variables and procedures*

When a procedure is imported from another module, its address is unavailable to the compiler. Instead, the procedure is identified by a number obtained from the imported module's symbol file. In place of the offset, the branch instruction holds

1. the number of the imported module,
2. the number of the imported procedure, and
3. a link in the list of BL instructions calling an external procedure.

This list is traversed by the linking loader, that computes the actual offset (fixup, see Chapter 6).

Imported variables are also referenced by a variable's number. In general, an access required 2 instructions. The first loads the static base register SB from a global table with the address of that module's data section. The module number of the imported variable serves as index. The second instruction loads the address of the variable, using the actual offset fixed up by the loader.

In the following example, modules P12a and P12b both export a procedure and a variable. They are referenced from the importing module Pattern12.

```
MODULE P12a;
  VAR k*: INT;
```

```
  PROC P*;
  BEGIN k := 1
  END P;
END P12a.
```

```
MODULE P12b;
  VAR x*: REAL;
```

```
  PROC Q*;
  BEGIN x := 1
  END Q;
END P12b.
```

```
MODULE Pattern12;
  IMPORT P12a, P12b;
  VAR i: INT; y: REAL;
BEGIN
```

i := P12a.k;	8D10xxxx	LDR SB 1 link	P12a
	80D00000	LDR R0 SB 0	P12a.k
	8D00xxxx	LDR SB 0 link	Pattern12
	A0D00000	STR R0 SB 0	Pattern12.i
y := P12b.x;	8D20xxxx	LDR SB 2 link	P12b
	80D00000	LDR R0 SB 0	P12b.x
	8D00xxxx	LDR SB 0 link	Pattern12
	A0D00004	STR R0 SB 4	Pattern12.y

END Pattern12.

#### 10.2.13 *RECORD extensions with pointers*

Fields of a record type  $R1$ , which is declared as an extension of a type  $R0$ , are simply appended to the fields of  $R0$ , i.e. their offsets are greater than those of the fields of  $R0$ . When a record is statically declared, its type is known by the compiler. If the record is referenced via a pointer, however, this is not the case. A pointer bound to a base type  $R0$  may well refer to a record of an extension  $R1$  of  $R0$ . Type tests (and type guards) allow to test for the actual type. This requires that a type can be identified at the time of program execution. Because the language defines name equivalence instead of structural equivalence of types, a type may be identified by a number. We use the address of a unique type descriptor for this purpose.

Therefore, type tests consist of a simple address comparison which is very fast. Type descriptors are stored in the module's area for data. Their address is called type tag. The tag of a (dynamically allocated) variable is stored as a prefix to its record (with offset -8).

A type descriptor contains - in addition to information stored for use by the garbage collector - a table of tags of all its base types. If, for instance, a type  $R2$  is an extension of  $R1$  which is an extension of  $R0$ , the descriptor of  $R2$

contains the tags of R1 and R0 as shown in Fig. 39. The table has a fixed number of 3 entries.




Figure 39: Type descriptors

A type test of the form  $p \text{ IS } T$  then, consists of a comparison of the type tag of  $p^\wedge$  at address  $p-8$  with the tag held in the descriptor of  $T$  at the extension level of the type of  $p^\wedge$ . A type guard  $p(T)$  is synonymous to the statement

```
IF ~(p IS T) THEN abort END
```

The following example features 3 record types with associated pointer types, and hence also 3 type descriptors. Each descriptor is 5 words long. Their addresses, and therefore their tags, are 0, 20, and 40 respectively.

```
0 00000020 FFFFFFFF FFFFFFFF FFFFFFFF FFFFFFFF
20 00000020 00014006 FFFFFFFF FFFFFFFF FFFFFFFF
40 00000020 00014005 00028001 FFFFFFFF FFFFFFFF

MODULE Pattern13;
```

## TYPE

```

P0 = POINTER TO R0;
P1 = POINTER TO R1;
P2 = POINTER TO R2;
R0 = RECORD x: INT END ;
R1 = RECORD (R0) y: INT END ;
R2 = RECORD (R1) z: INT END ;

```

## VAR

```

p0: P0;           60
p1: P1;           64
p2: P2;           68

```

## BEGIN

```

p0.x := 0;          LDR R0 SB 60
                    MOV R1 R0 0    p0.x
                    STR R1 R0 0    no type check

p1.y := 1;          LDR R0 SB 64
                    MOV R1 R0 1
                    STR R1 R0 4    p1.y

p0(P1).y := 3;      LDR R0 SB 60    p0
                    LDR R1 R0 -8    tag(p0)
                    LDR R1 R1 4
                    ADD R2 SB 20    TD P1
                    CMP R3 R2 R1
                    BLNE R12
                    MOV R1 R0 3
                    STR R1 R0 4    p0.z

p0(P2).z := 5;      LDR R0 SB 60    p0
                    LDR R1 R0 -8    tag(p0)

```

```

                                LDR R1 R1 8
                                ADD R2 SB 40    TD P2
                                CMP R3 R2 R1
                                BLNE R12
                                MOV R1 R0 5
                                STR R1 R0 8      p0.z
IF p1 IS P2 THEN              LDR R0 SB 64      p1
                                LDR R1 R0 -8     tag(p1)
                                LDR R1 R1 8
                                ADD R2 SB 40    TD P2
                                CMP R3 R2 R1
                                BNE 2
                                p0 := p2
                                LDR R0 SB 68
                                STR R0 SB 60
END
END Pattern13.

```

#### 10.2.14 *RECORD extensions as VAR parameters*

Records occurring as VAR-parameters may also require a type test at program execution time. This is because VAR-parameters effectively constitute hidden pointers. Type tests and type guards on VAR-parameters are handled in the same way as for variables referenced via pointers, with a slight difference, however. Statically declared record variables may be used as actual parameters, and they are not prefixed by a type tag. Therefore, the tag has to be supplied together with the variable's address when the procedure is called, i.e.

when the actual parameter is established. Record structured VAR-parameters therefore consist of address and type tag. This is similar to dynamic array descriptors consisting of address and length.

```
0 00000020 FFFFFFFF FFFFFFFF FFFFFFFF FFFFFFFF
20 00000020 00014006 FFFFFFFF FFFFFFFF FFFFFFFF
```

```
MODULE Pattern14;
```

```
  TYPE
```

```
    R0 = RECORD a, b, c: INT END ;
```

```
    R1 = RECORD (R0) d, e: INT END ;
```

```
  VAR
```

```
    r0: R0;                                40
```

```
    r1: R1;                                52
```

```
  PROC P(VAR r: R0);
```

```
  BEGIN
```

```
    ...
```

```
    r.a := 1;          LDR R1 SP 4    r
```

```
                      STR R0 R1 0    r.a
```

```
    r(R1).d := 2      LDR R0 SP 8    tag(r)
```

```
                      LDR R0 R0 4
```

```
                      ADD R1 SB 20   R1
```

```
                      CMP R2 R1 R0
```

```
                      BLNE R12
```

```
                      MOV R0 R0 2
```

```
                      LDR R1 SP 4    r
```

```
                      STR R0 R1 12   r.d
```



```

END P;          ...

BEGIN          ...
  P(r0);        ADD R0 SB 40    r0
                ADD R1 SB 0     tag(R0)
                BL P
  P(r1)         ADD R0 SB 52    r1
                ADD R1 SB 20    tag(R1)
                BL P
END Pattern14.  ...

```

### 10.2.15 *ARRAY assignments and strings*

```

MODULE Pattern15;
  VAR s0, s1: ARRAY 32 OF CHAR;

  PROC P(x: ARRAY OF CHAR);
  END P;
BEGIN
  s0 := "ABCDEF";  ADD R0 SB 0    @s0
                   ADD R1 SB 64   @"ABCDEF"
                   LDR R2 R1 0
                   ADD R1 R1 4
                   STR R2 R0 0
                   ADD R0 R0 4
                   ASR R2 R2 24    test for 0X
                   BNE -6
  s0 := s1;        ADD R0 SB 0    @s0

```

```

                                ADD R1 SB 32   @s1
                                MOV R2 R0 8    len
                                LDR R3 R1 0
                                ADD R1 R1 4
                                STR R3 R0 0
                                ADD R0 R0 4
                                SUB R2 R2 1
                                BNE -6
P(s1);                          ADD R0 SB 32   @s1
                                MOV R1 R0 32    len
                                BL -38          P
P("012345");                   ADD R0 SB 72    @"012345"
                                MOV R1 R0 7     len (incl 0X)
                                BL -42          P
P("%")                          ADD R0 SB 80    @"%"
                                MOV R1 R0 2     len
                                BL -46          P

END Pattern15.

```

#### 10.2.16 *Predeclared procedures*

```

MODULE Pattern16;
  VAR m, n: INT; x: REAL; u: SET;
      a, b: ARRAY 10 OF INT;
      s, t: ARRAY 16 OF CHAR;
BEGIN
  INC(m);                      ADD R0 SB 0      @m
                                LDR R1 R0 0

```

	ADD R1 R1 1	
	STR R1 R0 0	
DEC(n, 10);	ADD R0 SB 4	@n
	LDR R1 R0 0	
	SUB R1 R1 10	
	STR R1 R0 0	
INCL(u, 3);	ADD R0 SB 12	@u
	LDR R1 R0 0	
	OR R1 R1 8	{3}
	STR R1 R0 0	
EXCL(u, 7);	ADD R0 SB 12	@u
	LDR R1 R0 0	
	AND R1 R1 -129	-{7}
	STR R1 R0 0	
ASSERT(m < n);	LDR R0 SB 0	
	LDR R1 SB 4	
	CMP R0 R0 R1	
	BLGE R12	
UNPK(x, n);	LDR R0 SB 8	x
	ASR R1 R0 23	
	SUB R1 R1 127	
	STR R1 SB 4	n
	LSL R1 R1 23	
	SUB R0 R0 R1	
	STR R0 SB 8	x
PACK(x, n);	LDR R0 SB 8	x
	LDR R1 SB 4	n
	LSL R1 R1 23	

	ADD R0 R0 R1	
	STR R0 SB 8	x
s := "0123456789";	ADD R0 SB 96	@s
	ADD R1 SB 128	adr of string
	LDB R2 R1 0	loop
	ADD R1 R1 4	
	STB R2 R0 0	
	ADD R0 R0 4	
	ASR R2 R2 24	
	BNE -6	
IF s < t THEN	ADD R0 SB 96	@s
	ADD R1 SB 112	@t
	LDB R2 R0 0	loop
	ADD R0 R0 1	
	LDB R3 R1 0	
	ADD R1 R1 1	
	CMP R4 R2 R3	
	BNE 2	
	CMP R4 R2 0	
	BNE -8	
	BGE 3	
m := 1	MOV R0 R0 1	
	STR R0 SB 0	m
END		
END Pattern16.		

10.2.17 *Predeclared functions*

```

MODULE Pattern17;
  VAR m, n: INT;
      x, y: REAL;
      b: BOOL; ch: CHAR;
BEGIN
  n := ABS(m);      LDR R0 SB 0          m
                    CMP R0 R0 0
                    BGE 2
                    MOV R1 R0 0
                    SUB R0 R1 R0
                    STR R0 SB 4          n
  y := ABS(x);      LDR R0 SB 8          x
                    LSL R0 R0 1
                    ROR R0 R0 1
                    STR R0 SB 12        y
  b := ODD(n);      LDR R0 SB 4          n
                    AND R0 R0 1
                    BEQ 2
                    MOV R0 R0 1
                    B 1
                    MOV R0 R0 0
                    STB R0 SB 16        b
  n := ORD(ch);     LDB R0 SB 17        ch
                    STR R0 SB 4          n
  n := FLOOR(x);    LDR R0 SB 8          x
                    MOV' R1 R0 4B00H

```

```

                                FAD" R0 R0 R1      floor
                                STR R0 SB 4        n
    y := FLT(m);               LDR R0 SB 0        m
                                MOV' R1 R0 4B00H
                                FAD' R0 R0 R1      float
                                STR R0 SB 12       y
    n := LSL(m, 3);           LDR R0 SB 0        m
                                LSL R0 R0 3
                                STR R0 SB 4        n
    n := ASR(m, 8);           LDR R0 SB 0
                                ASR R0 R0 8
                                STR R0 SB 4
    m := ROR(m, n);           LDR R0 SB 0
                                LDR R1 SB 4
                                ROR R0 R0 R1
                                STR R0 SB 0
END Pattern17.

```

### 10.3 INTERNAL DATA STRUCTURES AND MODULE INTERFACES

#### 10.3.1 *Data structures*

In §10.1 it was explained that declarations inherently constitute context-dependence of the translation process. Although parsing still proceeds on the basis of a context-free syntax and relies on contextual information only in a few isolated instances, information provided by declarations af-

fects the generated code significantly. During the processing of declarations, their information is transferred into the "*symbol table*", a data structure of considerable complexity, from where it is retrieved for the generation of code.

This dynamic data structure is defined in module ORB in terms of 2 record types called *Object* and *Struct*. These types pervade all other modules with the exception of the scanner. They are therefore explained before further details of the compiler are discussed (see module ORB below).

For each declared identifier an instance of type *Object* is generated. The record holds the identifier and the properties associated with the identifier given in its declaration. Since Oberon is a statically typed language, every object has a type. It is represented in the record by its *typ* field, which is a pointer to a record of type *Struct*. Since many objects may be of the same type, it is appropriate to record the type's attributes only once and to refer to them via a pointer. The properties of type *Struct* will be discussed below.

The kind of object which a table entry represents is indicated by the field *class*. Its values are denoted by declared integer constants: *Var* indicates that the entry describes a variable, *Con* a constant, *Fld* a record field, *Par* a VAR-parameter, and *Proc* a procedure. Different kinds of entries carry different attributes. A variable or a parameter carries an address, a constant has a value, a record field has an offset, and a procedure has an entry address, a list of parameters, and a

result type. For each class the introduction of an extended record type would seem advisable. This was not done, however, for 3 reasons:

- 1<sup>st</sup> , the compiler was first formulated in (a subset of) Modula-2 which does not feature type extension.
- 2<sup>nd</sup> , not making use of type extensions would make it simpler to translate the compiler into other languages for porting the language to other computers. And
- 3<sup>rd</sup> , all extensions were known at the time the compiler was planned.

Hence extensibility provided no argument for the introduction of a considerable variety of types. The simplest solution lies in using the multi-purpose fields `val` and `dsc` for class-specific attributes. For example, `val` holds

- an address for variables, parameters, and procedures,
- an offset for record fields, and
- a value for constants.

The definition of a type yields a record of type `Struct`, regardless of whether it occurs within a type declaration, in which case also a record of type `Object` (`class=Typ`) is generated, or in a variable declaration, in which case the type remains anonymous. All types are characterized by a form



and a size. A type is either a basic type or a constructed type. In the latter case it refers to one or more other types. Constructed types are arrays, records, pointers, and procedural types. The attribute form refers to this classification. Its value is an integer.

Just as different object classes are characterized by different attributes, different forms have different attributes. Again, the introduction of extensions of type `Struct` was avoided. Instead, some of the fields of type `Struct` remain unused in some cases, such as for basic types, and others are used for form-specific attributes. For example, the attribute `base` refers to the element type in the case of an array, to the result type in the case of a procedural type, to the type to which a pointer is bound, or to the base type of a (extended) record type. The attribute `dsc` refers to the parameter list in the case of a procedural type, or to the list of fields in the case of a record type.

As an example, consider the following declarations:

```
CONST N = 100;
TYPE  Ptr = POINTER TO Rec;
      Rec = RECORD n: INTEGER; p, q: Ptr END ;
VAR   k: INTEGER;
      a: ARRAY N OF INTEGER;
PROC  P(x: INTEGER): INTEGER;
```

The corresponding data structure is shown in Fig. 40. For details, the reader is referred to the program listing of module ORB and the respective explanations.

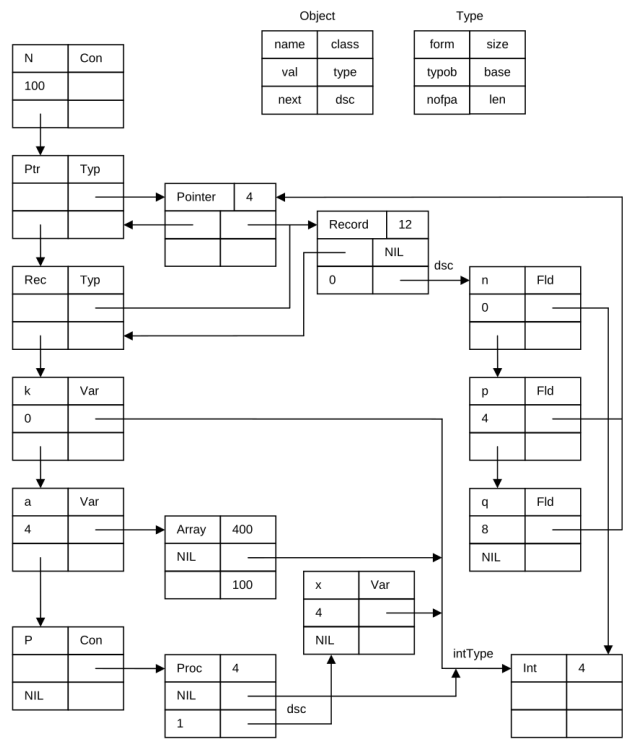


Figure 40: Representation of declarations

Only entries representing constructed types are generated during compilation. An entry for each basic type is established by the compiler's initialization. It consists of an Object holding the standard type's identifier and a Struct indicating its form, denoted by one of the values Byte, Bool,

Char, Int, Real, or Set. The object records of the basic types are anchored in global pointer variables in module ORB (which actually should be regarded as constants).

Not only are entries created upon initialization for basic types, but also for all standard procedures. Therefore, every compilation starts out with a symbol table reflecting all standard, pervasive identifiers and the objects they stand for.

We now return to the subject of Objects. Whereas objects of basic classes (Const, Var, Par, Fld, Typ, SProc, SFunc and Mod) directly reflect declared identifiers and constitute the context in which statements and expressions are compiled, compilations of expressions typically generate anonymous entities of additional, non-basic modes. Such entities reflect selectors, factors, terms, etc., i.e. constituents of expressions and statements. As such, they are of a transitory nature and hence are not represented by records allocated on the heap. Instead, they are represented by record variables local to the processing procedures and are therefore allocated on the stack. Their type is called *Item* and is a slight variation of the type *Object*. Items are not referenced via pointers.

Let us assume, for instance, that a term  $x*y$  is parsed. This implies that the operator and both factors have been parsed already. The factors  $x$  and  $y$  are represented by 2 variables of type *Item* of *Var* mode. The resulting term is again described by an item, and since the product is transitory,

i.e. has significance only within the expression of which the term is a constituent, it is to be held in a temporary location, in a register. In order to express that an item is located in a register, a new, non-basic mode `Reg` is introduced.

Effectively, all non-basic modes reflect the target computer's architecture, in particular its addressing modes. The more addressing modes a computer offers, the more item modes are needed to represent them. The additional item modes required by the RISC processor are. They are declared in module `ORG`:

<code>Reg</code>	direct register mode
<code>RegI</code>	indirect register mode
<code>Cond</code>	condition code mode

The use of the types `Object`, `Item`, and `Struct` for the various modes and forms, and the meaning of their attributes are explained in the following tables: Items have an attribute called `lev` which is part of the address of the item. Positive values denote the level of nesting of the procedure in which the item is declared; `lev = 0` implies a global object. Negative values indicate that the object is imported from the module with number `-lev`. The 3 types `Object`, `Item`, and `Struct` are defined in module `ORB`, which also contains procedures for accessing the symbol table.

Objects			Items		
class	val		a	b	r
0	Undef				
1	Const	val	val		
2	Var	adr	adr	base	
3	Par	adr	adr	off	
4	Fld	off	off		
5	Typ	TDadr	TDadr	modno	
6	SProc	num			
7	SFunc	num			
8	Mod				
10	Reg				regno
11	RegI		off		regno
12	Cond		Tjmp	Fjmp	condition code

Structures					
	form	nofpar	len	dsc	base
7	Pointer				base type
10	ProcTyp	nofpar		param	result type
12	Array		nofel		element type
13	Record	ext lev	desc adr	fields	extension type

10.3.2 Module interfaces

Before embarking on a presentation of the compiler’s main module, the parser, an overview of its remaining modules is

given in the form of their interfaces. The reader is invited to refer to them when studying the parser.

The interface of the scanner module ORS is simple. It defines the numeric values of all symbols. But its chief constituent is procedure `Get`. Each call yields the next symbol from the source text, identified by an integer. Global variables represent attributes of the read symbol in certain cases. If a number was read, `ival` or `rval` hold its numeric value. If an identifier or a string was read, `str` holds the ASCII values of the characters read.

Procedure `Mark` serves to generate a diagnostic output indicating a brief diagnostic and the scanner's current position in the source text. This procedure is located in the scanner, because only the scanner has access to its current position. `Mark` is called from all other modules.

```
DEFINITION ORS; (*Scanner*)  
  IMPORT Texts, Oberon;  
  
  TYPE Ident = ARRAY 32 OF CHAR;  
  VAR ival, slen: INT;  
       rval: REAL;  
       id: Ident;  
       str: ARRAY 256 OF CHAR;  
       errcnt: BOOL;
```

```
PROC Mark (msg: ARRAY OF CHAR);  
PROC Get (VAR sym: INT);  
PROC Init (source: Texts.Text; pos: INT);  
END ORS.
```

Module ORB defines the basic data structures representing declared objects and their types. It also contains procedures for accessing these structures. `NewObj` serves to insert a new identifier, and it returns a pointer to the allocated object. `ThisObj` returns the pointer to the object whose name equals the global scanner variable `ORS.id`. `Thisimport` and `thisfield` deliver imported objects and record fields with names equal to `ORS.id`.

Procedure `Import` serves to read the specified symbol file and to enter its identifier in the symbol table (`class=Mod`). Finally, `Export` generates the symbol file of the compiled module, containing descriptions of all objects and structures marked for export.

```
DEFINITION ORB; (*Base table handler*)  
TYPE  
  Object = POINTER TO ObjDesc;  
  Type = POINTER TO TypeDesc;  
  ObjDesc = RECORD  
    class, lev, expn: INT;  
    expo, rdo: BOOL;  
    next, dsc: Object;
```

```

        type: Type;
        name: ORS.Ident;
        val: INT
    END ;
    TypeDesc = RECORD
        form, ref, mno: INT; (*ref for import/export only*)
        nofpar: INT; (*for records: extension level*)
        len: INT; (*for records: address of descriptor*)
        dsc, typobj: Object;
        base: Type;
        size: INT
    END ;
    VAR topScope: Object;
        byteType, boolType, charType, intType, realType,
        setType, nilType, noType, strType: Type;

    PROC Init;
    PROC Close;
    PROC NewObj (VAR obj: Object; id: ORS.Ident; class: INT
    PROC thisObj (): Object;
    PROC thisimport (mod: Object): Object;
    PROC thisfield (rec: Type): Object;
    PROC OpenScope;
    PROC CloseScope;
    PROC Import (VAR modid, modid1: ORS.Ident);
    PROC Export (VAR modid: ORS.Ident;
        VAR newSF: BOOL; VAR key: INT);
    END ORB.

```



Module ORG contains the procedures for code generation. The names of these procedures indicate the respective constructs for which code is to be produced. Note that an individual code generator procedure is provided for every standard, predefined procedure. This is necessary, because they generate in-line code.

DEFINITION ORG;

```
CONST WordSize* = 4;
```

```
TYPE Item* = RECORD
```

```
  mode*: INT;
```

```
  type*: ORB.Type;
```

```
  a*, b*, r: INT;
```

```
  rdo*: BOOL (*read only*)
```

```
END ;
```

```
VAR pc: INT;
```

```
PROC MakeConstItem*(VAR x: Item; typ: ORB.Type; val: INT);
```

```
PROC MakeRealItem*(VAR x: Item; val: REAL);
```

```
PROC MakeStringItem*(VAR x: Item; len: INT);
```

```
PROC MakeItem*(VAR x: Item; y: ORB.Object; curlev: INT);
```

```
PROC Field*(VAR x: Item; y: ORB.Object); (* x := x.y *)
```

```
PROC Index*(VAR x, y: Item); (* x := x[y] *)
```

```
PROC DeRef*(VAR x: Item);
```

```
PROC BuildTD*(T: ORB.Type; VAR dc: INT);
```

```
PROC TypeTest*(VAR x: Item; T: ORB.Type; varpar, isguard: BOOL);
```

```
PROC Not*(VAR x: Item); (* x := ~x, Boolean operators *)
```

```

PROC And1*(VAR x: Item); (* x := x & *)
PROC And2*(VAR x, y: Item);
PROC Or1*(VAR x: Item); (* x := x OR *)
PROC Or2*(VAR x, y: Item);

```

```

PROC Neg*(VAR x: Item); (* x := -x, arithmetic operators *)
PROC AddOp*(op: LONGINT; VAR x, y: Item); (* x := x +- y *)
PROC MulOp*(VAR x, y: Item); (* x := x * y *)
PROC DivOp*(op: INT; VAR x, y: Item); (* x := x op y *)
PROC RealOp*(op: INT; VAR x, y: Item); (* x := x op y *)

```

```

PROC Singleton*(VAR x: Item); (* x := {x}, set operators *)
PROC Set*(VAR x, y: Item); (* x := {x .. y} *)
PROC In*(VAR x, y: Item); (* x := x IN y *)
PROC SetOp*(op: INT; VAR x, y: Item); (* x := x op y *)

```

```

PROC IntRelation*(op: INT; VAR x, y: Item); (* x := x < y *)
PROC SetRelation*(op: INT; VAR x, y: Item); (* x := x < y *)
PROC RealRelation*(op: INT; VAR x, y: Item); (* x := x < y *)
PROC StringRelation*(op: INT; VAR x, y: Item); (* x := x < y *)

```

```

PROC StrToChar*(VAR x: Item); (*assignments*)
PROC Store*(VAR x, y: Item); (* x := y *)
PROC StoreStruct*(VAR x, y: Item); (* x := y *)
PROC CopyString*(VAR x, y: Item); (*from x to y*)

```

```

PROC VarParam*(VAR x: Item; ftype: ORB.Type); (*parameters*)
PROC ValueParam*(VAR x: Item);

```

```

PROC OpenArrayParam*(VAR x: Item);
PROC StringParam*(VAR x: Item);

PROC For0*(VAR x, y: Item); (*For Statements*)
PROC For1*(VAR x, y, z, w: Item; VAR L: LONGINT);
PROC For2*(VAR x, y, w: Item);

(* Branches, procedure calls, procedure prolog and epilog *)
PROC Here*(): LONGINT;
PROC FJump*(VAR L: LONGINT);
PROC CFJump*(VAR x: Item);
PROC BJump*(L: LONGINT);
PROC CBJump*(VAR x: Item; L: LONGINT);
PROC Fixup*(VAR x: Item);
PROC PrepCall*(VAR x: Item; VAR r: LONGINT);
PROC Call*(VAR x: Item; r: LONGINT);
PROC Enter*(parblksize, locblksize: LONGINT; int: BOOL);
PROC Return*(form: INT; VAR x: Item; size: LONGINT; int: BOOL);

(* In-line code procedures*)
PROC Increment*(upordown: LONGINT; VAR x, y: Item);
PROC Include*(inorex: LONGINT; VAR x, y: Item);
PROC Assert*(VAR x: Item);
PROC New*(VAR x: Item);
PROC Pack*(VAR x, y: Item);
PROC Unpk*(VAR x, y: Item);
PROC Led*(VAR x: Item);
PROC Get*(VAR x, y: Item);

```

```
PROC Put*(VAR x, y: Item);
PROC Copy*(VAR x, y, z: Item);
PROC LDPSR*(VAR x: Item);
PROC LDREG*(VAR x, y: Item);

(*In-line code functions*)
PROC Abs*(VAR x: Item);
PROC Odd*(VAR x: Item);
PROC Floor*(VAR x: Item);
PROC Float*(VAR x: Item);
PROC Ord*(VAR x: Item);
PROC Len*(VAR x: Item);
PROC Shift*(fct: LONGINT; VAR x, y: Item);
PROC ADC*(VAR x, y: Item);
PROC SBC*(VAR x, y: Item);
PROC UML*(VAR x, y: Item);
PROC Bit*(VAR x, y: Item);
PROC Register*(VAR x: Item);
PROC H*(VAR x: Item);
PROC Adr*(VAR x: Item);
PROC Condition*(VAR x: Item);

PROC Open*(v: INT);
PROC SetDataSize*(dc: LONGINT);
PROC Header*;
PROC Close*(VAR modid: ORS.Ident; key, nofent: LONGINT);
END ORG.
```

## 10.4 THE PARSER

The main module ORP constitutes the parser. Its single command `Compile` - at the end of the program listing - identifies the source text according to the Oberon command conventions. It then calls procedure `Module` with the identified source text as parameter. The command forms are:

<code>ORP.Compile @</code>	The most recent selection identifies the beginning of the source text.
<code>ORP.Compile ^</code>	The most recent selection identifies the name of the source file.
<code>ORP.Compile f0 f1 ...</code>	<code>f0, f1, ...</code> are the names of source files

File names and the characters `@` and `^` may be followed by an option specification `/s`. Option `s` enables the compiler to overwrite an existing symbol file, thereby invalidating clients.

The parser is designed according to the proven method of top-down, recursive descent parsing with a look-ahead of a single symbol. The last symbol read is represented by the global variable `sym`. Syntactic entities are mirrored by procedures of the same name. Their goal is to recognize the specified construct in the source text. The start symbol and corresponding procedure is `Module`. The principal parser procedures are shown in Fig. 41, which also exhibits their

calling hierarchy. Loops in the diagram indicate recursion in the syntactic definition.

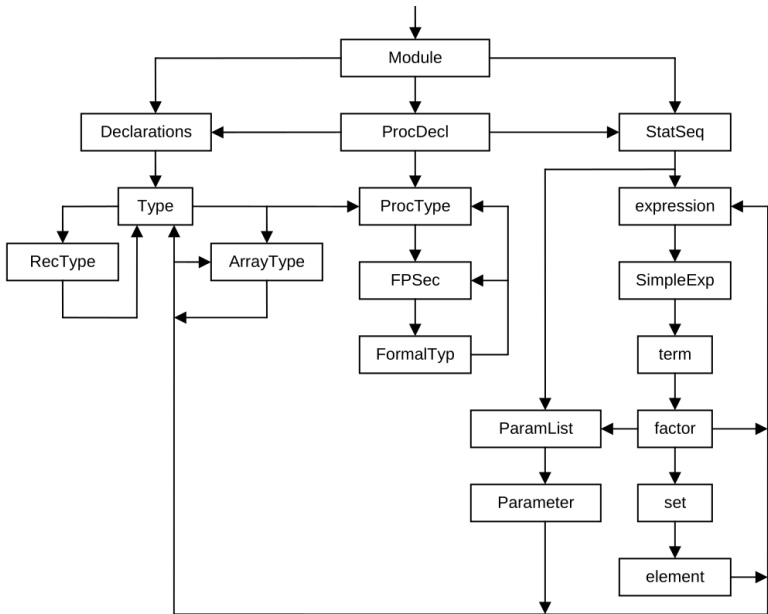


Figure 41: Parser procedure hierarchy

The rule of parsing strictly based on a single-symbol look-ahead and without reference to context is violated in three places. The prominent violation occurs in statements. If the first symbol of a statement is an identifier, the decision of whether an assignment or a procedure call is to be recognized is based on contextual information, namely the class of the identified object. The second violation occurs in qualident; if the identifier  $x$  preceding a period denotes a module, it is recognized together with the subsequent identifier as

a qualified identifier. Otherwise  $x$  supposedly denotes a record variable. The third violation is made in procedure selector; if an identifier is followed by a left parenthesis, the decision of whether a procedure call or a type guard is to be recognized is again made on the basis of contextual information, namely the mode of the identified object.

A fairly large part of the program is devoted to the discovery of errors. Not only should they be properly diagnosed. A much more difficult requirement is that the parsing process should continue on the basis of a good guess about the structure that the text should most likely have. The parsing process must continue with some assumption and possibly after skipping a short piece of the source text. Hence, this aspect of the parser is mostly based on heuristics. Incorrect assumptions about the nature of a syntactic error lead to secondary error diagnostics. There is no way to avoid them. A reasonably good result is obtained by the fact that procedure `ORS.Mark` inhibits an error report, if it lies less than 10 characters ahead of the last one. Also, the language Oberon is designed with the property that most large constructs begin with a unique symbol, such as `IF`, `WHILE`, `CASE`, `RECORD`, etc. These symbols facilitate the recovery of the parsing process in the erroneous text. More problematic are open constructs which neither begin nor end with key symbols, such as types, factors, and expressions. Relying on heuristics, the source text is skipped up to the first occurrence of a symbol which may begin a construct that follows

the one being parsed. The employed scheme may not be the best possible, but it yields quite acceptable results and keeps the amount of program devoted to the handling of erroneous texts within justifiable bounds.

Besides the parsing of text, the Parser also performs the checking for type consistency of objects. This is based on type information held in the global table, gained during the processing of declarations, which is also handled by the routines which parse. Thereby an unjustifiably large number of very short procedures is avoided. However, the strict target-computer independence of the parser is violated slightly: Information about variable allocation strategy including alignment, and about the sizes of basic types is used in the parser module. Whereas the former violation is harmless, because the allocation strategy is hardly controversial, the latter case constitutes a genuine target-dependence embodied in a number of explicitly declared constants. Mostly these constants are contained in the respective type definitions, represented by records of type `Type` initialized by ORB. The following procedures allocate objects and generate elements of the symbol table:

An inherently nasty subject is the treatment of forward references in a single-pass compiler. In Oberon, there are 2 such cases:

1. Forward declarations of procedures. They have been eliminated from the revision of the Oberon language



Declarations	Object(Con), Object(Typ), Object(Var)
ProcedureDeclaration	Object(xProc)
FormalType	Object(Var), Object(Par)
ORB.Import	Object(Mod)
RecordType	Object(Fld), Type(Record)
ArrayType	Type(Array)
ProcedureType	Type(ProcTyp)
Type	Type(Pointer)
FormalType	Type(Array)

in the year 2007 as they should be avoided if ever possible. If it is impossible, a remedy is to declare a variable of the given procedure type, and assign the procedure to be forwarded to this variable. The nastiness of procedure forward declarations originates in the necessity to specify parameters and result type in the forward declaration. These must be repeated in the actual procedure declaration, and one expects that a compiler verifies the equality (or equivalence) of the 2 declarations. This is a heavy burden for a case that very rarely occurs.

2. Forward declarations of pointer types also constitutes a nasty exception, but its exclusion would be difficult to justify. If in a pointer declaration the base type (to which the pointer is bound) is not found in the symbol table, a forward reference is therefore automatically assumed. An entry for the pointer type is generated

anyway (see procedure `Type`) and an element is inserted in the list of pointer base types to be fixed up. This list is headed by the global variable `pbsList`. When later in the text a declaration of a record type is encountered with the same identifier, the forward entry is recognized and the proper link is established (see procedure `Declarations`).

The compiler must check for undefined forward references when the current declaration scope is closed. This check is performed at the end of procedure `Declarations`.

The `with` statement had been eliminated from the language in its revision of 2007. Here it reappears in the form of a case statement, whose cases are not labelled by integers, but rather by types. What formerly was written as

```
IF x IS T1 THEN
  WITH x: T1 DO ... x ... END
ELSIF x IS T2 THEN
  WITH x: T2 DO ... x ... END
ELSIF ... END
```

is now written more simply and more efficiently as

```
CASE x OF
  T1: ... x ... |
  T2: ... x ... |
  ... END
```

where T1 and T2 are extensions of the type T0 of the case variable x. Compilation of this form of case statement merges the regional type guard of the former with statements with the type test of the former if statements. This case statement represents the only case where a symbol table entry - the type of x - is modified during compilation of statements. When the end of the with statement is reached, the change must be reverted.

## 10.5 THE SCANNER

The scanner module ORS embodies the lexicographic definitions of the language, i.e. the definition of abstract symbols in terms of characters. The scanner's substance is procedure Get, which scans the source text and, for each call, identifies the next symbol and yields the corresponding integer code. It is most important that this process be as efficient as possible. Procedure Get recognizes letters indicating the presence of an identifier (or reserved word), and digits signalling the presence of a number. Also, the scanner recognizes comments and skips them. The global variable ch stands for the last character read.

A sequence of letters and digits may either denote an identifier or a key word. In order to determine which is the case, a search is made in a table containing all key words for each would-be identifier. This table is sorted alphabetically and

according to the length of reserved words. It is initialized when the compiler is loaded.

The presence of a digit signals a number. Procedure Number first scans the subsequent digits (and letters) and stores them in a buffer. This is necessary, because hexadecimal numbers are denoted by the postfix letter H (rather than a prefix character). The postfix letter X specifies that the digits denote a character.

There exists one case in the language Oberon, where a look-ahead of a single character does not suffice to identify the next symbol. When a sequence of digits is followed by a period, this period may either be the decimal point of a real number, or it may be the first element of a range symbol ( .. ). Fortunately, the problem can be solved locally as follows: If, after reading digits and a period, a second period is present, the number symbol is returned, and the look-ahead variable *ch* is assigned the special value 7FX. A subsequent call of *Get* then delivers the range symbol. Otherwise the period after the digit sequence belongs to the (real) number.

## 10.6 SEARCHING THE SYMBOL TABLE, AND HANDLING SYMBOL FILES

### 10.6.1 *The structure of the symbol table*

The symbol table constitutes the context in which statements and expressions are parsed. Each procedure establishes a

scope of visibility of local identifiers. The records registering identifiers belonging to a scope are linked as a linear list. They are of type `Object`. Each object has a type.

Types are represented by records of type `Type`. These 2 types pervade the entire compiler, and they are defined as follows:

```
TYPE Object = POINTER TO ObjDesc;  
    Type    = POINTER TO TypeDesc;
```

```
ObjDesc = RECORD  
    class, lev, exno: INT;  
    expo, rdo: BOOL; (*exported / read-only*)  
    next, dsc: Object;  
    type: Type;  
    name: ORS.Ident;  
    val: INT  
END ;
```

```
TypeDesc = RECORD  
    form, ref, mno: INT; (*ref is only used for import/export*)  
    nofpar: INT; (*for procedures; extension level for records*)  
    len: INT; (*for arrays, len < 0 => open array; for records: address*)  
    dsc, typobj: Object;  
    base: Type; (*for arrays, records, pointers*)  
    size: INT; (*in bytes; always multiple of 4, except for Byte, Boolean*)  
END ;
```

Procedures for generating and searching the lists are contained in module ORB. If a new identifier is to be added, procedure NewObj first searches the list, and if the identifier is already present, a double definition is diagnosed. Otherwise the new element is appended, thereby preserving the order given by the source text.

Procedures, and therefore also scopes, may be nested. Each scope is represented by the list of its declared identifiers, and the list of the currently visible scopes are again connected as a list. Procedure OpenScope appends an element and procedure CloseScope removes it. The list of scopes is anchored in the global variable topScope and linked by the field dsc. It is treated like a stack. It consists of elements of type Object, each one being the header (class = Head) of the list of declared entities. As an example, the procedure for searching an object (with name ORS.id) is shown here:

```
PROC thisObj*(): Object;
  VAR s, x: Object;
BEGIN s := topScope;
  REPEAT x := s.next;
    WHILE (x # NIL) & (x.name # ORS.id) DO x := x.next EN
    s := s.dsc
  UNTIL (x # NIL) OR (s = NIL);
  RETURN x
END thisObj;
```

A snapshot of a symbol table for an example with nested scopes is shown in Fig. 42. It is taken when the following declarations are parsed and when the statement S is reached.

```
VAR x: INT;
```

```
PROC P(u: INT);  
BEGIN ... END P;
```

```
PROC Q(v: INT);  
  PROC R(w: INT);  
  BEGIN S END R;  
BEGIN ... END Q;
```

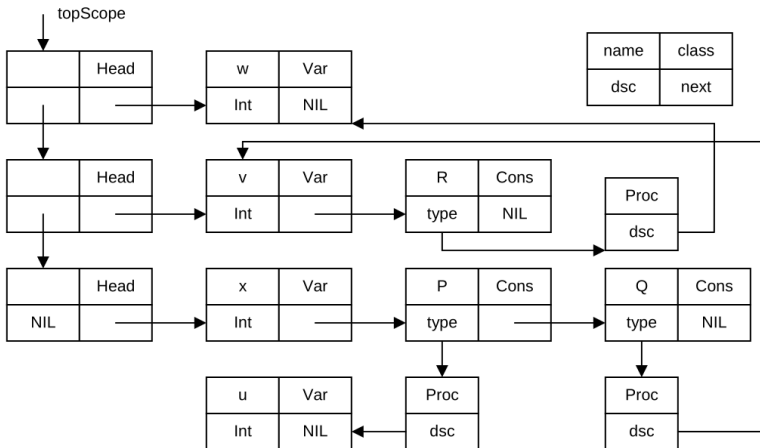


Figure 42: Snapshot of a symbol table

A search of an identifier proceeds first through the scope list, and for each header its list of object records is scanned. This

mirrors the scope rule of the language and guarantees that if several entities carry the same identifier, the most local one is selected. The linear list of objects represents the simplest implementation by far. A tree structure would in many cases be more efficient for searching, and would therefore seem more recommendable. Experiments have shown, however, that the gain in speed is marginal. The reason is that the lists are typically quite short. The superiority of a tree structure becomes manifest only when a large number of global objects is declared. We emphasize that when a tree structure is used for each scope, the linear lists must still be present, because the order of declarations is sometimes relevant in interpretation, e.g. in parameter lists.

Not only procedures, but also record types establish their own local scope. The list of record fields is anchored in the type record's field dsc, and it is searched by procedure `thisField`. If a record type `R1` is an extension of `R0`, then `R1`'s field list contains only the fields of the extension proper. The base type `R0` is referenced by the `BaseTyp` field of `R1`. Hence, a search for a field may have to proceed through the field lists of an entire sequence of record base types.

#### 10.6.2 *Symbol files*

The major part of module ORB is devoted to input and output of symbol files. A symbol file is a linearized form of an excerpt of the symbol table containing descriptions of all



exported (marked) objects. All exports are declared in the global scope. Procedure Export traverses the list of global objects and outputs them to the symbol file.

The structure of a symbol file is defined by the syntax specified below. The following terminal symbols are class and form specifiers or reference numbers for basic types with fixed values:

Classes: Con = 1, Var = 2, Par = 3, Fld = 4; Typ = 5

Forms: Byte = 1, Bool = 2, Char = 3, Int = 4, LInt = 5, Set = 6, Pointer = 7, NoTyp = 9, ProcTyp = 10, Array = 12, Record = 13

Syntax:

```
SymFile = null key name versionkey {object}.
```

```
object = (CON name type (value | exno) | TYP name type [{fix} 0]
```

```
type = ref (PTR type | ARR type len | REC type {field} 0 | PRO ty
```

```
field = FLD name type offset.
```

```
param = (VAR | PAR) type.
```

A procedure type description is contains a parameter list. Similarly, a record type description with form specifier Record contains the list of field descriptions. Note that a procedure is considered as a constant of a procedure type.

Objects have types, and types are referenced by pointers. These cannot be written on a file. The straight-forward so-

lution would be to use the type identifiers as they appear in the program to denote types. However, this would be rather crude and inefficient, and second, there are anonymous types, for which artificial identifiers would have to be generated.

An elegant solution lies in consecutively numbering types. Whenever a type is encountered the first time, it is assigned a unique reference number. For this purpose, records (in the compiler) of type `Type` contain the field `ref`. Following the number, a description of the type is then written to the symbol file. When the type is encountered again during the traversal of the data structure, only the reference number is issued, with negative sign. The global variable `ORB.Ref` functions as the running reference number.

When reading a symbol file, a positive reference number is followed by the type's description. A pointer to the type read is assigned to the global table `typtab` with the reference number as index. When a negative reference number is read (it is not followed by a type description), then the type is identified by `typtab[-ref]` (see procedure `InType`). In the following example, types are identified by their reference number (e.g. `R #14`), and later referenced by this number (`^14`).

```
MODULE A;
```

```
  CONST Ten* = 10; Dollar* = "\$";
```

```

TYPE R* = RECORD u*: INT; v*: SET END ;
      S* = RECORD w*: ARRAY 4 OF R END ;
      P* = POINTER TO R;
      A* = ARRAY 8 OF INT;
      B* = ARRAY 4, 5 OF REAL;
      C* = ARRAY 10 OF S;
      D* = ARRAY OF CHAR;

VAR x*: INT;
PROC Q0*;
BEGIN END Q0;
PROC Q1*(x, y: INT): INT;
BEGIN RETURN x+y END Q1;
END A.

```

```

class = CON Ten [^4] 10
class = CON Dollar [^3] 36
class = TYP R [#14 form = REC [^9] exno = 1 extlev = 0 size = 8 { v
class = TYP S [#15 form = REC [^9] exno = 2 extlev = 0 size = 32 { v
class = TYP P [#16 form = PTR [^14]]()
class = TYP A [#17 form = ARR [^4] len = 8 size = 32]()
class = TYP B [#18 form = ARR [#0 form = ARR [^5] len = 5 size = 20
class = TYP C [#19 form = ARR [^15] len = 10 size = 320]()
class = TYP D [#20 form = ARR [^3] len = -1 size = 8]()
class = VAR x [^{4}] 3
class = CON Q0 [#0 form = PRO [^9]()] 4
class = CON Q1 [#0 form = PRO [^4]( class = VAR [^4] class = VAR [^

```

After a symbol file has been generated, it is compared with the file from a previous compilation of the same module, if one exists. Only if the 2 files differ and if the compiler's `s`-option is enabled, is the old file replaced by the new version. The comparison is made by comparing byte after byte without consideration of the file's structure. This somewhat crude approach was chosen because of its simplicity and yielded good results in practice.

A symbol file must not contain addresses (of variables or procedures). If they did, most changes in the program would result in a change of the symbol file. This must be avoided, because changes in the implementation (rather than the interface) of a module are supposed to remain invisible to the clients. Only changes in the interface are allowed to effect changes in the symbol file, requiring recompilation of all clients. Therefore, addresses are replaced by export numbers. The variable `exno` (global in ORP) serves as running number (see `ORP.Declarations` and `ORP.ProcedureDecl`). The translation from export number to address is performed by the loader. Every code file contains a list (table) of addresses (of variables and entry points for procedures). The export number serves as index in this table to obtain the requested address. Export numbers are generated by the parser.

Objects exported from some module `M1` may refer in their declaration to some other module `Mo` imported by `M1`. It would be unacceptable, if an import of `M1` would then also

require the import of  $M_0$ , i.e. imply the automatic reading of  $M_0$ 's symbol file. It would trigger a chain reaction of imports that must be avoided. Fortunately, such a chain reaction can be avoided by making symbol files self-contained, i.e. by including in every symbol file the description of entities that stem from other modules. Such entities are always types.

The inclusion of types imported from other modules seems simple enough to handle: type descriptions must include a reference to the module from which the type was imported. This reference is the name and key of the respective module. However, there exists one additional complication that cannot be ignored. Consider a module  $M_1$  importing a variable  $x$  from a module  $M_0$ . Let the type  $T$  of  $x$  be defined in module  $M_0$ . Also, assume  $M_1$  to contain a variable  $y$  of type  $M_0.T$ . Evidently,  $x$  and  $y$  are of the same type, and the compiler compiling  $M_2$  must recognize this fact. Hence, when importing  $M_0$  during compilation of  $M_1$ , the imported element  $T$  must not only be registered in the symbol table, but it must also be recognized as being identical to the  $T$  already imported from  $M_2$  directly. It is rather fortunate that the language definition specifies equivalence of types on the basis of names rather than structure, because it allows type tests at execution time to be implemented by a simple address comparison.

The measures to be taken to satisfy the new requirements are as follows:

1. Every type element in a symbol file is given a module number. Before a type description is emitted to the file.
2. If a type to be exported has a name and stems from another, imported module, then also the name and key of the module are attached, from which the type stems (see end of procedure `ORB.OutType` and end of `ORB.InType`).

An additional detail is worth being mentioned here: Hidden pointers. We recall that individual fields of exported record types may be hidden. If marked (by an asterisk) they are exported and therefore visible in importing modules. If not marked, they are not exported and remain invisible, and evidently seem to be omissible in symbol files. However, this is a fallacy. They need to be included in symbol files, although without name, because of meta information to be provided for garbage collection. This is elucidated as follows:

Assume that a module `M1` declares a global pointer variables of a type imported from module `Mo`.

```
MODULE M0;  
    TYPE Ptr = POINTER TO Rec0;
```

```
    Rec0* = RECORD p*, q: Ptr ... END ;  
END M0.  
  
MODULE M1;  
    VAR p: M0.Ptr;  
        r: RECORD f: M0.Ptr; ... END ;  
END M1.
```

Here  $p$  and  $r.f$  are roots of data structures that must be visited by the garbage collector. If they are not, they will not be marked, and therefore collected with disastrous and entirely unpredictable consequences. The crux is that not only exported pointers ( $p.p$ ) must be listed, but also hidden ones ( $p.q$ ), although they are not accessible in module  $M1$ .

We chose to include hidden pointers in symbol files without their names, but with their type being of the form `ORB.NilTyp`. This must be considered in procedure `ORG.FindPtrs`, where the condition `typ.form = ORB.Pointer` must be extended to `(typ.form = ORB.Pointer) OR (typ.form = ORB.NilTyp)`.

But the story does not end here. Assume that in the example above module  $M1$  declares a type `Rec1` as a  $n$  extension of `Mo.Reco`. This requires the generation of a type descriptor. And this descriptor must include not only field  $p$ , but also the hidden field  $q$ . This is achieved by also extending the condition `typ.form = ORB.Pointer` in `ORG.FindPtrFlds` to `(typ.form = ORB.Pointer) OR (typ.form = ORB.NilTyp)`.

## 10.7 THE CODE GENERATOR

The routines for generating instructions are contained in a single module: ORG. They are fairly numerous, and therefore the interface of ORG is quite large. It is a procedural interface. This implies that there is no "intermediate code" or "intermediate data structure" between parser and code generator. This is one reason for the compactness of the code generator. The other is the regularity and simplicity of the processor architecture. In order to understand the following material, the reader is supposed to be familiar with this architecture (Appendix 2) and the generated code patterns for individual language constructs (Section 12.2).

A distinguishing feature of this compiler is that parsing proceeds top-down according to the principle of recursive descent in the parsing tree. This implies that for every syntactic construct a specific procedure is called. It carries the same name as the construct. It also implies that properties of the parsed construct can be represented by parameters of the parsing procedures. Consider, for example, the construct of simple expression:

```
SimpleExpression = term {"+" term}.
```

The corresponding parsing procedure is

```
PROC SimpleExpression(VAR x: Item);
```



```

    VAR y: Item;
BEGIN term(x);
    WHILE sym = plus DO ORS.Get(sym); term(y); ORG.AddOp(x, y) END
END SimpleExpression

```

The generating procedure `AddOp` receives 2 parameters representing the operands, and returns the result through the first parameter. This scheme carries the invaluable advantage of using operands efficiently allocated on the stack rather than dynamically allocated on the heap and subject to automatic storage retrieval (garbage collection). Here the processed operands quietly disappear from the stack upon exit from the parser procedure.

The parameters representing syntactic constructs are of type `Item` defined in `ORG`. This data type is rather similar to the type `Object` (in `ORB`). After all, it serves the same purpose; but it represents internal items rather than declared objects.

```

TYPE Item = RECORD
    mode: INT;
    type: ORB.Type;
    a, b, r: INT;
    rdo: BOOL (*read only*)
END

```

The attribute class of `Object` is renamed `mode` in `Item`. In fact, in some sense different classes evoke different (corre-

sponding) addressing modes as featured by the processor architecture. According to the architecture, additional modes may have to be introduced. Thanks to the simplicity of RISC, only three are needed:

Reg = 10; The item *x* is located in register *x.r*

RegI = 11; The item *x* is addressed indirectly through register

Cond = 12; The item is represented by the condition bit reg

Instructions are emitted sequentially and emitted by the four procedures *Puto*, *Put1*, *Put2*, *Put3*. They directly correspond to the instruction formats of the RISC processor (see Chapter 11). The instructions are stored in the array *code* and the compiler variable *pc* serves as running index.

```
PROC Put0(op, a, b, c: INT);      format F0
PROC Put1(op, a, b, im: INT);    format F1
PROC Put2(op, a, b, off: INT);   format F2
PROC Put3(op, cond, off: INT);   format F3
```

### 10.7.1 *Expressions*

Expressions consist of operands and operators. They are evaluated and have a value. First, a number of make-procedures transform objects into items (see Section 12.3.2). The principal one is *MakeItem*. Typical objects are variables (class, mode = *Var*). Global variables are addressed with base register *SB* (*x.r* = 13), local variables with the stack

pointer SP ( $x.r = 14$ ). VAR-parameters are addressed indirectly; the address is on the stack (class, mode = Par, Ind).  $x.a$  is the offset from the stack pointer.

Before an operator can be applied to operands, these must first be transferred (loaded) into registers. This is because the RISC performs operations only on registers. The loading is achieved by procedure load (and loadAdr) in ORG. The resulting mode is Reg. In allocating registers, a strict stack principle is used, starting with R0, up to R11. This is certainly not an optimal strategy and provides ample room for improvement (usually called optimization). The compiler variable RH indicates the next free register (top of register stack).

Base address SB is, as the name suggests, static. But this holds only within a module. It implies that on every transfer to a procedure in another module, the static base must be adjusted. The simplest way is to load SB before every external call, and to restore it to its old value after return from the procedure. We chose a different strategy: loading on demand (see below: global variables).

If a variable is indexed, has a field selector, is dereferenced, or has a type guard, this is detected in the parser by procedure selector. It calls generators Index, Field, DeRef, or TypeTest accordingly (see §10.3.2 and §10.2.1 - §10.2.4). These procedures cause item modes to change as follows:

	mode transition	instructions emitted	construct
1.	Index(x, y) (y is loaded into y.r)		
	Var → RegI	ADD y.r, SP, y.r	array variable
	Par → RegI	LDR RH, SP, x.a	array parameter
		ADD y.r, RH, y.r	
	RegI → RegI	ADD x.r, x.r, y.r	indexed array
2.	Field(x, y) (y.mode = Fld, y.a = field offset)		
	Var → Var	none	field designator, add c
	RegI → RegI	none	add field offset to x.a
	Par → Par	none	add field offset to x.b
3.	DeRef(x)		
	Var → RegI	LDR RH, SP, x.a	dereferenced x^
	Par → RegI	LDR RH, SP, x.a	dereferenced parameter
		LDR RH, RH, x.b	
	RegI → RegI	LDR x.r, x.r, x.a	

A fairly large number of procedures then deal with individual operators. Specifically, they are:

operators	on
Not, And1, And2, Or1, Or2	Booleans
Neg, AddOp, MulOp, DivOp	integers
RealOp	real numbers
Singleton, Set, In, SetOp	Sets

And finally, following the same pattern, are the procedures for relations (comparisons):

`IntRelation, SetRelation, RealRelation, StringRelation.`

See Appendix for listing of ORG.

We note in particular that if all operands are constants, their evaluation is performed by the compiler and not delegated to run-time. This is an important efficiency factor.

### 10.7.2 *Relations*

RISC does not feature any compare instruction. Instead, subtraction is used, because an implicit comparison with 0 is performed along with any arithmetic (or load) instruction. Instead of  $x < y$  we use  $x - y < 0$ . This is possible, because in addition to the computed difference deposited in a register, also the result of the comparison is deposited in the condition flags N (difference negative) and Z (difference zero). Relations therefore yield a result item  $x$  with mode `Cond.x.r` ( $= \text{relmap}[\text{sym}]$ ) identifies the relation. Branch instructions (jumps) are executed conditionally depending on these flags. The value  $x.r$  is then used when generating branch instructions. For example, the relation  $x < y$  is translated simply into

```
LDR R0, SP, x
```

```
LDR R1, SP, y
CMP R0, R0, R1
```

and the resulting item mode is  $x.mode = Cond$ ,  $x.r := "less"$ . (The mnemonic CMP is synonymous with SUB). More about relations and Boolean expressions will be explained in §12.7.6.

### 10.7.3 Set operations

The type SET represents sets of small integers in the range from 0 to 31. Bit  $i$  being 1 signals that  $i$  is an element of the set. This is a convenient representation, because the logical instructions directly mirror the set operations: AND implements set intersection, OR set union, and XOR the symmetric set difference. This representation also allows a simple and efficient implementation of membership tests. The instructions for the expression  $n \text{ IN } s$  is generated by procedure In. Assuming the value  $n$  in register R0, and the set  $s$  in R1, we obtain

```
ADD R0, R0, 1
ROR R1, R1, R0    rotate s by i+1 position, the
                  relevant bit moving to the sign bit
```

The resulting item mode is Cond with  $x.r = "minus"$ .

Of some interest are the procedures for generating sets, i.e. for processing  $\{m\}$ ,  $\{m \dots n\}$ , and  $\{m, n\}$ , where  $m, n$  are integer expressions.

We start with  $\{m\}$ . It is generated by procedure Singleton using a shift instruction. Assuming  $m$  in  $R0$ , the resulting code is

```
MOV R1, 0, 1
LSL R0, R1, R0    shift 1 by m bit positions to the left
```

Somewhat more sophisticated is the generation of  $\{m \dots n\}$  by procedure Set. Assuming  $m$  in  $R0$ , and  $n$  is  $R1$ , the resulting code is

```
MOV R2, 0, -2
LSL R1, R2, R1    shift -2 by n bit positions to the left
MOV R2, 0, -1
LSL R0, R2, R0    shift -1 by m bit positions to the left
XOR R0, R0, R1
```

The set  $\{m, n\}$  is generated as the union of  $\{m\}$  and  $\{n\}$ . If any of the element values is a constant, several possibilities of code improvement are possible. For details, the reader is referred to the source code of ORG.

#### 10.7.4 *Assignments*

Statements have an effect, but no result like expressions. Statements are executed, not evaluated. Assignments alter the value of variables through store instructions. The computation of the address of the affected variable follows the same scheme as for loading. The value to be assigned must be in a register.

Assignments of arrays (and records) are an exceptional case in so far as they are performed not by a single store instruction, but by a repetition. Consider  $y := x$ , where  $x$ , and  $y$  are both arrays of  $n$  integers. Assuming that the address of  $y$  is in register  $R0$ , that of  $x$  in  $R1$ , and the value  $n$  in  $R2$ . Then the resulting code is

```
L   LDR R3, R1, 0    source
      ADD R1, R1, 4
      STR R3, R0, 0    destination
      ADD R0, R0, 4
      SUB R2, R2, 1    counter
      BNE L
```

#### 10.7.5 *Conditional and repetitive statements*

These statements are implemented using branch instructions (jumps) as shown in §10.2.5 - §10.2.7. In all repetitive statements, backward jumps occur. Here, at the point of



return the value of the global variable `ORG.pc` is saved in a local (!) variable of the involved parsing procedure. It is retrieved when the backward jump is emitted. We note that branch instructions use a displacement rather than an absolute destination address. It is the difference between the branch instruction and the destination of the jump.

A difficulty, however, arises in the case of forward jumps, a difficulty inherent in all single-pass compilers: When the branch is issued, its destination is still unknown. It follows that the branch displacement must be later inserted when it becomes known, when the destination is reached. This is called a *fixup*. Here the method of fixup lists is used. The place of the instruction with still unknown destination is held in a variable `L` local to the respective parsing procedure. If several branches have the same destination, `L` is the heading of a list of the instructions to be fixed up, with its links placed in the instructions themselves in the place of the eventual jump displacement. This shown for the if statement by an excerpt of `ORP.StatSequence` with local variable `L0`:

```
ELSIF sym = ORS.if THEN
  ORS.Get(sym); ORG.FJump(L0); ORG.Fixup(x); expression(x);
  ORG.CFJump(x); Check(ORS.then, "no THEN"); StatSequence
WHILE sym = ORS.elsif DO
  ORS.Get(sym); expression(x); ORG.CFJump(x);
  StatSequence; L0 := 0;
END ;
```

```

    IF sym = ORS.else THEN ORS.Get(sym); ORG.FJump(L0); ORG
    ELSE ORG.Fixup(x)
    END ;
    ORG.FixLink(L0);

```

where in module ORG:

```

PROC CFJump(VAR x: Item); (*conditional forward jump*)
BEGIN
    IF x.mode # Cond THEN loadCond(x) END ;
    Put3(BC, negated(x.r), x.a); FixLink(x.b); x.a := pc-1
END CFJump;

PROC FJump(VAR L: LONGINT); (*unconditional forward jump*)
BEGIN Put3(BC, 7, L); L := pc-1
END FJump;

PROC fix(at, with: LONGINT);
BEGIN code[at] := code[at] DIV C24 * C24 + (with MOD C24)
END fix;

PROC FixLink(L: LONGINT);
    VAR L1: LONGINT;
BEGIN invalSB;
    WHILE L # 0 DO L1 := code[L] MOD 40000H; fix(L, pc-L-1); L := L1
END FixLink;

PROC Fixup(VAR x: Item);

```

```
BEGIN FixLink(x.a)
END Fixup;
```

In while-, repeat-, and for- statements essentially the same technique is used with the support of the identical procedures in ORG.

#### 10.7.6 *Boolean expressions*

In the case of arithmetic expressions, our compilation scheme results in a conversion from infix to postfix notation ( $x+y \Rightarrow xy+$ ). This is not applicable for Boolean expressions, because the operators & and OR are defined as follows:

```
x & y --> if x then y else FALSE
x OR y --> if x then TRUE else y
```

This entails that depending on the value of  $x$ ,  $y$  must not be evaluated. As a consequence, jumps may have to be taken across the code for  $y$ . Therefore, the same technique of conditional evaluation must be used as for conditional statements. In the case of an expression  $x \& y$  ( $x \text{ OR } y$ ), procedure `ORG.And1` resp. `ORG.Or1` must be called just after parsing  $x$  (see `ORP.term` resp. `ORP.SimpleExpression`). Only after parsing also  $y$  can the generators `ORG.And2` resp. `ORG.Or2` be called, providing the necessary fixups of forward jumps.

```
PROC And1(VAR x: Item); (* x := x & *)
```

```
BEGIN
```

```
  IF x.mode # Cond THEN loadCond(x) END ;
```

```
  Put3(BC, negated(x.r), x.a); x.a := pc-1; FixLink(x.b); x
```

```
END And1;
```

```
PROC And2(VAR x, y: Item);
```

```
BEGIN
```

```
  IF y.mode # Cond THEN loadCond(y) END ;
```

```
  x.a := merged(y.a, x.a); x.b := y.b; x.r := y.r
```

```
END And2;
```

A negative consequence of this scheme having condition flags in the processor is that when an item with mode Cond has to be transferred into mode Reg, as in a Boolean assignment, an unpleasantly complex instruction sequence must be generated. Fortunately, this case occurs quite rarely.

#### 10.7.7 *Procedures*

Before embarking on an explanation of procedure calls, entries and exits, we need to know how recursion is handled and how storage for local variables is allocated. Procedure calls cause a sequence of frames to be allocated in a stack fashion. These frames are the storage space for local variables. Each frame is headed by a single word containing the return address of the call. This address is deposited in R15 by the call instructions (BL, branch and link). The com-

piler "knows" the size of the frame to be allocated, and thus merely decrements the stack pointer SP (R14) by this amount. Upon return, SP is incremented by the same amount, and PC is restored by a branch instruction. In the following example, a procedure P is called, calling itself Q, and Q calling P again (recursion). The stack then contains 3 frames (see Fig. 43).




Figure 43: Stack frames

Scheme and layout determine the code sequences for call, entry and exit of procedures. Here is an example of a procedure P with 2 parameters:

```
Call:  LDR R0, param0
       LDR R1, param1
       BL  P
```

```

Prolog: SUB SP , SP, size    decrement SP
        STR LNK, SP, 0      push return adr
        STR R0 , SP, 4      push parameter 0
        STR R1 , SP, 8      push parameter1 ...

Epilog: LDR LNK, SP, 0      pop return adr
        ADD SP , SP, size   increment SP
        BR  LNK

```

When the call instruction is executed, parameters reside in registers, starting with  $R_0$ . For function procedures, the result is passed in register  $R_0$ . This scheme is very efficient; storing the parameters occurs only in a single place, namely at procedure entry, and not before each call. However, it has severe consequences for the entire register allocation strategy. Throughout the compiler, registers must be allocated in strict stack fashion. Furthermore, parameter allocation must start with  $R_0$ . This is a distinct drawback for function calls. If registers are occupied by other values loaded prior to the call, they must be cleared, i.e. the parameters must be saved and reloaded after return. This is rather cumbersome (see procedures `ORG.SaveRegisters` and `ORG.RestoreRegisters`).

```

F(x)                no register saving
x + F(x)
F(F(x))
(x+1) + F(x)        register saving necessary

```

### 10.7.8 *Type extension*

Static typing is an important principle in programming languages. It implies that every constant, variable or function is of a certain data type, and that this type can be derived by reading the program text without executing it. It is the key principle to introduce important redundancy in languages in such a form that a compiler can detect inconsistencies. It is therefore the key element for reducing the number of errors in programs.

However, it also acts as a restriction. It is, for example, impossible to construct data structures (arrays, trees) with different types of elements. In order to relax the rule of strictly static typing, the notion of *type extension* was introduced in Oberon. It makes it possible to construct inhomogeneous data structures without abandoning type safety. The price is that the checking of type consistency must in certain instances be deferred to runtime. Such checks are called *type tests*. The challenge is to defer to run-time as few checks as possible and as many as needed.

The solution in Oberon is to introduce families of types, and compatibility among their members. Their members are thus related, and a family forms a hierarchy. The principle idea is the following: Any record type  $T_0$  can be extended into a new type  $T_1$  by additional record fields (attributes).  $T_1$  is then called an extension of  $T_0$ , which in turn is said to be  $T_1$ 's *base type*.  $T_1$  is then type compatible with  $T_0$ , but not

vice-versa. This property ensures that in many cases static type checking is still possible. Furthermore, it turns out that run-time tests can be made very efficient, thus minimizing the overhead for maintaining type safety.

For example, given the declarations

```
TYPE R0 = RECORD u, v: INT END ;  
      R1 = RECORD (R0) w: INT END
```

we say that R1 is an *extension* of R0. R0 has the fields u and v, R1 has u, v, and w. The concept becomes useful in combination with pointers. Let

```
TYPE P0 = POINTER TO R0;  
      P1 = POINTER TO R1;  
VAR p0: P0; p1: P1;
```

Now it is possible to assign p1 to p0 (because a P1 is always also a P0), but not p0 to p1, because a P0 need not be a P1. This has the simple consequence that a variable of type P0 may well point to an extension of R0. Therefore, data structures can be declared with a base type, say P0, as common element type, but in fact they can individually differ, they can be any extension of the base type.

Obviously, it must be possible to determine the actual, current type of an element even if the base type is statically



fixed. This is possible through a *type test*, syntactically a Boolean factor:

$p0 \text{ IS } P1$  (short for  $p0^{\wedge} \text{ IS } R1$ )

Furthermore, we introduce the *type guard*. In the present example, the designator  $p0.w$  is illegal, because there is no field  $w$  in a record of type  $R0$ , even if the current value of  $p0^{\wedge}$  is a  $R1$ . As this case occurs frequently, we introduce the short notation  $p0(P1).w$ , implying a test  $p0 \text{ IS } P1$  and an abort if the test is not met.

It is important to mention that this technique also applies to formal variable parameters of record type, as they also represent a pointer to the actual parameter. Its type may be any extension of the type specified for the formal parameter in the procedure heading.

How are type test and type guard efficiently implemented? Our first observation is that they must consist of a single comparison only, similar to index checks. This in turn implies that types must be identified by a single word. The solution lies in using the unique address of the type descriptor of the (record) type. Which data must this descriptor hold? Essentially, type descriptors (TD) must identify the base types of a given type. Consider the following hierarchy:

TYPE T = RECORD ... END ;

```

T0  = RECORD ( T ) ... END ;    extension level 1
T1  = RECORD ( T ) ... END ;
T00 = RECORD ( T0 ) ... END ;    extension level 2
T01 = RECORD ( T0 ) ... END ;
T10 = RECORD ( T1 ) ... END ;
T11 = RECORD ( T1 ) ... END ;

```




Figure 44: A type hierarchy

In the symbol table, the field *base* refers to the ancestor of a given record type. Thus base of the type representing T11 points to T1, etc. Run-time checks, however, must be fast, and hence cannot proceed through chains of pointers. Instead, each TD contains an array with references to the ancestor TDs (including itself). For the example above, the TDs are as follows:

```

TD(T ) = [ T      ]
TD(T0 ) = [ T, T0  ]
TD(T1 ) = [ T, T1  ]
TD(T00) = [ T, T0, T00]
TD(T01) = [ T, T0, T01]
TD(T10) = [ T, T1, T10]
TD(T11) = [ T, T1, T11]

```

Evidently, the first element can be omitted, as it always refers to the common base of the type hierarchy. The last element always points to the TD's owner. TDs are allocated in the data area, the area for variables.

References to TDs are called *type tags*. They are required in 2 cases.

1<sup>st</sup> is for records referenced by pointers. Such dynamically allocated records carry an additional, hidden field holding their type tag. (A second additional word is reserved for use by the garbage collector. The offset of the tag field is therefore -8).

2<sup>nd</sup> case is that of record-typed VAR-parameters. In this case the type tag is explicitly passed along with the address of the actual parameter. Such parameters therefore require 2 words/registers.

A type test then consists of a test for equality of 2 type tags. In  $p \text{ IS } T$  the first tag is that of the  $n$ 'th entry of the TD of  $p^{\wedge}$ , where  $n$  is the extension level of  $T$ . The second tag is that of type  $T$ . This is shown in §10.2.12 (see also Fig. 39). The test then is as follows:

$$p^{\wedge}.\text{tag}^{\wedge}[n] = \text{adr}(T) \quad n \text{ is the extension level to } T$$

When declaring a record type, it is not known how many extensions, nor how many levels will be built on this type.

Therefore TD's should actually be infinite arrays. We decided to restrict them to 3 levels only. The first entry, which is never used for checking, is replaced by the size of the record.

#### 10.7.9 *Import and export, global variables*

Addresses of imported objects are not available to the compiler. Their computation must be left to the module loader (see Chapter 6). Similar to handling addresses of forward jumps, the compiler puts the necessary information in place of the actual address into the instruction itself. In the case of procedure calls, this is quite feasible, because the BL instruction features an offset field of 24 bits. The information consists of the module number and the export number of the imported object. In addition, there is a link to the previous instruction referring to an imported procedure. The origin of the list of procedure call fixups is rooted in the compiler variable `fixorgP`, and of the 24 bits in each BL instruction 4 bits are used for the module number, 8 bits for the object's export number, and 12 for the link. The loader need only scan this list to fix up the addresses (jump offsets).

Matters are more complex in the case of data. Object records in the symbol table have a field `lev`. It indicates the nesting level of variables local to procedures. It is also used for the module number in the case of variables of imported modules. Note that when importing, objects designating

modules are inserted in the symbol table, and the list of their own objects are attached in the field `dsc`. In this latter case, the module numbers have an inverted sign (are negative). Such imported objects are static, i.e. have a fixed address. In principle their absolute address could be computed (fixed) by the module loader. However, this is not practicable, because RISC instructions have an address offset of 16 bits only. It is therefore necessary in the general case to use a base address in conjunction with the offset. We use a single register for holding the static base (SB, R13). This register need be reloaded for every access to an imported variable. However, the compiler keeps track of external accesses; if a variable is to be accessed from the same module as the previous case, then reloading is avoided (see procedure `GetSB` and global compiler variable `curSB`).

This base address is fetched from a table global to the entire system. This module table contains one entry for every module loaded, namely the address of the module's data section. The address of the table is permanently in register `MT(=R12)`. An access to an imported variable therefore always requires 2 instructions:

```
LDR SB, MT, modno*4    base address of data section
LDR R0, SB, offset     offset computed by the loader from object's
```

Considering the fact that references to external variables are (or should be) rare, this circumstance is of no great concern.

(Note also that such accesses are read-only). More severe is the fact that we also treat global variables contained in the same module by the same technique. Their level number is 0. One might use a specific base register for the base of the current module. Its content would then have to be reloaded upon every procedure call and after every return. This is common technique, but we have here chosen to reload only when necessary, i.e. only when an access is at hand. This strategy rewards the programmer who sensibly uses global variables rarely.

#### 10.7.10 *Traps*

This compiler provides an extensive system of safeguard by providing run-time checks (aborts) in several cases:

trap number	trap cause
1	array index out of range
2	type guard failure
3	array or string copy overflow
4	access via NIL pointer
5	illegal procedure call
6	integer division by zero
7	assertion violated

These checks are implemented very efficiently in order not to downgrade a program's performance. Involved is typically a single compare instruction, plus a conditional branch (BLR

MT). It is assumed that entry 0 of the module table contain not a base address (module numbers start with 1), but a branch instruction to an appropriate trap routine. The trap number is encoded in bits 4:7 of the branch instruction.

The predefined procedure `Assert` generates a conditional trap with trap number 7. For example, the statement `Assert(m = n)` generates

```
LDR R0, m
LDR R1, n
CMP R0, R0, R1
BLR 1 , 7CH    branch and link if unequal through R12, trap #7
```

Procedure `New`, representing the operator `NEW`, has been implemented with the aid of the trap mechanism. (This is in order to omit in `ORG` any reference to module `Kernel`, which contains the allocation procedure `New`). The generated code for the statement `NEW(p)` is

```
ADD R0, SP, p    address of p
ADD R1, SB, tag  type tag
BLR 7 , 0CH      branch and link unconditionally through R12(MT)
```





## RISC PROCESSOR IMPLEMENTATION

---

### 11.1 INTRODUCTION

The design of the processor to be described here in detail was guided by two intentions. The first was to present an architecture that is distinct in its regularity, minimal in the number of features, yet complete and realistic. It should be ideal to present and explain the main principles of processors. In particular, it should connect the subjects of architectural and compiler design, of hardware and software, which are so closely interconnected.

Clearly “real”, commercial processors are far more complex than the one presented here. We concentrate on the fundamental concepts rather than on their elaboration. We strive for a fair degree of completeness of facilities, but refrain from their “optimization”. In fact, the dominant part of the vast size and complexity of modern processors and software is due to speed-up called optimization. It is the main culprit in obfuscating the basic principles, making them hard, if

not impossible to study. In this light, the choice of a RISC (Reduced Instruction Set Computer) is obvious.

The use of an FPGA provides a substantial amount of freedom for design. Yet, the hardware designer must be much more aware of availability of resources and of limitations than the software developer. Also, timing is a concern that usually does not occur in software, but pops up unavoidably in circuit design. Nowadays circuits are no longer described in terms of elaborate diagrams, but rather as a formal text. This lets circuit and program design appear quite similar. The circuit description language – we here use Verilog – appears almost the same as a programming language. But one must be aware that differences still exist, the main one being that in software we create mostly sequential processes, whereas in hardware everything “runs” concurrently. However, the presence of a language – a textual definition – is an enormous advantage over graphical schemata. Even more so are systems (tools) that compile such texts into circuits, taking over the arduous task of placing components and connecting them (routing). This holds in particular for FPGAs, where components and wires connecting them are limited, and routing is a very difficult and time-consuming matter.

The development of this RISC progressed through several stages. The first was the design of the architecture itself, (more or less) independent of subsequent implementation considerations. Then followed a first implementation called

RISC-o. For this a Harvard Architecture was chosen, implying that two distinct memories are used for program and for data. For both chip-internal block RAMs were used. The Harvard architecture allows for a neat separation of the arithmetic from the control unit.

But these blocks of RAM are relatively small on the used Spartan-3 development board (1 - 4K words). This board, however, provides also an FPGA-external static RAM with a capacity of 1 MByte. In a second effort, the BRAM for data was replaced by this SRAM. Both instructions and data are placed into the SRAM, resulting in a von Neumann architecture.

The RISC hardware is characterized by three interfaces. The 1<sup>st</sup> is the programmer's interface, the architecture, that is, those aspects that are relevant to the programmer, in particular, the instruction set. It is described in Appendix A2. The 2<sup>nd</sup> is the hardware interface between the processor core and its environment, described here. The 3<sup>rd</sup> is that which connects the environment with physical devices such as memory, keyboard and display. This is described in [12](#).

```
module RISC5(
    input clk, rst, stallX,
    input [31:0] inbus, codebus,
    output [19:0] adr, // memory and device addresses
    output rd, wr, ben, // read, write, byte enable
```

```
// control signals for memory
output [31:0] outbus);
```




Figure 45: The processor's interface

The main parts of the hardware interface are three busses, the data input and output busses, the code bus, and the address bus. Signals *rd* and *wr* indicate, whether a read or a write operation is to be performed. *ben* indicates a byte (rather than word) access. The entire processor operates synchronously on the clock *clk* (25 MHz on Spartan-3), *rst* is the reset signal (from a push button on the development board), and *stall* is the input to stall the processor.

First we concentrate on the implementation of the processor core, its realization in the form of circuits. They are divided into two parts, the Arithmetic/Logic Unit(ALU) processing data, and the control unit determining the flow of instructions.

## 11.2 THE ARITHMETIC AND LOGIC UNIT

The ALU features a bank of 16 registers with 32 bit words. Arithmetic and logical operations, represented by instructions, always operate on these registers. Data can be transferred between memory and registers by separate load and store instructions. This is an important characteristic of RISC architectures, developed between 1975 and 1985. It contrasts with the earlier CISC architectures (Complex Instruction Set): Memory is largely decoupled from the processor. A second important characteristic is that most instructions take a single clock cycle (25 MHz) for their execution. The exceptions are access to memory, multiplication and division. More about this will be presented later. This single-cycle rule makes such processors predictable in performance. The number of cycles and the time required for executing any instruction sequence is precisely defined. Predictability is essential in all real-time applications.

The data processing unit consisting of ALU and registers is shown in Figure 46. Evidently, data cycle from registers through the ALU, where an operation is performed, and the result is deposited back into a register. The ALU embodies the circuits for arithmetic operations, logical operations, and shifts. The operations available are listed below. They are described in more detail in Appendix A2. The operand  $n$  is either a register or a part of the instruction itself.

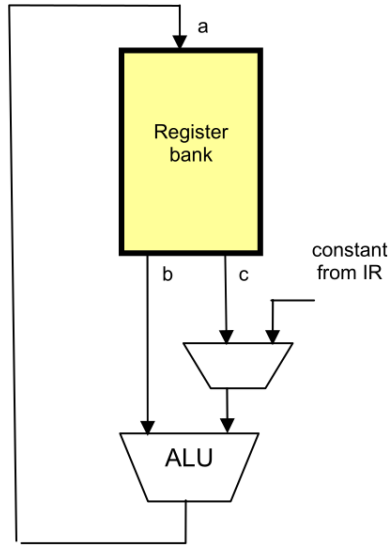


Figure 46: Processor core with ALU and registers

```

0 MOV a,n    R.a:=n
1 LSL a,b,n  R.a:=R.b ← n (shift L by n bits)
2 ASR a,b,n  R.a:=R.b → n (shift R by n bits
                           with sign extension)
3 ROR a,b,n  R.a:=R.b rot n (rotate R by n bits)

4 AND a,b,n  R.a:=R.b & n (logical
5 ANN a,b,n  R.a:=R.b & ~n operations)
6 IOR a,b,n  R.a:=R.b or n (inclusive or)
7 XOR a,b,n  R.a:=R.b xor n (exclusive or)

8 ADD a,b,n  R.a:=R.b + n (
9 SUB a,b,n  R.a:=R.b - n integer
  
```

```

10 MUL a,b,n R.a:=R.b x n arithmetic
11 DIV a,b,n R.a:=R.b div n )

12 FAD a,b,c R.a:=R.b + R.c (
13 FSB a,b,c R.a:=R.b - R.c floating-point
14 FML a,b,c R.a:=R.b x R.c arithmetic
15 FDV a,b,c R.a:=R.b / R.c )

```

The following excerpt describes the essence of the ALU circuits. It is written in the HDL Verilog and refers to the following wires and registers.

```

wire [31:0] IR; // instruction field(s):
wire p, q, u, v, w;
    // IR[31], IR[30], IR[29], IR[28], IR[16]
wire [3:0] op, ira, irb, irc;
    // IR[19:16], IR[27:24], IR[23:20], IR[3:0]
wire [15:0] imm; // IR[15:0]

wire [31:0] A, B, C0, C1, regmux;
wire [31:0] s3, t3, quotient, fsum, fprod, fquot;
wire [32:0] aluRes;
wire [63:0] product;

reg [31:0] R [0:15]; // array of 16 registers
reg N, Z, C, OV; // condition flags

```

$B$  and  $C0$  are the outputs from the register bank, and  $A$  is its input. The register numbers  $ira$  for port  $A$ ,  $irb$  for port  $B$ , and  $irc$  for port  $C0$  are taken from 4-bit fields of the instruction register  $IR$ .  $C1$  is the multiplexer selecting among the register output  $C0$  and the immediate field  $imm$ .  $s3$  and  $t3$  are outputs of the shift units (Sect. 16.2.1).  $product$  is the output of the multiplier (16.2.2),  $quotient$  and remainder those of the divider (16.2.3),  $fsum$  that of the floating-point adder (16.2.4),  $fprod$  that of the floating-point multiplier (16.2.5), and  $fquot$  the output of the floating-point divider (16.2.6).

```

assign A  = R[ira];
assign B  = R[irb];
assign C0 = R[irc];
assign C1 = q ? {{16{v}}}, imm} : C0;

```

The following represents the main instruction decoding and selection of results. The opcodes refer to specific values of fields  $p$  and  $op$  of  $IR$ . Note that if  $x$  then  $y$  else  $z$  is denoted in Verilog by  $x?y:z$ .

```

assign aluRes = MOV
    ? (q ? (~u ? {{16{v}}}, imm)
        : {imm, 16'b0})
    : (~u ? C0
        : (~irc[0] ? H

```



```

                                : {N, Z, C, OV, 20'b0,
                                8'b01010000}})))
:  LSL  ? t3           // L-shift unit output
:  (ASR
   |ROR) ? s3           // R-shift unit output
:  AND  ? B & C1
:  ANN  ? B & ~C1
:  IOR  ? B | C1
:  XOR  ? B ^ C1
:  ADD  ? B + C1 + (u & C)
:  SUB  ? B - C1 - (u & C)
:  MUL  ? product [31:0] // multiplier output
:  DIV  ? quotient
:  (FAD
   |FSB) ? fsum
:  FML  ? fprod
:  FDV  ? fquot : 0;

```

The input to the register bank, *regmux*, is selected from either *aluRes*, *inbus* (for LDR instructions), or the program address *nipc* (for branch and link instructions). The signal *regwr* determines, whether data are to be stored (written) into the register bank. Details must be gathered from the respective program listing RISC.v.

```

always @ (posedge clk) begin
    R[ira] <= regwr      ? regmux      : A ;

```

```

N      <= regwr      ? regmux[31] : N ;
Z      <= regwr      ? (regmux==0) : Z ;
C      <= (ADD|SUB) ? aluRes[32] : C ;
OV     <= (ADD|SUB) ? aluRes[32]
                        ^ aluRes[31] : OV;
end

```

Whenever a register is written, the condition flags are also affected. They are *N* (*aluRes* negative), *Z* (*aluRes* zero), *C* (carry), and *OV* (overflow). The latter apply only to addition and subtraction.

#### 11.2.1 Shifters

Shifters are multi-way multiplexers. For a 32-bit word, the simplest solution would be 32 32-way multiplexers. But this is hardly economical. On the FPGA used here, 4-way muxes are basic cells. It is therefore beneficial, to compose a shifter out of 4-way muxes. Now the obvious solution is to use 3 levels of muxes through which data flow. The first level shifts by amounts of 0, 1, 2, or 3, the second by amounts of 0, 4, 8, 12, and the third by 0 or 16. This scheme is programmed as follows for left shifts (instruction LSL) with *B* as input, *sc0* = *C1*[1 : 0] and *sc1* = *C1*[3 : 2] as shift counts, and *t3* as output:

```

assign t1 = (sc0 == 3) ? { B[28:0], 3'b0} :

```

```

(sc0 == 2) ? { B[29:0], 2'b0} :
(sc0 == 1) ? { B[30:0], 1'b0} : B ;
assign t2 = (sc1 == 3) ? {t1[19:0], 12'b0} :
(sc1 == 2) ? {t1[23:0], 8'b0} :
(sc1 == 1) ? {t1[27:0], 4'b0} : t1;
assign t3 = C1[4] ? {t2[15:0], 16'b0} : t2;

```

The solution for right shifts is analogous. An additional level of multiplexing is required, shifting in either the sign bit (ASR with sign propagation) or bits from the low end of the word (ROR), making a barrel shifter. This selection is controlled by the instruction bit  $w = IR[16]$ .

### 11.2.2 *Multiplication*

Multiplication is an inherently more complex operation than addition and subtraction. After all, multiplication can be composed (of a sequence) of additions. There are many methods to implement multiplication, all – of course – based on the same concept of a series of additions. They show the fundamental problem of trade-off between time and space (circuitry). Some solutions operate with a minimum of circuitry, namely a single adder used for all 32 additions executed sequentially (in time). They obviously sacrifice speed. The other extreme is multiplication in a single cycle, using 32 adders in series (in space). This solution is fast, but the amount of required circuitry is high.

Before we present the sequential solution, let us briefly recapitulate the basics of a multiplication  $p := xy$ . Here  $p$  is the product,  $x$  the multiplier, and  $y$  the multiplicand. Let  $x$  and  $y$  be unsigned integers. Consider  $x$  in binary form.

$$x = x_{31}2^{31} + x_{30}2^{30} + \dots + x_12^1 + x_02^0$$

Evidently, the product is the sum of 32 terms of the form  $x_k2^k y$ , i.e. of  $y$  left shifted by  $k$  positions multiplied by  $x_k$ . Since  $x_k$  is either 0 or 1, the product is either 0 or  $y$  (shifted). Multiplication is thus performed by an adder and a selector. The selector is controlled by  $x_k$ , a bit of the multiplier. Instead of selecting this bit among  $x_0 \dots x_{31}$ , we right shift  $x$  by one bit in each step. Then the selection is always according to  $x_0$ . The add-shift step then is

```
IF ODD(x) THEN p := p + y END;
y := 2*y; x := x DIV 2
```

whereby multiplication by 2 is done by a left shift, and division by 2 by a right shift: As an example, consider the multiplication of two 4-bit integers  $x = 5$  and  $y = 3$ , requiring 4 steps:

The shifting of  $x$  to the right also suggests that instead of shifting  $y$  to the left in each step, we keep  $y$  in the same position and shift the partial sum  $p$  to the right. We notice

	p	x	y
	0000'0000	0101	0000'0011
add y to p	0000'0011	0101	0000'0011
shift	0000'0011	0010	0000'0110
add o to p	0000'0011	0010	0000'0110
shift	0000'0011	0001	0000'1100
add y to p	0000'1111	0001	0000'1100
shift	0000'1111	0000	0001'1000
add o to p	0000'1111	0000	0001'1000
shift	0000'1111	0000	0011'0000
			p = 15

that the size of  $x$  decreases by 1 in each step, whereas the size of  $p$  increases by 1. This allows to pack  $p$  and  $x$  into a single double register  $\langle B, A \rangle$  with a shifting border line. At the end, it contains the product  $p = xy$ .

	p	x
	0000	0101
add y to p	0011	0101
shift	00011	010
add o to p	00011	010
shift	000011	01
add y to p	001111	01
shift	0001111	0
add o to p	0001111	0
shift	00001111	
		p = 15

$$p = \{B[31:0], A\{31:[32-k]\}\},$$



advances as long as the input `MUL` is active (high). `MUL` indicates that the current operation is a multiplication, and the signal is stable until the processor advances to the next instruction. This happens when step 31 is reached (Figure 48).




Figure 48: Generating *stall*

The details of the simple multiplier are listed below:

```
module Multiplier(
    input CLK, MUL, u,
    output stall,
    input [31:0] x, y,
    output [63:0] z);

    reg [4:0] S; // state
    reg [31:0] B2, A2; // high and low parts of partial product
    wire [32:0] B0, B00, B01;
    wire [31:0] B1, A0, A1;

    assign stall = MUL & ~(S == 31);
    assign B00 = (S == 0) ? 0 : {B2[31] & u, B2};
```

```
assign B01 = A0[0] ? {y[31] & u, y} : 0;
assign B0 = ((S == 31) & u) ? B00 - B01 : B00 + B01;
assign B1 = B0[32:1];
assign A0 = (S == 0) ? x : A2;
assign A1 = {B0[0], A0[31:1]};
assign z = {B1, A1};

always @ (posedge(CLK)) begin
    B2 <= B1; A2 <= A1;
    S <= MUL ? S+1 : 0;
end
endmodule
```

Implementing multiplication in hardware made the operation about 30 times faster than its solution by software. A significant factor! As multiplication is a relatively rare operation – at least in comparison with addition and subtraction – early RISC designs (MIPS, SPARC, ARM) refrained from its full implementation in hardware. Instead, an instruction called multiply step was provided, performing a single add-shift step in one clock cycle. A multiplication was then programmed by a sequence of 32 step instructions, typically provided as a subroutine. This measure of economy was abandoned, when hardware became faster and cheaper.

The FPGA used on the Spartan-3 board features a welcome facility for speeding up multiplication, namely fast 18 x 18 bit multiplier units. These are made available as basic



cells of the FPGA, and they multiply in a single clock cycle. Considering an operand  $x = x_12^{16} + x_0$ , the product is obtained as the sum of only 4 terms:

$$p = xy = x_1y_12^{32} + (x_0y_1 + x_1y_0)2^{16} + x_0y_0$$

Thereby multiplication of two 32-bit integers can be performed in 2 cycles only, one for multiplications, one for addition. Four multipliers are needed. For details, the reader is referred to the program listing (module Multiplier1).

### 11.2.3 *Division*

Division is similar to multiplication in structure, but slightly more complicated. We present its implementation by a sequence of 32 shift-subtract steps, the complement of add-shift. We here discuss division of unsigned integers only.

$$q = x \text{ DIV } y$$

$$r = x \text{ MOD } y$$

$q$  is the quotient,  $r$  the remainder. These are defined by the invariants

$$x = qy + r \text{ with } 0 \leq r < y$$

Both  $q$  and  $r$  are held in registers. Initially we set  $r$  to  $x$ , the dividend, and then subtract multiples of  $y$  (the divisor) from it, each time checking that the result is not negative. This shift-subtract step is

```

r := 2*r; q := 2*q;
IF r - y >= 0 THEN r := r - y END

```

As an example, consider the division of the 8-bit integer  $x = 14$  by the 4-bit integer  $y = 4$ , where multiplication and division by 2 are done by shifts:

	r	q	y	
	0000'1110	0000	0001'1000	
shift	0000'1110	0000	0001'1000	$r < y$
sub 0 from r	0000'1110	0000	0000'1100	
shift	0000'1110	0000	0000'1100	$r \geq y$
sub y from r	0000'0010	0001	0000'1100	
shift	0000'0010	0010	0000'0110	$r < y$
sub 0 from r	0000'0010	0010	0000'0110	
shift	0000'0010	0100	0000'0011	$r < y$
sub y from r	0000'0010	0100	0000'0011	$q = 4, r = 2$

As with multiplication this arrangement may be simplified by putting  $r$  and  $q$  into a double-length shift register, and by shifting  $r$  to the left instead of  $y$  to the right. This results in

This scheme is represented by the circuit shown in Figure 49.

	r	q
	0000'1110	
shift	0001'110	0 $r < Y$
sub o from r	0001'110	0
shift	0011'10	00 $r \geq Y$
sub y from r	0000'10	01
shift	0001'0	010 $r < Y$
sub o from r	0001'0	010
shift	0010	0100 $r < Y$
sub o from r	0010	0100 $q = 4, r = 2$

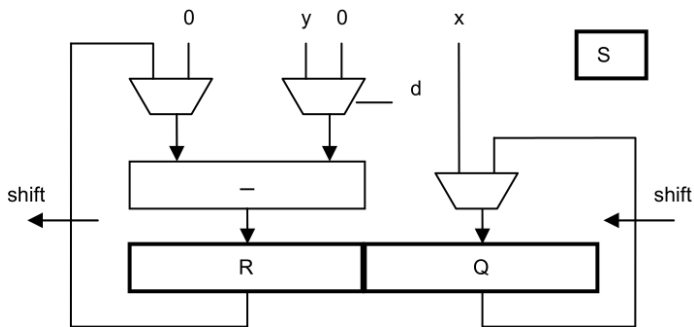


Figure 49: Schematic of divider

Stall generation is the same as for the multiplier. A division takes 32 clock cycles. Further details are shown in the subsequent program listing.

```

module Divider(
    input clk, DIV,
    output stall,

```

```

    input [31:0] x, y,
    output [31:0] quot, rem);

reg [4:0] S; // state
reg [31:0] r3, q2;
wire [31:0] r0, r1, r2, q0, q1, d;

assign stall = DIV & ~(S == 31);
assign r0 = (S == 0) ? 0 : r3;
assign d = r1 - y;
assign r1 = {r0[30:0], q0[31]};
assign r2 = d[31] ? r1 : d;
assign q0 = (S == 0) ? x : q2;
assign q1 = {q0[30:0], ~d[31]};
assign rem = r2;
assign quot = q1;

always @ (posedge(clk)) begin
    r3 <= r2; q2 <= q1;
    S <= DIV ? S+1 : 0;
end
endmodule

```

### 11.3 FLOATING-POINT ARITHMETIC

The RISC uses the IEEE Standard for representing REAL (floating-point) numbers with 32 bits. The word is divided

into 3 fields: s for the sign, e for the exponent, and m for the mantissa. The value is

$$x = (-1)^s 2^{-127} 1.m \text{ with } 1.0 \leq m < 2.0 \text{ (normalized form)}$$

Numbers are represented in sign-magnitude form. This implies that for sign inversion only the sign bit must be inverted, and exponent and mantissa remain unchanged.

Zero is a special case represented by 32 0-bits, and therefore has to be treated separately. Furthermore, e = 255 denotes "not a number". It is generated in the case of arithmetic overflow.




Figure 50: IEEE standard floating-point representation of REAL numbers

### 11.3.1 Floating-point addition

If two numbers are to be added, they must have the same exponent. This implies that the summand with the smaller exponent must be denormalized. m is shifted to the right and e is incremented accordingly. That is, if d is the difference of the two exponents, m is multiplied by 2<sup>d</sup>, and e is incremented by d. After the addition, the sum must be

rounded and post-normalized.  $m$  is shifted to the left and  $e$  is decremented accordingly. The shift amount is determined by the position of the leftmost one-bit. This results in the scheme shown in Figure 16.7, and the module's interface is

```
module FPAdder(  
    input clk, run, u, v,  
    input [31:0] x, y,  
    output stall,  
    output [31:0] z);
```

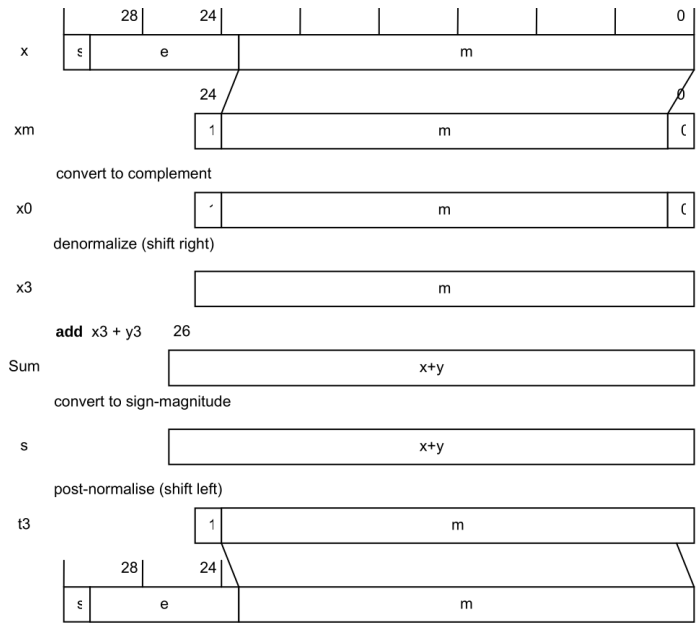


Figure 51: Steps of floating-point addition

It is important to achieve proper rounding. This is done by extending the mantissa of both operands by a guard bit, initialized to 0. A one is added (effectively 0.5) and at the end the guard bit is discarded.

The two predefined conversion functions FLT and FLOOR are conveniently implemented as additions. A denormalized 0 is added to the argument, effecting the proper shift. In the case of FLT (modifier bit  $u = 1$ ), denormalization is omitted (no 1-bit inserted), and in the case of FLOOR (modifier bit  $v = 1$ ), post-normalization is suppressed.

### 11.3.2 *Floating-point multiplication*

A product is given by the equation

$$p = xy = (2^{x_e} x_m)(2^{y_e} y_m) = 2^{x_e + y_e} (x_m * y_m)$$

$$p = (x_s, x_e, x_m)(y_s, y_e, y_m) = (x_s \text{ xor } y_s, x_e + y_e, x_m y_m)$$

That is, exponents are added, mantissas multiplied. Denormalization is not needed. Post-normalization is a right shift of at most one bit, because if  $1.0 \leq x_m, y_m < 2.0$ , the result satisfies  $1.0 \leq x_m * y_m < 4.0$ . The sign of the product is the exclusive or of the signs of the arguments. The multiplier module's interface is

```
module FPMultiplier(
```

```

input clk, run,
input [31:0] x, y,
output stall,
output [31:0] z);

```

### 11.3.3 Floating-point division

A quotient is given by the equation

```

module FPDivider(
    input clk, run,
    input [31:0] x, y,
    output stall,
    output [31:0] z);

```

## 11.4 THE CONTROL UNIT

$$q = x/y = (2^{x_e} x_m) / (2^{y_e} y_m) = 2^{x_e - y_e} (x_m / y_m)$$

$$q = (sx, ex, mx) / (sy, ey, my) = (sx \text{ xor } sy, ex - ey, mx / my)$$

That is, exponents are subtracted, mantissas divided. Denormalization is not needed. Post-normalization requires a left shift by at most a single bit, because if  $1.0 \leq x_m, y_m < 2.0$ , the result satisfies  $0.5 \leq x_m / y_m < 2.0$ . The sign of the product is the exclusive or of the signs of the arguments. The divider module's interfaces is



The control unit determines the sequence of executed instructions. It contains two registers, the program counter PC holding the address of the current instruction, and the current instruction register IR holding the instruction currently being interpreted. Instructions are obtained from memory through the codebus (see interface), from where the decoding signals emanate. Mostly, the arithmetic unit and the control unit operate concurrently (in parallel). While the arithmetic unit performs the operation held in register IR and data signals flow through the ALU, the control unit fetches in the same clock cycle the next instruction from memory in the location with the address held in PC. Next address and next instruction are latched in the registers at the end of a cycle. This scheme constitutes a one-element pipeline of instructions.

The principal task of the control unit is to generate the address of the next instruction. There are essentially only four cases:

1. Zero on reset.
2. The next instructions address is  $PC+1$  (all instructions except branches)
3. The branch target  $PC+1 + \text{offset}$ . (Branch instructions).
4. It is taken from a data register. (This is used for returning from procedures).

This is reflected by the following program text, and shown in Figure 52.

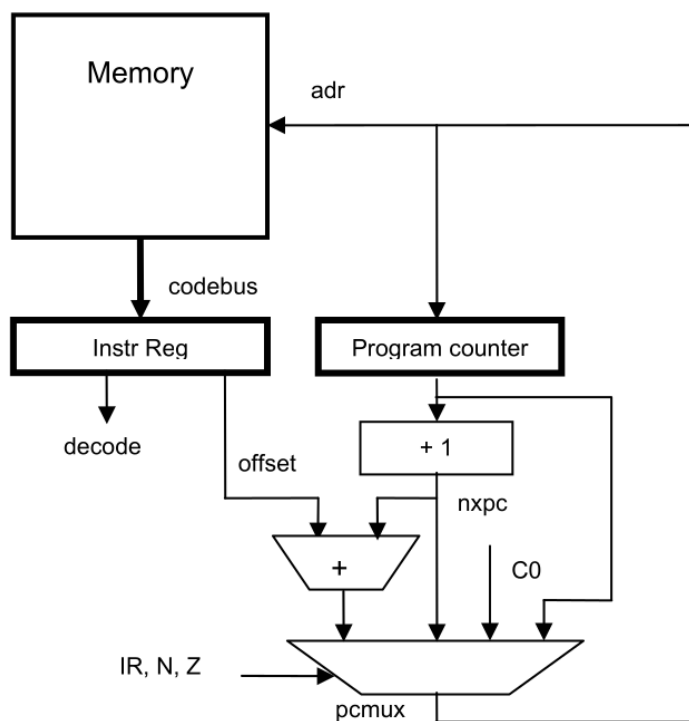


Figure 52: The control unit

```

reg [17:0] PC;
reg [31:0] IRBuf;
wire [31:0] IR;
wire [31:0] pmout;
wire [17:0] pcmux, nxc;
wire cond;

```

```

IR = codebus;
nxcpc = PC + 1;
pcmux = (~rst) ? 0 :
    (stall) ? PC : // stall
    (BR & cond & u) ? off + nxcpc :
    (BR & cond & ~u) ? C0[19:2] :
    nxcpc;

always @ (posedge clk) PC <= pcmux; end

```

Branches are the only conditional instructions. Whether a branch is taken or not, is determined by the combination of the condition flags selected by the condition code field of the branch instruction. IR[27] is the condition sense inversion bit.

```

reg N, Z, C, OV; // condition flags
wire S;
assign S = N ^ OV;
assign cond = IR[27] ^
    ((cc == 0) & N | // MI, PL
    (cc == 1) & Z | // EQ, NE
    (cc == 2) & C | // CS, CC
    (cc == 3) & OV | // VS, VC
    (cc == 4) & (C|Z) | // LS, HI
    (cc == 5) & S | // LT, GE
    (cc == 6) & (S|Z) | // LE, GT
    (cc == 7)); // T, F

```

There is, unfortunately, a complication obfuscating the simple scheme presented so far. It stems from the necessity to initialize the processor. Only registers and memory blocks (BRAM) can be initialized and loaded by the available FPGA-tools. How, then, is a program (in our case the boot loader) moved into memory, the chip-external SRAM? The following scheme has been chosen:

The initial program is loaded into a BRAM (1K x 32). This block is memory-mapped into high-end addresses in the range of the data stack. On startup, the flag PMsel is set and IR is loaded from pmout (from the BRAM) at StartAdr. At the end of the program (boot loader), a branch instruction with destination 0 jumps to the beginning of the program that had just been loaded into SRAM by the boot loader. This is, presumably, but not necessarily, the operating system. The following changes and additions are required:

```
localparam StartAdr = 18'b1111111000000000000; // 0FE000H
```

```
reg PMsel; // memory select for instruction fetch
```

```
reg [31:0] IRBuf;
```

```
dbram32 PM ( // BRAM
```

```
    .clk (clk),
```

```
    .rdb (pmout), // output port
```

```
    .ab (pcmux[10:0])); // address
```

```
assign IR = PMsel ? pmout : IRBuf;

always @ (posedge clk) begin
    PMsel <= ~rst | (pcmux[17:11] == 7'b1111111);
    IRBuf <= stall ? IRBuf : codebus;
    ...
end;
```



## PROCESSOR'S ENVIRONMENT

---

The RISC processor is embedded in an environment (module RISCTop.v) connecting it with elements that are FPGA-chip external, but whose are provided on the Spartan development board (Figure 53). The environment consists of an address decoder, a data multiplexer, and interfaces to the memory and peripheral devices.

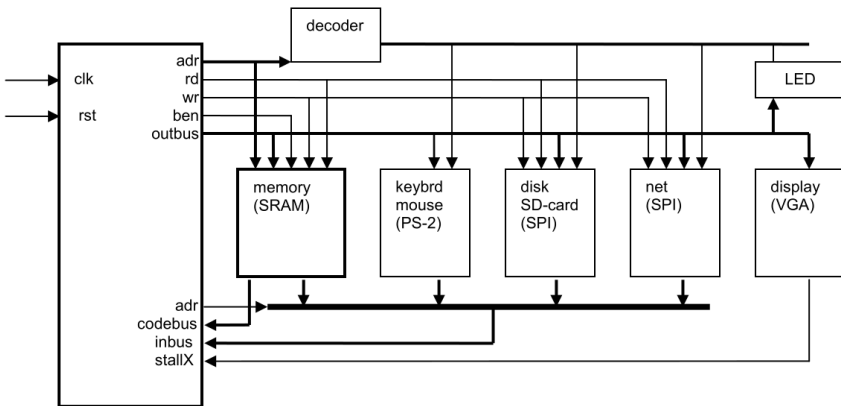


Figure 53: The RISC configuration

The decoder for output and the multiplexer for input determine the various addresses of devices:

	adr	input	output
0	oFFFFFFFCoH	millisecond counter	reserved
4	oFFFFFFFC <sub>4</sub> H	switches	LEDs
8	oFFFFFFFC <sub>8</sub> H	RS-232 data	RS-232 data
12	oFFFFFFFCCH	RS-232 status	RS-232 control
16	oFFFFFFFD <sub>0</sub> H	SPI data(SD-card,net)	SPI data(SD...)
20	oFFFFFFFD <sub>4</sub> H	SPI status	SPI control
24	oFFFFFFFD <sub>8</sub> H	PS/2 keyboard	
28	oFFFFFFFDCH	mouse	

The circuitry connecting with the SRAM is part of this module, whereas the drivers for the other devices are described in separate modules. Note: The signals to and from devices must be listed in the heading of the top module, which is not imported by any other module. Their pin numbers are specified in a configuration file (.ucf). For details, the reader is referred to the program listing, as several items are rather dependent on the given Spartan-3 board.

### 12.1 THE SRAM MEMORY

The design of the circuitry around a static RAM is quite straight forward. The only controls are a read (SRoe) and a write enable signal (SRwe). Since the SRAM multiplexes data lines for input and output, a tri-state driver (SRbuf) must be used on the FPGA. This is shown schematically in Figure 54.



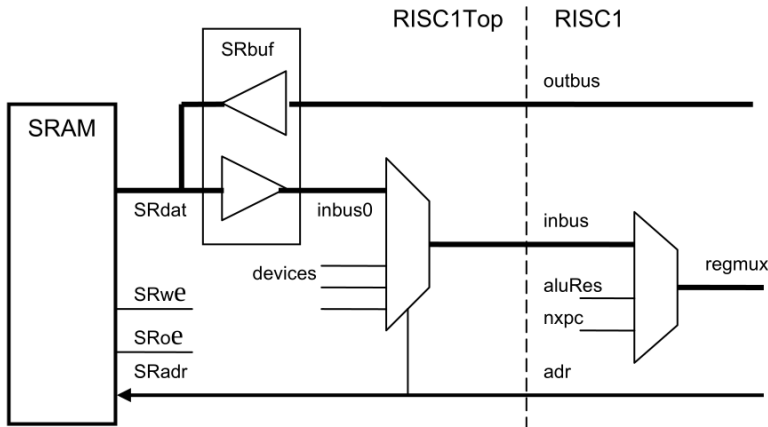


Figure 54: Connections between processor and SRAM

However, there is a complication: the feature of byte-wise access. After all, the present RAM is 32 bits wide (actually there are two 512K x 16-bit chips in parallel). Evidently, some multiplexing is unavoidable. The task is significantly eased by the chip's feature of four separate write enables, one for each byte of a word. The selection of the byte affected is determined by address bits 0 and 1 (which are ignored in the case of word-access). This scheme is shown in Figure 55. The codebus bypasses the multiplexers.

## 12.2 PERIPHERAL INTERFACES

Each of the interfaces to external media is implemented as a separate module and can therefore easily be exchanged. Modules are connected with the processor by the input and the output bus, and by enable signals wr and rd.




Figure 55: Multiplexers for SRAM byte access

#### 12.2.1 *The PS/2 interface for the keyboard*

PS/2 is mostly used for input devices. It uses 2 wires (apart from ground), one for data, one for the clock. It uses a synchronous transmission, and the clock is driven by the device. Here it is used for the keyboard and the mouse (see Sect. 17.2.5). Transmission occurs in packets of 8 bits. An optional third wire serves for output. It is not used in this application. The interface is very simple and consists of an 8-bit buffer register. The following describes the interface for the keyboard.

A bit is shifted into the data register whenever the clock shows a falling edge, i.e. the clock signal  $Q_0$  is low and the clock delayed by one cycle  $Q_1$  is high.

In the driver for the keyboard a 16-byte fifo buffer is inserted, forming a queue. This is necessary in order to avoid loss of characters when the processor is tied up in computation.




Figure 56: The PS/2 configuration

```

module PS2(
    input clk, rst,
    input done, // "byte has been read"
    output rdy, // "byte is available"
    output shift, // shift in, transmitter
    output [7:0] data,
    input PS2C, // serial input
    input PS2D); // clock

    reg Q0, Q1; // synchronizer and falling edge detector
    reg [10:0] shreg;
    reg [3:0] inptr, outptr;
    reg [7:0] fifo [15:0]; // 16 byte buffer
    wire endbit;

    assign endbit = ~shreg[0]; //start bit reached correct pos
    assign shift = Q1 & ~Q0;
    assign data = fifo[outptr];
    assign rdy = ~(inptr == outptr);

    always @ (posedge clk) begin
        Q0 <= PS2C; Q1 <= Q0;
    
```

```

shreg <= (\~rst | endbit) ? 11'h7FF :
shift ? {PS2D, shreg[10:1]} : shreg;
outptr <= \~rst ? 0 : rdy \& done ? outptr+1 : outptr;
inptr <= \~rst ? 0 : endbit ? inptr+1 : inptr;
if (endbit) fifo[inptr] <= shreg[8:1];
end
endmodule

```

### 12.2.2 *The Mouse*

Subsequently we present two Mouse interfaces. The first (MouseP) is based on the PS/2 Standard and caters for most commercially available mice. The second (MouseX) is included here for historical reasons. It was used by the computer Lilith in 1979, and used the same Mouse as its ancestor Alto (at PARC, 1975). It is distinguished by a very simple hardware without its own microprocessor, which is currently contained in most mice. This goes at a cost of a 9-wire cable. But today, microprocessors are cheaper than cables. We include this interface here, because it allows for a simple explanation of the principle of pointing devices.

The first interface uses the PS/2 Standard, that is, a 2-wire cable (not counting ground and power). It complies with the commercial standard of pointing devices. Details are shown on module MouseP.v.

```

module MouseP (input rst, clk,

```

```
inout PS2C, PS2D,  
output [27:0] out);  
endmodule
```

The second interface described here is not based on any standard, but it features the same interface to the software environment. Its principles are very simple and easily explained, and it refrains from the use of a mouse-internal processor. The price for this simplicity is a cable with 7 wires (plus 2 for power and ground), namely 3 for 3 buttons, and 2 for each direction,  $x$  (left/right) and  $y$  (up/down).

Let us first explain how signals indicating movements are derived. The key reason for the solution's simplicity is that these signals are directly mirrored by the position of a cursor on the display. The human user simply moves the Mouse until the cursor has reached the desired position (for example, at a displayed object). Thereby, the human eye and hand are included in the feedback loop providing the desired precision. This represents a very clever symbiosis between man and computer.

An actual movement is recognized by a simple light sensor (we will restrict our observation to a single coordinate  $x$ ). The movement is transmitted to a wheel consisting of a transparent disc with intransparent spokes. A light beam shines through the disk and is received by the light sensor. Each time a spoke passes, the light is blocked. Any change

in the sensor output signals a movement (see Figure 57). Unfortunately, this scheme does not allow to recognize the direction of the movement (left or right). A second light and sensor solve this problem. The distance between the two lights is half the distance of adjacent spokes.




Figure 57: Wheel with spokes and sensors

The signal pair  $x_0$ ,  $x_1$  originating from a movement (with constant speed) to the left or to the right is shown in Figure 58.




Figure 58: Signals resulting from movements

The logic equations for movements to the left and right (or up and down) are derived from this signal pair. For each signal a register records the state. Therefore it can be determined whether a move to the left, or to the right, or no move had occurred. The sampling frequency is irrelevant,

as long as it is high enough. Let  $x_{01}$  be  $x_{00}$  delayed by one clock cycle, and  $x_{11}$  be  $x_{10}$  delayed by one cycle.

$x_{00}$	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1
$x_{01}$	0	0	0	0	1	1	1	1	0	0	0	0	1	1	1	1
$x_{10}$	0	0	1	1	0	0	1	1	0	0	1	1	0	0	1	1
$x_{11}$	0	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
right	0	1	0	0	0	0	0	1	1	0	0	0	0	0	1	0
left	0	0	1	0	1	0	0	0	0	0	0	1	0	1	0	0

Every active right signal causes the 10-bit  $x$  counter to be incremented, and every left to be decremented.

An identical circuit is used for the up/down direction, with its wheel set perpendicular to the first wheel. Finally, the output is packed into a single word. 3 bits are taken by the keys, and 10 by each of the two counters.

```

module MouseX(
    input clk,
    input [6:0] in,
    output [27:0] out);

    reg x00, x01, x10, x11, y00, y01, y10, y11;
    reg ML, MM, MR; // keys
    reg [9:0] x, y; // counters

    wire xup, xdn, yup, ydn;

```

```

assign xup = ~x00&~x01&~x10&x11 | ~x00&x01&x10&~x11;
assign yup = ~y00&~y01&~y10&y11 | ~y00&y01&y10&~y11;
assign xdn = ~x00&~x01&x10&~x11 | ~x00&x01&~x10&x11;
assign ydn = ~y00&~y01&y10&~y11 | ~y00&y01&~y10&y11;
assign out = {1'b0, ML, MM, MR, 2'b0, y, 2'b0, x};

always @ (posedge clk) begin
    x00 <= in[3]; x01 <= x00; x10 <= in[2]; x11 <= x10;
    y00 <= in[1]; y01 <= y00; y10 <= in[0]; y11 <= y10;
    MR <= ~in[4]; MM <= ~in[5]; ML <= ~in[6];
    x <= xup ? x+1 : xdn ? x-1 : x;
    y <= yup ? y+1 : ydn ? y-1 : y;
end
endmodule

```

### 12.2.3 *The SPI interface for the SD-card (disk) and the Net*

SPI (Standard Peripheral Interface) is similar to PS/2, and also synchronous. However, there may be many participants. They are configured in a loop as shown in Figure 59, and the clock is provided by a master, namely the RISC. SPI requires 3 wires (apart from ground).

Here, however, no use is made of SPI's ring topology. Instead, One master interface is serving both the disk and the net. The connection is determined in module RISC5Top. The packet (and thus the shift register) is 32 bits long






Figure 59: SPI-configuration as a ring

Transmission frequency is 0.4 MHz at startup (as required by the SD-card), and then is raised to 8.33 MHz. Details are shown in the respective program listing.

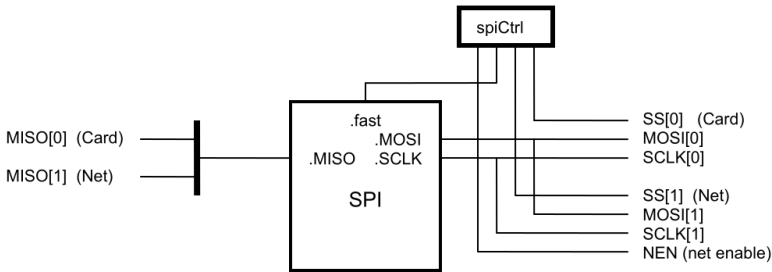


Figure 60: Connections between SPI, SD-card, and Net  
(see RISCTop5.v)

// Motorola Serial Peripheral Interface (SPI) PDR 23.3.12 /  
16.10.13 // transmitter / receiver of words (fast, clk/3) or  
bytes (slow, clk/64) // e.g 8.33MHz or 400kHz respectively  
at 25MHz (slow needed for SD-card init) // note: bytes are  
always MSbit first; but if fast, words are LSByte first

```
module SPI(
  input clk, rst,
  input start, fast,
```

```

input [31:0] dataTx,
output [31:0] dataRx,
output reg rdy,
input MISO,
output MOSI, SCLK);
endmodule

```

The SPI specifications postulate that bytes are sent with the most significant bit first. This results in a somewhat twisted scheme for shifting bits (see Figure 61).




Figure 61: Shifting with MSB first

#### 12.2.4 The display controller

A controller for a raster scan display feeds data from memory to the display. The data area in memory is called frame buffer. It contains a fixed number of bits for each pixel on the screen. In this case, there is exactly one bit per pixel, signalling black or white. For a 1024 x 768 pixel display area, 96 Kbyte are required.

The pixel position on the display is not determined by an address. Instead, data are received by the display purely sequentially, and the position is indirectly determined by

two synchronization signals, hsync (for horizontal sync) at the end of each line, and vsync (for vertical sync) at the end of every frame. This scheme originates from cathode ray tube (CRT) monitors, where an electron beam is sweeping the screen. It is deflected by magnetic fields, which require some time to sweep back. The timing with retrace periods was retained for LCD displays as a legacy.

The heart of the controller consists of a data buffer (32 bits) fed from memory and shifted out bit by bit to the display, and of two counters hcnt and vcnt, representing the horizontal and vertical coordinates. The memory word address is derived from hcnt and vcnt:

$$\text{vidadr} = (\text{hcnt} \text{ DIV } 32) + (\text{vcnt} * 32) + \text{org}$$

Every line consists of 1024 pixels (32 words). The challenge is to find a design with as few registers and comparators as possible. There are two signals for suppressing video data: hblank, vblank. They are needed for turning the light off during retrace.




Figure 62: Synchronization and blanking signals

Let us generalize this scheme to displays of  $w$  pixels per line and  $h$  lines per frame. Also, let  $w'$  be the number of pixels per line including those of the retrace time, and  $h'$  be the number of lines including the vertical retrace. Also, let the number of displayed frames per second be  $n$ . Then the pixel frequency is

$$f = w' \times h' \times n.$$

This will in all probability be different from the system clock's frequency. Therefore the need arises for a different pixel clock. It is generated by the FPGA's built-in digital clock manager (dcm). It multiplies and divides the system clock by selectable factors. Note that the refresh rate may vary within certain bounds for all brands of monitors. Therefore, a simple factor may be chosen for division and multiplication. Examples:

$$(1024 \times 768) (1280 \times 1024)$$

$$1182 \times 791 \times 60 = 56'097'720 \quad 1536 \times 1280 \times 60 = 117'964'800$$

rounded up to rounded up to

$$60 \text{ MHz } 125 \text{ MHz}$$

The pixel buffer is fed from the video buffer driven by the system clock, and it is shifted and read by the pixel clock. This makes a double-buffering necessary, as shown in Figure 63. Also the counters are driven by this pixel clock. The

numbers for `hcnt` and `vcnt` shown are, of course, device-specific (see Figure 62).




Figure 63: Buffering the video output

```

module VID(
    input clk, clk25, inv,
    input [31:0] viddata,
    output reg req, // read request
    output hsync, vsync, // to display
    output [17:0] vidadr,
    output [2:0] RGB);
localparam Org = 18'b1101\_1111\_1111\_0000\_00; // DFF00
reg [9:0] hcnt, vcnt;
reg [4:0] hword; // from hcnt, but latched in the clk domain
reg [31:0] vidbuf, pixbuf;
reg hblank;
endmodule
  
```

Both the display controller and the processor access memory directly. It therefore becomes necessary to arbitrate in the case where both require access simultaneously, that is, to decide which has priority. The decision is simple, because the

display controller is time-critical and must not be delayed. The processor, on the other hand, can easily be delayed by the already present stalling scheme. The signal (wire) `dspreq` stalls the processor (`stallX`) and decides whether the memory address (`SRadr`) should be taken from the processor (`adr`) or the display controller (`vidadr`). The following multiplexer is placed in module `RISCTop`:

```
assign SRadr = dspreq ? vidadr : adr[19:2];
```

#### 12.2.5 *The RS-232 interface*

RS-232 is an old standard for serial data transmission (see also Sect. 9.4). We chose to describe it here in detail because of its frequent use and inherent simplicity. RS-232 uses 2 wires (apart from ground), one for input (`RxD`) and one for output (`TxD`) as shown in Figure 64. Data are transmitted in packets of a fixed length, here of length 8, i.e. byte-wise. Since there is no clock wire, bytes are transmitted asyn-




Figure 64: RS-232 configuration

chronously. Their beginning is marked by a start-bit, and at the end a stop-bit is appended. Hence, a packet is 10 bits

long (see Figure 65). Within a packet, transmission is synchronous, i.e. with a fixed clock rate, on which transmitter and receiver agree. The packet length is short enough to admit slight deviations. The standard defines several packet lengths and many clock rates. Here we use a rate of 19200 or 115200 bit/s.




Figure 65: RS-232 packet format

The input signal start triggers the state machine by setting register run. The transmitter has 2 counters and a shift register. Counter tick runs from 0 to 1302, yielding a frequency of  $25'000 / 1302 = 19.2$  KHz, the transmission rate for bits. The signal endtick advances counter bitcnt, running from 0 to 9 (the number of bits in a packet). Signal endbit resets run and the counter to 0. Signal rdy indicates whether or not a next byte can be loaded and sent.

```
module RS232T(
    input clk, rst, // system clock, 25 MHz
    input start, // request to accept and send a byte
    input [7:0] data,
```

The receiver is structured very similarly with 2 counters and a shift register. The state machine is triggered by an incoming start bit at RxD. The state rdy is set when the last data bit has been received, and it is reset by the done signal, generated when reading a byte. The line RxD is sampled in the middle of the bit period rather than at the end, namely when  $\text{midtick} = \text{endtick}/2$ .

```

    output rdy, // status
    output TxD); // serial data

wire endtick, endbit;
reg run;
reg [11:0] tick;
reg [3:0] bitcnt;
reg [8:0] shreg;

assign endtick = tick == 1302;
assign endbit = bitcnt == 9;
assign rdy = ~run;
assign TxD = shreg[0];

always @ (posedge clk) begin
    run <= (~rst | endtick & endbit) ? 0 : start ? 1 : run;
    tick <= (run & ~endtick) ? tick + 1 : 0;
    bitcnt <= (endtick & ~endbit) ? bitcnt + 1 :
    (endtick & endbit) ? 0 : bitcnt;
    shreg <= (~rst) ? 1 : start ? {data, 1'b0} :

```



```

    endtick ? {1'b1, shreg[8:1]} : shreg;
end
endmodule

```

```

module RS232R(
    input clk, rst,
    input done, // "byte has been read"
    input RxD,
    output rdy,
    output [7:0] data);

wire endtick, midtick;
reg run, stat;
reg [11:0] tick;
reg [3:0] bitcnt;
reg [7:0] shreg;

assign endtick = tick == 1302;
assign midtick = tick == 651;
assign endbit = bitcnt == 8;
assign data = shreg;
assign rdy = stat;

always @ (posedge clk) begin
    run <= (~RxD) | (~rst | endtick & endbit) & run;
    tick <= (run & ~endtick) ? tick + 1 : 0;
    bitcnt <= (endtick & ~endbit) ? bitcnt + 1 :
        (endtick & endbit) ? 0 : bitcnt;
end

```

```
    shreg <= midtick ? {RxD, shreg[7:1]} : shreg;  
    stat <= (endtick & endbit) ? 1 : (~rst | done) ? 0 : sta  
end  
endmodule
```

SPEC





## OBERON PROGRAMMING LANGUAGE

---

Make it as simple as possible, but not simpler.

– A. Einstein

### A.1 INTRODUCTION

Oberon is a general-purpose programming language that evolved from Modula-2. Its principal new feature is the concept of *type extension*. It permits the construction of new data types on the basis of existing ones and to relate them.

This is not intended as a programmer's tutorial. It is intentionally kept concise. Its function is to serve as a reference for programmers, implementors, and manual writers. What remains unsaid is mostly intentionally left so, either because it is derivable from stated rules of the language, or because it would unnecessarily restrict the freedom of implementors.

This document describes the language defined in 1988/90 as revised in 2007/16.

### A.1.1 *Syntax*

A language is an infinite set of sentences, namely the sentences well formed according to its syntax. In Oberon, these sentences are called *compilation units*. Each unit is a finite sequence of symbols from a finite *vocabulary*. The vocabulary of Oberon consists of identifiers, numbers, strings, operators, delimiters, and comments. They are called *lexical symbols* and are composed of sequences of characters. (Note the distinction between symbols and characters.)

To describe the syntax, an *Extended Backus-Naur Formalism* (EBNF) is used. Brackets [ and ] denote optionality of the enclosed sentential form, and braces { and } denote its repetition (possibly zero times). Syntactic entities (non-terminal symbols) are denoted by English words expressing their intuitive meaning. Symbols of the language vocabulary (terminal symbols) are denoted by strings enclosed in quote marks or by words in capital letters.

### A.1.2 *Vocabulary*

The following lexical rules must be observed when composing symbols. Blanks and line breaks must not occur within symbols (except in comments, or blanks in strings). They are ignored unless they are essential to separate 2 consecutive symbols. Symbols are case-sensitive.

*Identifiers*

Sequences of letters and digits. The 1st must be a letter.

```
digit  = 0-9
letter = A-Z | a-z
id     = letter {letter | digit}
```

Examples:

```
x scan Oberon GetSymbol firstLetter
```

*Numbers*

(Unsigned) integers or real numbers:

- Integers are sequences of digits and may be followed by a suffix letter H, which indicates hexadecimal representation. No suffix represents the default decimal.
- Real numbers always contain the decimal point. Optionally, it may also contain a decimal scale factor: leaded by a E, which is pronounced as "times ten to the power of".

```
dec    = digit {digit}
hexdig = digit | A-F
```

```

hex    = digit {hexdig}
int     = dec | hex H
sign    = + | -
real    = [sign] dec . {digit} [E [sign] dec]
number = int | real

```

Examples:

```

1987          100H = 256
-12.3         4.567E8 = 456700000

```

### *Strings*

Sequences of characters enclosed in a pair of quote marks ("). Strings cannot contain any delimiting quote mark. Alternatively, a single-character string may be specified by the ordinal number of the character in hexadecimal notation followed by an X. The number of characters in a string is called the *length* of the string.

```

string = " {char} " | hex X

```

Examples:

```

"OBERON"      "Don't worry!"      22X

```



*Operators & Delimiters*

Special characters, character pairs, or *reserved words* listed below. Reserved words consist exclusively of capital letters and cannot be used in the role of identifiers.

( )	+ - * /	< = >	. , ;	:	:=	..
[ ]	&   ~ ^	<= # >=	ARRAY	IMPORT		
{ }	DIV	CASE	BEGIN	MODULE		
	END	ELSE	CONST	RECORD		
BY IS	FOR	PROC	ELSIF	REPEAT		
DO OF	MOD	THEN	FALSE	RETURN		
IF OR	NIL	TRUE	UNTIL			
IN TO	VAR	TYPE	WHILE	POINTER		

*Comments*

Arbitrary character sequences opened by the bracket (\* and closed by \*), which may NOT be nested. They may be inserted between any 2 symbols in a program and do not affect its meaning.

## A.2 DECLARATIONS

Every identifier occurring in a program must be introduced by a declaration, unless it is a predefined identifier. Declarations also serve to specify certain permanent properties of

an object, such as whether it is a constant, a type, a variable, or a procedure.

The identifier is then used to refer to the associated object. This is possible in those parts of a program only which are within the scope of the declaration. No identifier may denote more than one object within a given scope. The scope extends textually from the point of the declaration to the end of the block (procedure or module) to which the declaration belongs and hence to which the object is local.

In its declaration, an identifier in the module's scope may be followed by an export mark (\*) to indicate that it be exported from its declaring module. In this case, the identifier may be used in other modules, if they import the declaring module. The identifier is then prefixed by the identifier designating its module (see [A.6](#)). The prefix and the identifier are separated by a period (.) and together are called a *Qualified Identifier*.

```
qid = [id.]id
xid = id[*]
```

The following identifiers are predefined; they are defined in [A.2.2](#) or [A.5.2](#):

ABS	FLT	NEW	SET	INCL	FLOOR
ASR	INC	ODD	BYTE	PACK	ASSERT

CHR	LEN	ORD	CHAR	REAL	BOOL
DEC	LSL	ROR	EXCL	UNPK	INT

### A.2.1 *Constant*

A constant declaration associates an identifier with a constant value.

```
const = xid=expr
```

A constant expression can be evaluated by a *mere textual scan* without actually executing the program. Its operands (see [A.3](#)) are constants. Examples of constant declarations:

```
N      = 100
limit  = 2*N -1
all    = {0 .. WordSize-1}
name   = "Oberon"
```

### A.2.2 *Type*

A data type determines the set of values which variables of that type may assume, and the operators that are applicable. A type declaration is used to associate an identifier with a type. The types define the structure of variables of this type and, by implication, the operators that are applicable to components. There are 2 different data structures, namely *arrays* and *records*, with different component selectors.

```
typedef = xid = type
type = pre | arr | rec | ptr | pro | fun
pre = BOOL | BYTE | CHAR | INT | REAL | SET
```

Examples:

```
Table = ARRAY N OF REAL
Tree  = POINTER TO Node
Node  = RECORD key: INT;
        left, right: Tree
      END

CenterNode = RECORD (Node)
        name: ARRAY 32 OF CHAR;
        subnode: Tree
      END

Function  = PROC(x: INT): INT
```

### *Basic Types*

The following basic types are denoted by predeclared identifiers. The associated operators are defined in [A.3.2](#), and the predeclared function procedures in [A.5.2](#). The values of a given basic type are the following:

```
BOOL TRUE and FALSE
BYTE integers between 0 and 255
```

CHAR characters in a standard charset  
 INT integers  
 REAL real numbers  
 SET sets of integers between 0 and an  
     implementation-dependent max

BYTE is compatible with INT, and vice-versa.

## ARRAY

Structure consisting of a fixed number (*length*) of elements which are all of the same type (*element type*), and designated by indices between 0 and  $length - 1$ :

arr = ARRAY len{, len} OF type

## Declaration

ARRAY N0, N1, ..., Nk OF T

is an abbreviation of

ARRAY N0 OF  
     ARRAY N1 OF  
     ...  
     ARRAY Nk OF T

Examples:

```

ARRAY N OF INT
ARRAY 10, 20 OF REAL

```

### *RECORD*

Structure consisting of a fixed number of elements of any possible type, specifying for each element (*field*) its type and identifier denoting the field, whose scope is the record itself, but is also visible within field designators (see [A.3.1](#)) referring to elements of record variables:

```

rec = RECORD [(base)] [fields] END
base = qid

fields = field {; field}
field = xids: type
xids = xid{, xid}

```

If a record type is exported, field identifiers that are to be visible outside the declaring module must be marked (*public fields*); unmarked ones are *private fields*.

Record types are extensible, i.e. a record type can be defined as an extension of another record type. In the examples above, `CenterNode` (directly) extends `Node`, which is the (di-

rect) base type of `CenterNode`. More specifically, `CenterNode` extends `Node` with the fields `name` and `subnode`.

**Definition:** A type  $T$  *extends* a type  $T_0$ , if it equals  $T_0$ , or directly extends an extension of  $T_0$ . Conversely, a type  $T_0$  is a *base type* of  $T$ , if it equals  $T$ , or is the direct base type of a base type of  $T$ .

Examples:

```
RECORD day, month, year: INT
END
```

```
RECORD
  name, firstname: ARRAY 32 OF CHAR;
  age: INT;
  salary: REAL
END
```

## *POINTER*

Variables of a pointer type  $P$  assume as values pointers to variables of some type  $T$ . It must be a record type. The pointer type  $P$  is said to be bound to  $T$ , and  $T$  is the pointer base type of  $P$ . Pointer types inherit the extension relation of their base types, if there is any. If a type  $T$  is an extension of  $T_0$  and  $P$  is a pointer type bound to  $T$ , then  $P$  is also an extension of  $P_0$ , the pointer type bound to  $T_0$ .

```
ptr = POINTER TO type
```

If a type  $P$  is defined as `POINTER TO  $T$` , the identifier  $T$  can be declared textually following the declaration of  $P$ , but (if so) it must lie within the same scope.

If  $p$  is a variable of type  $P = \text{POINTER TO } T$ , then a call of the predefined procedure `NEW( $p$ )` has the following effect (see [A.5.2](#)): A variable of type  $T$  is allocated in free storage, and a pointer to it is assigned to  $p$ . This pointer  $p$  is of type  $P$  and the referenced variable  $p^{\wedge}$  is of type  $T$ . Failure of allocation results in  $p$  obtaining the value `NIL`. Any pointer variable may be assigned to `NIL`, which points to no variable at all.

## *PROC*

Variables of a procedure type  $T$  have a procedure (or `NIL`) as value. If a procedure  $P$  is assigned to a procedure variable of type  $T$ , the (types of the) parameters of  $P$  must be the same as those indicated in the parameters of  $T$ . The same holds for the result type in the case of a function procedure (see [A.5.1](#)).  $P$  must not be declared local to another procedure, and neither can it be a standard procedure.

```
pro = PROC [params]  
fun = PROC [params]: qid
```



### A.2.3 *Variable*

Variable declarations serve to introduce variables and associate them with identifiers that must be unique within the given scope. They also serve to associate fixed data types with the variables.

```
field = xids: type
```

Variables whose identifiers appear in the same list are all of the same type.

Examples of variable declarations (see examples in [A.2.2](#)):

```
i, j, k: INT
  x, y: REAL
  p, q: BOOL
    s: SET
    f: Function
    a: ARRAY 100 OF REAL
    w: ARRAY 16 OF RECORD
                                ch: CHAR;
                                count: INT
                                END
  t: Tree
```

### A.3 EXPRESSIONS

Constructs denoting rules of computation whereby constants and current values of variables are combined to derive other values by the application of operators and function procedures. Expressions consist of *operands* and *operators*, whose specific associations may be changed by using parentheses.

#### A.3.1 *Operands*

With the exception of sets and literal constants, i.e. numbers and strings, operands are denoted by designators. A *designator* consists of an identifier referring to the constant, variable, or procedure to be designated. This identifier may possibly be qualified by module identifiers (see [A.2](#) and [A.6](#)), and it may be followed by selectors, if the designated object is a structure element.

If  $A$  designates an array, then  $A[E]$  denotes that element of  $A$  whose index is the current value of the expression  $E$ . The type of  $E$  must be of `INT`. A designator of the form  $A[E_1, E_2, \dots, E_n]$  stands for  $A[E_1][E_2] \dots [E_n]$ . If  $p$  designates a pointer variable,  $p^\wedge$  denotes the variable which is referenced by  $p$ . If  $r$  designates a record, then  $r.f$  denotes the field  $f$  of  $r$ . If  $p$  designates a pointer,  $p.f$  denotes the field  $f$  of the record  $p^\wedge$ , i.e. the dot implies dereferencing and  $p.f$  stands for  $p^\wedge.f$ .

The type guard  $v(T_0)$  asserts that  $v$  is of type  $T_0$ , i.e. it aborts program execution, if it is not of type  $T_0$ . The guard is applicable, if

1.  $T_0$  is an extension of the declared type  $T$  of  $v$ , and
2.  $v$  is a variable parameter of record type, or  $v$  is a pointer.

```
des = qid {sel}
sel = . id | '[' exprs ']' | ^ | (qid)
exprs = expr{, expr}
```

If the designated object is a variable, then the designator refers to the variable's current value. If the object is a procedure, a designator without parameter list refers to that procedure. If it is followed by a (possibly empty) parameter list, the designator implies an activation of the procedure and stands for the value resulting from its execution. The (types of the) arguments must correspond to the parameters as specified in the procedure's declaration (see [A.5](#)).

Examples of designators (see examples in [A.2.3](#)):

<code>i</code>	<code>(INT)</code>
<code>a[i]</code>	<code>(REAL)</code>
<code>w[3].ch</code>	<code>(CHAR)</code>
<code>t.key</code>	<code>(INT)</code>

```

t.left.right          (Tree)
t(CenterNode).subnode (Tree)

```

### A.3.2 Operators

The syntax of expressions distinguishes between 4 classes of operators with different precedences (binding strengths). The operator  $\sim$  has the highest precedence, followed by multiplication operators, addition operators, and relations. Operators of the same precedence associate from left to right. For example,  $x-y-z$  stands for  $(x-y)-z$ .

```

expr = exp [rel exp]
rel  = = | # | < | <= | > | >= | IN | IS
exp  = [sign] term {add term}
add  = sign | OR
term = factor {mul factor}
mul  = * | / | DIV | MOD | &
factor = number | string | NIL | TRUE | FALSE
        | set | pc | (expr) | ~factor
set   = '{' [seg {, seg}] '}'
seg   = expr[.. expr]
args  = ( [exprs] )

```

The set  $m..n$  denotes  $m, m+1, \dots, n-1, n$  if  $m \leq n$ , else the empty set. The available operators are listed in the following tables. In some instances, several different operations are

designated by the same operator symbol. In these cases, the actual operation is identified by the type of the operands.

### *Logical Operators*

symbol	result
OR	logical disjunction
&	logical conjunction
~	negation

These operators apply to BOOL operands and yield a BOOL result.

p OR q stands for "if p then TRUE, else q"  
 p & q stands for "if p then q, else FALSE"  
 ~ p stands for "not p"

### *Arithmetic Operators*

symbol	result
+	sum
-	difference
*	product
/	quotient
DIV	int quotient
MOD	modulus

The operators  $+$ ,  $-$ ,  $*$ , and  $/$  apply to operands of numeric types. Both operands must be of the same type, which is also the type of the result. When used as unary operators,  $-$  denotes sign inversion and  $+$  denotes the identity operation.

The operators DIV and MOD apply to integer operands only. Let  $q = x \text{ DIV } y$ , and  $r = x \text{ MOD } y$ . Then quotient  $q$  and remainder  $r$  are defined by the equation:

$$x = qy + r \text{ where } 0 \leq r < y$$

### *Set Operators*

symbol	result
$+$	union
$-$	difference
$*$	intersection
$/$	symmetric set difference

When used with a single operand of type SET, the minus sign denotes the set complement.

### *Relations*

Relations are Boolean. The ordering relations  $<$ ,  $<=$ ,  $>$ ,  $>=$  apply to the numeric types, CHAR, and CHAR arrays. The relations  $=$  and  $\#$  also apply to B00L, SET, pointer and procedure types.

symbol	relation
=	equal
#	unequal
<	less
<=	less or equal
>	greater
>=	greater or equal
IN	set membership
IS	type test

$x \text{ IN } s$  stands for " $x$  is an element of  $s$ ".  $x$  must be of INT, and  $s$  of SET.

$v \text{ IS } T$  stands for " $v$  is of type  $T$ ", which is called *type test*. It is applicable, if

1.  $T$  is an extension of the declared type  $T_0$  of  $v$ , and
2.  $v$  is a variable parameter of record type or  $v$  is a pointer.

Assuming, for instance, that  $T$  is an extension of  $T_0$  and that  $v$  is a designator declared of  $T_0$ , then the test  $v \text{ IS } T$  determines whether the actually designated variable is (not only a  $T_0$ , but also) a  $T$ . The value of  $\text{NIL IS } T$  is undefined.

Examples of expressions (refer to examples in [A.2.3](#)):

<code>i DIV 3</code>	<code>(INT)</code>
<code>~p OR q</code>	<code>(BOOL)</code>
<code>(i+j) * (i-j)</code>	<code>(INT)</code>
<code>s - {8, 9, 13}</code>	<code>(SET)</code>
<code>a[i+j] * a[i-j]</code>	<code>(REAL)</code>
<code>(0&lt;=i) &amp; (i&lt;100)</code>	<code>(BOOL)</code>
<code>t.key = 0</code>	<code>(BOOL)</code>
<code>k IN {i .. j-1}</code>	<code>(BOOL)</code>
<code>t IS CenterNode</code>	<code>(BOOL)</code>

#### A.4 STATEMENTS

Denote actions. There are statements of:

**ELEMENTARY** Simplest, not composed of other ones:

- the assignment, and
- procedure call.

**STRUCTURED** Composed of others, expressing sequencing and execution of:

- conditional,
- selective, and
- repetitive.

**EMPTY** No actions, relaxing punctuation rules in [A.4.1](#).



`s = [assign | pc | if | cases | while | repeat | for]`

#### A.4.1 *Statement Sequences*

Denote the sequence of actions specified by component statements, separated by semicolons:

`ss = s{; s}`

#### A.4.2 *Assignments*

Assign variables new values specified by the expression following the assignment operator, `:=`, read as "becomes":

`assign = des:=expr`

If a value parameter is structured (of array or record type), no assignment to it or its elements is permitted. Neither may assignments be made to imported variables.

The type of the expression must be the same as of the designator. The following exceptions hold:

1. The constant NIL can be assigned to variables of any pointer or procedure type.
2. Strings can be assigned to any array of CHAR, provided the number of characters in the string is less than that of the array. (A null character is appended).

Single-character strings can also be assigned to variables of type CHAR.

3. In the case of records, the type of the source must be an extension of the one of the destination.
4. An open array may be assigned to an array of equal base type.

Examples of assignments (see examples in [A.2.3](#)):

```
i := 0
p := i = j
x := FLT(i + 1)
k := (i + j) DIV 2
f := log2
s := {2, 3, 5, 7, 11, 13}
a[i] := (x+y) * (x-y)
t.key := i
w[i+1].ch := "A"
```

#### A.4.3 *Procedure Calls*

A procedure call activates a procedure. It may contain arguments to be substituted in place of their corresponding parameters defined in the procedure declaration (see [A.5](#)). The correspondence is established by their positions in the lists of arguments/parameters, respectively. There are 2 kinds of parameters:

**VARIABLE** arguments must be designators denoting variables. If it designates an element of a structured variable, selectors are evaluated before procedure execution.

**VALUE** the corresponding argument must be an expression, evaluated prior to the procedure activation. The resulting value is assigned to the parameter constituting a local variable (see also [A.5.1](#)).

```
pc = des [args]
```

Examples of procedure calls: (see [A.5](#))

```
ReadInt(i)
WriteInt(2*j + 1, 6)
INC(w[k].count)
```

#### A.4.4 IF

Specify the conditional execution of statements guarded by preceding boolean expressions, which is called *guards*. They are evaluated in sequence of occurrence, until one is TRUE, whereafter its guarded statement sequence will be executed. If no guards satisfied, the ELSE one is executed (if there is).

```
if = IF expr THEN ss
```

```
{ELSIF expr THEN ss}  
      [ELSE ss]  
END
```

Example:

```
IF (ch >= "A") & (ch <= "Z") THEN readIdentifier  
ELSIF (ch >= "0") & (ch <= "9") THEN readNumber  
ELSIF ch = 22X THEN readString  
END
```

#### A.4.5 CASE

Specify the selection and execution of a statement sequence according to the value of an expression. First the case expression is evaluated, then the statement sequence is executed whose case label list contains the obtained value. If the case expression is of INT or CHAR, all labels must be integers or single-character strings, respectively.

```
cases = CASE expr OF case {'|' case} END  
case  = [ranges: ss]  
  
ranges = range {, range}  
range  = label [.. label]  
label  = int | string | qid
```

Example:

```

CASE k OF 0: x := x + y
          | 1: x := x - y
          | 2: x := x * y
          | 3: x := x / y
END

```

The type  $T$  of the case expression (case variable) may also be a record or pointer type. Then the case labels must be extensions of  $T$ , and in the statements  $S_i$  labelled by  $T_i$ , the case variable is considered as of type  $T_i$ . Example:

```

TYPE R = RECORD      a: INT END;
  R0 = RECORD (R) b: INT END;
  R1 = RECORD (R) b: REAL END;
  R2 = RECORD (R) b: SET END;
  P = POINTER TO R;
  P0 = POINTER TO R0;
  P1 = POINTER TO R1;
  P2 = POINTER TO R2;
VAR p: P;

CASE p OF
  P0: p.b := 10 |
  P1: p.b := 2.5 |
  P2: p.b := {0, 2}
END

```

#### A.4.6 WHILE

Specify repetition: if any of boolean expressions (*guards*) yields TRUE, the corresponding statement sequence will be executed. It ends while NONE yields TRUE:

```
while = WHILE expr DO ss
        {ELSIF expr DO ss} END
```

Examples:

```
WHILE j > 0 DO
  j := j DIV 2; i := i+1
END
```

```
WHILE (t # NIL) & (t.key # i) DO
  t := t.left
END
```

```
WHILE m > n DO m := m - n
ELSIF n > m DO n := n - m
END
```

#### A.4.7 REPEAT

Specify repeated execution of a statement sequence until conditions satisfied:

repeat = REPEAT ss UNTIL expr

The sequence is executed at least once.

#### A.4.8 FOR

Specify repeated execution of a statement sequence for the *control variable* be progressed at the specified (integer) T0 value:

for = FOR id := expr T0 expr [BY expr] DO ss END

FOR v := beg T0 end BY step DO ss END

is equivalent to

```
v := beg;
WHILE v <= end DO ss; v := v + step END
```

when  $step > 0$ . While  $step < 0$ , it is equivalent to

```
v := beg;
WHILE v >= end DO ss; v := v + step END
```

$v$ ,  $beg$ ,  $end$ ,  $step$  must all be of INT. Moreover, the  $step$  must NOT be zero. Its default is 1, if not specified.

## A.5 PROCEDURES

Procedure declarations consist of:

- a heading, specifies the
  - procedure identifier,
  - parameters, and
  - result type (if any).
- a body, contains
  - declarations, and
  - statements.

A repeated procedure identifier following END at the end finalizes the whole declaration.

There are 2 kinds of procedures, namely:

1. proper procedures, and
2. function procedures.

The latter are activated by a function designator as a constituent of an expression, and yield a result that is an operand in the expression. Proper procedures are activated by a procedure call. A function procedure is distinguished in the declaration by indication of the type of its result follow-



ing the parameter list. Its body must end with a RETURN clause which defines the result of the function procedure.

All constants, variables, types, and procedures declared within a procedure body are local to the procedure. The values of local variables are undefined upon entry to the procedure. Since procedures may be declared as local objects too, procedure declarations may be nested.

In addition to its parameters and locally declared objects, the objects declared globally are also visible in the procedure.

The use of the procedure identifier in a call within its declaration implies recursive activation of the procedure.

```

proc = prop | func
prop = PROC xid [params]; decls
      [BEGIN ss] END id
func = PROC xid [params]: qid; decls
      [BEGIN ss] [RETURN expr] END id
decls = [CONST {const;}]
        [TYPE {typedef;}]
        [VAR {field;}]
        {proc;}

```

#### A.5.1 *Parameters*

Parameters are identifiers which denote arguments specified in the procedure call. The correspondence between param-

eters and arguments is established when the procedure is called. There are two kinds of parameters, namely value and variable parameters. A variable parameter corresponds to an argument that is a variable, and it stands for that variable. A value parameter corresponds to an argument that is an expression, and it stands for its value, which cannot be changed by assignment. However, if a value parameter is of a basic type, it represents a local variable to which the value of the actual expression is initially assigned.

The kind of a parameter is indicated in the parameter list: Variable parameters are denoted by the symbol VAR and value parameters by the absence of a prefix.

A function procedure without parameters must have an empty parameter list. It must be called by a function designator whose argument list is empty too.

Parameters are local to the procedure, i.e. their scope is the program text which constitutes the procedure declaration.

```
params = ( [param{; param}] )  
param = [VAR] id{, id}: {ARRAY OF} qid
```

The type of each parameter is specified in the parameter list. For variable parameters, it must be identical to the corresponding argument's type, except in the case of a record, where it must be a base type of the corresponding argument's type.

If the parameter's type is specified as

```
ARRAY OF T
```

the parameter is said to be an open array, and the corresponding argument may be of arbitrary length.

If a parameter specifies a procedure type, then the corresponding argument must be either a procedure declared globally, or a variable (or parameter) of that procedure type. It cannot be a predefined procedure. The result type of a procedure can be neither a record nor an array.

Examples of procedure declarations:

```
PROC ReadInt(VAR x: INT);
  VAR i : INT; ch: CHAR;
BEGIN i := 0; Read(ch);
  WHILE ("0" <= ch) & (ch <= "9") DO
    i := 10*i + (ORD(ch)-ORD("0")); Read(ch)
  END ;
  x := i
END ReadInt
```

```
PROC WriteInt(x: INT); (* 0 <= x < 10^5 *)
  VAR i: INT;
  buf: ARRAY 5 OF INT;
BEGIN i := 0;
```

```

    REPEAT buf[i] := x MOD 10; x := x DIV 10;
        INC(i) UNTIL x = 0;
    REPEAT DEC(i); Write(CHR(buf[i] + ORD("0")))
        UNTIL i = 0
END WriteInt

PROC log2(x: INT): INT;
    VAR y: INT; (* assume x>0 *)
BEGIN y := 0;
    WHILE x > 1 DO x := x DIV 2; INC(y) END;
    RETURN y
END log2

```

### A.5.2 *Predefined*

The following table lists the predefined functions. Some are generic, i.e. they apply to several types of operands.  $v$  stands for a variable,  $x$  and  $n$  for expressions, and  $T$  for a type.

Name	ArgType	ResType	Function
ABS( $x$ )	numeric	type( $x$ )	absolute value
ODD( $x$ )	INT	BOOL	$x \bmod 2 = 1$
LEN( $v$ )	array	INT	length
LSL( $x, n$ )	INT	INT	logical shift left
ASR( $x, n$ )	INT	INT	signed shift right
ROR( $x, n$ )	INT	INT	rotated right

Type conversion functions:

Name	ArgType	ResType	Function
FLOOR(x)	REAL	INT	round down
FLT(x)	INT	REAL	identity
ORD(x)	CHAR   BOOL   SET	INT	ordinal # of
CHR(x)	INT	CHAR	character of

Prop procedures:

Name	ArgType	Procedure
INC(v)	INT	$v := v + 1$
INC(v,n)	INT	$v := v + n$
DEC(v)	INT	$v := v - 1$
DEC(v,n)	INT	$v := v - n$
INCL(v,x)	SET,INT	$v := v + x$
EXCL(v,x)	SET,INT	$v := v - x$
NEW(v)	pointer type	allocate $v^{\wedge}$
ASSERT(b)	BOOL	abort if $\sim b$
PACK(x,n)	REAL,INT	pack into x
UNPK(x,n)	REAL,INT	unpack x

The FLOOR(x) yields the largest int not greater than x:

$$\text{FLOOR}(1.5) = 1$$

$$\text{FLOOR}(-1.5) = -2$$

The parameter  $n$  of `PACK` represents the exponent of  $x$ . `PACK( $x$ ,  $y$ )` is equivalent to  $x := x * 2^y$ . `UNPK` is the reverse operation. The resulting  $x$  is normalized, such that  $1.0 \leq x < 2.0$ .

## A.6 MODULES

A module is a collection of declarations of constants, types, variables, and procedures, and a sequence of statements for the purpose of assigning initial values to the variables. A module typically constitutes a text that is compilable as a unit.

```

module = MODULE id;
        [IMPORT import{, import};]
        decls
        [BEGIN ss]
        END id.
import = id[:= id]

```

The import list specifies the modules of which the module is a client. If an identifier  $x$  is exported from a module  $M$ , and if  $M$  is listed in a module's import list, then  $x$  is referred to as  $M.x$ . If the form " $M := M1$ " is used in the import list, an exported object  $x$  declared within  $M1$  is referenced in the importing module as  $M.x$ .

Identifiers that are to be visible in client modules, i.e. which are to be exported, must be marked by an asterisk (export

mark) in their declaration. Variables are always exported in read-only mode.

The statement sequence following the symbol `BEGIN` is executed when the module is added to a system (loaded). Individual (parameterless) procedures can thereafter be activated from the system, and these procedures serve as commands. Example:

```
MODULE Out; (*exported: Write, WriteInt, WriteLn*)
  IMPORT Texts, Oberon;
  VAR W: Texts.Writer;

  PROC Write*(ch: CHAR);
  BEGIN Texts.Write(W, ch)
  END Write;

  PROC WriteInt*(x, n: INT);
    VAR i: INT; a: ARRAY 16 OF CHAR;
  BEGIN i := 0;
    IF x < 0 THEN Texts.Write(W, "-"); x := -x END ;
    REPEAT a[i] := CHR(x MOD 10 + ORD("0"));
      x := x DIV 10; INC(i) UNTIL x = 0;
    REPEAT Texts.Write(W, " "); DEC(n) UNTIL n <= i;
    REPEAT DEC(i); Texts.Write(W, a[i]) UNTIL i = 0
  END WriteInt;

  PROC WriteLn*;
```

```
BEGIN Texts.WriteLine(W);  
    Texts.Append(Oberon.Log, W.buf)  
END WriteLine;  
BEGIN Texts.OpenWriter(W)  
END Out.
```

#### A.6.1 *SYSTEM*

The optional module *SYSTEM* contains definitions that are necessary to program low-level operations referring directly to resources particular to a given computer and/or implementation.

These include for example facilities for accessing devices that are controlled by the computer, and perhaps facilities to break the data type compatibility rules otherwise imposed by the language definition.

There are 2 reasons for providing facilities in *SYSTEM*:

1. Their value is implementation-dependent, i.e., it is not derivable from the language's definition, and
2. they may corrupt a system (e.g. PUT).

It is strongly recommended to restrict their use to specific low-level modules, as such modules are inherently non-portable and not "type-safe". However, they are easily recognized due to the identifier *SYSTEM* appearing in the module's



import lists. The subsequent definitions are generally applicable. However, individual implementations may include in their module `SYSTEM` additional definitions that are particular to the specific, underlying computer. In the following,  $v$  stands for a variable,  $x$ ,  $a$ , and  $n$  for expressions.

Name	ArgType	ResType	Function
ADR( $v$ )	any	INT	address of variable
SIZE( $T$ )	any	INT	size in bytes
BIT( $a, n$ )	INT	BOOL	bit $n$ of <code>mem[a]</code>

Name	ArgType	Procedure
GET( $a, v$ )	INT,	$v := \text{mem}[a]$
PUT( $a, x$ )	any basic	<code>mem[a] := x</code>
COPY( $src$ , $dst, n$ )	INT	copy $n$ consecutive words from $src$ to $dst$

The following are additional functions and procedures accepted by the compiler for the RISC processor:

Name	ArgType	ResType	Function
VAL( $T, n$ )	scalar	T	identity
ADC( $m, n$ )	INT	INT	add with carry C
SBC( $m, n$ )	INT	INT	subtract with carry C
UML( $m, n$ )	INT	INT	unsigned multiplication
COND( $n$ )	INT	BOOL	IF Cond( $n$ ) THEN ...

Name	ArgType	Procedure
LED(n)	INT	display n on LEDs

SYNTAX OF OBERON

---

```
digit  = 0-9
hexdig = digit|A-F
letter = A-Z|a-z

id      = letter{letter|digit}
qid     = [id.]id
xid     = id[*]

dec     = digit{digit}
hex     = digit{hexdig}
int     = dec | hex H
sign    = +|-
real    = [sign]dec.{digit}[E[sign]dec]
number = int|real
string = " {char} " | hex X

const   = xid=expr

typedef = xid=type
type    = pre | arr | rec | ptr | pro | fun
```

```

pre      = BOOL|BYTE|CHAR|INT|REAL|SET
arr      = ARRAY len{, len} OF type
len      = expr
rec      = RECORD[ (base)][ fields] END
base     = qid
fields   = field{; field}
field    = xids: type
xids     = xid{, xid}
ptr      = POINTER TO type
pro      = PROC[ params]
fun      = PROC[ params]: qid

expr     = exp [rel exp]
rel      = =|<|<=|>|>=|IN|IS
exp      = [sign]term {add term}
add      = sign|OR
term     = factor {mul factor}
mul      = *|/|DIV|MOD|&
factor   = number | string | NIL | TRUE | FALSE
          | set | pc | (expr) | ~factor
des      = qid{sel}
sel      = .id|['exprs']|^|(qid)
set      = '{' [seg{, seg}] '}'
seg      = expr[..expr]
exprs    = expr{, expr}
args     = ([exprs])

s        = [assign|pc|if|cases|while|repeat|for]

```

```

assign = des:=expr
pc      = des[args]
ss      = s{; s}
if      =      IF expr THEN ss
          {ELSIF expr THEN ss}
          [ELSE ss] END
cases   = CASE expr OF case {'|' case} END
case    = [ranges: ss]
ranges  = range{, range}
range   = label[..label]
label   = int|string|qid
while   = WHILE expr DO ss {ELSIF expr DO ss} END
repeat  = REPEAT ss UNTIL expr
for      = FOR id:=expr TO expr[ BY expr] DO ss END

proc    = prop | func
prop    = PROC xid [params]; decls
          [BEGIN ss] END id
func    = PROC xid [params]: qid; decls
          [BEGIN ss] [RETURN expr] END id
decls   = [CONST {const;}] [TYPE {typdef;}] [VAR {field;}]
          {proc;}
params  = ([param{; param}])
param   = [VAR] id{, id}: {ARRAY OF} qid

module  = MODULE id;[ IMPORT import{, import};] decls
          [BEGIN ss] END id.
import  = id[:= id]

```